



Confluent Developer Training: Building Kafka Solutions

Kafka Fundamentals

Chapter 3



Course Contents

01: Introduction

02: The Motivation for Apache Kafka

>>> 03: Kafka Fundamentals

04: Kafka's Architecture

05: Developing With Kafka

06: More Advanced Kafka Development

07: Schema Management In Kafka

08: Kafka Connect for Data Movement

09: Basic Kafka Administration

10: Kafka Stream Processing

11: Conclusion

Appendix A: Installation Recommendations

Kafka Fundamentals


- **In this chapter you will learn:**
 - How Producers write data to a Kafka cluster
 - How data is divided into Partitions, and then stored on Brokers
 - How Consumers read data from the cluster

Kafka Fundamentals

- **An Overview of Kafka**
- *Kafka Producers*
- *Kafka Brokers*
- *Kafka Consumers*
- *Kafka's Use of ZooKeeper*
- *Comparisons with Traditional Message Queues*
- *Hands-On Exercise: Using Kafka's Command-Line Tools*
- *Chapter Review*


Reprise: A Very High-Level View of Kafka

- *Producers* send data to the Kafka cluster
- *Consumers* read data from the Kafka cluster
- *Brokers* are the main storage and messaging components of the Kafka cluster



Kafka Messages

- The basic unit of data in Kafka is a *message*
 - *Message* is sometimes used interchangeably with *record*
 - Producers write messages to Brokers
 - Consumers read messages from Brokers

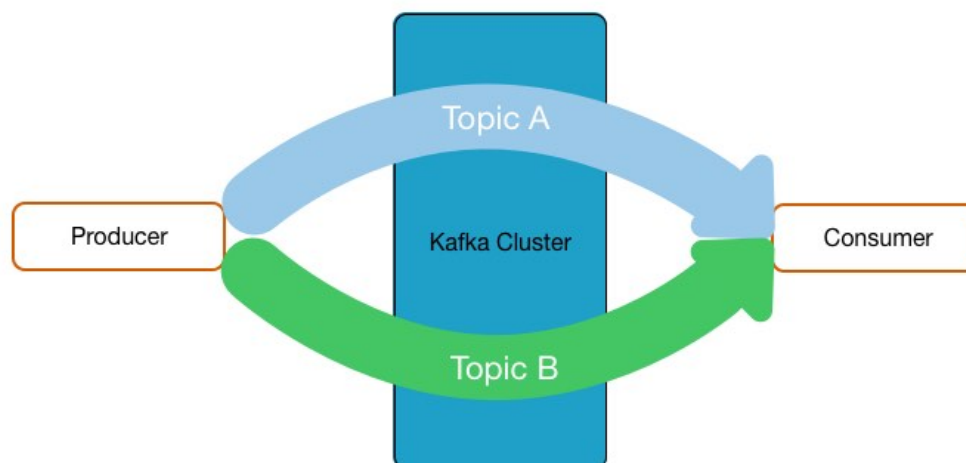


Key-Value Pairs

- **A message is a key-value pair**
 - All data is stored in Kafka as byte arrays
 - Producer provides serializers to convert the key and value to byte arrays
 - Key and value can be any data type

Topics


- **Kafka maintains streams of messages called *Topics***
 - Logical representation
 - They categorize messages into groups



- **Developers decide which Topics exist**
 - By default, a Topic is auto-created when it is first used
- **One or more Producers can write to one or more Topics**
- **There is no limit to the number of Topics that can be used**

Partitioned Data Ingestion

- **Producers shard data over a set of *Partitions***
 - Each Partition contains a subset of the Topic's messages
 - Each Partition is an ordered, immutable log of messages
- **Partitions are distributed across the Brokers**
- **Typically, the message key is used to determine which Partition a message is assigned to**



Load Balancing and Semantic Partitioning

- **Producers use a partitioning strategy to assign each message to a Partition**
- **Having a partitioning strategy serves two purposes**
 - Load balancing: shares the load across the Brokers
 - Semantic partitioning: user-specified key allows locality-sensitive message processing
- **The partitioning strategy is specified by the Producer**
 - Default strategy is a hash of the message key
 - `hash(key) % number_of_partitions`
 - If a key is not specified, messages are sent to Partitions on a round-robin basis
- **Developers can provide a custom partitioner class**

Kafka Components

- **There are four key components in a Kafka system**
 - Producers
 - Brokers
 - Consumers
 - ZooKeeper
- **We will now investigate each of these in turn**

Kafka Fundamentals

- *An Overview of Kafka*
- **Kafka Producers**
- *Kafka Brokers*
- *Kafka Consumers*
- *Kafka's Use of ZooKeeper*
- *Comparisons with Traditional Message Queues*
- *Hands-On Exercise: Using Kafka's Command-Line Tools*
- *Chapter Review*

Producer Basics


- **Producers write data in the form of *messages* to the Kafka cluster**
- **Producers can be written in any language**
 - Native Java, C/C++, Python, Go, .NET, JMS clients are supported by Confluent
 - Clients for many other languages exist
 - Confluent develops and supports a REST (REpresentational State Transfer) server which can be used by clients written in any language
- **A command-line Producer tool exists to send messages to the cluster**
 - Useful for testing, debugging, etc.

Kafka Fundamentals

- *An Overview of Kafka*
- *Kafka Producers*
- **Kafka Brokers**
- *Kafka Consumers*
- *Kafka's Use of ZooKeeper*
- *Comparisons with Traditional Message Queues*
- *Hands-On Exercise: Using Kafka's Command-Line Tools*
- *Chapter Review*

Broker Basics

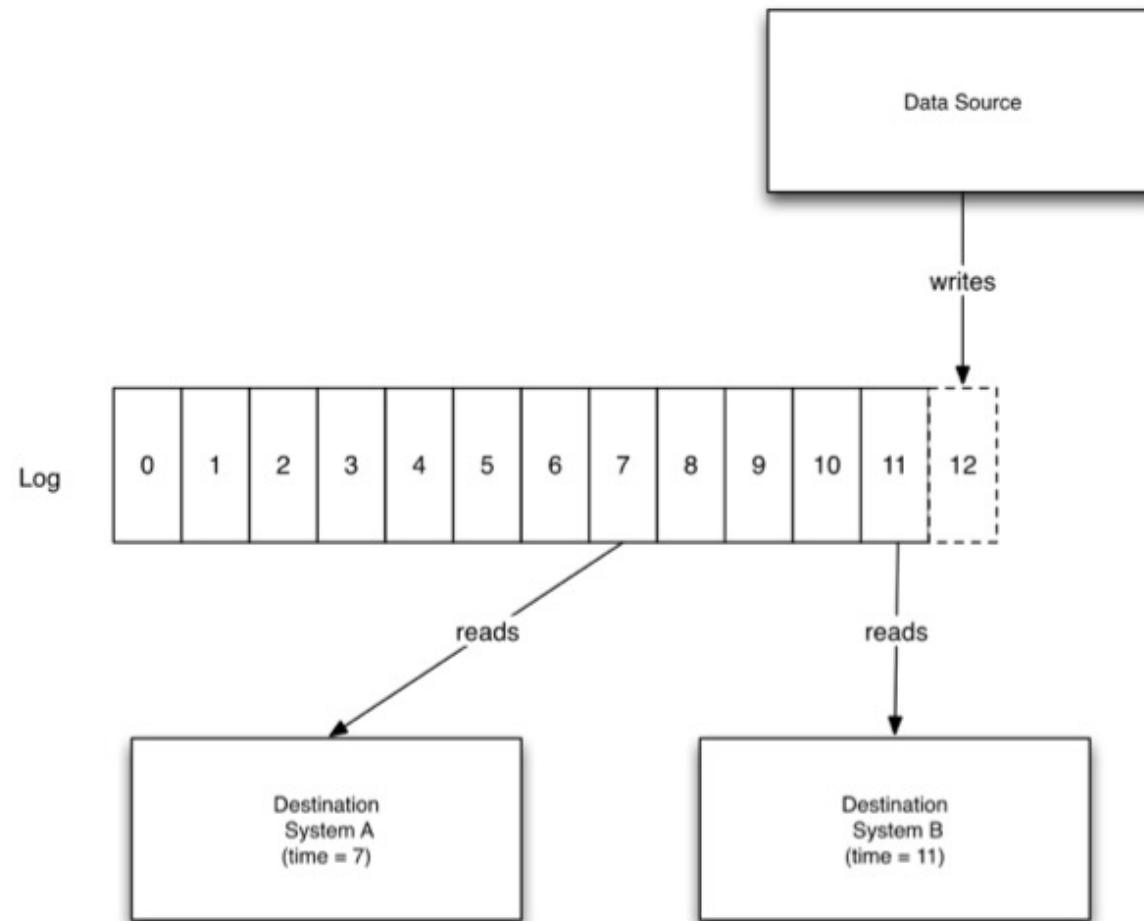
- **Brokers receive and store messages when they are sent by the Producers**
- **A production Kafka cluster will have three or more Brokers**
 - Each can handle hundreds of thousands, or millions, of messages per second



Brokers Manage Partitions

- **Messages in a Topic are spread across Partitions in different Brokers**
- **Typically, a Broker manages multiple Partitions**
- **Each Partition is stored on the Broker's disk as one or more log files**
 - Not to be confused with log4j files used for monitoring
- **Each message in the log is identified by its *offset* number**
 - A monotonically increasing value
- **Kafka provides a configurable retention policy for messages to manage log file growth**
 - Retention policies can be configured per Topic

Messages are Stored in a Persistent Log

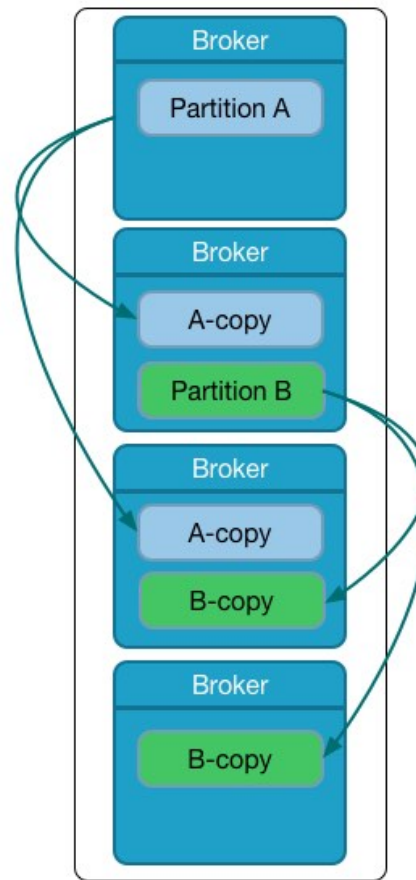


Metadata in a Kafka Message

- **A Kafka Message contains data and metadata**
 - Key/value pair
 - Offset
 - Timestamp
 - Compression type
 - Magic byte
 - Optional message headers API
 - Application teams can add custom key-value paired metadata to messages
 - Additional fields to support batching, exactly once semantics, replication protocol
- **Latest message format:**
<http://kafka.apache.org/documentation.html#messageformat>

Fault Tolerance via a Replicated Log

- Partitions can be replicated across multiple Brokers
- Replication provides fault tolerance in case a Broker goes down
 - Kafka automatically handles the replication



Kafka Fundamentals


- *An Overview of Kafka*
- *Kafka Producers*
- *Kafka Brokers*
- **Kafka Consumers**
- *Kafka's Use of ZooKeeper*
- *Comparisons with Traditional Message Queues*
- *Hands-On Exercise: Using Kafka's Command-Line Tools*
- *Chapter Review*

Consumer Basics


- **Consumers pull messages from one or more Topics in the cluster**
 - As messages are written to a Topic, the Consumer will automatically retrieve them
- **The *Consumer Offset* keeps track of the latest message read**
 - If necessary, the Consumer Offset can be changed
 - For example, to reread messages
- **The Consumer Offset is stored in a special Kafka Topic**
- **A command-line Consumer tool exists to read messages from the cluster**
 - Useful for testing, debugging, etc.

Distributed Consumption

- **Different Consumers can read data from the same Topic**
 - By default, each Consumer will receive all the messages in the Topic
- **Multiple Consumers can be combined into a *Consumer Group***
 - Consumer Groups provide scaling capabilities
 - Each Consumer is assigned a subset of Partitions for consumption



The Result: A Linearly Scalable Firehose



Kafka Fundamentals

- *An Overview of Kafka*
- *Kafka Producers*
- *Kafka Brokers*
- *Kafka Consumers*
- **Kafka's Use of ZooKeeper**
- *Comparisons with Traditional Message Queues*
- *Hands-On Exercise: Using Kafka's Command-Line Tools*
- *Chapter Review*


What is ZooKeeper?

- **ZooKeeper is a centralized service that can be used by distributed applications**
 - Open source Apache project
 - Enables highly reliable distributed coordination
 - Maintains configuration information
 - Provides distributed synchronization
- **Used by many projects**
 - Including Kafka and Hadoop
- **Typically consists of three or five servers in an *ensemble***
 - This provides resiliency should a machine fail



How Kafka Uses ZooKeeper

- **Kafka Brokers use ZooKeeper for a number of important internal features**
 - Cluster management
 - Failure detection and recovery
 - Access Control List (ACL) storage



Quiz: Question

- **Provide the correct relationship — 1:1, 1:N, N:1, or N:N**
 - Broker to Partition — ?
 - Key to Partition — ?
 - Producer to Topic — ?
 - Consumer Group to Topic — ?
 - Consumer (in a Consumer Group) to Partition — ?

Quiz: Answer

- **Provide the correct relationship — 1:1, 1:N, N:1, or N:N**
 - Broker to Partition — N:N
 - Key to Partition — N:1
 - Producer to Topic — N:N
 - Consumer Group to Topic — N:N
 - Consumer (in a Consumer Group) to Partition — 1:N

Kafka Fundamentals

- *An Overview of Kafka*
- *Kafka Producers*
- *Kafka Brokers*
- *Kafka Consumers*
- *Kafka's Use of ZooKeeper*
- **Comparisons with Traditional Message Queues**
- *Hands-On Exercise: Using Kafka's Command-Line Tools*
- *Chapter Review*

Kafka's Benefits over Traditional Message Queues

- **Durability and Availability**

- Messages are replicated on multiple machines for reliability
- Cluster can handle Broker failures

- **Excellent scalability**

- Even a small cluster can process a large volume of messages
 - Tests have shown that three low-end machines can easily deal with two million writes per second

- **Very high throughput**

- **Supports both real-time and batch consumption**

- **Data retention**

Multi-Subscription and Scalability

- **Multi-subscription provides easy distribution of data**
 - Once the data is in Kafka, it can be read by multiple different Consumers
 - For instance, a Consumer which writes the data to the Hadoop Distributed File System (HDFS), another to do real-time analysis on the data, etc.
- **Multiple Brokers, multiple Topics, and Consumer Groups provide very high scalability**
 - Kafka was designed as a distributed system from the ground up
 - Enables parallelism

The Advantages of a Pull Architecture


- **Kafka Consumers pull messages from the Brokers**
 - This is in contrast to some other systems, which use a *push* design
- **The advantages of pulling, rather than pushing, data, include:**
 - The ability to add more Consumers to the system without reconfiguring the cluster
 - The ability for a Consumer to go offline and return later, resuming from where it left off
 - No problems with the Consumer being overwhelmed by data
 - It can pull, and process, the data at whatever speed it needs to
 - A slow Consumer will not affect Producers

The Page Cache for High Performance

- Unlike some systems, Kafka itself does not require a lot of RAM
- Logs are held on disk and read when required
- Kafka makes use of the operating system's page cache to hold recently-used data
 - Typically, recently-Produced data is the data which Consumers are requesting
- A Kafka Broker running on a system with a reasonable amount of RAM for the OS to use as cache will typically be able to swamp its network connection
 - In other words the network, not Kafka itself, will be the limiting factor on the speed of the system

Speeding Up Data Transfer

- Kafka uses zero-copy data transfer (Broker → Consumer)



Kafka Fundamentals

- *An Overview of Kafka*
- *Kafka Producers*
- *Kafka Brokers*
- *Kafka Consumers*
- *Kafka's Use of ZooKeeper*
- *Comparisons with Traditional Message Queues*
- **Hands-On Exercise: Using Kafka's Command-Line Tools**
- *Chapter Review*

Hands-On Exercise: Using Kafka's Command-Line Tools

- In this Hands-On Exercise you will use Kafka's command-line tools to Produce and Consume data
- Please refer to the Hands-On Exercise Manual

Kafka Fundamentals

- *An Overview of Kafka*
- *Kafka Producers*
- *Kafka Brokers*
- *Kafka Consumers*
- *Kafka's Use of ZooKeeper*
- *Comparisons with Traditional Message Queues*
- *Hands-On Exercise: Using Kafka's Command-Line Tools*
- **Chapter Review**

Chapter Review

- **A Kafka system is made up of Producers, Consumers, and Brokers**
 - ZooKeeper provides co-ordination services for the Brokers
- **Producers write messages to Topics**
 - Topics are broken down into partitions for scalability
- **Consumers read data from one or more Topics**

Kafka's Architecture

Chapter 4



Course Contents

01: Introduction

02: The Motivation for Apache Kafka

03: Kafka Fundamentals

>>> 04: Kafka's Architecture

05: Developing With Kafka

06: More Advanced Kafka Development

07: Schema Management In Kafka

08: Kafka Connect for Data Movement

09: Basic Kafka Administration

10: Kafka Stream Processing

11: Conclusion

Appendix A: Installation Recommendations

Kafka's Architecture

- How Kafka's log files are stored on the Kafka Brokers
- How Kafka uses replicas for reliability
- How Consumer Groups and Partitions provide scalability

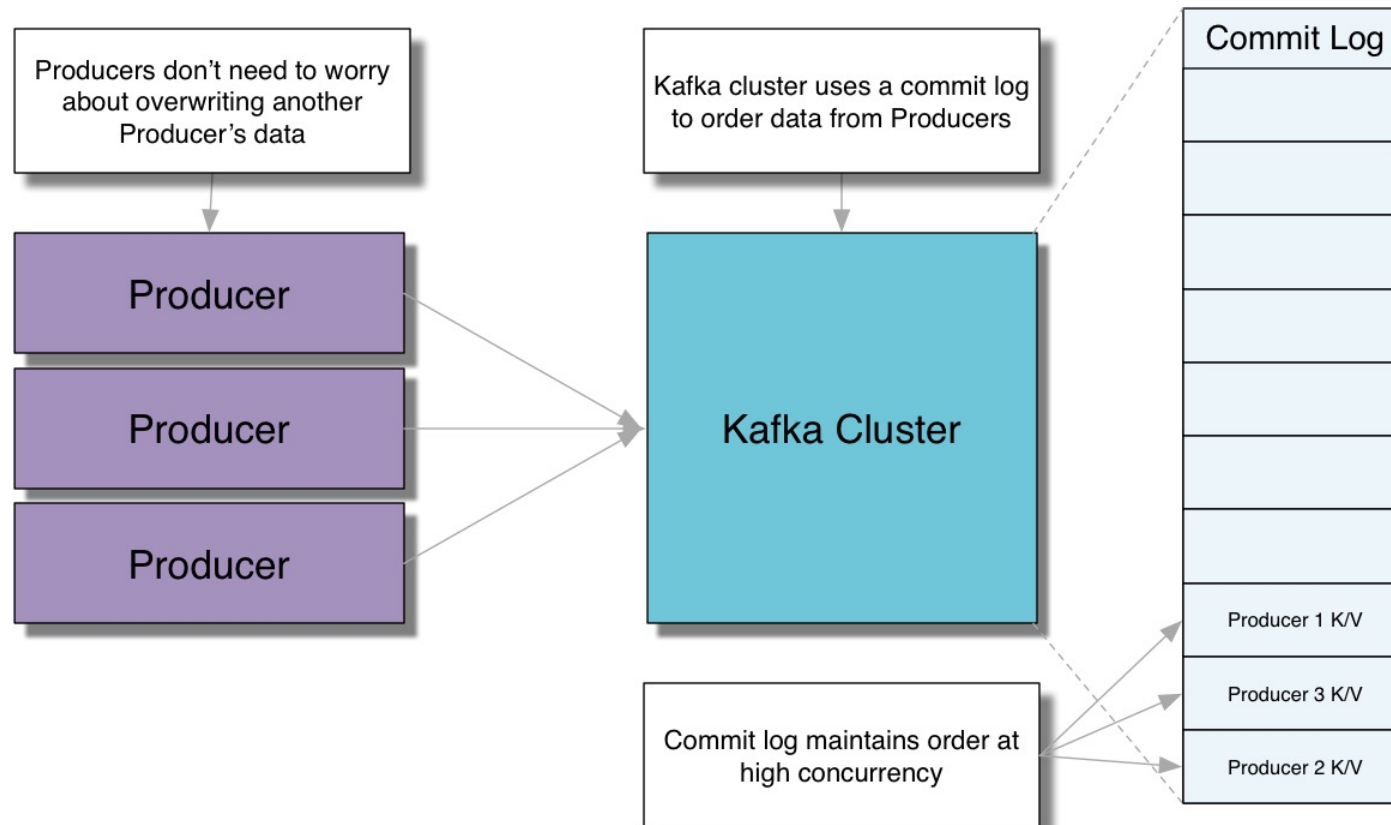
Kafka's Architecture

- **Kafka's Log Files**
- *Replicas for Reliability*
- *Partitions and Consumer Groups for Scalability*
- *Hands-On Exercise: Consuming from Multiple Partitions*
- *Chapter Review*

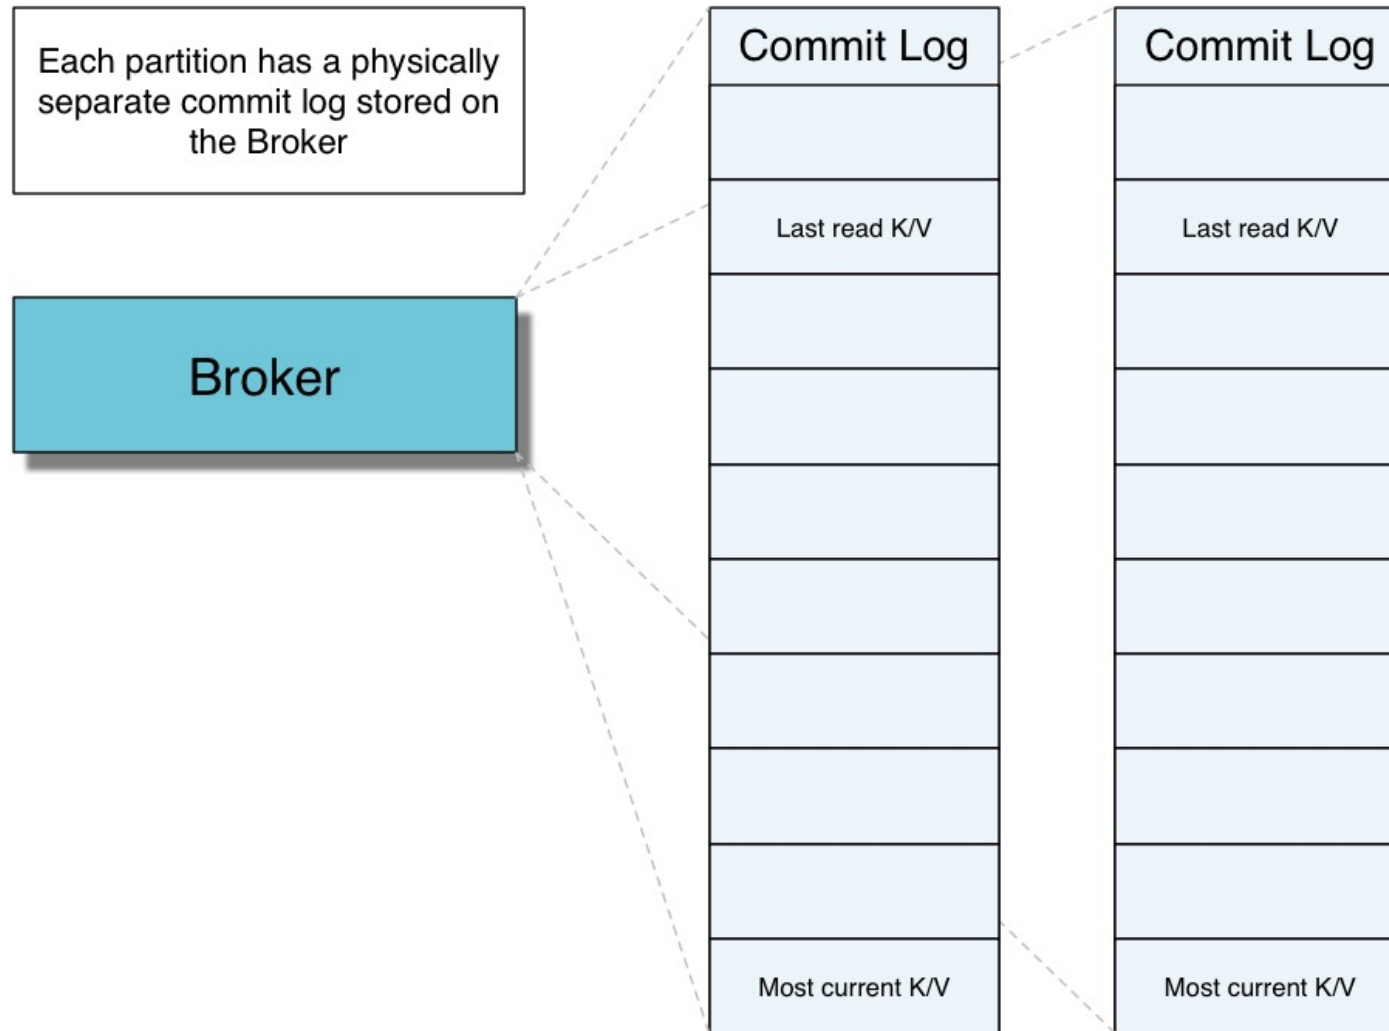
What is a Commit Log?

- A Commit Log is a way to keep track of changes as they happen
- Commonly used by databases to keep track of all changes to tables
- Kafka uses commit logs to keep track of all messages in a particular Topic
 - Consumers can retrieve previous data by backtracking through the commit log

The Commit Log for High Concurrency



Partitions Are Stored as Separate Logs



Kafka's Architecture


- *Kafka's Log Files*
- **Replicas for Reliability**
- *Partitions and Consumer Groups for Scalability*
- *Hands-On Exercise: Consuming from Multiple Partitions*
- *Chapter Review*

Problems With our Current Model

- **So far, we have said that each Broker manages one or more Partitions for a Topic**
- **This does not provide reliability**
 - A Broker failing would result in all of those Partitions being unavailable
- **Kafka takes care of this by replicating each partition**
 - The replication factor is configurable

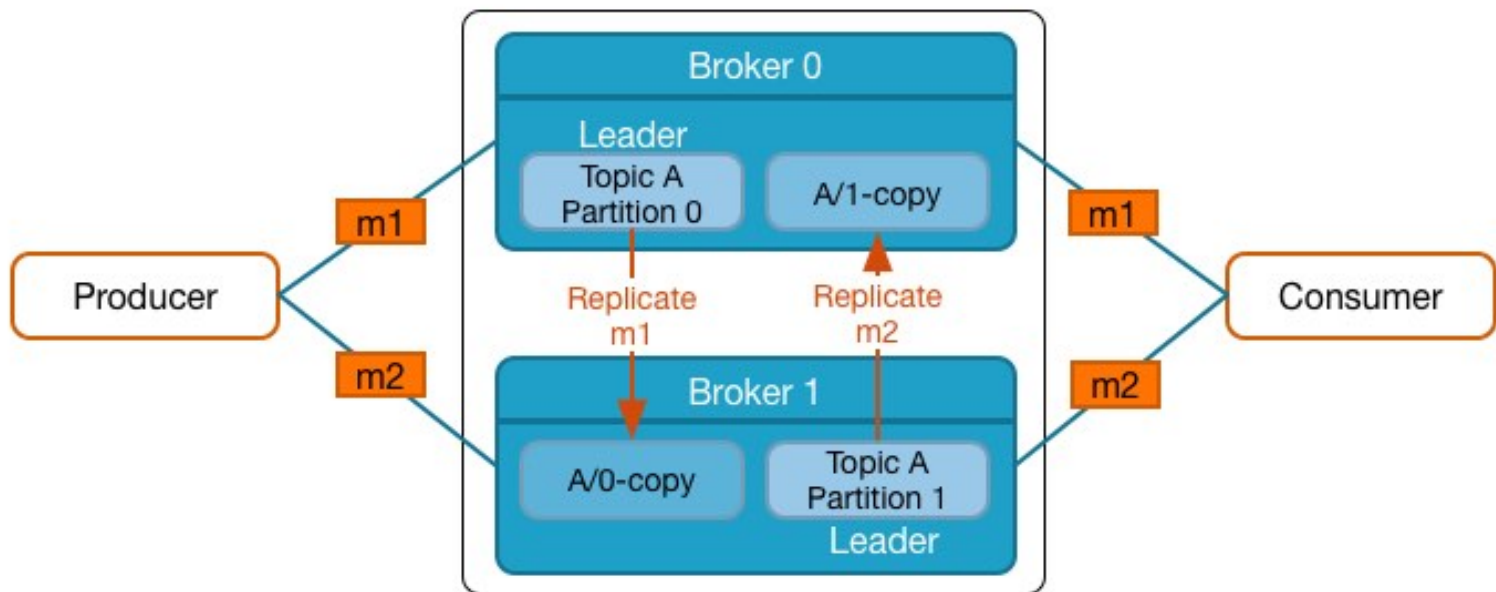
Replication of Partitions

- **Kafka maintains replicas of each partition on other Brokers in the cluster**
 - Number of replicas is configurable
- **One Broker is the leader for that Partition**
 - All writes and reads go to and from the leader
 - Other Brokers are followers



Important: Clients Do Not Access Followers

- It is important to understand that **Producers only write to the leader**
- Likewise, **Consumers *only* read from the leader**
 - They do not read from the replicas
 - Replicas only exist to provide reliability in case of Broker failure
- In the diagram below, m1 hashes to Partition 0 and m2 hashes to Partition 1



- If a leader fails, the Kafka cluster will elect a new leader from among the followers

In-Sync Replicas

- You may see information about “In-Sync Replicas” (ISR) from some Kafka command-line tools
- ISRs are replicas which are up-to-date with the leader
 - If the leader fails, it is the list of ISRs which is used to elect a new leader
- Although this is more of an administration Topic, it helps to be familiar with the term ISR

Managing Broker Failures

- **One Broker in the entire cluster is designated as the Controller**
 - Detects Broker failure/restart via ZooKeeper
- **Controller action on Broker failure**
 - Selects a new leader and updates the ISR list
 - Persists the new leader and ISR list to ZooKeeper
 - Sends the new leader and ISR list changes to all Brokers
- **Controller action on Broker restart**
 - Sends leader and ISR list information to the restarted Broker
- **If the Controller fails, one of the other Brokers will become the new Controller**

Kafka's Architecture

- *Kafka's Log Files*
- *Replicas for Reliability*
- **Partitions and Consumer Groups for Scalability**
- *Hands-On Exercise: Consuming from Multiple Partitions*
- *Chapter Review*

Scaling using Partitions

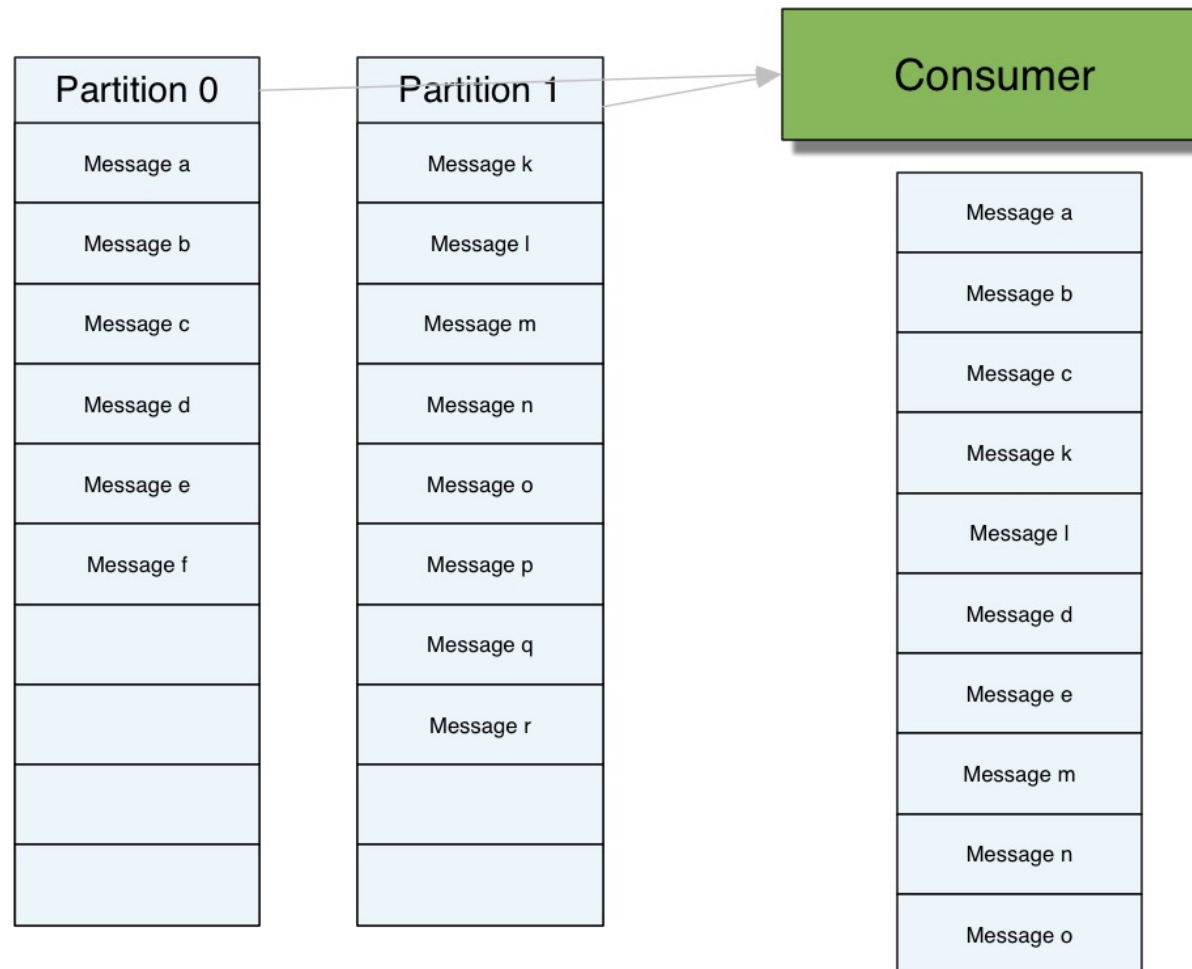
- **Recall: All Consumers read from the leader of a Partition**
 - No clients write to, or read from, followers
- **This can lead to congestion on a Broker if there are many Consumers**
- **Splitting a Topic into multiple Partitions can help to improve performance**
 - Leaders for different Partitions can be on different Brokers

Preserve Message Ordering

- **Messages with the same key, from the same Producer, are delivered to the Consumer in order**
 - Kafka hashes the key and uses the result to map the message to a specific Partition
 - Data within a Partition is stored in the order in which it is written
 - Therefore, data read from a Partition is read in order *for that partition*
- **If the key is null and the default Partitioner is used, the record is sent to a random Partition**

An Important Note About Ordering

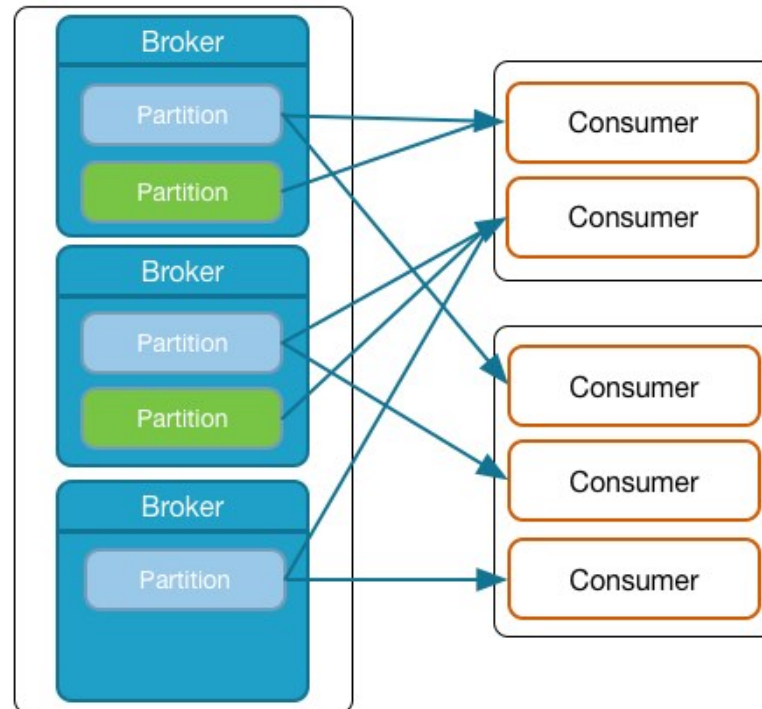
- If there are multiple Partitions, you will not get total ordering across all messages when reading data



- Question: how can you preserve message order if the application requires it?

Group Consumption Balances Load

- **Multiple Consumers in a *Consumer Group***
 - The `group.id` property is identical across all Consumers in the group
- **Consumers in the group are typically on separate machines**
- **Automatic failure detection and load rebalancing**



Partition Assignment within a Consumer Group


- **Partitions are 'assigned' to Consumers**

- A single Partition is consumed by only one Consumer in any given Consumer Group
- *i.e.*, you are guaranteed that messages with the same key will go to the same Consumer
 - Unless you change the number of partitions (see later)
- `partition.assignment.strategy` in the Consumer configuration

Partition Assignment Strategy: Range

■ Range (default)


- Topic/Partition: topics A and B, each with two Partitions 0 and 1
- Consumer Group: Consumers c1, c2, c3
- c1: {A-0, B-0} c2: {A-1, B-1} c3: {}



Partition Assignment Strategy: RoundRobin

- RoundRobin

- c1: {A-0, B-1} c2: {A-1} c3: {B-0}



- Partition assignment is automatically recomputed on changes in Partitions/Consumers

Partition Assignment Strategy: Sticky

- **Sticky assignment strategy was introduced in Kafka 0.11 (Confluent 3.3)**
- **Guarantees an assignment that is as balanced as possible**
- **Preserves existing Partition assignments to Consumers during reassignment to reduce overhead**
 - Kafka Consumers retain pre-fetched messages for Partitions assigned to them before a reassignment
 - Reduces the need to cleanup local Partition state between rebalances

Consumer Groups: Limitations

- **The number of useful Consumers in a Consumer Group is constrained by the number of Partitions on the Topic**
 - Example: If you have a Topic with three partitions, and ten Consumers in a Consumer Group reading that Topic, only three Consumers will receive data
 - One for each of the three Partitions

Consumer Groups: Caution When Changing Partitions

- **Recall: All messages with the same key will go to the same Consumer**
 - However, if you change the number of Partitions in the Topic, this may not be the case
 - Example: Using Kafka's default Partitioner, Messages with key *K1* were previously written to Partition 2 of a Topic
 - After repartitioning, new messages with key *K1* may now go to a different Partition
 - Therefore, the Consumer which was reading from Partition 2 may not get those new messages, as they may be read by a new Consumer

Kafka's Architecture

- *Kafka's Log Files*
- *Replicas for Reliability*
- *Partitions and Consumer Groups for Scalability*
- **Hands-On Exercise: Consuming from Multiple Partitions**
- *Chapter Review*

Hands-On Exercise: Consuming from Multiple Partitions

- In this Hands-On Exercise, you will create a Topic with multiple Partitions, write data to the Topic, then read the data back to see how ordering of the data is affected
- Please refer to the Hands-On Exercise Manual

Kafka's Architecture

- *Kafka's Log Files*
- *Replicas for Reliability*
- *Partitions and Consumer Groups for Scalability*
- *Hands-On Exercise: Consuming from Multiple Partitions*
- **Chapter Review**

Chapter Review

- **Kafka uses commit logs to store all its data**
 - These allow the data to be read back by any number of Consumers
- **Topics can be split into Partitions for scalability**
- **Partitions are replicated for reliability**
- **Consumers can be collected together in Consumer Groups**
 - Data from a specific Partition will go to a single Consumer in the Consumer Group
- **If there are more Consumers in a Consumer Group than there are Partitions in a Topic, some Consumers will receive no data**

Developing With Kafka

Chapter 5



Course Contents

01: Introduction

02: The Motivation for Apache Kafka

03: Kafka Fundamentals

04: Kafka's Architecture

>>> 05: Developing With Kafka

06: More Advanced Kafka Development

07: Schema Management In Kafka

08: Kafka Connect for Data Movement

09: Basic Kafka Administration

10: Kafka Stream Processing

11: Conclusion

Appendix A: Installation Recommendations

Developing With Kafka

- **In this chapter you will learn:**
 - How to write a Producer using the Java API
 - How to use the REST proxy to access Kafka from other languages
 - How to write a basic Consumer using the New Consumer API

Developing With Kafka

- **Programmatically Accessing Kafka**
- *Writing a Producer in Java*
- *Using the REST API to Write a Producer*
- *Hands-On Exercise: Writing a Producer*
- *Writing a Consumer in Java*
- *Using the REST API to Write a Consumer*
- *Hands-On Exercise: Writing a Basic Consumer*
- *Chapter Review*

The Kafka API

- **Since Kafka 0.9, Kafka includes Java clients in the `org.apache.kafka.clients` package**
 - These are intended to supplant the older Scala clients
 - They are available in a JAR which has as few dependencies as possible, to reduce code size
- **There are client libraries for many other languages**
 - The quality and support for these varies
- **Confluent provides and supports client libraries for C/C++, Python, Go, and .NET**
- **Confluent also maintains a REST Proxy for Kafka**
 - This allows any language to access Kafka via REST
 - (REpresentational State Transfer; essentially, a way to access a system by making HTTP calls)

Our Class Environment

- **During the course this week, we anticipate that you will be writing code either in Java...**
 - In which case, you will use Kafka's Java API
- **...or Python**
 - In which case, you will use the REST Proxy
- **If you wish to use some other programming language to access the REST proxy, you can do so**
 - Be aware that your instructor may not be familiar with your language of choice, though

Developing With Kafka

- *Programmatically Accessing Kafka*
- **Writing a Producer in Java**
- *Using the REST API to Write a Producer*
- *Hands-On Exercise: Writing a Producer*
- *Writing a Consumer in Java*
- *Using the REST API to Write a Consumer*
- *Hands-On Exercise: Writing a Basic Consumer*
- *Chapter Review*

The Producer API

- **To create a Producer, use the `KafkaProducer` class**
- **This is thread safe; sharing a single Producer instance across threads will typically be faster than having multiple instances**
- **Create a `Properties` object, and pass that to the Producer**
 - You will need to specify one or more Broker host/port pairs to establish the initial connection to the Kafka cluster
 - The property for this is `bootstrap.servers`
 - This is only used to establish the initial connection
 - The client will use all servers, even if they are not all listed here
 - Question: why not just specify a single server?

Important Properties Elements (1)

Name	Description
<code>bootstrap.servers</code>	List of Broker host/port pairs used to establish the initial connection to the cluster
<code>key.serializer</code>	Class used to serialize the key. Must implement the Serializer interface
<code>value.serializer</code>	Class used to serialize the value. Must implement the Serializer interface
<code>compression.type</code>	How data should be compressed. Values are none , snappy , gzip , lz4 . Compression is performed on batches of records

Important Properties Elements (2)

Name	Description
acks	Number of acknowledgment the Producer requires the leader to have before considering the request complete. This controls the durability of records. acks=0 : Producer will not wait for any acknowledgment from the server; acks=1 : Producer will wait until the leader has written the record to its local log; acks=all : Producer will wait until all in-sync replicas have acknowledged receipt of the record

Creating the Properties and KafkaProducer Objects

```
1 Properties props = new Properties;  
2 props.put("bootstrap.servers", "broker1:9092");  
3 props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");  
4 props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");  
5  
6 KafkaProducer<String, String> producer = new KafkaProducer<>(props);
```

- Other serializers available: ByteArraySerializer, IntegerSerializer, LongSerializer
- StringSerializer encoding defaults to UTF8
 - Can be customized by setting the property `serializer.encoding`

Helper Classes

- Kafka includes helper classes `ProducerConfig`, `ConsumerConfig`, and `StreamsConfig`
 - Provide predefined constants for commonly configured properties
- Examples

```
// With helper class
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "broker1:9092");
// Without helper class
props.put("bootstrap.servers", "broker1:9092");
```

```
// With helper class
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, StringSerializer.class);
// Without helper class
props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
```

```
// With helper class
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class);
// Without helper class
props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
```

Sending Messages to Kafka

```
1 String k = "mykey";  
2 String v = "myvalue";  
3 ProducerRecord<String, String> record = new ProducerRecord<String, String>("my_topic", k, v); ①  
4 producer.send(record); ②
```

① **ProducerRecord can take an optional timestamp if you don't want to use current system time**

② **Alternatively:**

```
producer.send(new ProducerRecord<String, String>("my_topic", k, v));
```


send() Does Not Block

- **The send() call is asynchronous**
 - It does not block; it returns immediately and your code continues
- **It returns a Future which contains a RecordMetadata object**
 - The Partition the record was put into and its offset
- **To force send() to block, call `producer.send(record).get()`**

When Do Producers Actually Send?

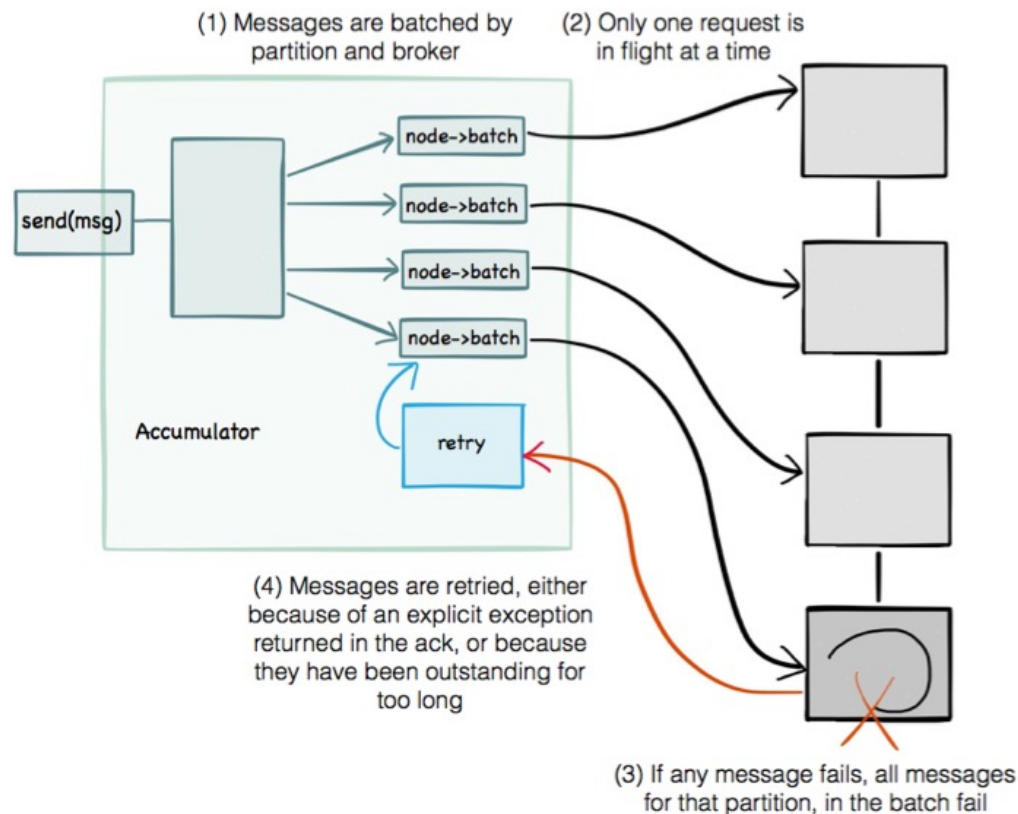
- **A Producer `send()` returns immediately after it has added the message to a local buffer of pending record sends**
 - This allows the Producer to send records in batches for better performance
- **Then the Producer flushes multiple messages to the Brokers based on batching configuration parameters**
 - You can also manually flush by calling the Producer method `flush()`
- **Do not confuse the Producer method `flush()` in a Producer context with the term `flush` in a Broker context**
 - `flush` in a Broker context refers to when messages get written from the page cache to disk

Producer Retries

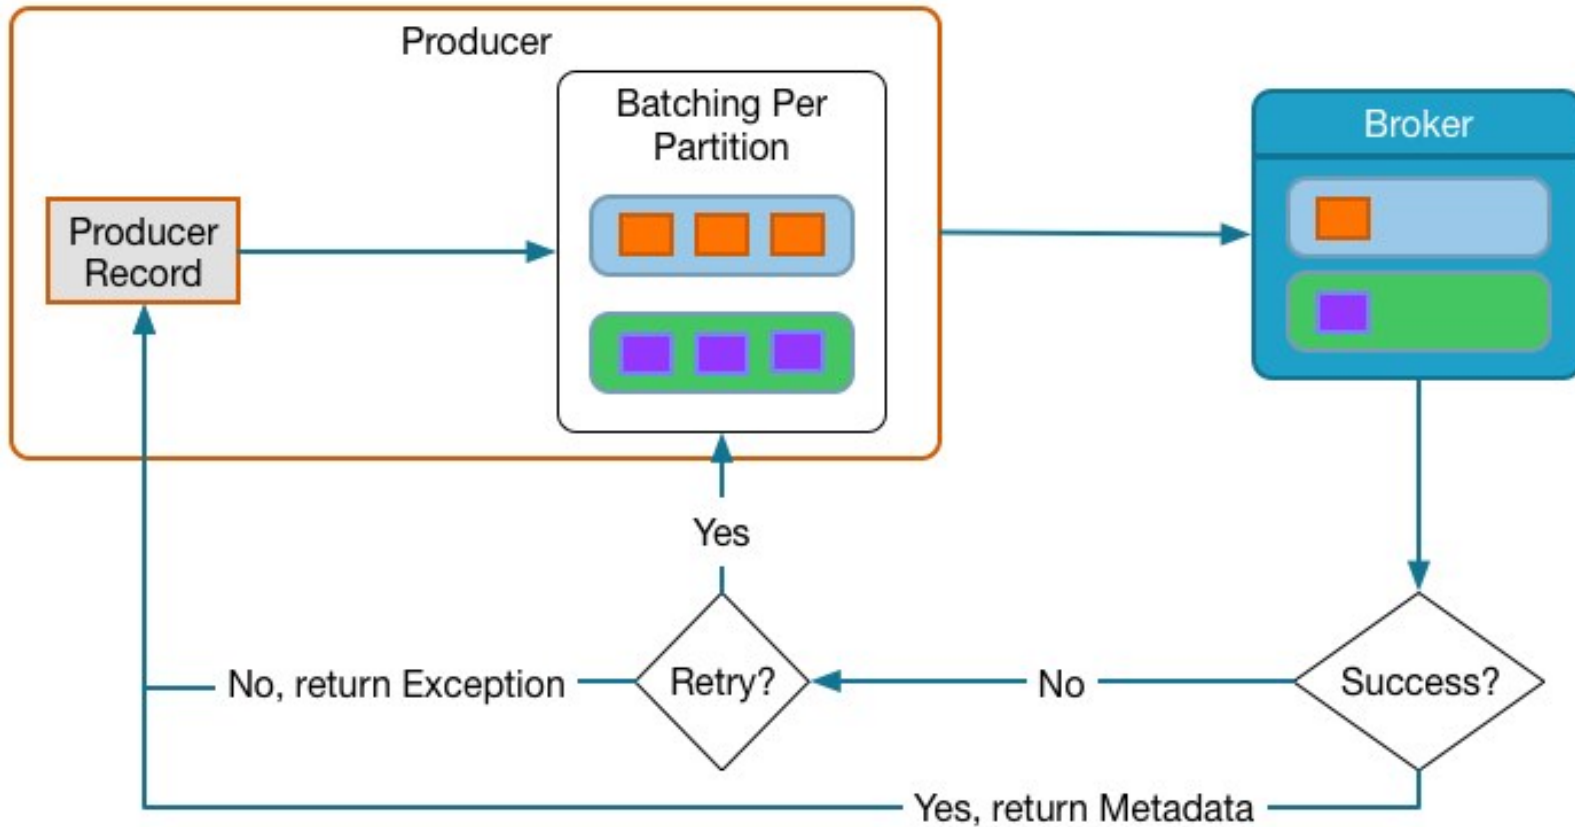
- **Developers can configure the retries configuration setting in the Producer code**
 - retries: how many times the Producer will attempt to retry sending records
 - Only relevant if acks is not 0
 - Hides transient failure
 - Ensure lost messages are retried rather than just throwing an error
 - retries: number times to retry (Default: 0)
 - `retry.backoff.ms`: pause added between retries (Default: 100)
 - For example, `retry.backoff.ms=100` and `retries=600` will retry for for 60 seconds

Preserve Message Send Order

- If `retries > 0`, message ordering could change
- To preserve message order, set `max.in.flight.requests.per.connection=1` (Default: 5)
 - May impact throughput performance because of lack of request pipelining



Batching and Retries



Performance Tuning Batching

- **Producers can adjust batching configuration parameters**
 - `batch.size` message batch size in bytes (Default: 16KB)
 - `linger.ms` time to wait for messages to batch together (Default: 0, *i.e.*, send immediately)
 - High throughput: large `batch.size` and `linger.ms`, or flush manually
 - Low latency: small `batch.size` and `linger.ms`
 - `buffer.memory` (Default: 32MB)
 - The Producer's buffer for messages to be sent to the cluster
 - Increase if Producers are sending faster than Brokers are acknowledging, to prevent blocking

send() and Callbacks (1)

- `send(record)` is equivalent to `send(record, null)`
- **Instead, it is possible to supply a Callback as the second parameter**
 - This is invoked when the send has been acknowledged
 - It is an Interface with an `onCompletion` method:

```
onCompletion(RecordMetadata metadata, java.lang.Exception exception)
```

- **Callback parameters**
 - metadata will be `null` if an error occurred
 - exception will be `null` if no error occurred

send() and Callbacks (2)

- Parameters correlated to a particular record can be passed into the Callback's constructor
- Example code, with lambda function and closure instead of requiring a separate class definition

```
1 producer.send(record, (recordMetadata, e) -> {  
2   if (e != null) {  
3     e.printStackTrace();  
4   } else {  
5     System.out.println("Message String = " + record.value() + ", Offset = " + recordMetadata  
6       .offset());  
7   }  
8 });
```


Closing the Producer

- **close():** blocks until all previously sent requests complete

```
1 producer.close();
```

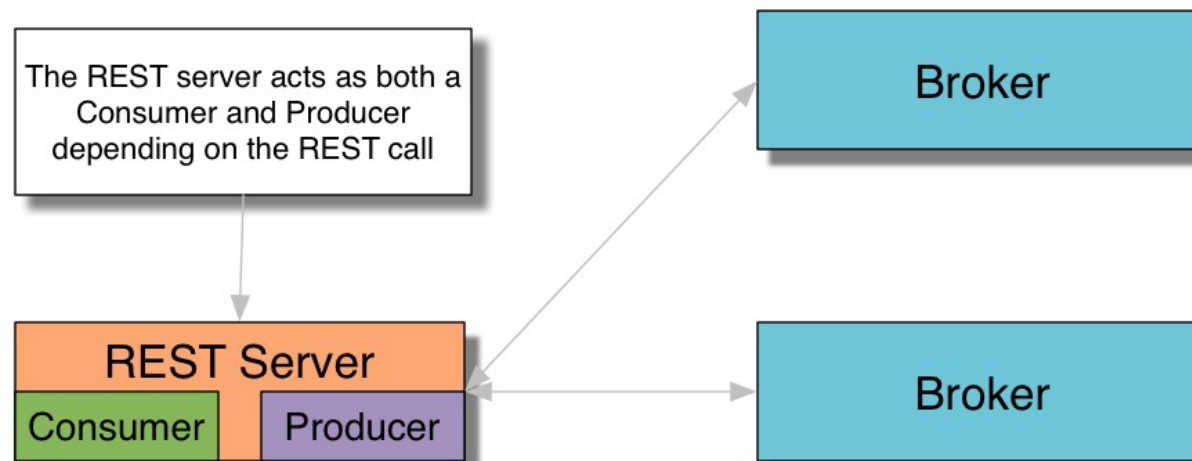
- **close(timeout, TimeUnit):** waits up to timeout for the producer to complete the sending of all incomplete requests
 - If the producer is unable to complete all requests before the timeout expires, this method will fail any unsent and unacknowledged records immediately

Developing With Kafka

- *Programmatically Accessing Kafka*
- *Writing a Producer in Java*
- **Using the REST API to Write a Producer**
- *Hands-On Exercise: Writing a Producer*
- *Writing a Consumer in Java*
- *Using the REST API to Write a Consumer*
- *Hands-On Exercise: Writing a Basic Consumer*
- *Chapter Review*

About the REST Proxy

- The REST Proxy allows you to use HTTP to perform actions on the Kafka cluster
- The REST calls are translated into native Kafka calls
- This allows virtually any language to access Kafka
- Uses POST to send data to Kafka
 - Embedded formats: JSON, base64-encoded JSON, or Avro-encoded JSON
- Uses GET to retrieve data from Kafka



A Python Producer Using the REST Proxy

```
1 #!/usr/bin/python
2
3 import requests
4 import base64
5 import json
6
7 url = "http://restproxy:8082/topics/my_topic"
8 headers = {
9     "Content-Type" : "application/vnd.kafka.binary.v1+json"
10 }
11 # Create one or more messages
12 payload = {"records":
13     [{
14         "key":base64.b64encode("firstkey"),
15         "value":base64.b64encode("firstvalue")
16     ]}
17 # Send the message
18 r = requests.post(url, data=json.dumps(payload), headers=headers)
19 if r.status_code != 200:
20     print "Status Code: " + str(r.status_code)
21     print r.text
```

Developing With Kafka

- *Programmatically Accessing Kafka*
- *Writing a Producer in Java*
- *Using the REST API to Write a Producer*
- **Hands-On Exercise: Writing a Producer**
- *Writing a Consumer in Java*
- *Using the REST API to Write a Consumer*
- *Hands-On Exercise: Writing a Basic Consumer*
- *Chapter Review*

Hands-On Exercise: Writing a Producer

- In this Hands-On Exercise, you will write a Kafka Producer either in Java or Python
- Please refer to the Hands-On Exercise Manual

Developing With Kafka

- *Programmatically Accessing Kafka*
- *Writing a Producer in Java*
- *Using the REST API to Write a Producer*
- *Hands-On Exercise: Writing a Producer*
- **Writing a Consumer in Java**
- *Using the REST API to Write a Consumer*
- *Hands-On Exercise: Writing a Basic Consumer*
- *Chapter Review*

Consumers and Offsets

- **Each message in a Partition has an offset**
 - The numerical value indicating where the message is in the log
- **Kafka tracks the Consumer Offset for each partition of a Topic the Consumer (or Consumer Group) has subscribed to**
 - It tracks these values in a special Topic
- **Consumer offsets are committed automatically by default**
 - We will see later how to manually commit offsets if you need to do that
- **Tip: the Consumer Offset is the value of the next message the Consumer will read, not the last message that has been read**
 - For example, if the Consumer Offset is 9, this indicates that messages 0 to 8 have already been processed, and that message 9 will be the next one sent to the Consumer

Important Consumer Properties

- Important Consumer properties include:

Name	Description
<code>bootstrap.servers</code>	List of Broker host/port pairs used to establish the initial connection to the cluster
<code>key.deserializer</code>	Class used to deserialize the key. Must implement the Deserializer interface
<code>value.deserializer</code>	Class used to deserialize the value. Must implement the Deserializer interface
<code>group.id</code>	A unique string that identifies the Consumer Group this Consumer belongs to.
<code>enable.auto.commit</code>	When set to true (the default), the Consumer will trigger offset commits based on the value of auto.commit.interval.ms (default 5000ms)

Creating the Properties and KafkaConsumer Objects

```
1 Properties props = new Properties();
2 props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "broker1:9092");
3 props.put(ConsumerConfig.GROUP_ID_CONFIG, "samplegroup");
4 props.put(ConsumerConfig.KEY_SERIALIZER_CLASS_CONFIG, StringDeserializer.class);
5 props.put(ConsumerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringDeserializer.class);
6
7 KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props);
8 consumer.subscribe(Arrays.asList("my_topic", "my_other_topic")); ① ②
```

- ① The Consumer can subscribe to as many Topics as it wishes. Note that this call is not additive; calling subscribe again will remove the existing list of Topics, and will only subscribe to those specified in the new call
- ② You may also use regular expressions (*i.e.*, Pattern) for topic subscription

Reading Messages from Kafka with `poll()`

```
1 while (true) { ①
2   ConsumerRecords<String, String> records = consumer.poll(100); ②
3   for (ConsumerRecord<String, String> record : records)
4     System.out.printf("offset = %d, key = %s, value = %s\n", record.offset(), record.key(),
5     record.value());
6 }
```

① Loop forever

② Each call to `poll` returns a (possibly empty) list of messages.

Controlling the Number of Messages Returned

- **By default, `poll()` fetches available messages from multiple partitions across multiple Brokers**
 - Up to the maximum data size *per partition*
 - `max.partition.fetch.bytes`: default value 1048576
 - A Topic with many partitions could result in extremely large amounts of data being returned
- **The optional timeout parameter in the call to `poll()` is across all fetch requests**
 - This controls the maximum amount of time in milliseconds that the Consumer will block if no new records are available
 - If records are available, it will return immediately
- **In Kafka 0.10.0, `max.poll.records` was introduced**
 - Property limits the total number of records retrieved in a single call to `poll()`

Performance Tuning Consumer Fetch Requests

- **Consumption goal: high throughput or low latency?**
 - High throughput: send more data at a given time
 - Low latency: send immediately when data is available
- **`fetch.min.bytes`, `fetch.max.wait.ms`**
 - `fetch.min.bytes` (Default: 1)
 - Broker waits for messages to accumulate to this size batch before responding
 - `fetch.max.wait.ms` (Default: 500)
 - Broker will not wait longer than this duration before returning a batch
- **High throughput**
 - Large `fetch.min.bytes`, reasonable `fetch.wait.max.ms`
- **Low latency**
 - `fetch.min.bytes=1`

Message Size Limit

- **Try not to change the maximum message size unless it is unavoidable**
 - Kafka is not optimized for very large messages
- **If you must change the maximum size for a batch of messages that the Broker can receive from a Producer**
 - Broker: `message.max.bytes` (Default: 1MB)
 - Topic override: `max.message.bytes` (Default: 1MB)

Preventing Resource Leaks

- It is good practice to wrap the code in a `try{ }` block, and close the `KafkaConsumer` object in a `finally{ }` block to avoid resource leaks
- It is important to note that `KafkaConsumer` is not thread-safe

```
1 try {
2     while (true) { ①
3         ConsumerRecords<String, String> records = consumer.poll(100); ②
4         for (ConsumerRecord<String, String> record : records)
5             System.out.printf("offset = %d, key = %s, value = %s\n", record.offset(), record.key(),
record.value());
6     }
7 } finally {
8     consumer.close();
9 }
```

Developing With Kafka

- *Programmatically Accessing Kafka*
- *Writing a Producer in Java*
- *Using the REST API to Write a Producer*
- *Hands-On Exercise: Writing a Producer*
- *Writing a Consumer in Java*
- **Using the REST API to Write a Consumer**
- *Hands-On Exercise: Writing a Basic Consumer*
- *Chapter Review*

A Python Consumer Using the REST API (1)

```
1 #!/usr/bin/python
2
3 import requests
4 import base64
5 import json
6 import sys
7
8 # Base URL for interacting with REST server
9 baseurl = "http://restproxy:8082/consumers/group1" ①
10
11 # Create the Consumer instance
12 print "Creating consumer instance"
13 payload = {
14     "format": "binary"
15 }
16 headers = {
17     "Content-Type" : "application/vnd.kafka.v1+json"
18 }
```

① We are creating a Consumer instance in a Consumer Group called group1

A Python Consumer Using the REST API (2)

```
19 r = requests.post(baseurl, data=json.dumps(payload), headers=headers)
20
21 if r.status_code != 200:
22     print "Status Code: " + str(r.status_code)
23     print r.text
24     sys.exit("Error thrown while creating consumer")
25
26 # Base URI is used to identify the consumer instance
27 base_uri = r.json()["base_uri"]
```

A Python Consumer Using the REST API (3)

```
28 # Get the message(s) from the Consumer
29 headers = {
30     "Accept" : "application/vnd.kafka.binary.v1+json"
31 }
32
33 # Request messages for the instance on the Topic
34 r = requests.get(base_uri + "/topics/my_topic", headers=headers, timeout=20)
35
36 if r.status_code != 200:
37     print "Status Code: " + str(r.status_code)
38     print r.text
39     sys.exit("Error thrown while getting message")
```

A Python Consumer Using the REST API (4)

```
40 # Output all messages
41 for message in r.json():
42     if message["key"] is not None:
43         print "Message Key:" + base64.b64decode(message["key"])
44         print "Message Value:" + base64.b64decode(message["value"])
45
46 # When we're done, delete the Consumer
47 headers = {
48     "Accept" : "application/vnd.kafka.v1+json"
49 }
50
51 r = requests.delete(base_uri, headers=headers)
52
53 if r.status_code != 204:
54     print "Status Code: " + str(r.status_code)
55     print r.text
```

Developing With Kafka

- *Programmatically Accessing Kafka*
- *Writing a Producer in Java*
- *Using the REST API to Write a Producer*
- *Hands-On Exercise: Writing a Producer*
- *Writing a Consumer in Java*
- *Using the REST API to Write a Consumer*
- **Hands-On Exercise: Writing a Basic Consumer**
- *Chapter Review*

Hands-On Exercise: Writing a Basic Consumer

- In this Hands-On Exercise, you will write a basic Kafka Consumer
- Please refer to the Hands-On Exercise Manual

Developing With Kafka

- *Programmatically Accessing Kafka*
- *Writing a Producer in Java*
- *Using the REST API to Write a Producer*
- *Hands-On Exercise: Writing a Producer*
- *Writing a Consumer in Java*
- *Using the REST API to Write a Consumer*
- *Hands-On Exercise: Writing a Basic Consumer*
- **Chapter Review**

Chapter Review

- **The Kafka API provides Java clients for Producers and Consumers**
- **Client libraries for other languages are available**
 - Confluent provides and supports client libraries for C/C++, Python, Go, and .NET
- **Confluent's REST Proxy allows other languages to access Kafka without the need for native client libraries**

More Advanced Kafka Development

Chapter 6



Course Contents

01: Introduction

02: The Motivation for Apache Kafka

03: Kafka Fundamentals

04: Kafka's Architecture

05: Developing With Kafka

>>> 06: More Advanced Kafka Development

07: Schema Management In Kafka

08: Kafka Connect for Data Movement

09: Basic Kafka Administration

10: Kafka Stream Processing

11: Conclusion

Appendix A: Installation Recommendations

More Advanced Kafka Development

- **In this chapter you will learn:**

- How to enable exactly once semantics (EOS)
- How to specify the offset to read from
- How to manually commit reads from the Consumer
- How to create a custom Partitioner
- How to control message delivery reliability

More Advanced Kafka Development

- **Exactly Once Semantics**
- *Specifying Offsets*
- *Consumer 'Liveness' and Rebalancing*
- *Manually Committing Offsets*
- *Partitioning Data*
- *Message Durability*
- *Hands-On Exercise: Accessing Previous Data*
- *Chapter Review*

Motivation for Exactly Once Semantics

- **Write real-time, mission-critical streaming applications that require guarantees that data is processed “exactly once”**
 - Complements the “at most once” or “at least once” semantics that exist today
- **Exactly Once Semantics (EOS) bring strong transactional guarantees to Kafka**
 - Prevents duplicate messages from being processed by client applications
 - Even in the event of client retries and Broker failures
- **Use cases: tracking ad views, processing financial transactions, etc.**

EOS in Kafka

- EOS was introduced in Kafka 0.11 (Confluent 3.3)
- Supported with the new Java Producer and Consumer
- Supported with the Kafka Streams API
- Transaction Coordinator manages transactions and assign Producer IDs
- Updated log format with sequence numbers and Producer IDs

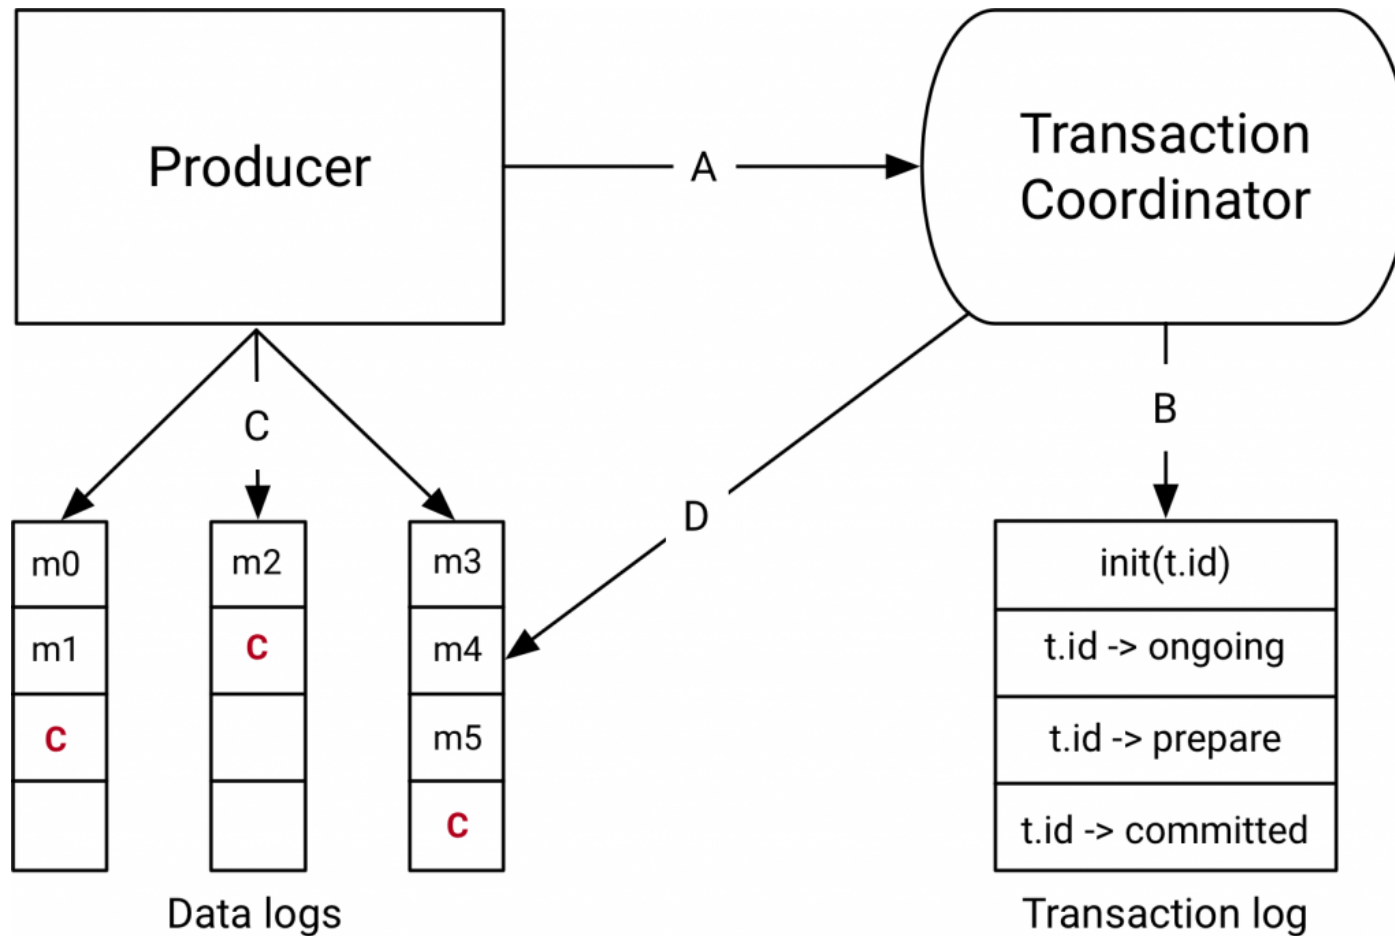
Enabling EOS on the Producer Side (1)

- **Configure the Producer send operation to be idempotent**
 - An idempotent operation is one which can be performed many times without causing a different effect than only being performed once
 - Removes the possibility of duplicate messages being delivered to a particular Topic Partition due to Producer or Broker errors
 - Set `enable.idempotence` to `true` (Default: `false`)
- **A Producer writes atomically across multiple Partitions using "transactional" messages**
 - Either all or no messages in a batch are visible to Consumers
 - Use the new transactions API
- **Allow reliability semantics to span multiple producer sessions**
 - Set `transactional.id` to guarantee that transactions using the same transactional ID have completed prior to starting new transactions

Enabling EOS on the Producer Side (2)

```
1 producer.initTransactions();
2 try {
3     producer.beginTransaction();
4     producer.send(record1);
5     producer.send(record2);
6     producer.commitTransaction();
7 } catch(ProducerFencedException e) {
8     producer.close();
9 } catch(KafkaException e) {
10     producer.abortTransaction();
11 }
```


Producer Transactions across Partitions



Enabling EOS on the Consumer Side

- **Configure the Consumer to read committed transactional messages**
 - Set `isolation.level` to `read_committed` (Default: `read_uncommitted`)
 - `read_committed`: reads only committed transactional messages and all non-transactional messages
 - `read_uncommitted`: reads all transactional messages and all non-transactional messages

More Advanced Kafka Development

- *Exactly Once Semantics*
- **Specifying Offsets**
- *Consumer 'Liveness' and Rebalancing*
- *Manually Committing Offsets*
- *Partitioning Data*
- *Message Durability*
- *Hands-On Exercise: Accessing Previous Data*
- *Chapter Review*

Determining the Offset When a Consumer Starts

- **The Consumer property `auto.offset.reset` determines what to do if there is no valid offset in Kafka for the Consumer's Consumer Group**
 - When a particular Consumer Group starts the first time
 - If the Consumer offset is less than the smallest offset
 - If the Consumer offset is greater than the last offset
- **The value can be one of:**
 - `earliest`: Automatically reset the offset to the earliest available
 - `latest`: Automatically reset to the latest offset available
 - `none`: Throw an exception if no previous offset can be found for the ConsumerGroup
- **The default is `latest`**

Changing the Offset Within the Consumer (1)

- **The KafkaConsumer API provides ways to view offsets and dynamically change the offset from which the Consumer will read**
- **View offsets**
 - `position(TopicPartition)` provides the offset of the next record that will be fetched
 - `offsetsForTimes(Map<TopicPartition, Long> timestampsToSearch)` looks up the offsets for the given Partitions by timestamp
- **Change offsets**
 - `seekToBeginning(Collection<TopicPartition>)` seeks to the first offset of each of the specified Partitions
 - `seekToEnd(Collection<TopicPartition>)` seeks to the last offset of each of the specified Partitions
 - `seek(TopicPartition, offset)` seeks to a specific offset in the specified Partition

Changing the Offset Within the Consumer (2)

- For example, to seek to the beginning of all Partitions that are being read by a Consumer for a particular Topic, you might do something like:

```
1 consumer.subscribe(Arrays.asList("my_topic"));
2 consumer.poll(0);
3 consumer.seekToBeginning(consumer.assignment()); ①
```

- ① `assignment()` returns a list of all `TopicPartitions` currently assigned to this Consumer

Changing the Offset Within the Consumer (3)

- For example, to reset the offset to a particular timestamp, you could do something like:

```
1 for (TopicPartition topicPartition : partitionSet) {
2     timestampsToSearch.put(topicPartition, MY_TIMESTAMP); ①
3 }
4
5 Map<TopicPartition, OffsetAndTimestamp> result = consumer.offsetsForTimes(timestampsToSearch);
6 ②
7 for (Map.Entry<TopicPartition, OffsetAndTimestamp> entry : result.entrySet()) {
8     consumer.seek(entry.getKey(), entry.getValue().offset()); ③
9 }
```

- ① Add each Partition and timestamp to the HashMap
- ② Get the offsets for each Partition
- ③ Seek to the specified offset for each Partition

Question: what is a use case where this would be useful?

More Advanced Kafka Development

- *Exactly Once Semantics*
- *Specifying Offsets*
- **Consumer 'Liveness' and Rebalancing**
- *Manually Committing Offsets*
- *Partitioning Data*
- *Message Durability*
- *Hands-On Exercise: Accessing Previous Data*
- *Chapter Review*

Consumer 'Liveness'

- **The Consumer communicates to the Brokers regularly to let them know that it is still alive**
- **If a Consumer is believed to be dead, it is removed from the Consumer Group and a Consumer rebalance is performed**
 - Reallocates all Partitions of a Topic to live members of the Consumer Group

Consumer Heartbeats for 'Liveness'

- **As of Kafka 0.10.1.0, heartbeats are sent in a background thread, separate from calls to poll**
- **New defaults**
 - `session.timeout.ms` now defaults to 10 seconds
 - `max.poll.records` now defaults to 500
- **To avoid issues where a Consumer stops processing records but the background thread is still running, poll must still be called periodically**
 - `max.poll.interval.ms`: maximum allowable interval between calls to poll (Default: 5 minutes)

Pausing Record Retrieval

- **There are situations where Consumers may wish to temporarily stop consumption of new messages**
 - The Consumer can pause and resume consumption on a per-Partition basis
 - To maintain liveness, `poll` must be called periodically, even in Kafka 0.10.1.0 and above
- **Use cases**
 - Prior to 0.10.1.0, continue heartbeating during long message processing times
 - Processing and joining data from two topics, where one lags behind another
- **How pause works**
 - Each pause call takes a Collection of `TopicPartition` objects
 - While paused, calls to `poll` return no new messages for those Partitions
- **Caution: In this scenario, you will probably need to manually manage topic offsets**

Partition Assignments for Consumers

- We have said that all the data from a particular Partition will go to the same Consumer
- This is true *as long as* the number of Consumers in a Consumer Group does not change
- If the number of Consumers changes, a Partition rebalance occurs
 - Partition ownership is moved around between the Consumers to spread the load evenly
 - No guarantees that a Consumer will get the same or different Partition
- Consumers cannot consume messages during the rebalance, so this results in a short pause in message consumption

Consumer Rebalancing

- **Consumer rebalances are initiated when**
 - A Consumer leaves the Consumer group (either by failing to send a timely heartbeat or by explicitly requesting to leave)
 - A new Consumer joins the Consumer Group
 - A Consumer changes its Topic subscription
 - The Consumer Group notices a change to the Topic metadata for any subscribed Topic (e.g. an increase in the number of Partitions)
- **Rebalancing does not occur if**
 - A Consumer calls pause
- **During rebalance, consumption is paused**
- **Question: could adding a Consumer to a Consumer Group cause Partition assignment to change?**

The Case For and Against Rebalancing

- **Typically, Partition rebalancing is a good thing**
 - Allows you to add more Consumers to a Consumer Group dynamically, without having to restart all the other Consumers in the group
 - Automatically handles situations where a Consumer in the Consumer Group fails
- **However a rebalance is like a fresh restart**
 - Consumers may or may not get a different set of Partitions
 - If your Consumer is relying on getting all data from a particular Partition, this could be a problem
 - A previously-paused Partition is no longer paused
- **Managing rebalances where Partition assignment matters**
 - Option 1: only have a single Consumer for the entire topic
 - Option 2: provide a `ConsumerRebalanceListener` when calling `subscribe()`
 - Implement `onPartitionsRevoked` and `onPartitionsAssigned` methods

More Advanced Kafka Development

- *Exactly Once Semantics*
- *Specifying Offsets*
- *Consumer 'Liveness' and Rebalancing*
- **Manually Committing Offsets**
- *Partitioning Data*
- *Message Durability*
- *Hands-On Exercise: Accessing Previous Data*
- *Chapter Review*

Default Behavior: Automatically Committed Offsets

- **By default, `enable.auto.commit` is set to `true`**
 - Consumer offsets are periodically committed in the background
 - This happens during the `poll()` call
- **This is typically the desired behavior**
- **However, there are times when it may be problematic**

Problems With Automatically Committed Offsets

- **By default, automatic commits occur every five seconds**
- **Imagine that two seconds after the most recent commit, a rebalance is triggered**
 - After the rebalance, Consumers will start consuming from the latest committed offset position
- **In this case, the offset is two seconds old, so all messages that arrived in those two seconds will be processed twice**
 - This provides “at least once” delivery

Manually Committing Offsets

- **A Consumer can manually commit offsets to control the committed position**
 - Disable automatic commits: set `enable.auto.commit` to `false`
- **`commitSync()`**
 - Blocks until it succeeds, retrying as long as it does not receive a fatal error
 - For “at most once” delivery, call `commitSync()` immediately after `poll()` and then process the messages
 - Consumer should ensure it has processed all the records returned by `poll()` or it may miss messages
- **`commitAsync()`**
 - Returns immediately
 - Optionally takes a callback that will be triggered when the Broker responds
 - Has higher throughput since the Consumer can process next message batch before commit returns
 - Trade-off is the Consumer may find out later the commit failed

Manually Committing Offsets From the REST Proxy

- It is possible to manually commit offsets from the REST Proxy

```
1 payload = {
2     "format": "binary",
3     # Manually/Programmatically commit offset
4     "auto.commit.enable": "false"
5 }
6
7 headers = {
8     "Content-Type" : "application/vnd.kafka.v1+json"
9 }
10
11 r = requests.post(baseurl, data=json.dumps(payload), headers=headers)
12
13 # Commit the offsets
14 if shouldCommit() == True:
15     r = requests.post(base_uri + "/offsets", headers=headers, timeout=20)
16     if r.status_code != 200:
17         print "Status Code: " + str(r.status_code)
18         print r.text
19         sys.exit("Error thrown while committing")
20     print "Committed"
```

Aside: What Offset is Committed?

- **Note: The offset committed (whether automatically or manually) is the offset of the next record to be read**
 - Not the offset of the last record which was read

Storing Offsets Outside of Kafka

- **By default, Kafka stores offsets in a special Topic**
 - Called `__consumer_offsets`
- **In some cases, you may want to store offsets outside of Kafka**
 - For example, in a database table
- **If you do this, you can read the value and then use `seek()` to move to the correct position when your application launches**

Quiz: Question

- A team responsible for consuming data from a Kafka cluster is reporting that the cluster is not properly managing offsets. They describe that when they add Consumers to a Consumer Group, all the Consumers see some previously-consumed messages. What is happening?

Quiz: Answer

- **When the Consumer Group membership changes, it triggers a rebalance**
 - Existing Partitions may be reassigned to different Consumers
- **Offsets are managed per Partition not per Consumer**
 - This works only if the Consumer code is written well and the offsets are committed upon consumption
- **Action:**
 - Review Consumer client code for committing offsets during consumption
 - Commit intervals
 - `commitSync()` vs `commitAsync()`
 - Consumers should implement `ConsumerRebalanceListener` code
 - Consumer can commit offsets in reaction to receiving an `onPartitionsRevoked` event

More Advanced Kafka Development

- *Exactly Once Semantics*
- *Specifying Offsets*
- *Consumer 'Liveness' and Rebalancing*
- *Manually Committing Offsets*
- **Partitioning Data**
- *Message Durability*
- *Hands-On Exercise: Accessing Previous Data*
- *Chapter Review*

Kafka's Default Partitioning Scheme

- Recall that by default, if a message has a key, Kafka will hash the key and use the result to map the message to a specific Partition
- This means that all messages with the same key will go to the same Partition
- If the key is null and the default Partitioner is used, the record will be sent to a random Partition (using a round-robin algorithm)
- You may wish to override this behavior and provide your own Partitioning scheme
 - For example, if you will have many messages with a particular key, you may want one Partition to hold just those messages

Creating a Custom Partitioner

- **To create a custom Partitioner, you should implement the Partitioner interface**
 - This interface includes `configure`, `close`, and `partition` methods, although often you will only implement `partition`
 - `partition` is given the `Topic`, `key`, `serialized key`, `value`, `serialized value`, and `cluster metadata`
- **It should return the number of the Partition this particular message should be sent to (0-based)**

Custom Partitioner: Example

- Assume we want to store all messages with a particular key in one Partition, and distribute all other messages across the remaining Partitions

```
1 public class MyPartitioner implements Partitioner {
2     public void configure(Map<String, ?> configs) {}
3     public void close() {}
4
5     public int partition(String topic, Object key, byte[] keyBytes,
6                          Object value, byte[] valueBytes, Cluster cluster) {
7         List<PartitionInfo> partitions = cluster.partitionsForTopic(topic);
8         int numPartitions = partitions.size();
9
10        if ((keyBytes == null) || (!(key instanceof String)))
11            throw new InvalidRecordException("Record did not have a string Key");
12
13        if (((String) key).equals("OurBigKey")) ①
14            return 0; // This key will always go to Partition 0
15
16        // Other records will go to the rest of the Partitions using a hashing function
17        return (Math.abs(Utils.murmur2(keyBytes)) % (numPartitions - 1)) + 1;
18    }
19 }
```

- ① This is the key we want to store in its own Partition

An Alternative to a Custom Partitioner

- It is also possible to specify the Partition to which a message should be written when creating the `ProducerRecord`
 - `ProducerRecord<String, String> record = new ProducerRecord<String, String>("my_topic", 0, key, value);`
 - Will write the message to Partition 0
- Discussion: Which method is preferable?

More Advanced Kafka Development

- *Exactly Once Semantics*
- *Specifying Offsets*
- *Consumer 'Liveness' and Rebalancing*
- *Manually Committing Offsets*
- *Partitioning Data*
- **Message Durability**
- *Hands-On Exercise: Accessing Previous Data*
- *Chapter Review*

Replication Factor Affects Message Durability

- Recall that Topics can be replicated for durability
- Default replication factor is 1
 - Can be specified when a Topic is first created, or modified later

Delivery Acknowledgment

- The `acks` configuration parameter determines the behavior of the Producer when sending messages
- Use this to configure the durability of messages being sent

<code>acks=0</code>	Producer will not wait for any acknowledgment. The message is placed in the Producer's buffer, and immediately considered sent
<code>acks=1</code>	The Producer will wait until the leader acknowledges receipt of the message
<code>acks=all</code>	The leader will wait for acknowledgment from all in-sync replicas before reporting the message as delivered to the Producer

- `min.insync.replicas` with `acks=all`
 - defines the minimum number of replicas in ISR needed to satisfy a produce request

Important Timeouts

- **`request.timeout.ms`**
 - Maximum time waiting for the response to a request
 - After timeout, the client will resend or throw an error if retries are exhausted
 - Recall that retries defaults to 0 and is only relevant if acks is not 0
 - Default: 30s
- **`max.block.ms`**
 - The `send()` call waits this time until the message can be successfully put into the bufferpool
 - After timeout, it will throw a `TimeoutException`
 - Default: 60s

More Advanced Kafka Development

- *Exactly Once Semantics*
- *Specifying Offsets*
- *Consumer 'Liveness' and Rebalancing*
- *Manually Committing Offsets*
- *Partitioning Data*
- *Message Durability*
- **Hands-On Exercise: Accessing Previous Data**
- *Chapter Review*

Hands-On Exercise: Accessing Previous Data

- In this Hands-On Exercise you will create a Consumer which will access data already stored in the cluster.
- Please refer to the Hands-On Exercise Manual

More Advanced Kafka Development

- *Exactly Once Semantics*
- *Specifying Offsets*
- *Consumer 'Liveness' and Rebalancing*
- *Manually Committing Offsets*
- *Partitioning Data*
- *Message Durability*
- *Hands-On Exercise: Accessing Previous Data*
- **Chapter Review**

Chapter Review

- Your Consumer can move through the data in the Cluster, reading from the beginning, the end, or any point in between
- You may need to take Consumer Rebalancing into account when you write your code
- It is possible to specify your own Partitioner if Kafka's default is not sufficient for your needs
- You can configure the reliability of message delivery by specifying different values for the acks configuration parameter

Schema Management In Kafka

Chapter 7



Course Contents

01: Introduction

02: The Motivation for Apache Kafka

03: Kafka Fundamentals

04: Kafka's Architecture

05: Developing With Kafka

06: More Advanced Kafka Development

>>> 07: Schema Management In Kafka

08: Kafka Connect for Data Movement

09: Basic Kafka Administration

10: Kafka Stream Processing

11: Conclusion

Appendix A: Installation Recommendations

Schema Management In Kafka

- **In this chapter you will learn:**

- What Avro is, and how it can be used for data with a changing schema
- How to write Avro messages to Kafka
- How to use the Schema Registry for better performance

Schema Management In Kafka

- **An Introduction to Avro**
- *Avro Schemas*
- *The Schema Registry*
- *Hands-On Exercise: Using Kafka with Avro*
- *Chapter Review*

What is Serialization?

- **Serialization is a way of representing data in memory as a series of bytes**
 - Needed to transfer data across the network, or store it on disk
- **Deserialization is the process of converting the stream of bytes back into the data object**
- **Java provides the `Serializable` package to support serialization**
 - Kafka has its own serialization classes in `org.apache.kafka.common.serialization`
- **Backward compatibility and support for multiple languages are a challenge for any serialization system**

The Need for a More Complex Serialization System

- So far, all our data has been plain text
- This has several advantages, including:
 - Excellent support across virtually every programming language
 - Easy to inspect files for debugging
- However, plain text also has disadvantages:
 - Data is not stored efficiently
 - Non-text data must be converted to strings
 - No type checking is performed
 - It is inefficient to convert binary data to strings

Avro: An Efficient Data Serialization System



- **Avro is an Apache open source project**
 - Created by Doug Cutting, the creator of Hadoop
- **Provides data serialization**
- **Data is defined with a self-describing schema**
- **Supported by many programming languages, including Java**
- **Provides a data structure format**
- **Supports code generation of data types**
- **Provides a container file format**
- **Avro data is binary, so stores data efficiently**
- **Type checking is performed at write time**

Schema Management In Kafka

- *An Introduction to Avro*
- **Avro Schemas**
- *The Schema Registry*
- *Hands-On Exercise: Using Kafka with Avro*
- *Chapter Review*

Avro Schemas

- **Avro schemas define the structure of your data**
- **Schemas are represented in JSON format**
- **Avro has three different ways of creating records:**
 - Generic
 - Write code to map each schema field to a field in your object
 - Reflection
 - Generate a schema from an existing Java class
 - Specific
 - Generate a Java class from your schema
 - This is the most common way to use Avro classes

Avro Data Types (Simple)

- Avro supports several simple and complex data types
 - Following are the most common

Name	Description	Java equivalent
boolean	True or false	boolean
int	32-bit signed integer	int
long	64-bit signed integer	long
float	Single-precision floating-point number	float
double	Double-precision floating-point number	double
string	Sequence of Unicode characters	java.lang.CharSequence
bytes	Sequence of bytes	java.nio.ByteBuffer
null	The absence of a value	null

Avro Data Types (Complex)

Name	Description
record	A user-defined field comprising one or more simple or complex data types, including nested records
enum	A specified set of values
union	Exactly one value from a specified set of types
array	Zero or more values, each of the same type
map	Set of key/value pairs; key is always a string , value is the specified type
fixed	A fixed number of bytes

- **record is the most important of these, as we will see**

Example Avro Schema (1)

```
{
  "namespace": "model",
  "type": "record",
  "name": "SimpleCard",
  "fields": [
    {
      "name": "suit",
      "type": "string",
      "doc" : "The suit of the card"
    },
    {
      "name": "card",
      "type": "string",
      "doc" : "The card number"
    }
  ]
}
```


Example Avro Schema (2)

- By default, the schema definition is placed in `src/main/avro`
 - File extension is `.avsc`
- The namespace is the Java package name, which you will import into your code
- `doc` allows you to place comments in the schema definition

Example Schema Snippet with **array** and **map**

```
{
  "name": "cards_list",
  "type" : {
    "type" : "array",
    "items": "string"
  },
  "doc" : "The cards played"
},
{
  "name": "cards_map",
  "type" : {
    "type" : "map",
    "values": "string"
  },
  "doc" : "The cards played"
},
}
```

Example Schema Snippet with enum

```
{  
  "name": "suit_type",  
  "type" : {  
    "type" : "enum",  
    "name" : "Suit",  
    "symbols" : ["SPADES", "HEARTS", "DIAMONDS", "CLUBS"]  
  },  
  "doc" : "The suit of the card"  
},
```

Schema Evolution

- **Avro schemas may evolve as updates to code happen**
 - Schema registry allows schema evolution
- **We often want compatibility between schemas**
 - Backward compatibility
 - Code with a new version of the schema can read data written in the old schema
 - Code that reads data written with the schema will assume default values if fields are not provided
 - Forward compatibility
 - Code with previous versions of the schema can read data written in a new schema
 - Code that reads data written with the schema ignores new fields
 - Full compatibility
 - Forward and Backward

Compatibility Examples

- Consider a schema written with the following fields

```
{ "name": "suit", "type": "string"},  
{ "name": "card", "type": "string"}
```

- Backward compatibility: Consumer is expecting the following schema and assumes default for omitted size field

```
{ "name": "suit", "type": "string"},  
{ "name": "card", "type": "string"},  
{ "name": "size", "type": "string", "default": "" }
```

- Forward compatibility: Consumer is expecting the following schema and ignores additional card field

```
{ "name": "suit", "type": "string"}
```

- Question: What could be an example of Full compatibility using the fields above?

Schema Management In Kafka

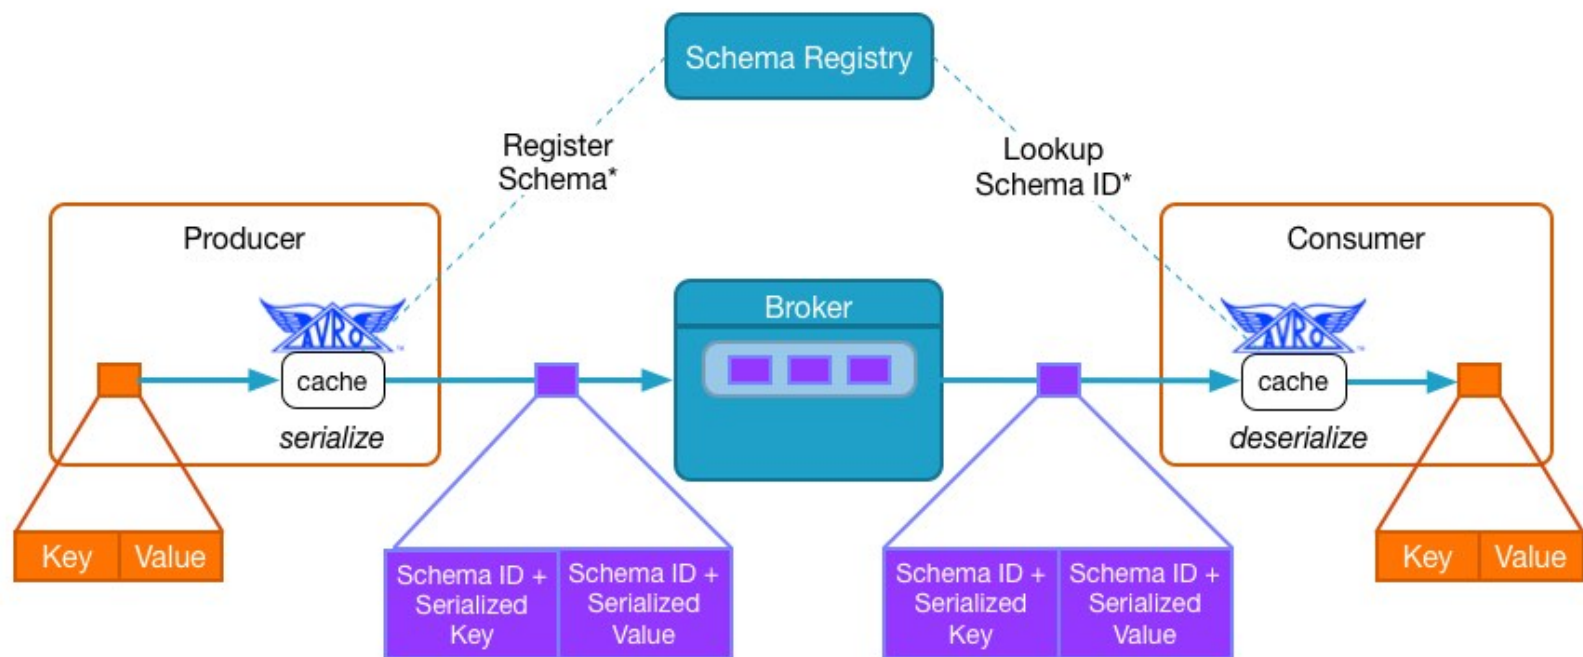
- *An Introduction to Avro*
- *Avro Schemas*
- **The Schema Registry**
- *Hands-On Exercise: Using Kafka with Avro*
- *Chapter Review*

What Is the Schema Registry?

- **Schema registry provides centralized management of schemas**
 - Stores a versioned history of all schemas
 - Provides a RESTful interface for storing and retrieving Avro schemas
 - Checks schemas and throws an exception if data does not conform to the schema
 - Allows evolution of schemas according to the configured compatibility setting
- **Sending the Avro schema with each message would be inefficient**
 - Instead, a globally unique ID representing the Avro schema is sent with each message
- **The Schema Registry stores schema information in a special Kafka topic**
- **The Schema Registry is accessible both via a REST API and a Java API**
 - There are also command-line tools, `kafka-avro-console-producer` and `kafka-avro-console-consumer`

Schema Registration and Dataflow

- The message key and value can be independently serialized
- Producers serialize data and prepend the schema ID
- Consumers use the schema ID to deserialize the data
- Schema Registry communication is only on the first message of a new schema
 - Producers and Consumers cache the schema/ID mapping for future messages



Client Support for the Schema Registry (1)

■ Java client

- New schemas automatically registered

```
Properties newProperties = new Properties();
newProperties.put(AbstractKafkaAvroSerDeConfig.SCHEMA_REGISTRY_URL_CONFIG,
"http://schemaregistry1:8081");
```

■ Python client (confluent-kafka-python)

- Support for Schema Registry introduced in Kafka 0.10.2 (Confluent 3.2)

```
value_schema = avro.load('ValueSchema.avsc')
key_schema = avro.load('KeySchema.avsc')
avroProducer = AvroProducer({'bootstrap.servers': 'broker101', 'schema.registry.url':
'http://schemaregistry1:9001'}, default_key_schema=key_schema, default_value_schema=value_schema)
```

Client Support for the Schema Registry (2)

■ Other clients

- Use the Schema Registry REST API to manually pre-register schemas and use the IDs in the requests

```
$ curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" \
  --data '{"schema": "{\"type\": \"string\"}"}' \
  http://schemaregistry1:8081/subjects/Kafka-key/versions
$ curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" \
  --data '{"schema": "{\"type\": \"string\"}"}' \
  http://schemaregistry1:8081/subjects/Kafka-value/versions
$ curl -X GET http://localhost:8081/schemas/ids/1
```

Different Versions of Schemas in the Same Topic

- Different versions of schemas in the same Topic can be Backward, Forward, or Full compatible
 - Default is BACKWARD
- If they are neither, set compatibility to NONE
 - Code has no assumptions on schema as long as it is valid Avro
 - Code has full burden to read and process data
- Configuring compatibility
 - Use REST API

```
$ curl -X PUT -i -H "Content-Type: application/vnd.schemaregistry.v1+json" \  
  --data '{"compatibility": "NONE"}' \  
  http://schemaregistry1:8081/config/my_topic
```

Java Avro Producer example

```
1 Properties props = new Properties();
2 props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "broker1:9092");
3 // Configure serializer classes
4 props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
5           io.confluent.kafka.serializers.KafkaAvroSerializer.class);
6 props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
7           io.confluent.kafka.serializers.KafkaAvroSerializer.class);
8 // Configure schema repository server
9 props.put(AbstractKafkaAvroSerDeConfig.SCHEMA_REGISTRY_URL_CONFIG,
10 "http://schemaregistry1:8081");
11 // Create the producer expecting Avro objects
12 KafkaProducer<Object, Object> avroProducer = new KafkaProducer<Object, Object>(props);
13 // Create the Avro objects for the key and value
14 CardSuit suit = new CardSuit("spades");
15 SimpleCard card = new SimpleCard("spades", "ace");
16 // Create the ProducerRecord with the Avro objects and send them
17 ProducerRecord<Object, Object> record = new
18   ProducerRecord<Object, Object>(
19     "my_avro_topic", suit, card);
20 avroProducer.send(record);
```

Java Avro Consumer Example

```
1 public class CardConsumer {
2     public static void main(String[] args) {
3         Properties props = new Properties();
4         props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "broker1:9092");
5         props.put(ConsumerConfig.GROUP_ID_CONFIG, "testgroup");
6         props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
7         "io.confluent.kafka.serializers.KafkaAvroDeserializer");
8         props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
9         "io.confluent.kafka.serializers.KafkaAvroDeserializer");
10        props.put(AbstractKafkaAvroSerDeConfig.SCHEMA_REGISTRY_URL_CONFIG,
11        "http://schemaregistry1:8081");
12        props.put(KafkaAvroDeserializerConfig.SPECIFIC_AVRO_READER_CONFIG, "true");
13
14        KafkaConsumer<CardSuit, SimpleCard> consumer = new KafkaConsumer<>(props);
15        consumer.subscribe(Arrays.asList("my_avro_topic"));
16
17        while (true) {
18            ConsumerRecords<CardSuit, SimpleCard> records = consumer.poll(100);
19            for (ConsumerRecord<CardSuit, SimpleCard> record : records) {
20                System.out.printf("offset = %d, key = %s, value = %s\n", record.offset(), record.
21                key().getSuit(), record.value().getCard());
22            }
23        }
24    }
25 }
```

Python Avro Producer Example

```
1 # Read in the Avro files
2 key_schema = open("my_key.avsc", 'rU').read()
3 value_schema = open("my_value.avsc", 'rU').read()
4
5 producerurl = "http://kafkarest1:8082/topics/my_avro_topic"
6 headers = {
7     "Content-Type" : "application/vnd.kafka.avro.v1+json"
8 }
9 payload = {
10     "key_schema": key_schema,
11     "value_schema": value_schema,
12     "records":
13     [{
14         "key": {"suit": "spades"},
15         "value": {"suit": "spades", "card": "ace"}
16     }]
17 # Send the message
18 r = requests.post(producerurl, data=json.dumps(payload), headers=headers)
19 if r.status_code != 200:
20     print "Status Code: " + str(r.status_code)
21     print r.text
```

Python Avro Consumer Example

```
1 # Get the message(s) from the consumer
2 headers = {
3     "Accept" : "application/vnd.kafka.avro.v1+json"
4 }
5 # Request messages for the instance on the topic
6 r = requests.get(base_uri + "/topics/my_avro_topic", headers=headers, timeout=20)
7 if r.status_code != 200:
8     print "Status Code: " + str(r.status_code)
9     print r.text
10    sys.exit("Error thrown while getting message")
11 # Output all messages
12 for message in r.json():
13     keysuit = message["key"]["suit"]
14     valuesuit = message["value"]["suit"]
15     valuecard = message["value"]["card"]
16    # Do something with the data
```

Command-line Consumer Example

```
$ kafka-avro-console-consumer --bootstrap-server broker1:9092 \  
--from-beginning --topic my_avro_topic
```


Schema Management In Kafka

- *An Introduction to Avro*
- *Avro Schemas*
- *The Schema Registry*
- **Hands-On Exercise: Using Kafka with Avro**
- *Chapter Review*

Hands-On Exercise: Using Kafka with Avro

- In this Hands-On Exercise, you will write and read Kafka data with Avro
- Please refer to the Hands-On Exercise Manual

Schema Management In Kafka

- *An Introduction to Avro*
- *Avro Schemas*
- *The Schema Registry*
- *Hands-On Exercise: Using Kafka with Avro*
- **Chapter Review**

Chapter Review

- Using a serialization format such as Avro makes sense for complex data
- The Schema Registry makes it easy to efficiently write and read Avro data to and from Kafka by centrally storing the schema