


# Big Stream Processing Systems

**Data Systems Group**

<https://bigdata.cs.ut.ee/>

**Big Data Management**

University of Tartu  
Fall 2019




An overview

---

<sup>1</sup>Slides are based on content from Cloudera and Confluent


# Motivation

Data pipelines start with a small number of systems to integrate. A single ETL (extract, transform, load) process move data from the source to the interested applications.




# Motivation

But data pipeline grow over time. Adding new system causes the need of new ETL process. The code-base grows together with data formats and services.




# Motivation

Things end up messy when sources and sinks are coupled!



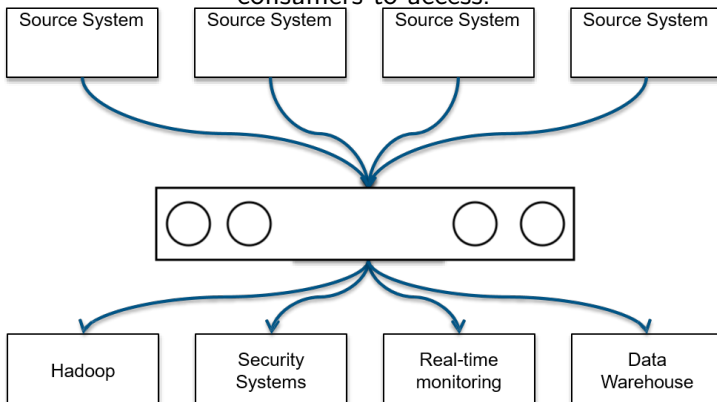
# An alternative: Publish/Subscribe

PubSubs decouple data sources and their consumers making communication asynchronous and processing scalable.




# An alternative: Publish/Subscribe

PubSubs organize messages logically so that it is easier for the interested consumers to access.



# Apache Kafka


Apache Kafka is an horizontally scalable, fault-tolerant, publish-subscribe system. It can process over 1 trillion messages without neglecting durability, i.e., it persists data on disk.



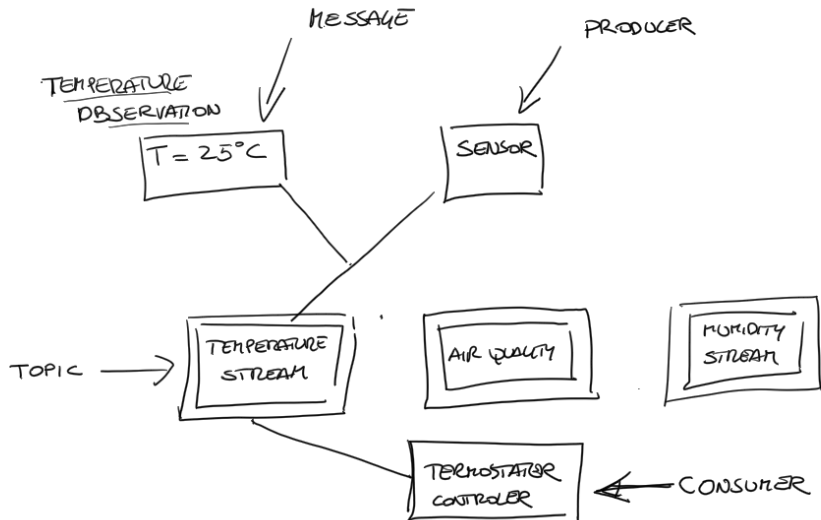


# Kafka Conceptual View

- **Messages**, the basic unit in Kafka, are organized in **Topics**
- **Producers** write messages to topics
- **Consumers** read messages from topics




# Kafka Conceptual View: Example



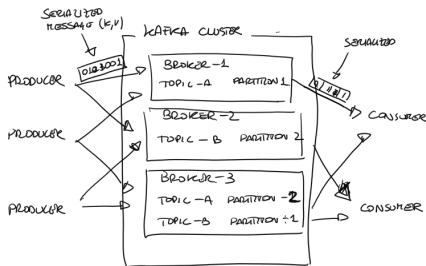
# Kafka Logical View

- **Messages** are key-value pairs
- **Brokers** are the main component inside the Kafka Cluster.
- **Producers** write messages to a certain broker
- **Consumers** read messages by from a certain broker

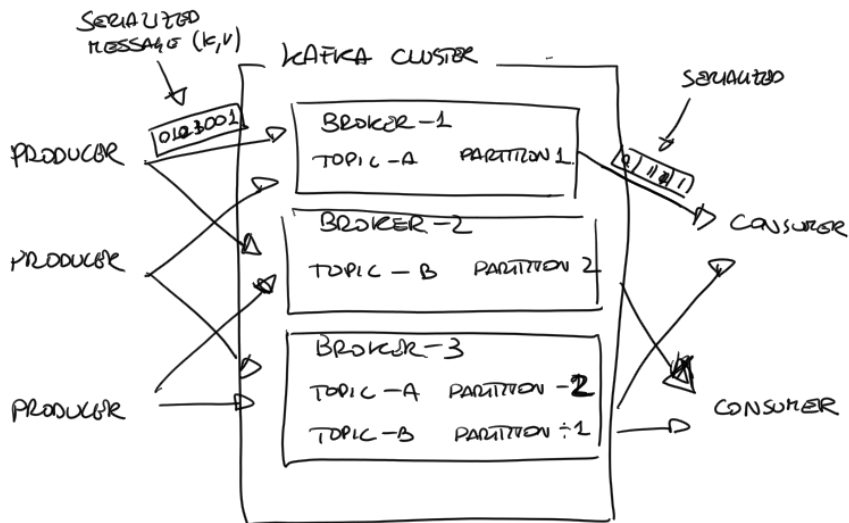


# Kafka Physical View

- **Topics** are partitioned across brokers using the message **Key**.
- Typically, **Producers** has the message key to determine the partition. Also they serialize the message
- **Consumers** read messages by from brokers and de-serialize them



# Kafka Physical View: Zoom In



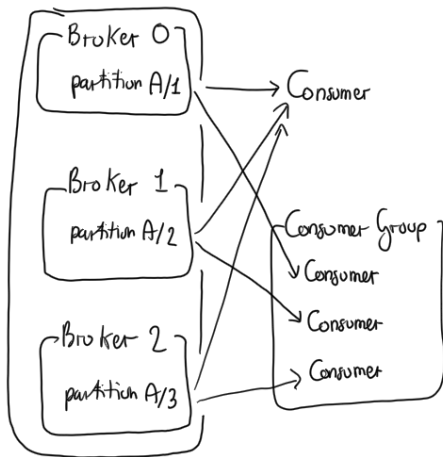
# Topics Partitions


Producers shard data over a set of Partitions

- Each Partition contains a subset of the Topic's messages
- Typically, the message key is used to determine which Partition a message is assigned to
- Each Partition is an ordered, immutable log of messages

# Topics Partitions and Distributed Consumption

- Different Consumers can read data from the same Topic
  - By default, each Consumer will receive all the messages in the Topic
- Multiple Consumers can be combined into a Consumer Group
  - Consumer Groups provide scaling capabilities
  - Each Consumer is assigned a subset of Partitions for consumption





Internals

---

<sup>2</sup>Slides are based on content from Cloudera and Confluent



# Messages and Metadata

Messages are Key-Value pairs and there is not restriction on what each of them can be.

Additionally, messages are enriched with metadata:

- Offset
- Timestamp
- Compression type
- Magic byte
- Optional message headers API
- Application teams can add custom key-value paired metadata to messages
- Additional fields to support batching, exactly once semantics, replication protocol

## Topics Partitions: Physical View

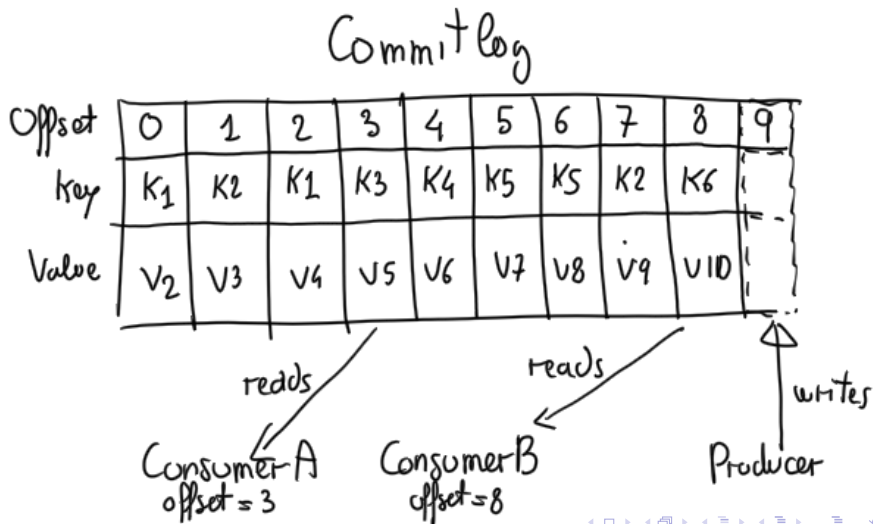
Each Partition is stored on the Broker's disk as one or more log files Each message in the log is identified by its offset number

Commit log

Offset	0	1	2	3	4	5	6	7	8	
key	K <sub>1</sub>	K <sub>2</sub>	K <sub>1</sub>	K <sub>3</sub>	K <sub>4</sub>	K <sub>5</sub>	K <sub>5</sub>	K <sub>2</sub>	K <sub>6</sub>	
Value	V <sub>2</sub>	V <sub>3</sub>	V <sub>4</sub>	V <sub>5</sub>	V <sub>6</sub>	V <sub>7</sub>	V <sub>8</sub>	V <sub>9</sub>	V <sub>10</sub>	

# Topics Partitions: Physical View

Messages are always appended. Consumers can consume from different offset. Brokers are single thread to guarantee consistency



# Topics Partitions: Load Balancing

Producers use a partition strategy to assign each message a partition

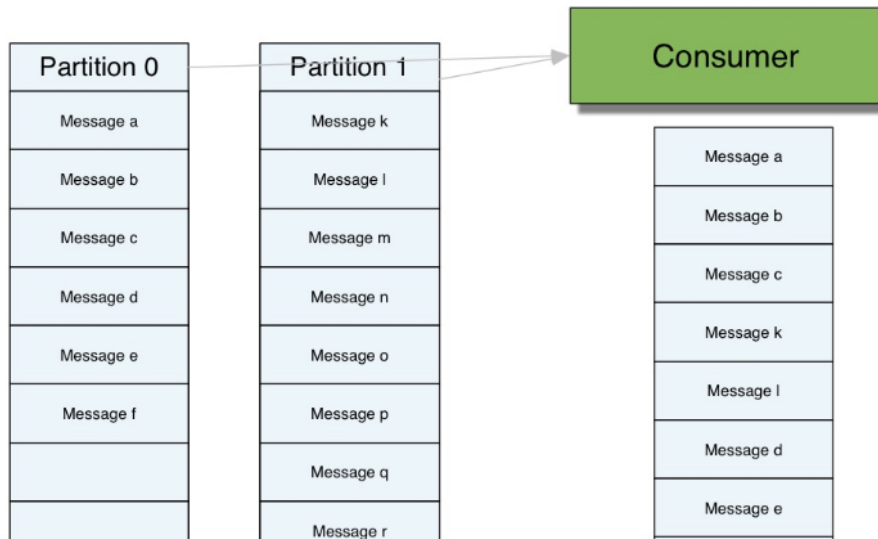
- To ensure load balancing across the Brokers
- To allow user-specified key

You can customize the partition strategy, but!

- it must ensure load balancing across the Brokers too, i.e.,  $\text{hash}(\text{key}) \ \% \ \text{number\_of\_partitions}$
- if key is not specified, messages are sent to Partitions on a round-robin basis

## Important: About Ordering

If there are multiple Partitions, you will not get total ordering across all messages when reading data



# Log Retention

- Duration default: messages will be retained for seven days
- Duration is configurable per Broker by setting
  - a time period
  - a size limit
- Topic can override a Broker's retention policy
- When cleaning up a log
  - the default policy is delete
  - An alternate policy is compact

# Log Compaction

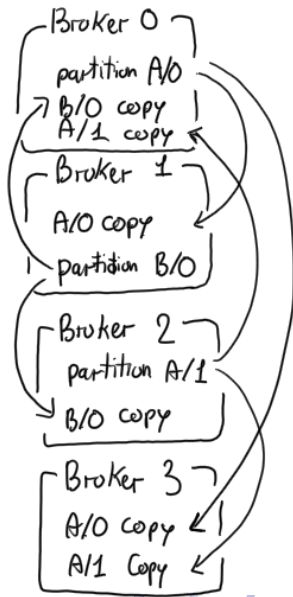
A compacted log retains at least the last known message value for each key within the Partition Before After

Offset	0	1	2	3	4	5	6	7	8	
key	K <sub>1</sub>	K <sub>2</sub>	K <sub>1</sub>	K <sub>3</sub>	K <sub>4</sub>	K <sub>5</sub>	K <sub>5</sub>	K <sub>2</sub>	K <sub>6</sub>	
Value	V <sub>2</sub>	V <sub>3</sub>	V <sub>4</sub>	V <sub>5</sub>	V <sub>6</sub>	V <sub>7</sub>	V <sub>8</sub>	V <sub>9</sub>	V <sub>10</sub>	

K <sub>1</sub>	K <sub>3</sub>	K <sub>4</sub>	K <sub>5</sub>	K <sub>2</sub>	K <sub>6</sub>
V <sub>4</sub>	V <sub>5</sub>	V <sub>6</sub>	V <sub>8</sub>	V <sub>9</sub>	V <sub>10</sub>

# Fault Tolerance via a Replicated Log

- Kafka maintains replicas of each partition on other Brokers in the cluster
  - Number of replicas is configurable
- One Broker is the leader for that Partition
  - All writes and reads go to and from the leader
  - Other Brokers are followers
- Replication provides fault tolerance in case a Broker goes down






# Important: Clients do not Access Followers

It is important to understand that Producers and Consumers only write/read to/from the leader

- Replicas only exist to provide reliability in case of Broker failure
- If a leader fails, the Kafka cluster will elect a new leader from among the followers



In the diagram, m1 hashes to Partition 0 and m2 hashes to Partition 1

# Delivery Semantics

- At least once
  - Messages are never lost but may be redelivered
- At most once
  - Messages are lost but never redelivered
- Exactly once
  - Messages are delivered once and only once

- ZooKeeper is a centralized service that stores configurations for distributed applications
- Kafka Brokers use ZooKeeper for a number of important internal features
  - Cluster management
  - Failure detection and recovery
  - Access Control List (ACL) storage

Provide the correct relationship - 1:1, 1:N, N:1, or N:N -

- Broker to Partition - ?
- Key to Partition - ?
- Producer to Topic - ?
- Consumer Group to Topic - ?
- Consumer (in a Consumer Group) to Partition - ?


# Quiz

Provide the correct relationship - 1:1, 1:N, N:1, or N:N -

- Broker to Partition - N:N
- Key to Partition - N:1
- Producer to Topic - N:N
- Consumer Group to Topic - N:N
- Consumer (in a Consumer Group) to Partition - 1:N

# Getting Exactly Once Semantics

- Must consider two components
  - Durability guarantees when publishing a message
  - Durability guarantees when consuming a message
- Producer
  - What happens when a produce request was sent but a network error returned before an ack?
  - Use a single writer per partition and check the latest committed value after network errors
- Consumer
  - Include a unique ID (e.g. UUID) and de-duplicate.
  - Consider storing offsets with data



Streams

---


<sup>3</sup>Slides are based on content from Cloudera and Confluent

# Stream Processing with Kafka


- As of version 0.10.0, Kafka streams library has been added to Kafka distribution
- It is no longer just a distributed message broker
- You can process messages in the different Kafka topics in real-time
- You can produce new messages to (other) topics



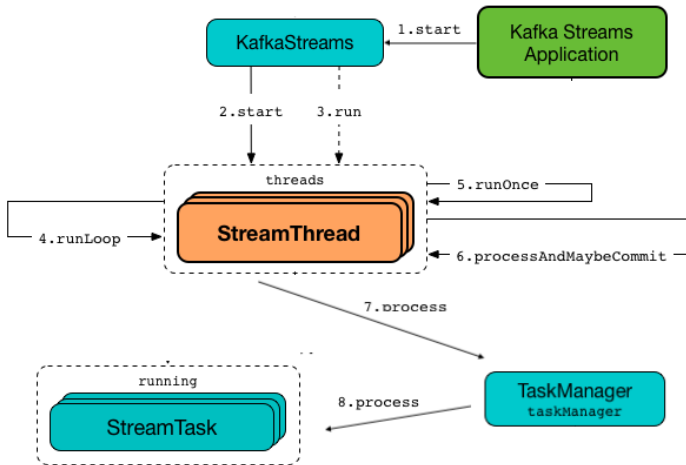
# Kafka Streams Library



# Kafka Streams: Closer Look




# Kafka Streams: Even Closer Look



# Kafka Streams: How

- Stream Application is the main abstraction for the User to interact
- Data streams are elicited from topics
- Stream Threads are stream processor threads (a Java Thread) that runs the main record processing loop when started
- Stream Task are build upon Producer and Consumer APIs

# Kafka Streams Model<sup>4</sup>



<sup>4</sup>Sax, Matthias J., et al. "Streams and tables: Two sides of the same coin." Proceedings of the International Workshop on Real-Time Business Intelligence and Analytics. 2018.


# Stream-Table Duality<sup>5</sup>

- Streams carry individual stateless events
  - An individual event can represent a state change, e.g., for a table
- Tables can be used as a state at a certain time (snapshot)
  - Accumulation of the individual events (stateful)

---

<sup>5</sup>Matthias J. Sax, Guozhang Wang, Matthias Weidlich, Johann-Christoph Freytag.  
*Streams and Tables: Two Sides of the Same Coin*. BIRTE 2018: 1:1-1:10

# Stream-Table: Chess Analogy




Streams record History  
"The sequence of moves."



Tables represent State  
"The state of the board at last move."

*Source: Micheal Noll*


# Example



*Source: Micheal Noll*



# Example




*Source: Micheal Noll*

# Example




*Source: Micheal Noll*

# Example




*Source: Micheal Noll*

# Example



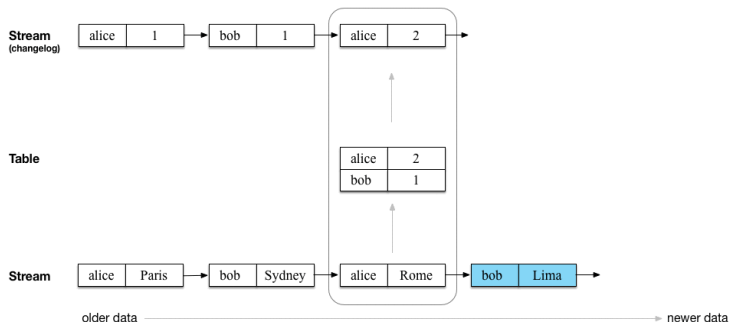
*Source: Micheal Noll*

# Example



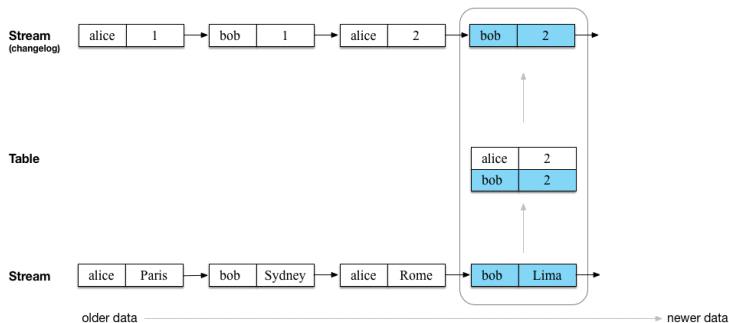
*Source: Micheal Noll*

# Example



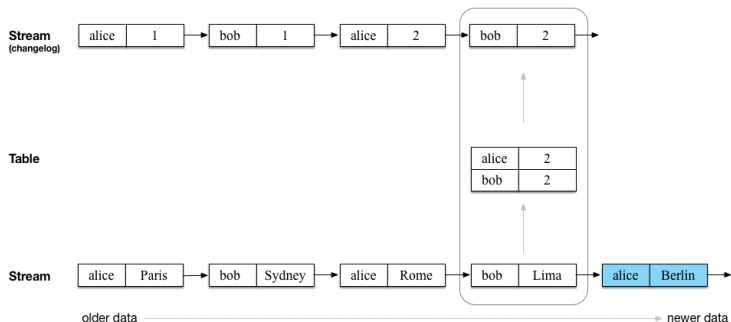
*Source: Micheal Noll*

# Example



*Source: Micheal Noll*

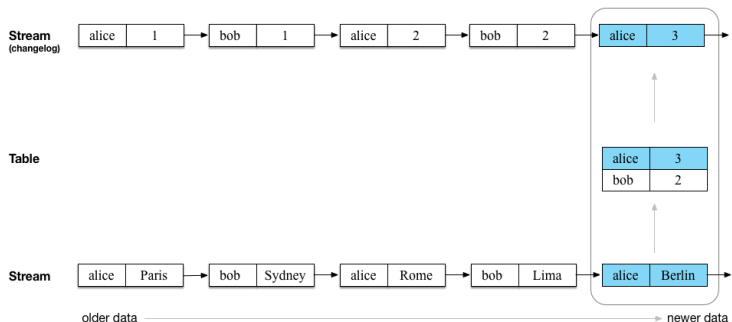
# Example



*Source: Micheal Noll*



# Example

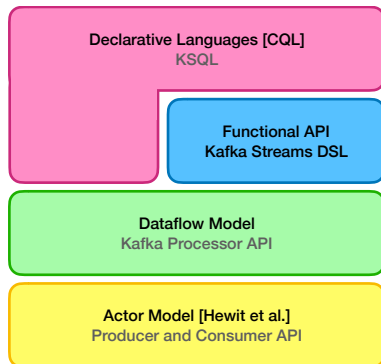


*Source: Micheal Noll*

# Programming with Streams

Stream processing frameworks hide execution details from the programmers, and manage them in the background.

There are different abstraction levels that a programmer can use to express streaming computations.



# Kafka Streams: APIs



Kafka allows writing streaming programs at different levels of abstraction:

- Kafka Processor API builds directly on top of Apache Kafka. It allows defining operators and organize them in topologies, i.e., Direct-Acyclic Graphs (DAG).
- Kafka Streams DSL builds on Kafka Processor API. It implements a number of functional operations and let the programmers design data pipeline that are automatically translated into DAGs.
- KSQL is yet another level on top of Kafka Streams. It provides a SQL-based declarative syntax that allows defining simple ETL (Extract, Transform, Load) jobs on top of Kafka Streams DSL.

- KStream
  - Record stream
  - Each record describes an event in the real world
  - Example: click stream
- KTable
  - Changelog stream
  - Each record describes a change to a previous record
  - Example: position report stream


# Processor Topology

- Close idea to Storm Topology
  - DAG in General
- Several topologies can be linked together
  - Achievable via writing back to Kafka




# Kafka Streams DSL

```
KStream<..> stream1 = builder.stream("topic1");
KStream<..> stream2 = builder.stream("topic2");
KStream<..> joined = stream1.leftJoin(stream2, ...);
KTable<..> aggregated = joined.aggregateByKey(...);
aggregated.to("topic3");
```



# Stateful Processing

- Stateful processors
  - Windowing
  - Joins
  - Aggregation
- Kafka provides a configurable local state store
  - Memory
  - Disk



# Notions of Time

- Recall we have
  - Event time: when the data was actually generated
  - Processing time: when the data was received/processed by the system
- Kafka provides a uniform `Timestamp Extractor`
  - Based on Kafka configuration `log.message.timestamp.type`, Kafka streams will read either the ingestion or the event time (default)
  - You can still create your own extractor



# Windowing

- Kafka Streams supports time-based windows only
  - **Tumbling**
  - Sliding (called hopping)
  - Session

```
KStream<String, GenericRecord> pageViews = ...;  
// Count page views per window, per user, with tumbling windows of size 5 minutes  
KTable<Windowed<String>, Long> windowedPageViewCounts = pageViews  
    .groupByKey(Grouped.with(Serdes.String(), genericAvroSerde))  
    .windowedBy(TimeWindows.of(Duration.ofMinutes(5))) .count();
```

# Windowing

- Kafka Streams supports time-based windows only
  - Tumbling
  - **Sliding (called hopping)**
  - Session

```
KStream<String, GenericRecord> pageViews = ...;  
// Count page views per window, per user, with hopping windows of size 5 minutes  
// that advance every 1 minute  
KTable<Windowed<String>, Long> windowedPageViewCounts = pageViews  
    .groupByKey(Grouped.with(Serdes.String(), genericAvroSerde))  
    .windowedBy(TimeWindows.of(Duration.ofMinutes(5)).advanceBy(Duration.ofMinutes(1))).count()
```

# Windowing

- Kafka Streams supports time-based windows only
  - Tumbling
  - Sliding (called hopping)
  - **Session**

```
KStream<String, GenericRecord> pageViews = ...;  
// Count page views per session, per user, with session windows that have  
// an inactivity gap of 5 minutes  
KTable<Windowed<String>, Long> sessionizedPageViewCounts = pageViews  
    .groupByKey(Grouped.with(Serdes.String(), genericAvroSerde))  
    .windowedBy(SessionWindows.with(Duration.ofMinutes(5))).count();
```

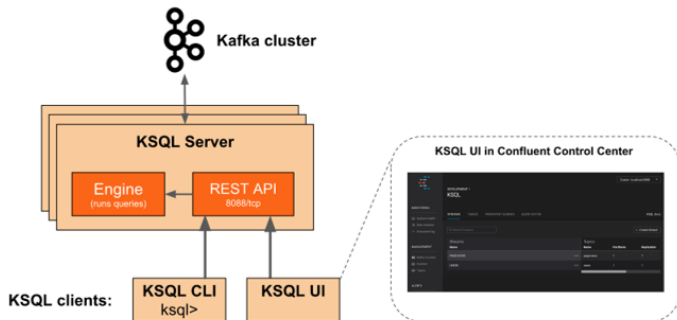
# Late Arrival

- By default, Kafka streams emits an updated result for each newly arriving record
  - No need for special handling for late arrival, simply a new result is emitted
  - Window retention is a configuration parameter in Kafka, default is one day
  - No watermark support
- Some times, you need to receive one result at the end of the window
  - You can do that in your code

```
KGroupedStream<UserId, Event> grouped = ...;
grouped.windowedBy(TimeWindows.of(Duration.ofHours(1)).grace(Duration.ofMinutes(10)))
.count()
.suppress(Suppressed.untilWindowCloses(unbounded()))
.filter((windowedUserId, count) -> count < 3)
.toStream()
.foreach((windowedUserId, count) -> sendAlert(windowedUserId.window(),
windowedUserId.key(), count));
```

- Brings SQL support to Kafka Streams
- Streaming ETL
  - DDL
  - Querying
  - Link streams to tables
    - Don't confuse it with KStream and KTable

# KSQL



# Create Stream

Stream is KSQL's wrapper for the data in a Kafka topic

```
CREATE STREAM ratings (  
  rating_idlong ,  
  user_idint ,  
  stars int ,  
  route_idint ,  
  rating_timelong ,  
  channel varchar ,  
  message varchar )  
WITH(  
  value_format='JSON', kafka_topic='ratings');
```

Content of  
messages

```
CREATE STREAM ratings (  
  rating_idlong,  
  user_idint,  
  stars int,  
  route_idint,  
  rating_timelong,  
  channel varchar,  
  message varchar)  
WITH(  
  value_format='JSON', kafka_topic='ratings');
```

# Selecting From the Stream

```
SELECT *  
FROM ratings  
WHERE stars <= 2  
AND lcase(channel) LIKE '%ios%'  
AND user_id > 0  
LIMIT 10;
```



# Selecting From the Stream

```
SELECT *  
FROM ratings  
WHERE stars <= 2  
AND lcase(channel) LIKE '%ios%'  
AND user_id > 0  
LIMIT 10;
```

We can derive another stream based on the query result



```
CREATE STREAM poor_ratings AS  
SELECT *  
FROM ratings  
WHERE stars <= 2  
AND lcase(channel) LIKE '%ios%';
```

# Create Table

```
CREATE TABLE users (  
  uidint ,  
  name varchar ,  
  elite varchar)  
WITH(  
  Key= 'uid' ,  
  value_format='JSON' , kafka_topic='mysql-users' );
```

# Enrich Stream with Table Data

```
CREATE STREAM vip-poor-ratings AS  
SELECT uid, name, elite ,  
stars, route_id, rating_time, message  
FROM poor-ratingsr LEFT JOIN users u ON r.user_id= u.uid  
WHERE u.elite= 'P';
```

# Aggregation and Windowing

```
SELECT uid, name, count(*) as rating_count  
FROM vip_poor_ratings  
WINDOW TUMBLING(size 2 minutes)  
GROUPBY uid, name;
```

# The End

Thank You  
Mahalo  
Kiitos  
Tack  
Grazie  
Obrigado  
Takk  
Gracias  
Merci  
Thanks  
Toda

