

# Data Engineering

LTAT.02.007

Ass Prof. Riccardo Tommasini

Assistants: **Fabiano Spiga, Mohamed Ragab, Hassan Eldeeb**

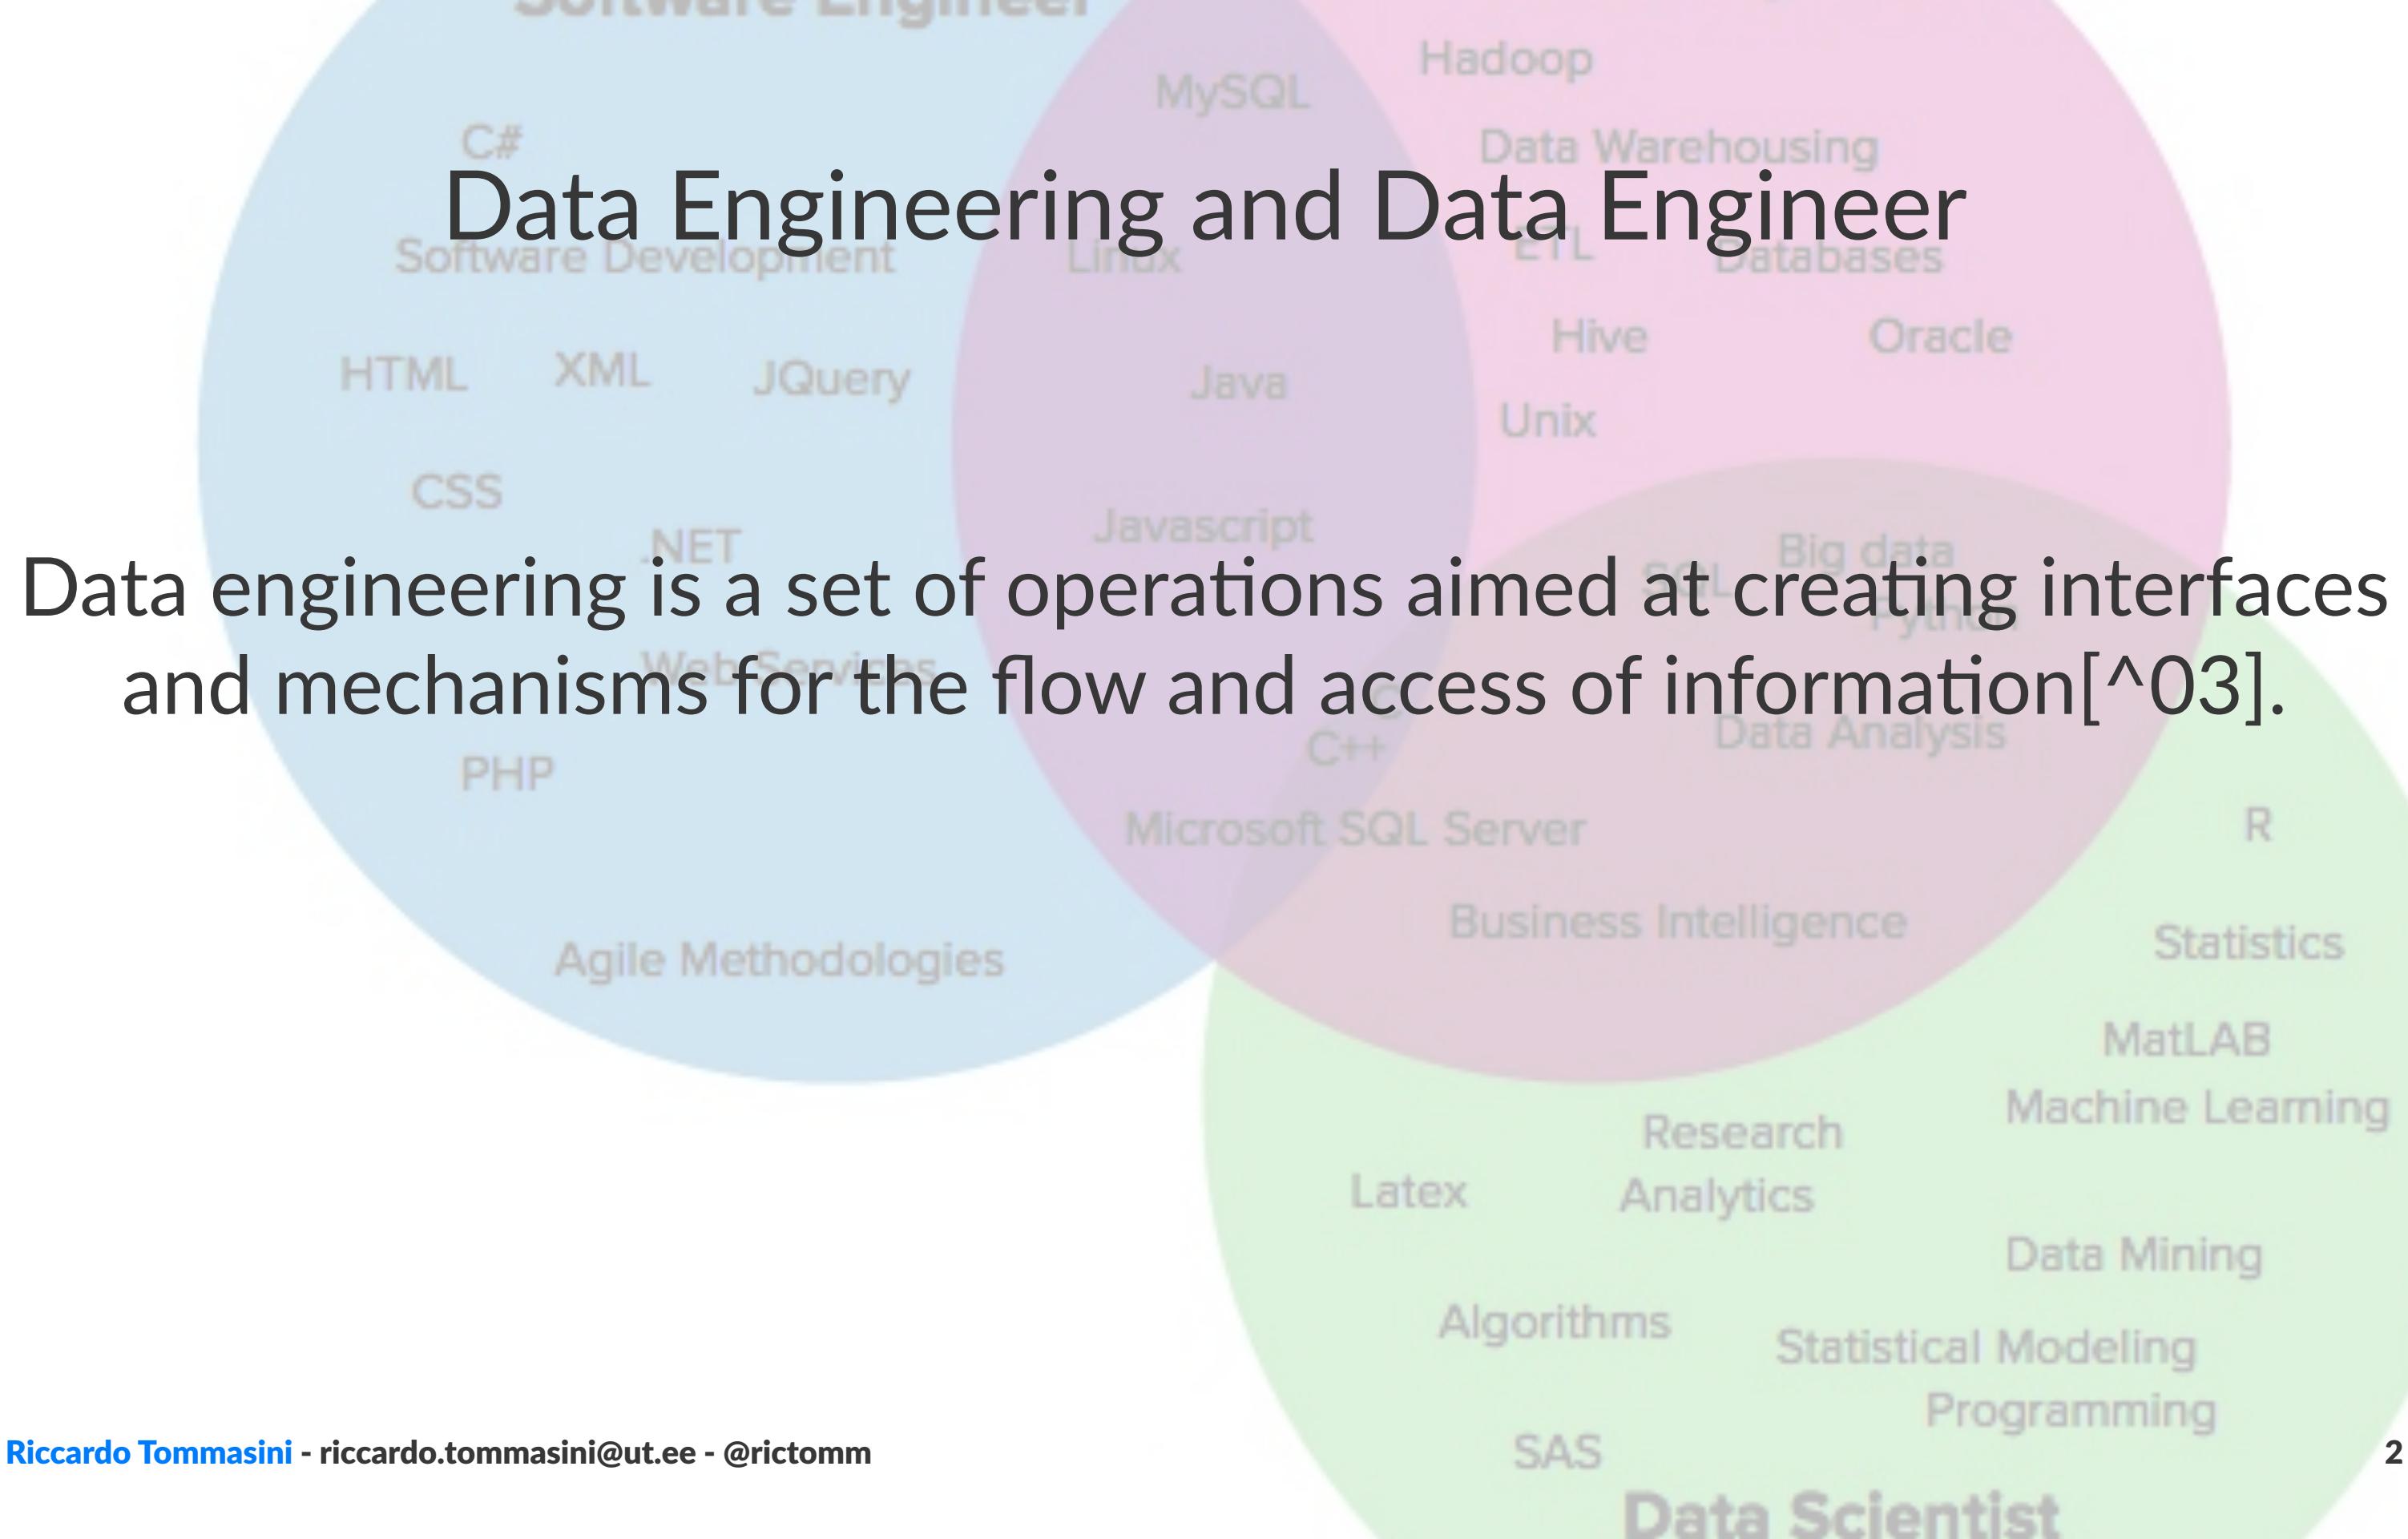


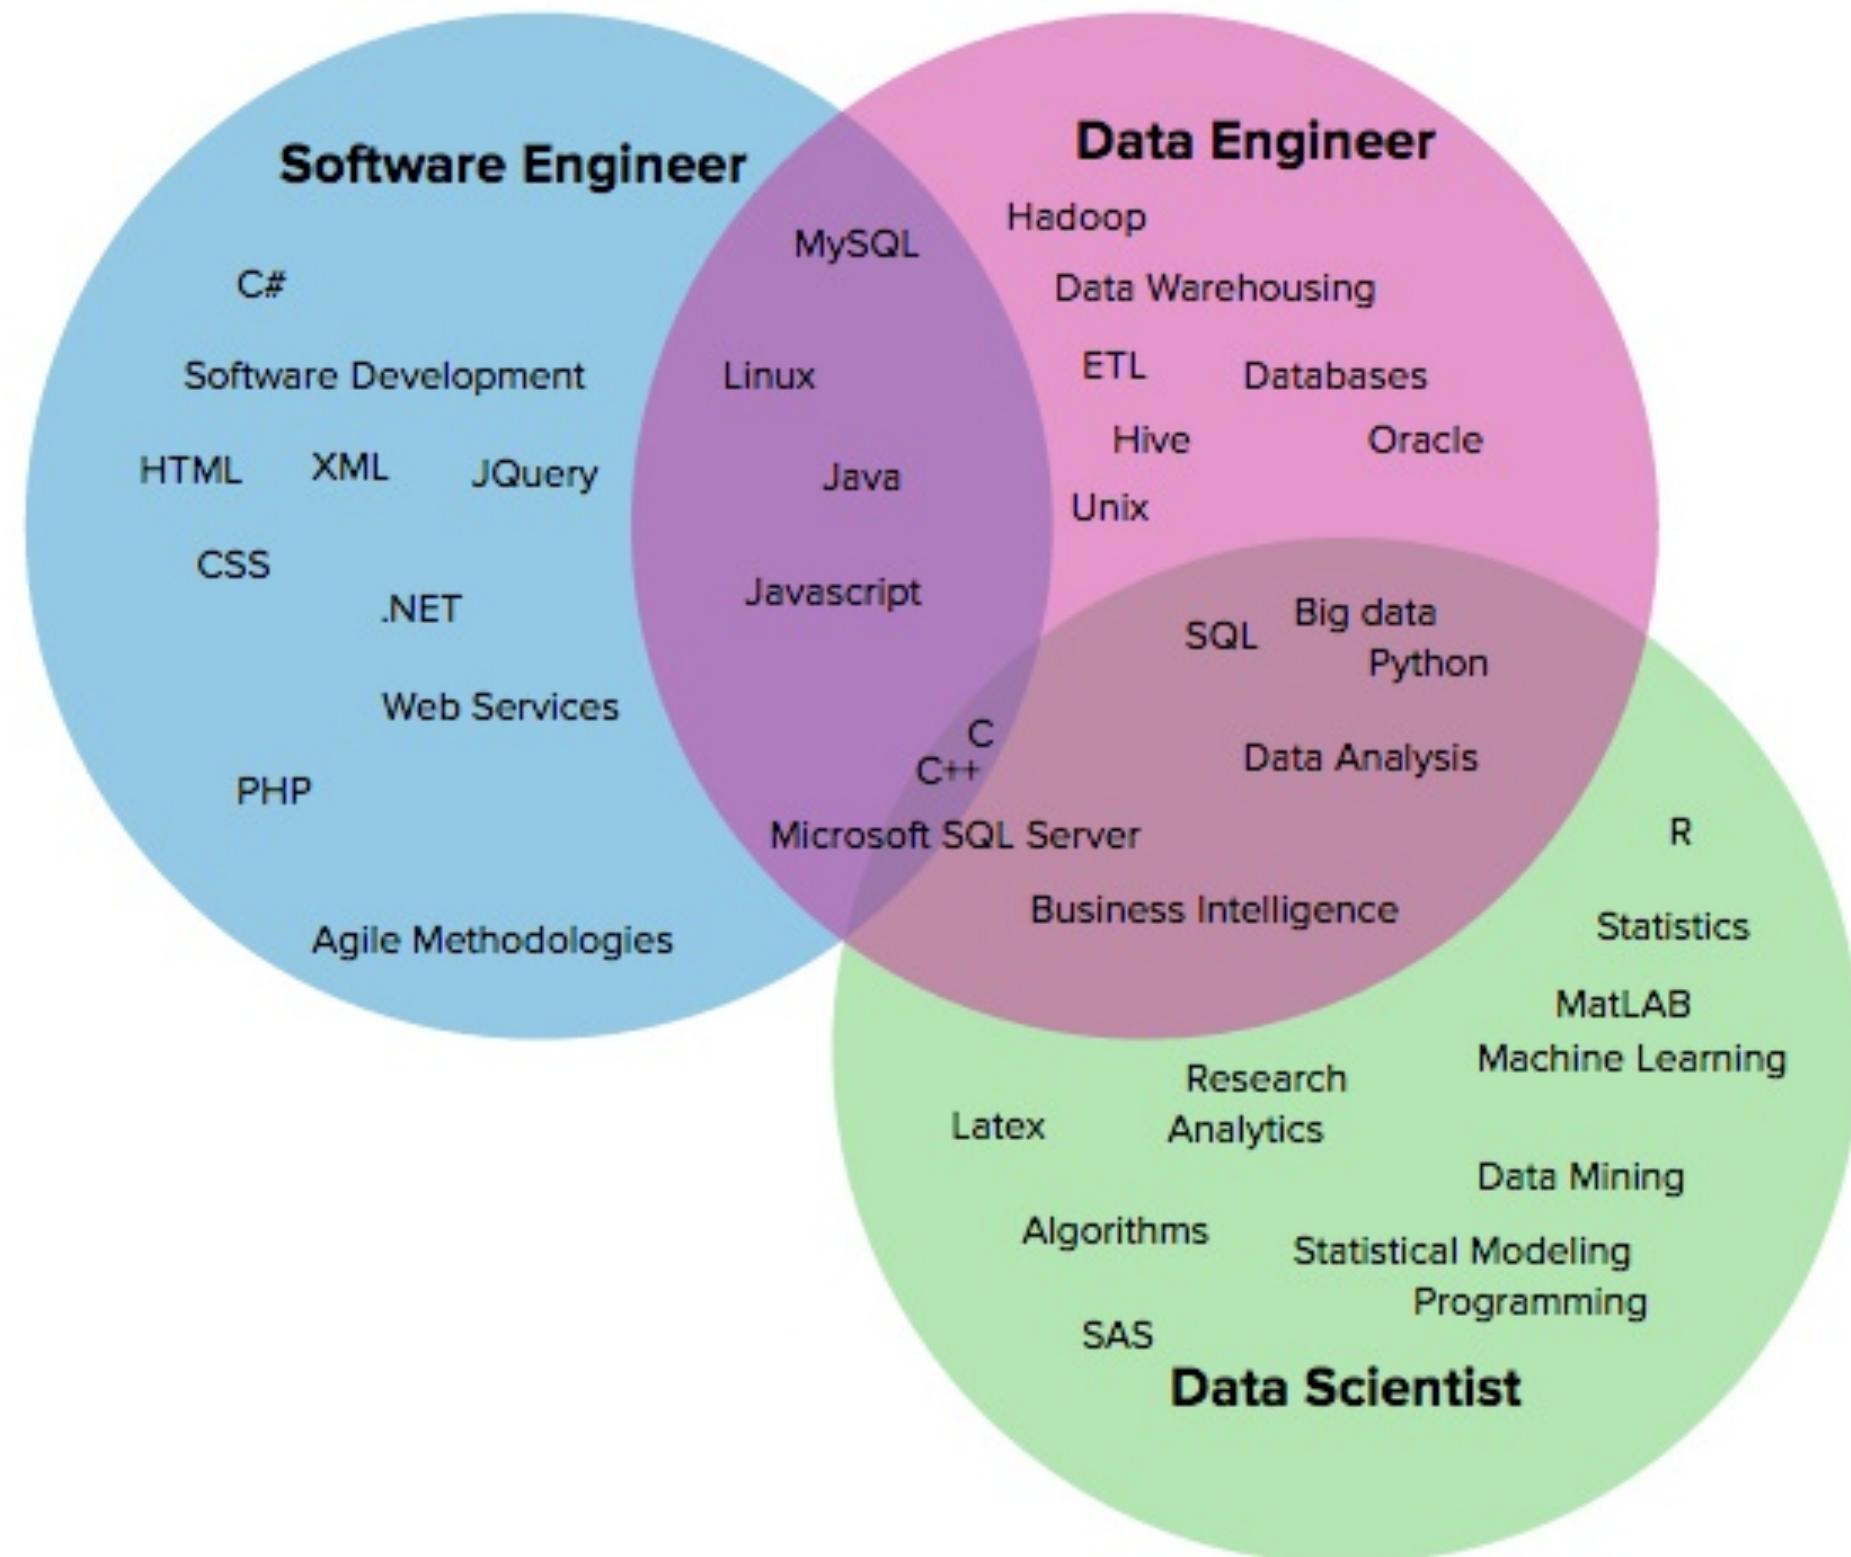
[https://courses.cs.ut.ee/2020/  
dataeng](https://courses.cs.ut.ee/2020/dataeng)

Forum

Moodle

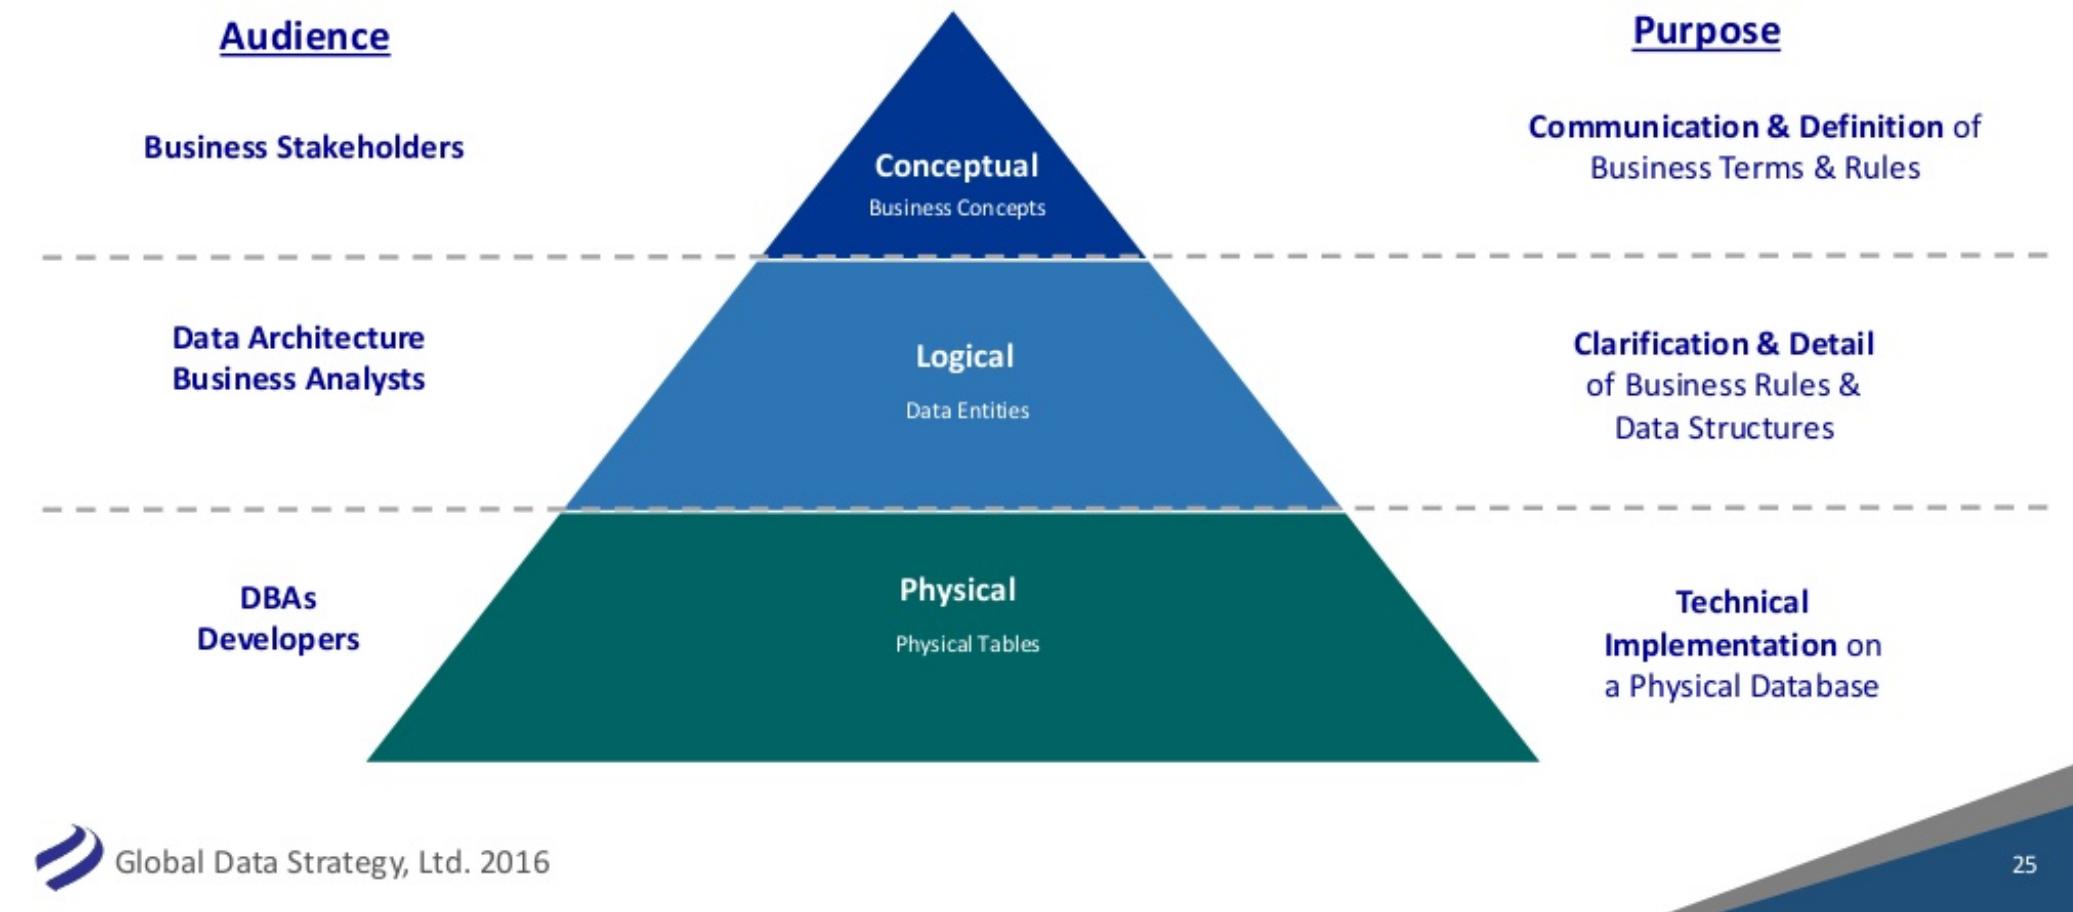






# Data Modelling

## Levels of Data Modeling



# CAP Theorem (Brewer's Theorem)

It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:

# CAP Theorem (Brewer's Theorem)

It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:

- **Consistency**: ???

# CAP Theorem (Brewer's Theorem)

It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:

- **Consistency**: ???
- **Availability**: ???

# CAP Theorem (Brewer's Theorem)

It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:

- **Consistency**: ???
- **Availability**: ???
- **Partition tolerance**: ???

# CAP Theorem (Brewer's Theorem)

It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:

- **Consistency**: all nodes see the same data at the same time
- **Availability**: ???
- **Partition tolerance**: ???

# CAP Theorem (Brewer's Theorem)

It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:

- **Consistency**: all nodes see the same data at the same time
- **Availability**: Node failures do not prevent other survivors from continuing to operate (a guarantee that every request receives a response whether it succeeded or failed)
- **Partition tolerance**: ???

# CAP Theorem (Brewer's Theorem)

It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:

- **Consistency**: all nodes see the same data at the same time
- **Availability**: Node failures do not prevent other survivors from continuing to operate (a guarantee that every request receives a response whether it succeeded or failed)
- **Partition tolerance**: the system continues to operate despite arbitrary partitioning due to network failures (e.g., message loss)

## Data Replication

Replication means keeping a copy of the same data on multiple machines that are connected via a network

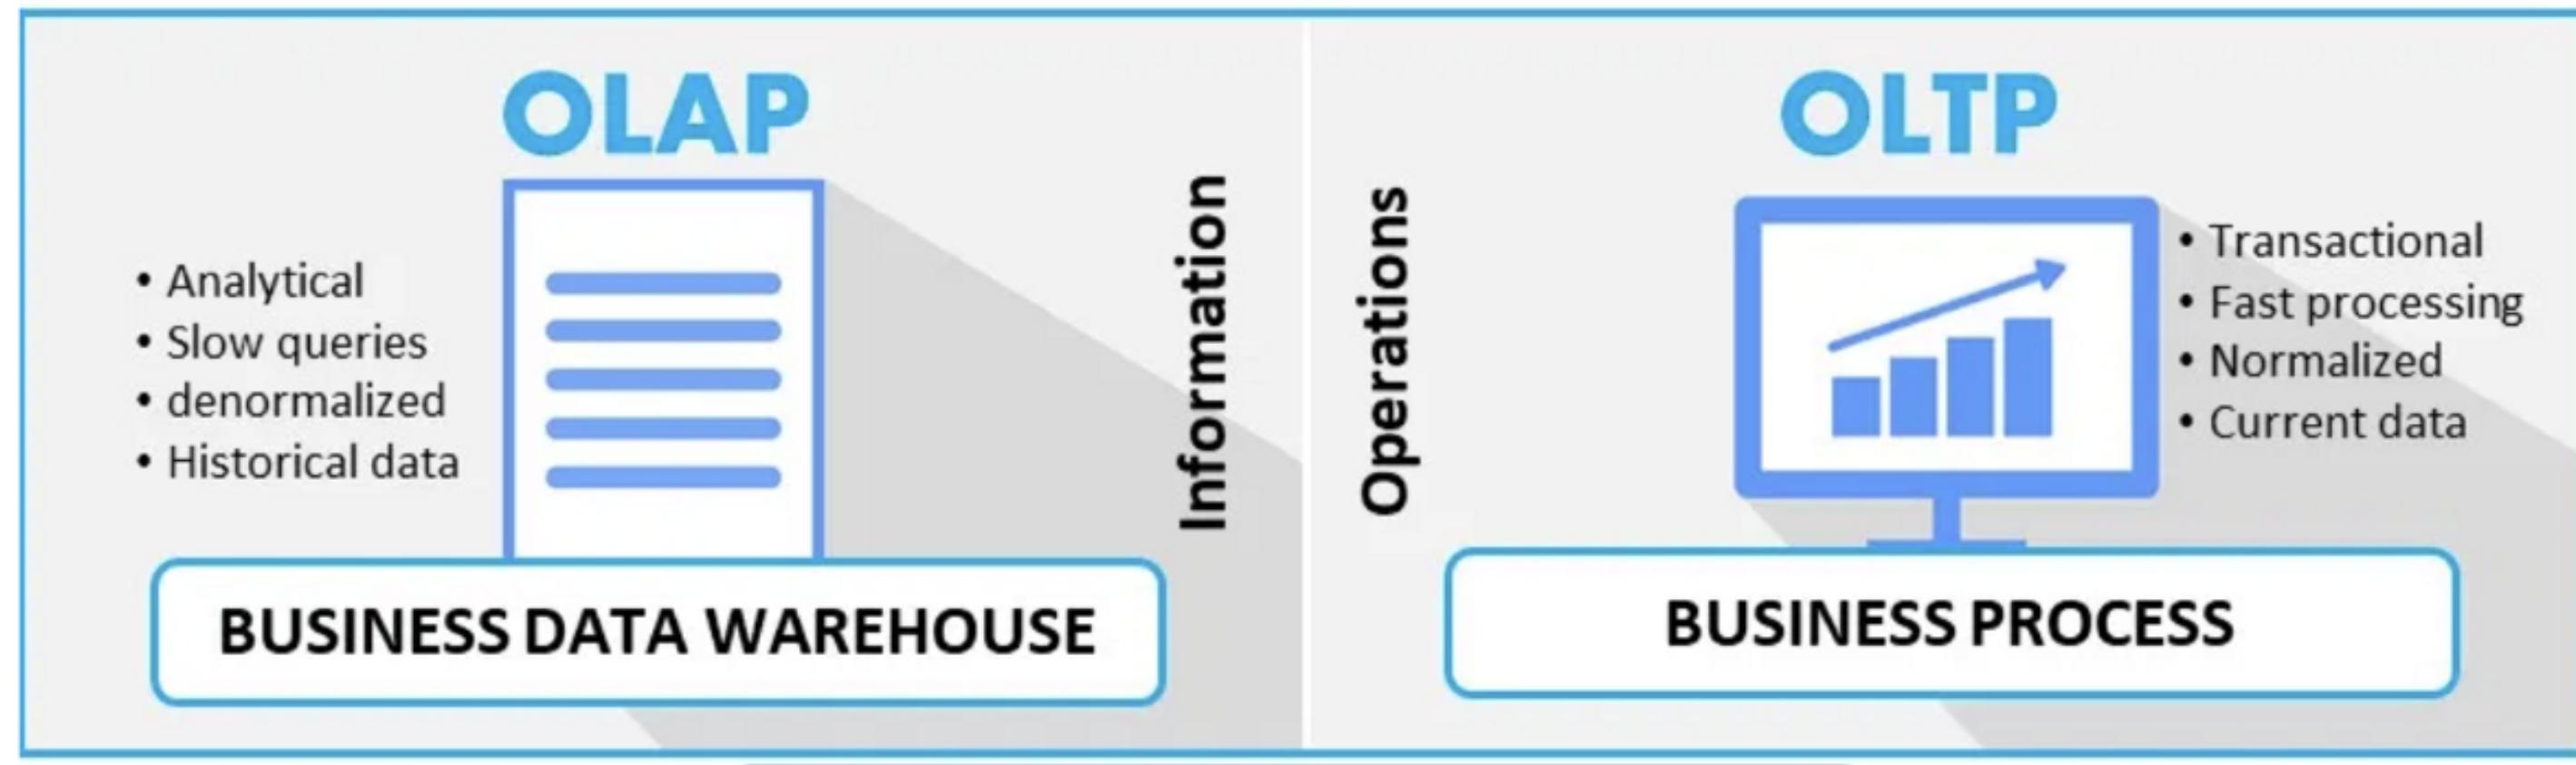


# Data Partitioning (Sharding)

breaking a large database down into smaller ones

# We Talked about Workloads

## OLAP Vs OLTP



# Data Modelling for Databases

## [Relational Algebra]

is a procedural language consisting of a six basic operations that take one or two relations as input and produce a new relation as their result:

## [Relational Algebra]

is a procedural language consisting of a six basic operations that take one or two relations as input and produce a new relation as their result:

- select:  $\sigma$

## [Relational Algebra]

is a procedural language consisting of a six basic operations that take one or two relations as input and produce a new relation as their result:

- select:  $\sigma$
- project:  $\Pi$

## [Relational Algebra]

is a procedural language consisting of a six basic operations that take one or two relations as input and produce a new relation as their result:

- select:  $\sigma$
- project:  $\Pi$
- union:  $\cup$

## [Relational Algebra]

is a procedural language consisting of a six basic operations that take one or two relations as input and produce a new relation as their result:

- select:  $\sigma$
- project:  $\Pi$
- union:  $\cup$
- set difference:  $-$

## [Relational Algebra]

is a procedural language consisting of a six basic operations that take one or two relations as input and produce a new relation as their result:

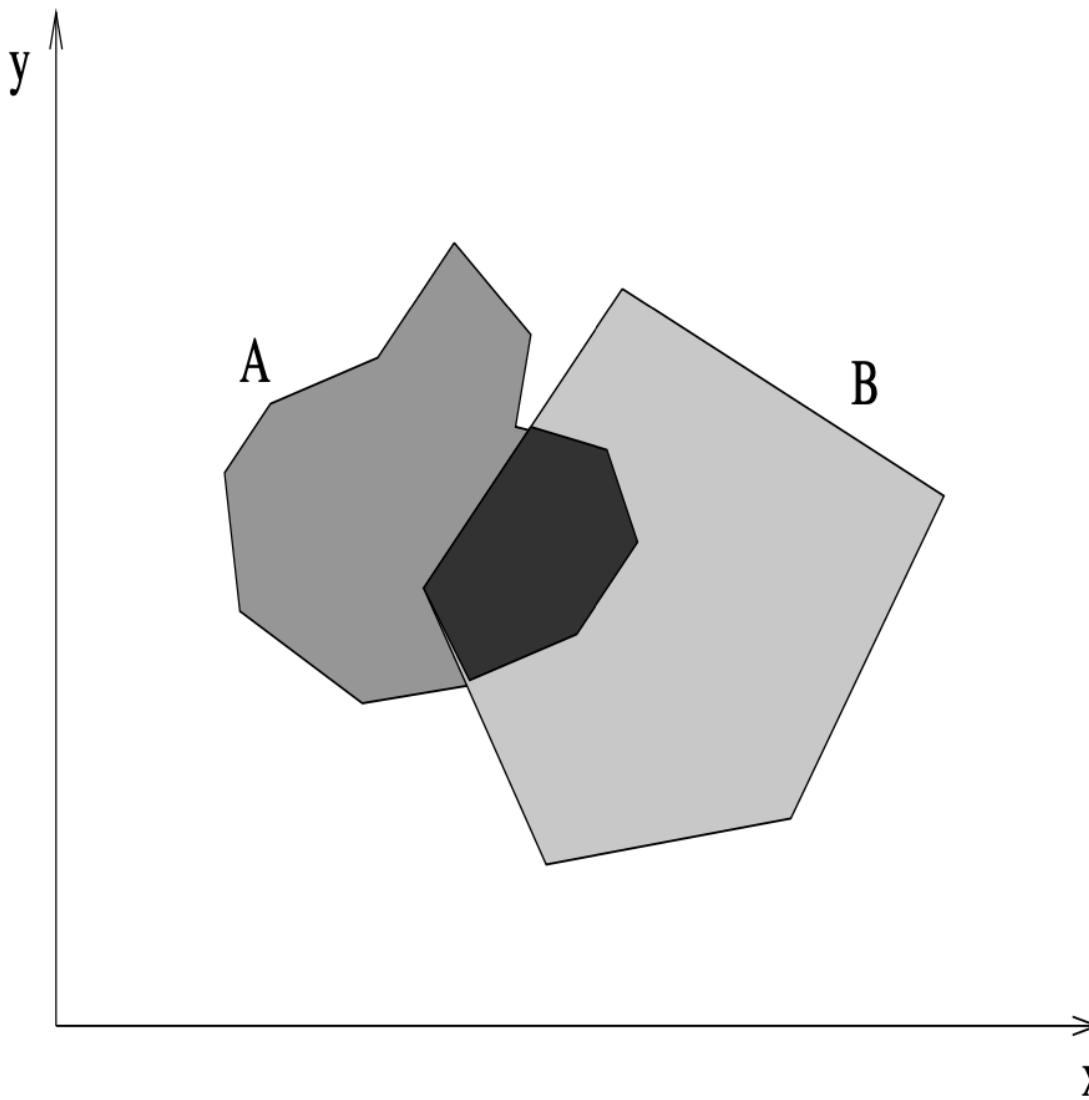
- select:  $\sigma$
- project:  $\Pi$
- union:  $\cup$
- set difference:  $-$
- Cartesian product:  $\times$

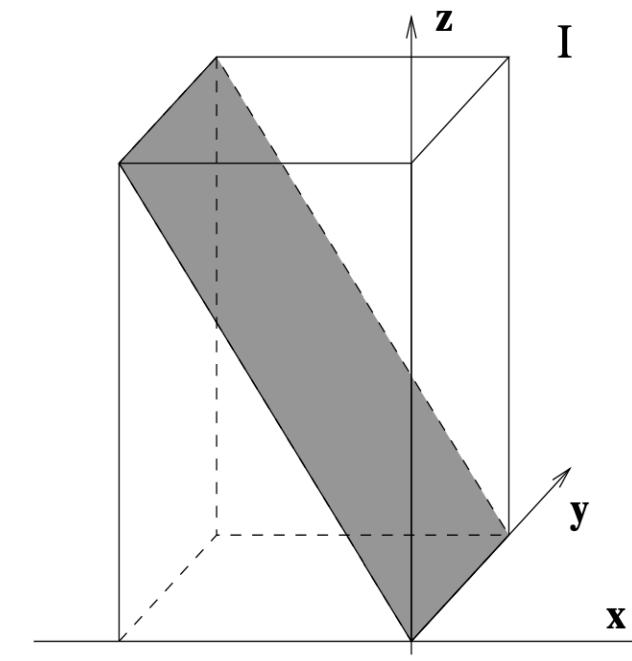
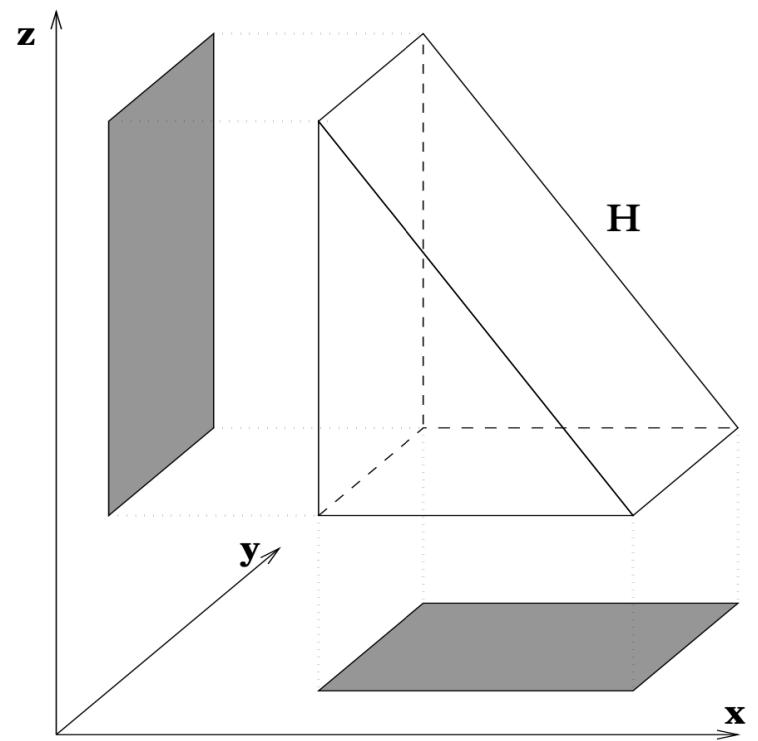
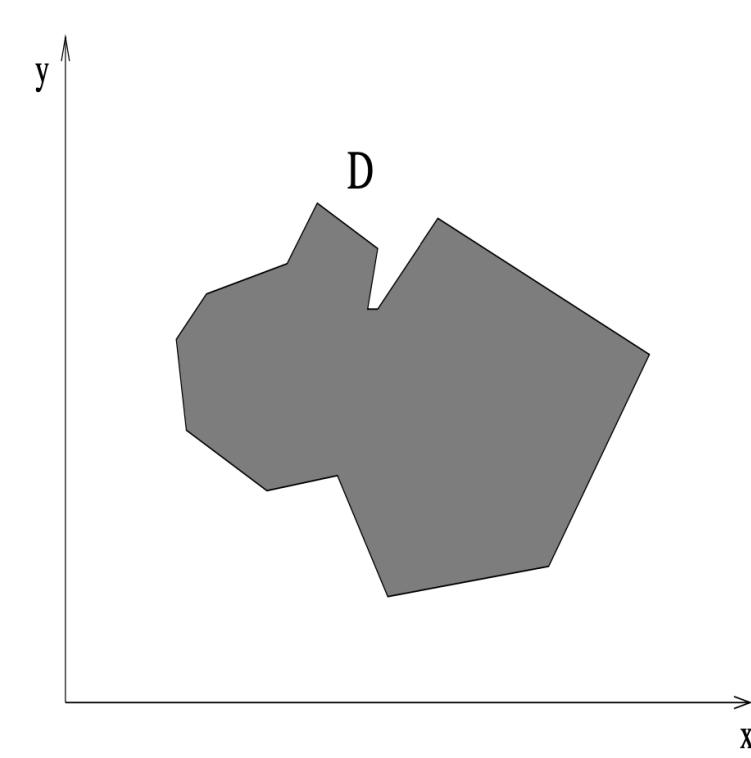
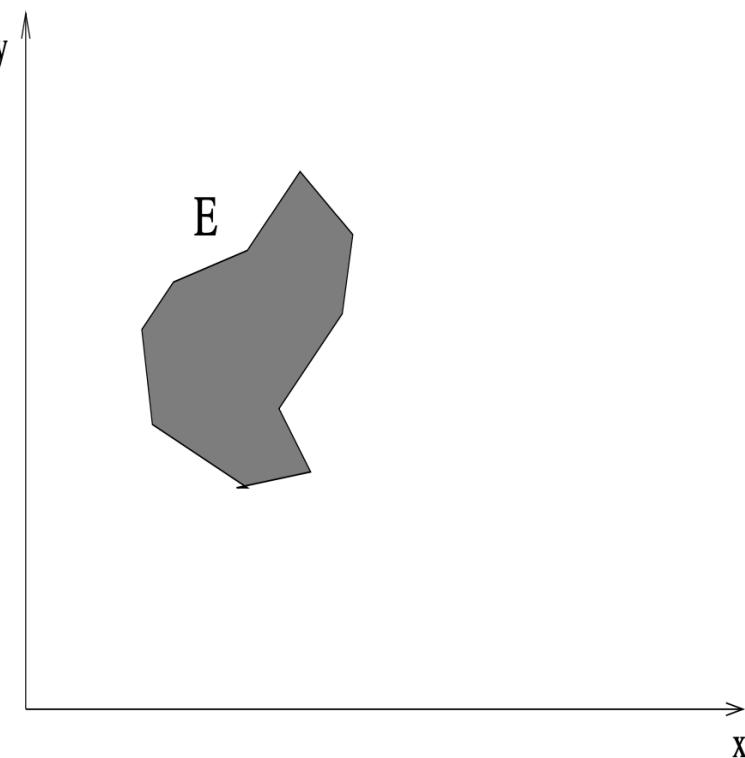
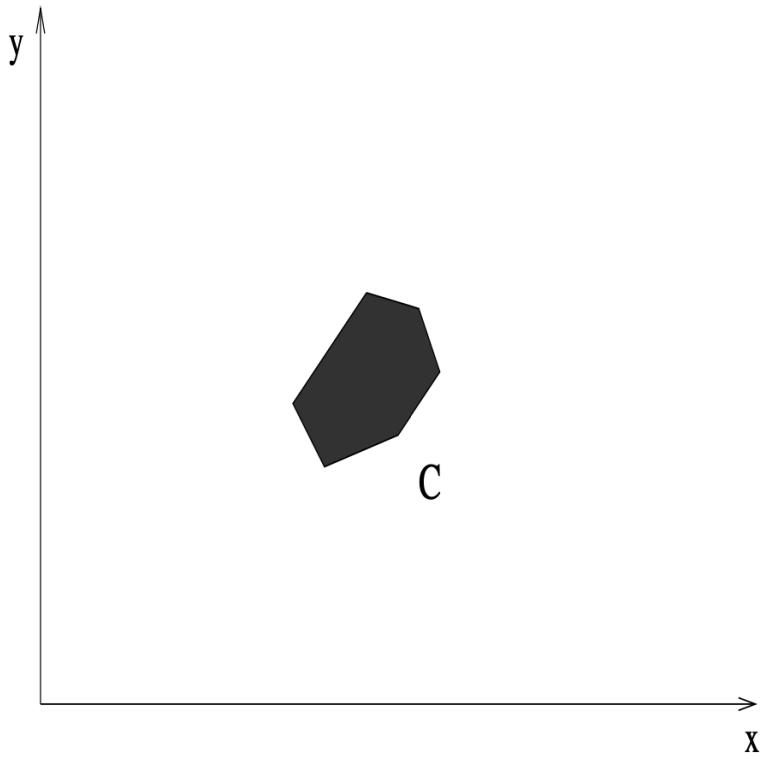
## [Relational Algebra]

is a procedural language consisting of a six basic operations that take one or two relations as input and produce a new relation as their result:

- select:  $\sigma$
- project:  $\Pi$
- union:  $\cup$
- set difference:  $-$
- Cartesian product:  $\times$
- rename:  $\rho$

# [[Relational Algebra]] Visualized





# Refresh on ACID Properties

# Refresh on ACID Properties

- ACID, which stands for Atomicity, Consistency, Isolation, and Durability[^11]

# Refresh on ACID Properties

- ACID, which stands for Atomicity, Consistency, Isolation, and Durability[^11]
- **Atomicity** refers to something that cannot be broken down into smaller parts.

# Refresh on ACID Properties

- ACID, which stands for Atomicity, Consistency, Isolation, and Durability[^11]
- **Atomicity** refers to something that cannot be broken down into smaller parts.
  - It is not about concurrency (which comes with the I)

# Refresh on ACID Properties

- ACID, which stands for Atomicity, Consistency, Isolation, and Durability[^11]
- **Atomicity** refers to something that cannot be broken down into smaller parts.
  - It is not about concurrency (which comes with the I)
- **Consistency** (overused term), that here relates to the data *invariants* (integrity would be a better term IMHO)

# Refresh on ACID Properties

- ACID, which stands for Atomicity, Consistency, Isolation, and Durability[^11]
- **Atomicity** refers to something that cannot be broken down into smaller parts.
  - It is not about concurrency (which comes with the I)
- **Consistency** (overused term), that here relates to the data *invariants* (integrity would be a better term IMHO)
- **Isolation** means that concurrently executing transactions are isolated from each other.

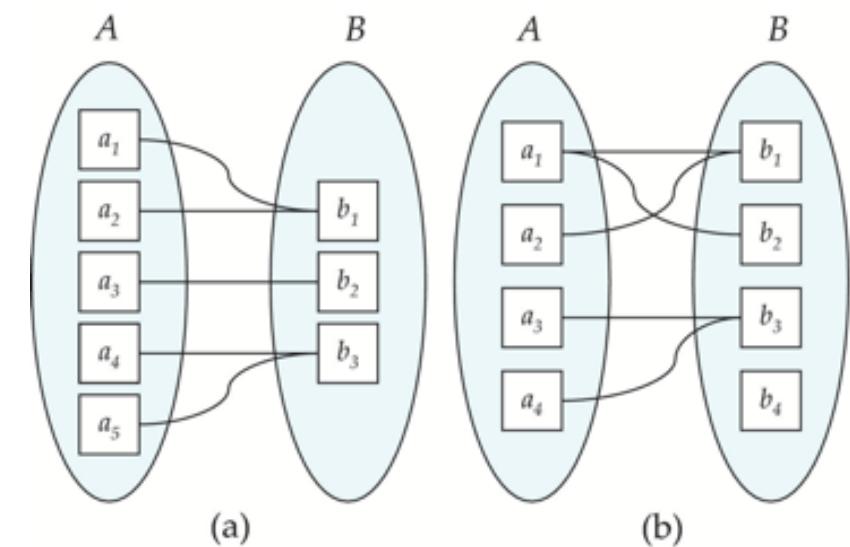
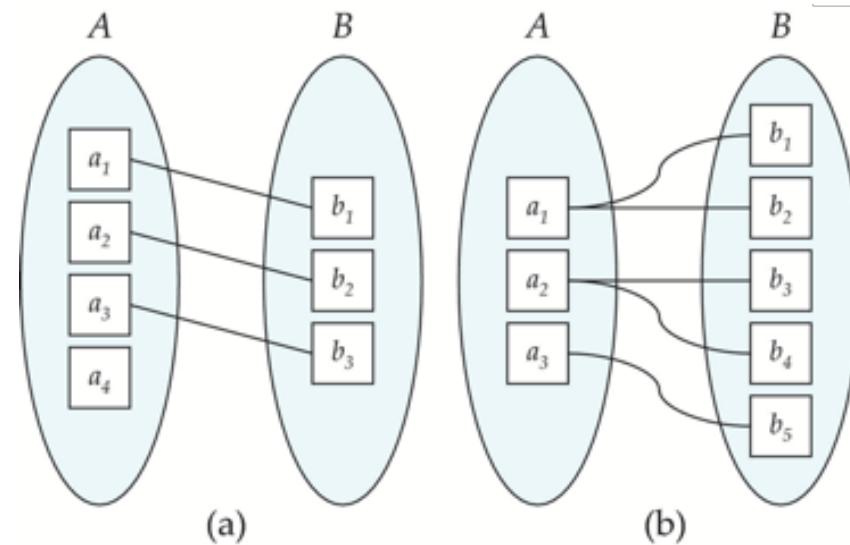
# Refresh on ACID Properties

- ACID, which stands for Atomicity, Consistency, Isolation, and Durability[^11]
- **Atomicity** refers to something that cannot be broken down into smaller parts.
  - It is not about concurrency (which comes with the I)
- **Consistency** (overused term), that here relates to the data *invariants* (integrity would be a better term IMHO)
- **Isolation** means that concurrently executing transactions are isolated from each other.
  - Typically associated with serializability, but there weaker options.

# Refresh on ACID Properties

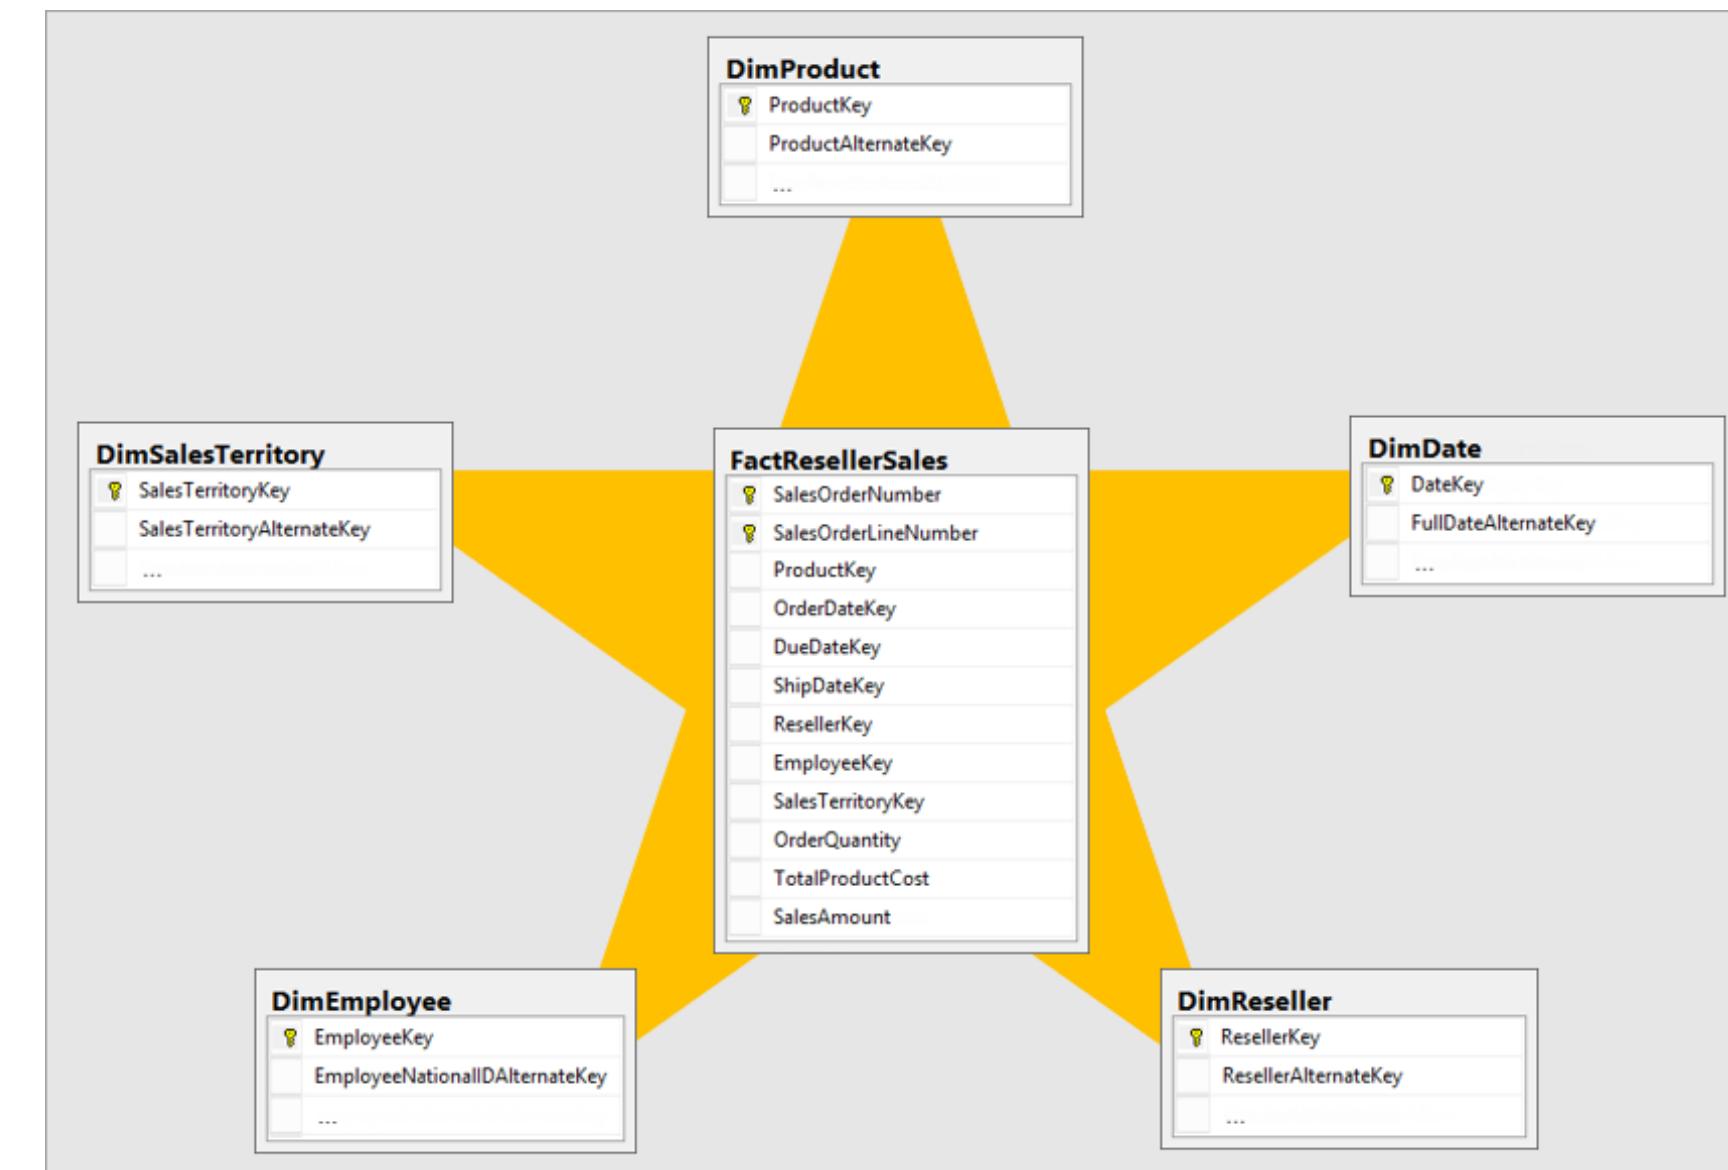
- ACID, which stands for Atomicity, Consistency, Isolation, and Durability[^11]
- **Atomicity** refers to something that cannot be broken down into smaller parts.
  - It is not about concurrency (which comes with the I)
- **Consistency** (overused term), that here relates to the data *invariants* (integrity would be a better term IMHO)
- **Isolation** means that concurrently executing transactions are isolated from each other.
  - Typically associated with serializability, but there weaker options.
- **Durability** means (fault-tolerant) persistency of the data, once the transaction is completed.

# Cardinality Visualized



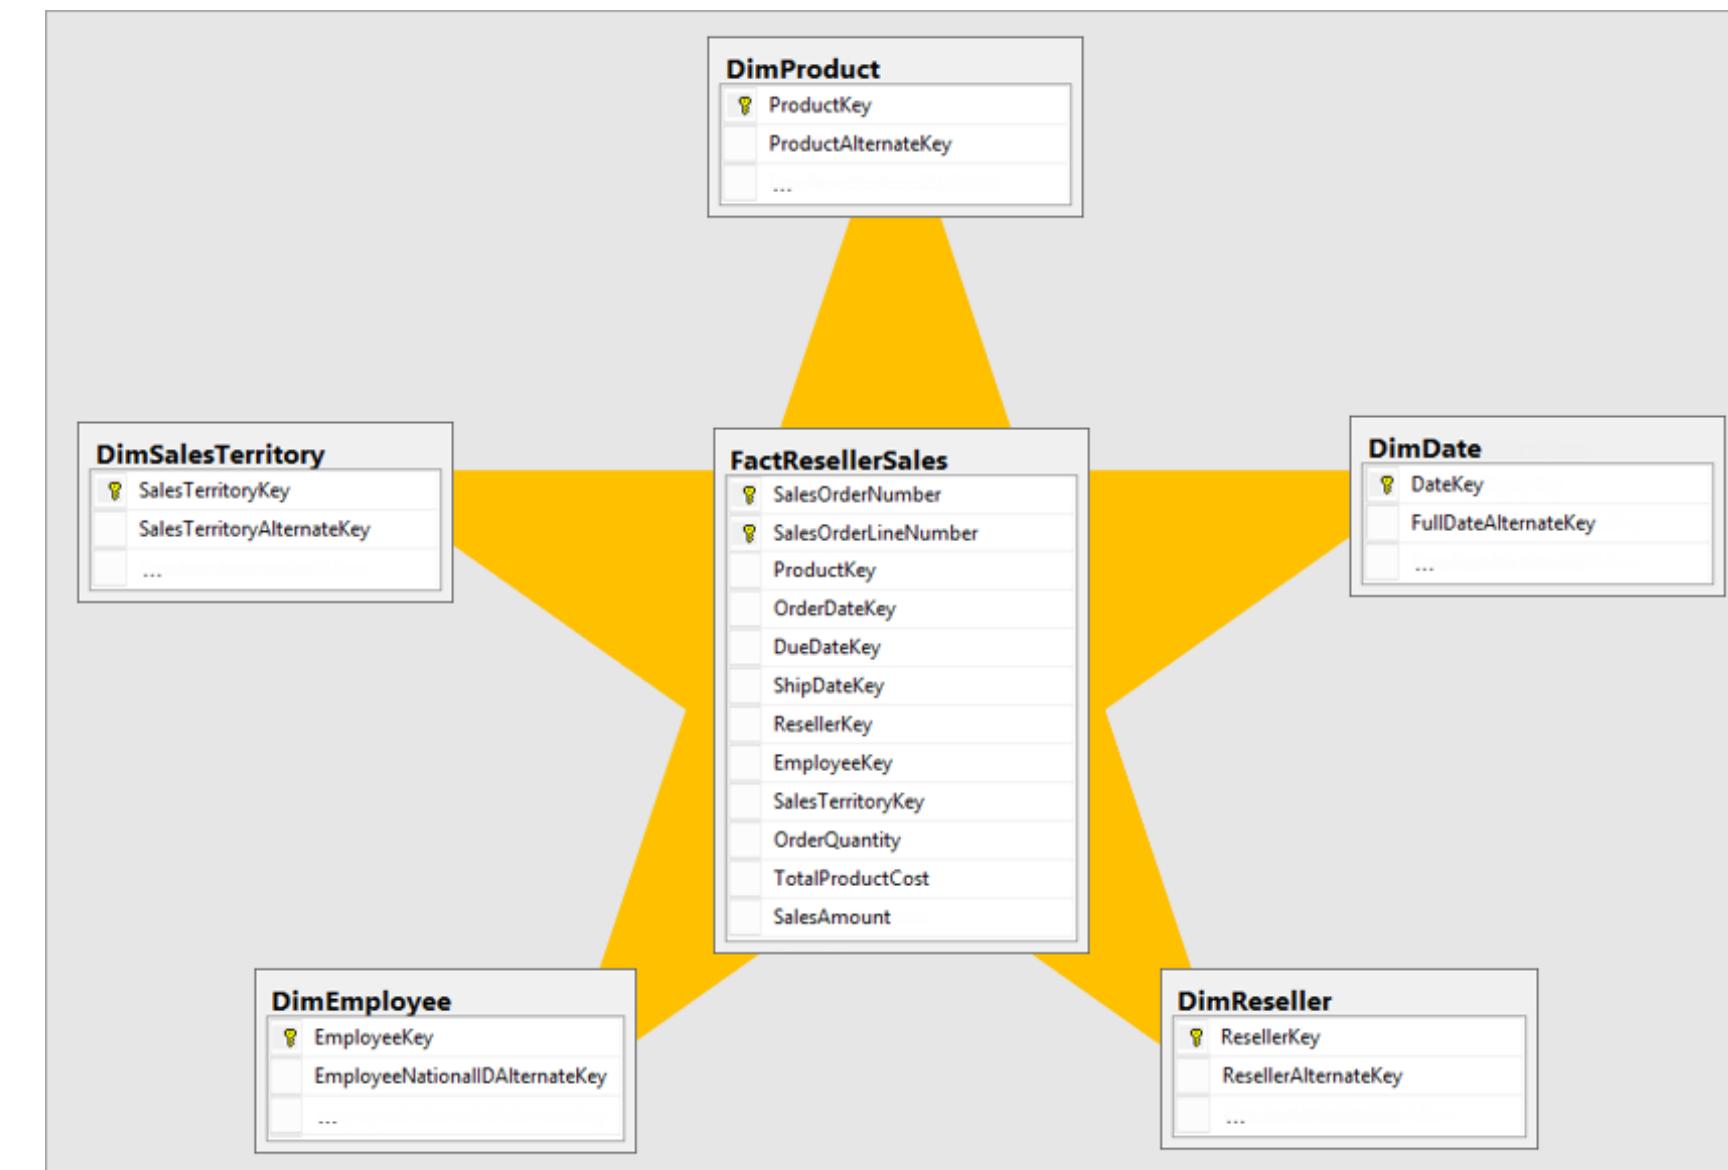
# Data Modelling for Data Warehouse (Briefly)

# Star Schema



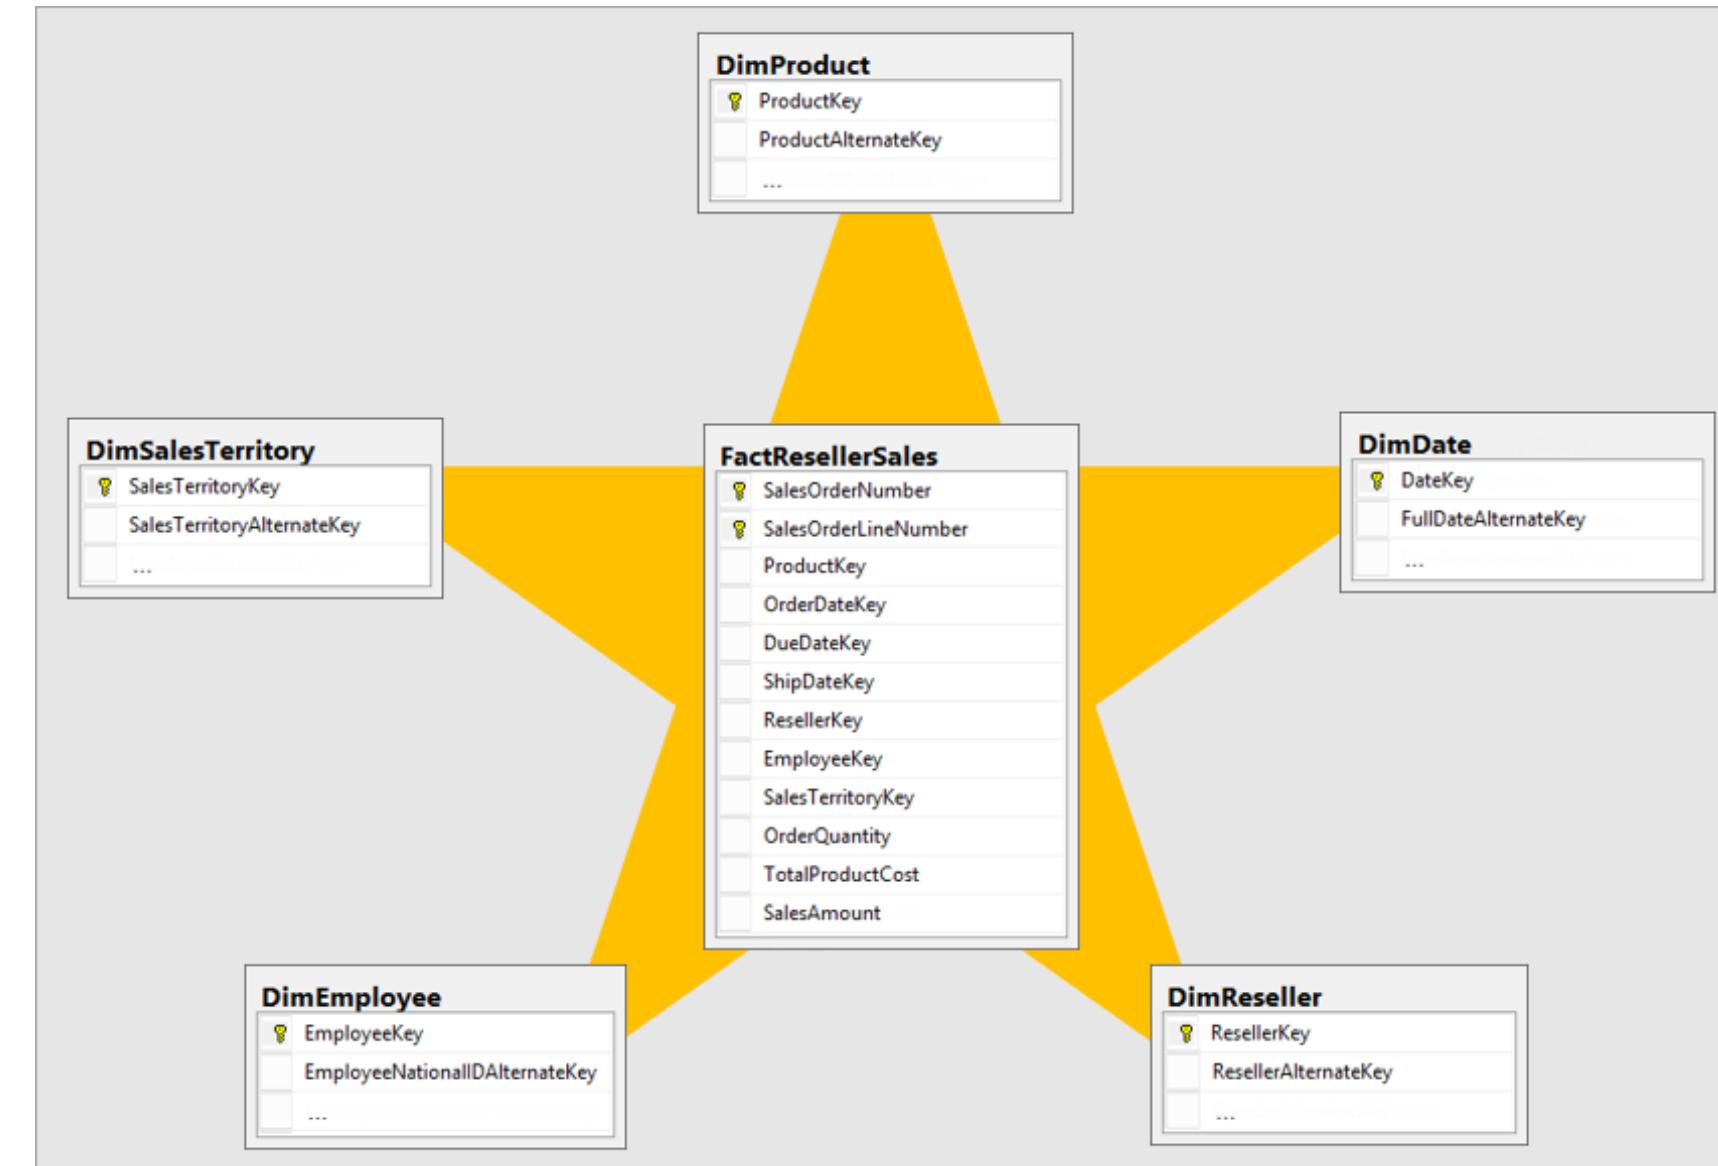
# Star Schema

- A **fact table** contains the numeric measures produced by an operational measurement event in the real world.



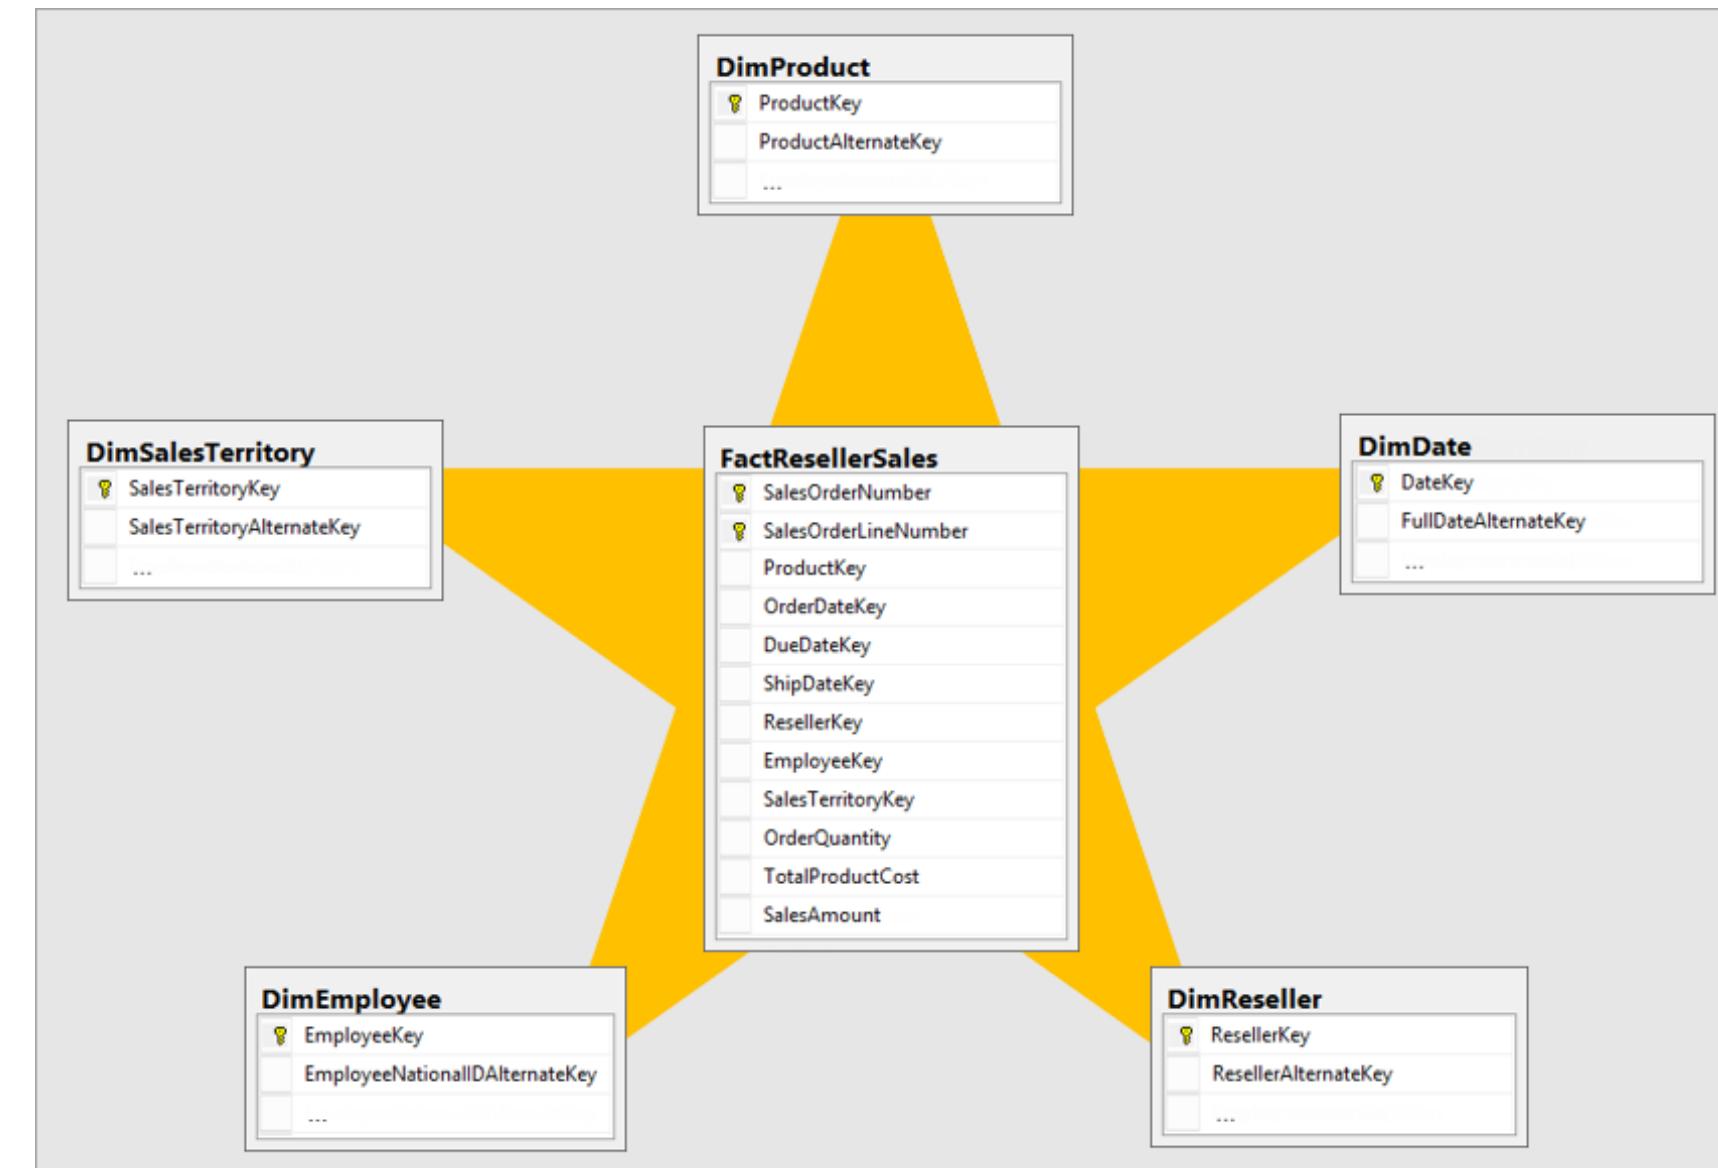
# Star Schema

- A **fact table** contains the numeric measures produced by an operational measurement event in the real world.
- A **single fact** table row has a one-to-one relationship to a measurement event as described by the fact table's grain.

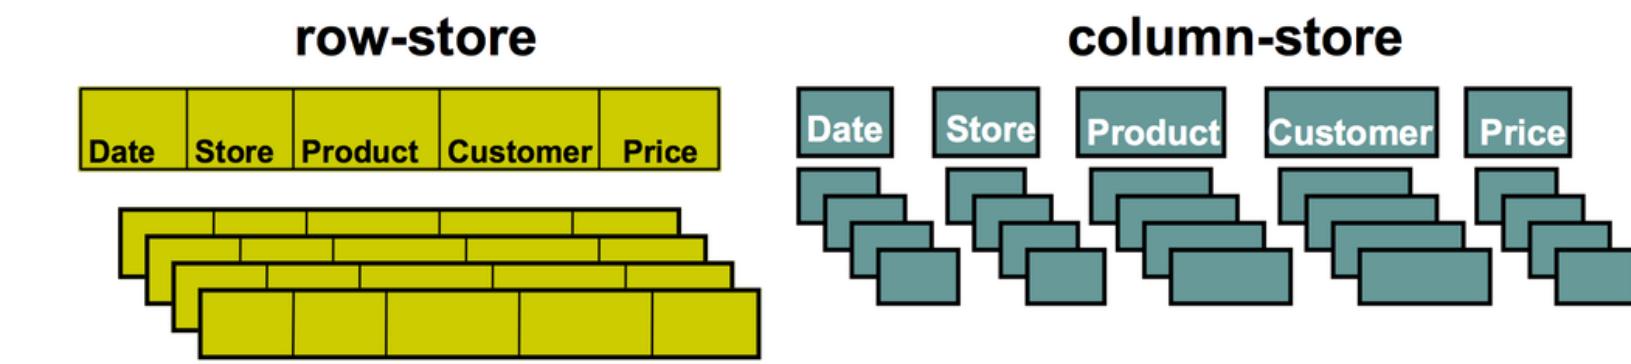


# Star Schema

- A **fact table** contains the numeric measures produced by an operational measurement event in the real world.
- A **single fact** table row has a one-to-one relationship to a measurement event as described by the fact table's grain.
- **Dimensions** provide context to business process events, e.g., who, what, where, when, why, and how.

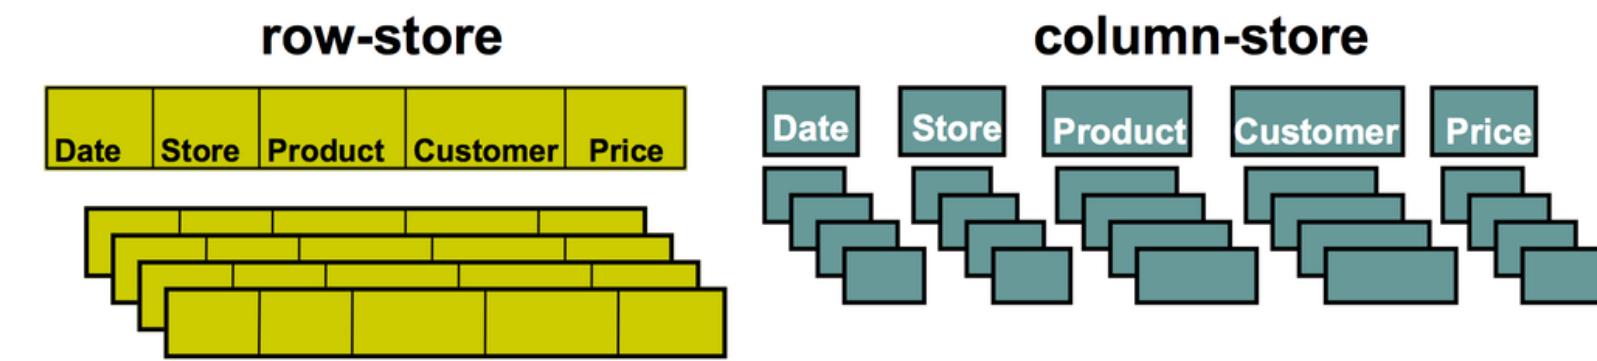


# A note on Storage



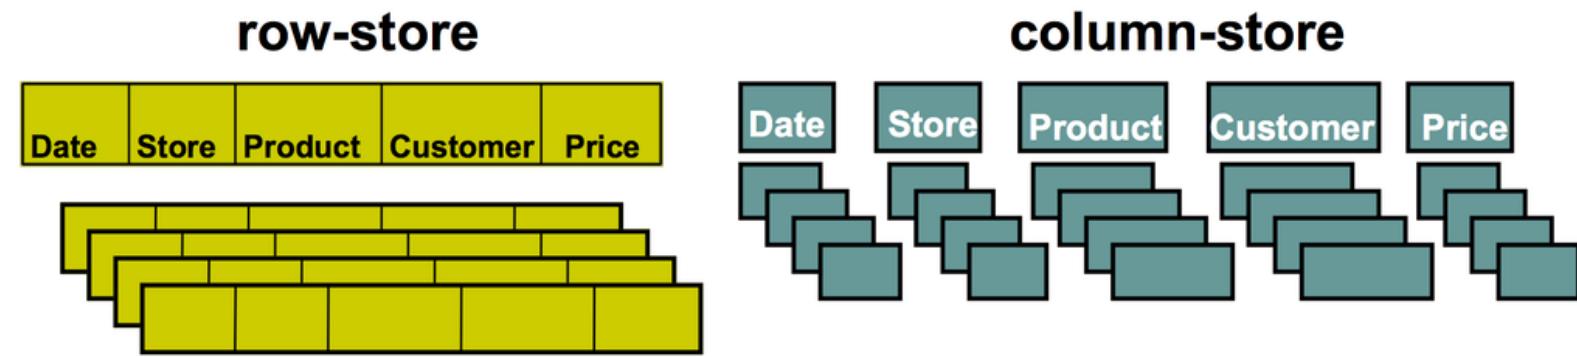
## A note on Storage

- Data warehouse typically interact with OLTP database to expose one or more OLAP system.



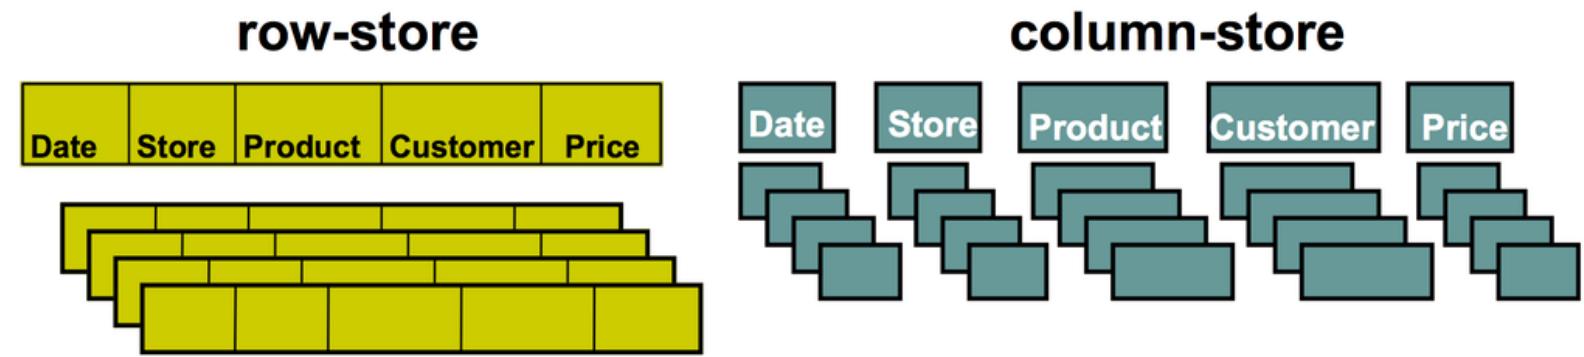
## A note on Storage

- Data warehouse typically interact with OLTP database to expose one or more OLAP system.
- Such OLAP system adopt storage optimized for analytics, i.e., Column Oriented



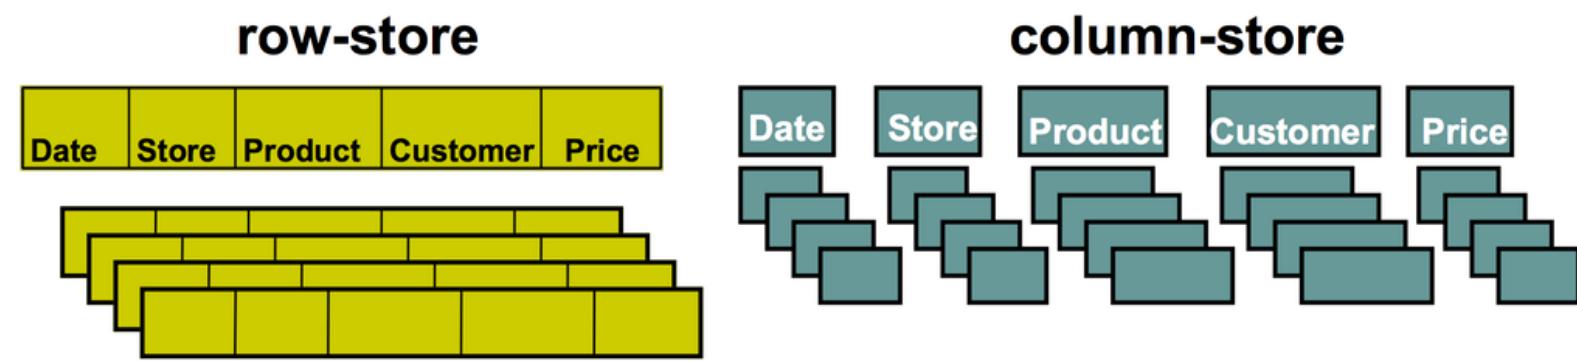
## A note on Storage

- Data warehouse typically interact with OLTP database to expose one or more OLAP system.
- Such OLAP system adopt storage optimized for analytics, i.e., Column Oriented
- The column-oriented storage layout relies on each column file containing the rows in the same order.



## A note on Storage

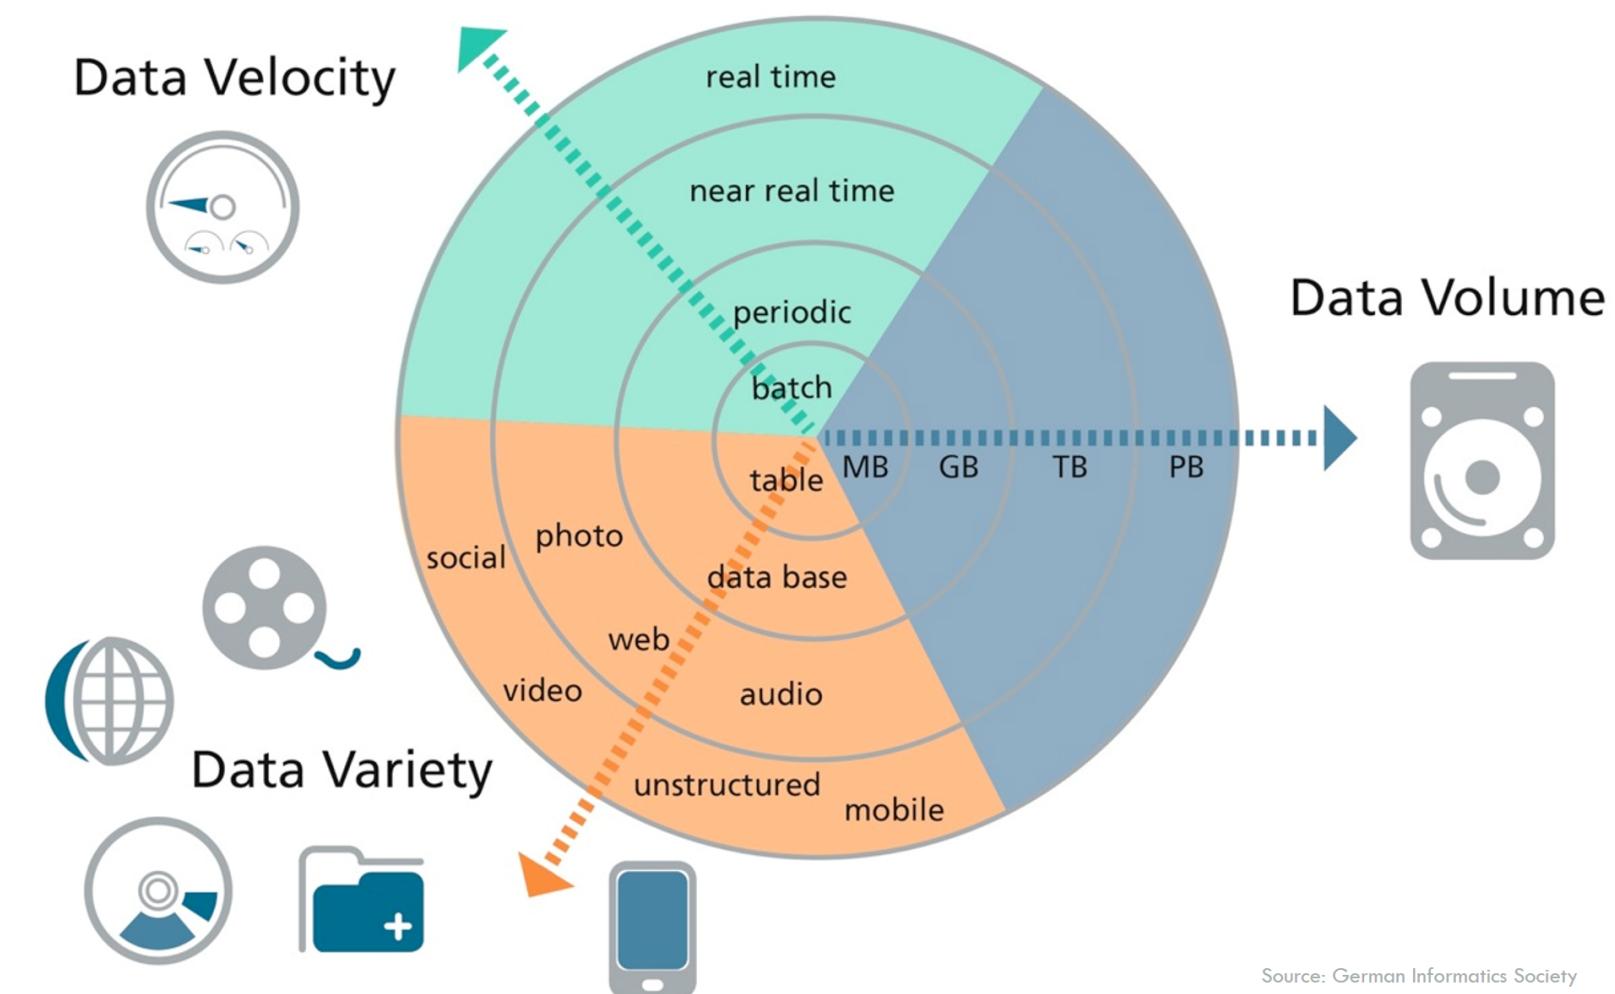
- Data warehouse typically interact with OLTP database to expose one or more OLAP system.
- Such OLAP system adopt storage optimized for analytics, i.e., Column Oriented
- The column-oriented storage layout relies on each column file containing the rows in the same order.
- Not just relational data, e.g., Apache Parquet



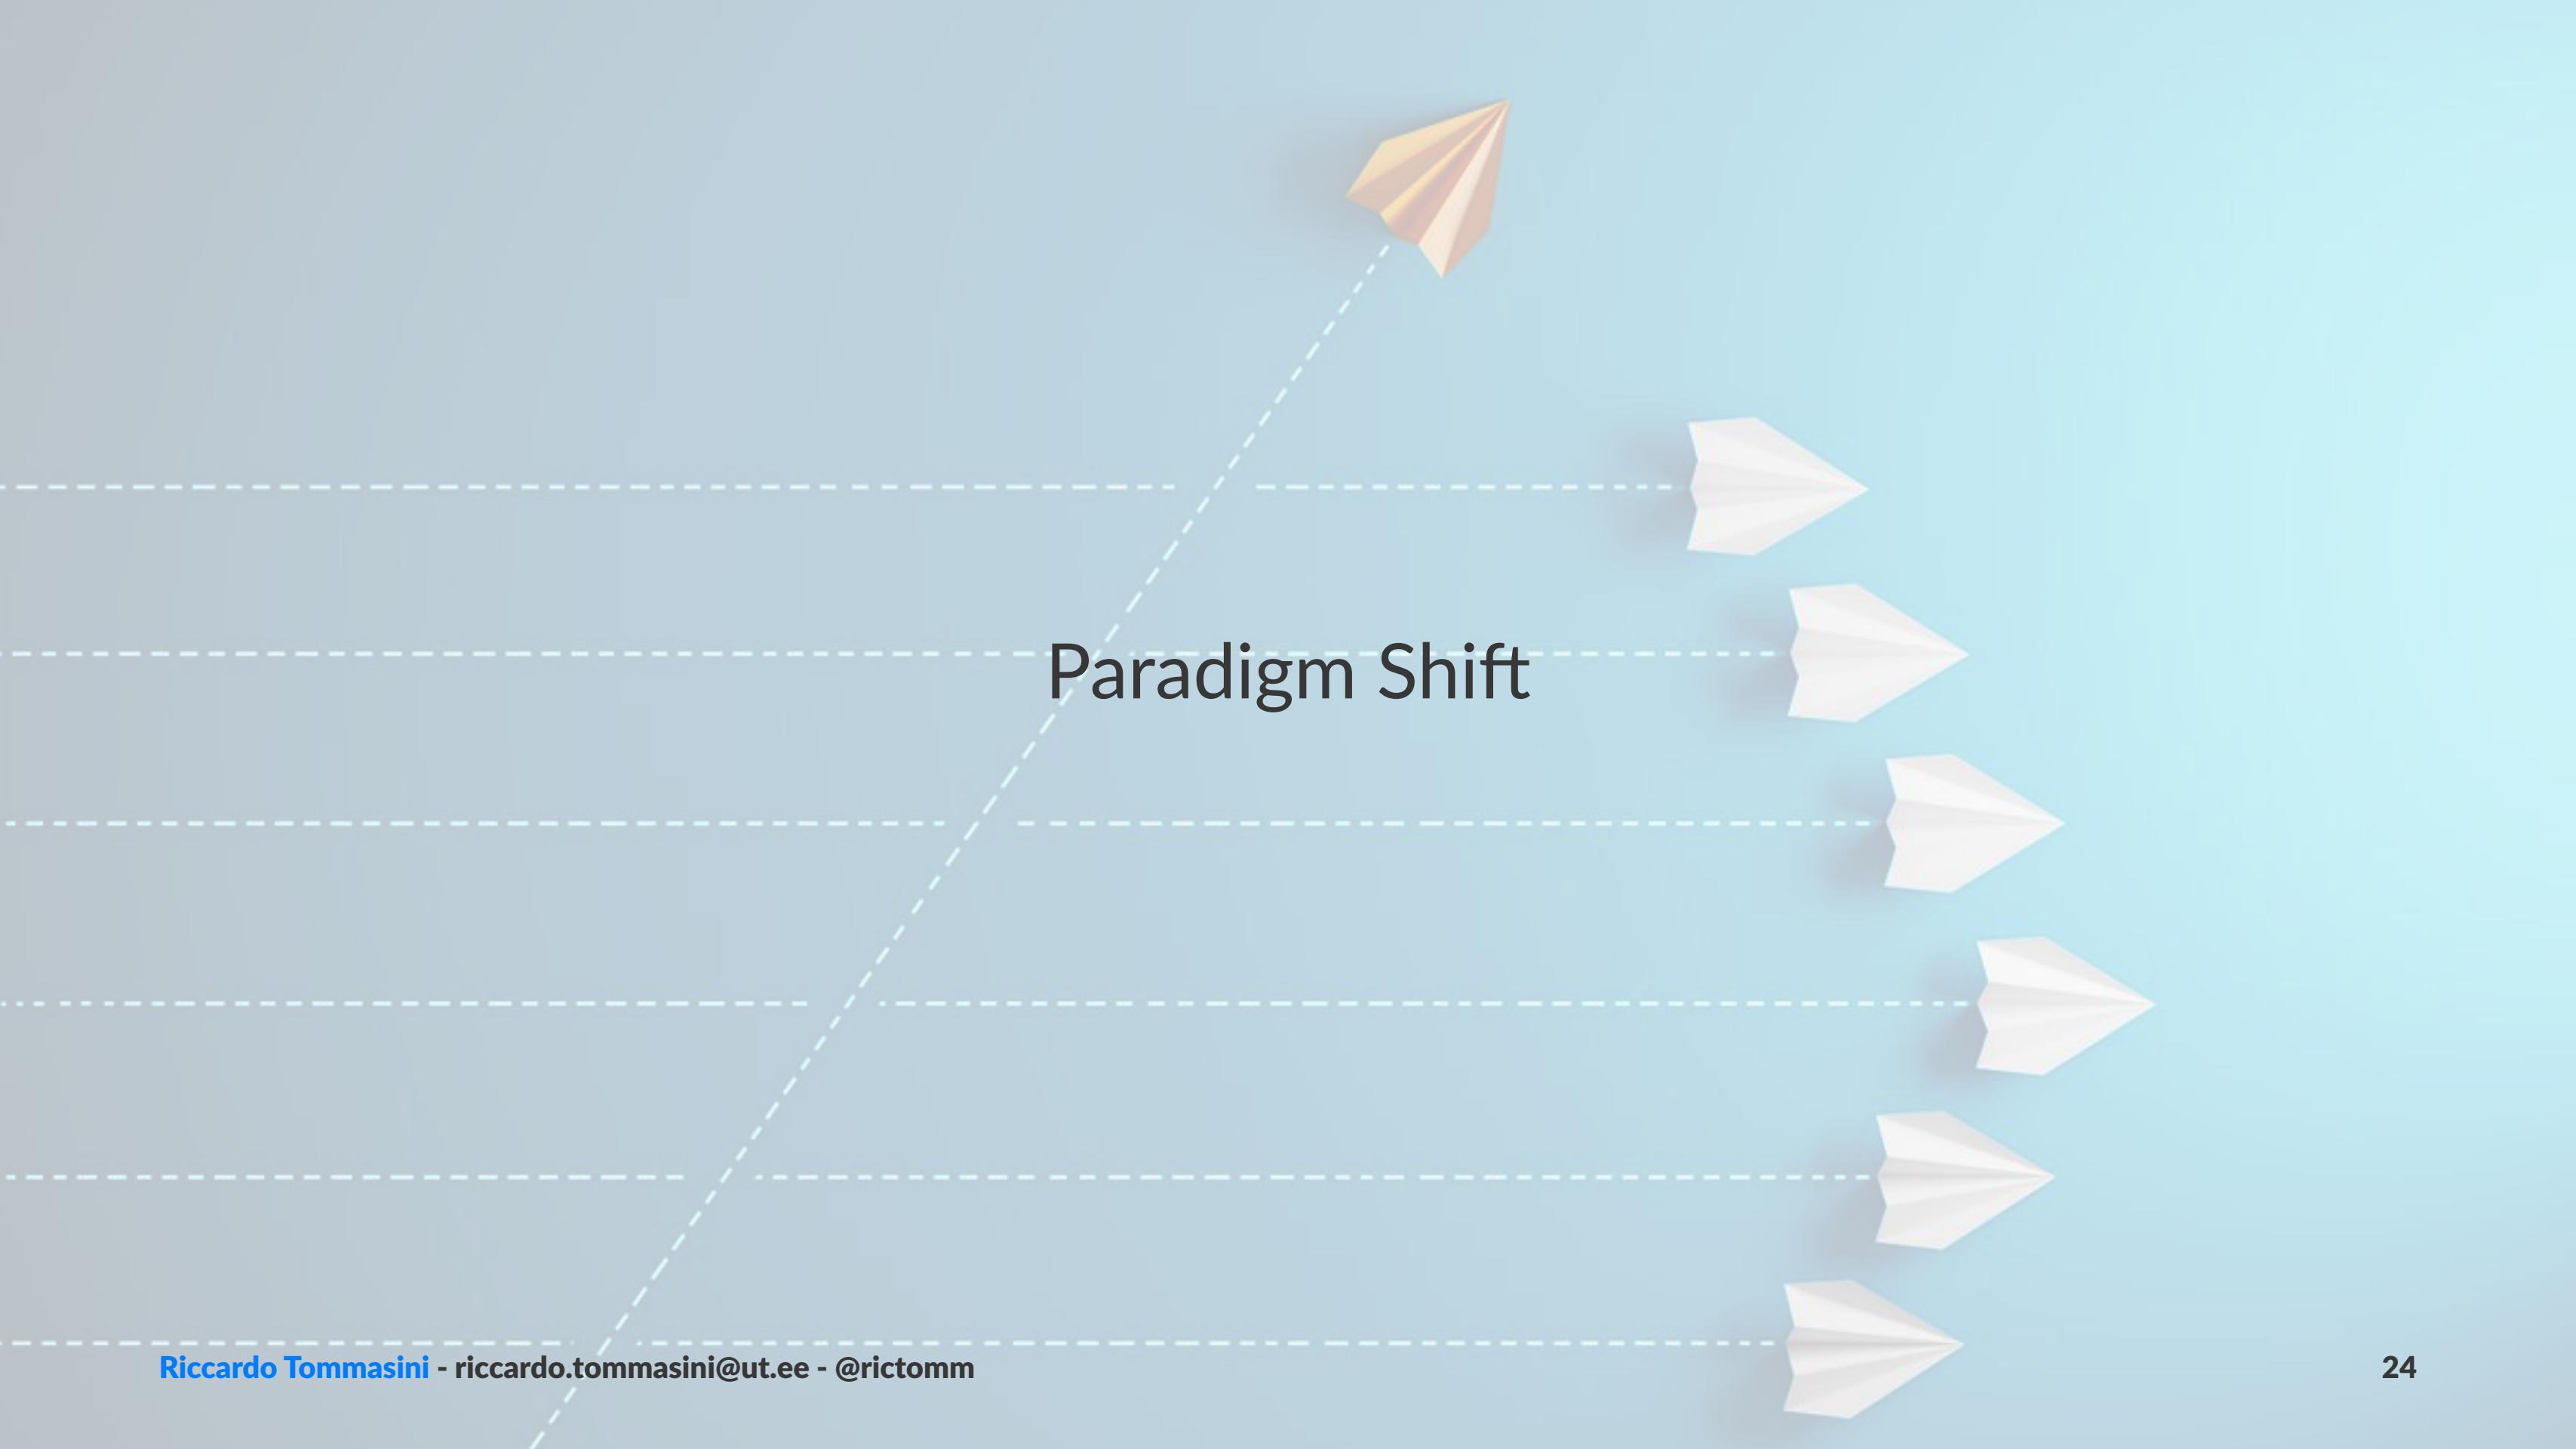
# Data Modelling for Big Data

# Big Data

# Challenges 014



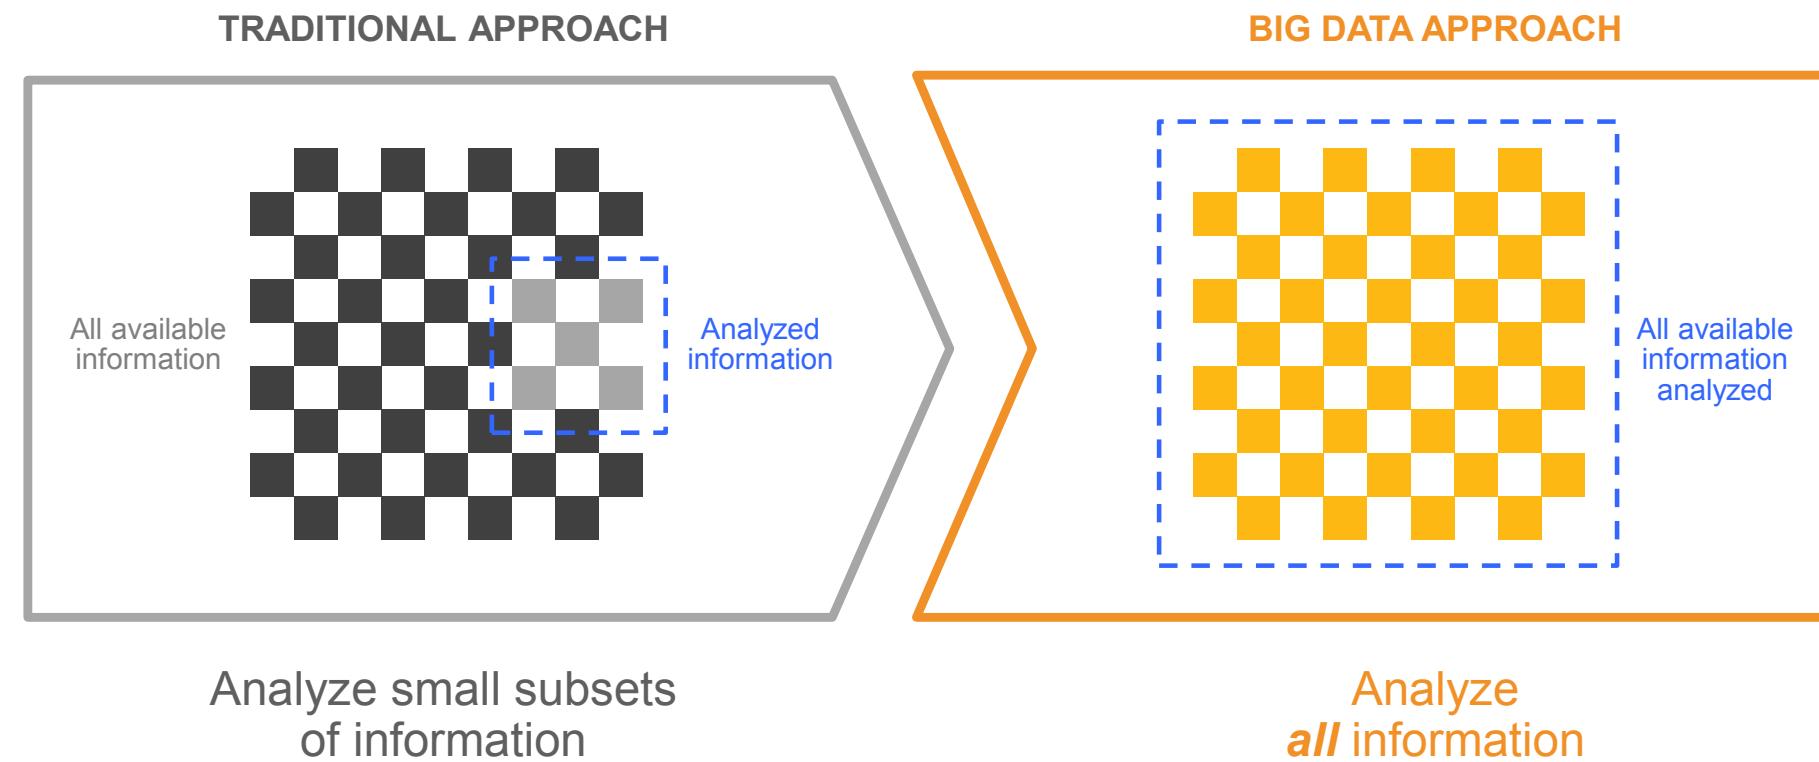
014 Lanely, 2001



Paradigm Shift

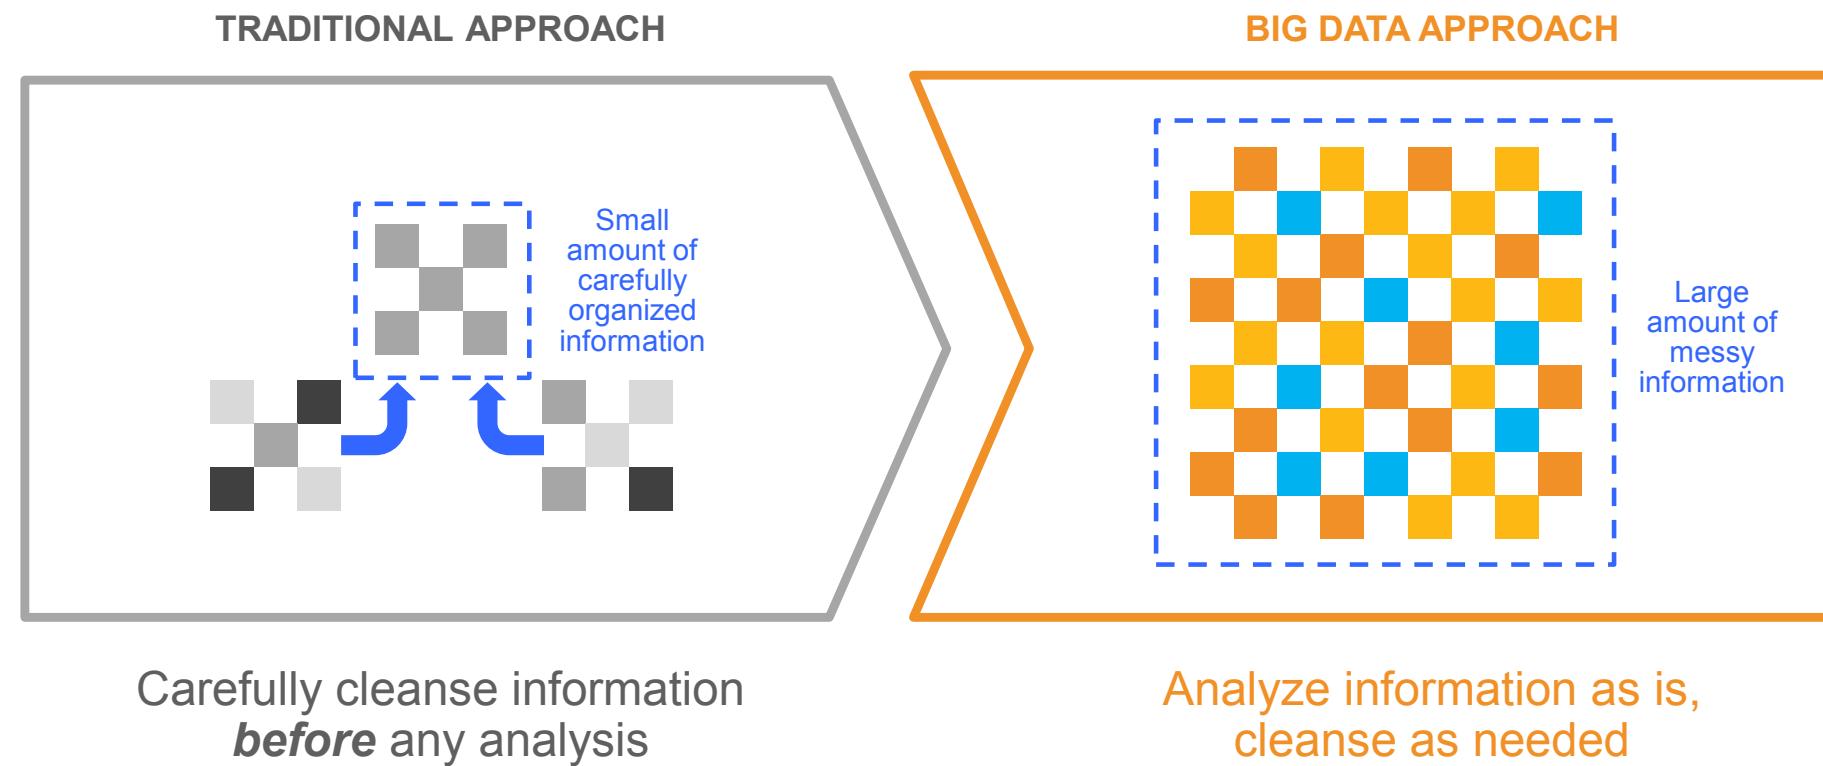
## Paradigm shifts enabled by big data

### Leverage more of the data being captured



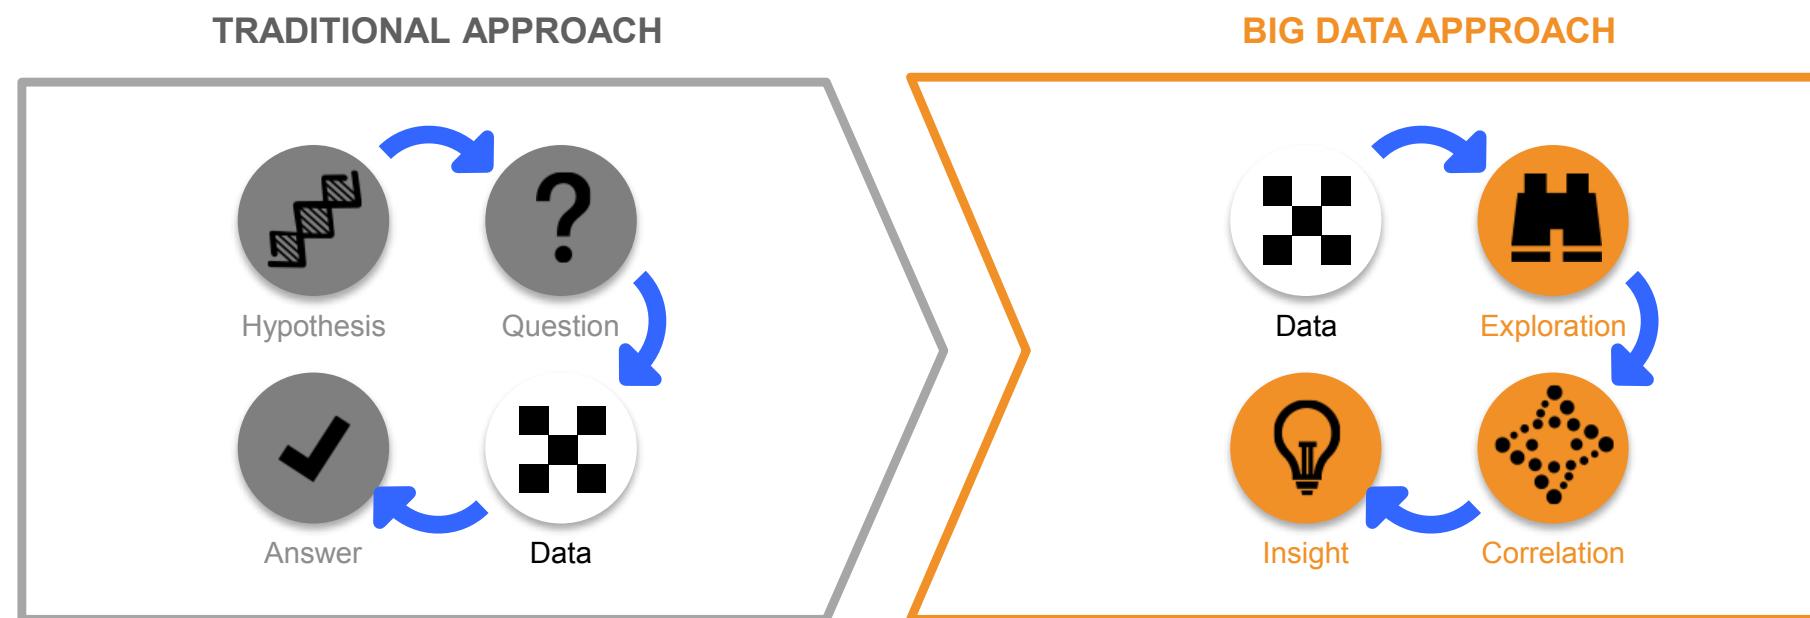
## Paradigm shifts enabled by big data

### Reduce effort required to leverage data



## Paradigm shifts enabled by big data

Data leads the way—and sometimes correlations are good enough



Start with hypothesis and test against selected data

Explore **all** data and identify correlations

## Paradigm shifts enabled by big data

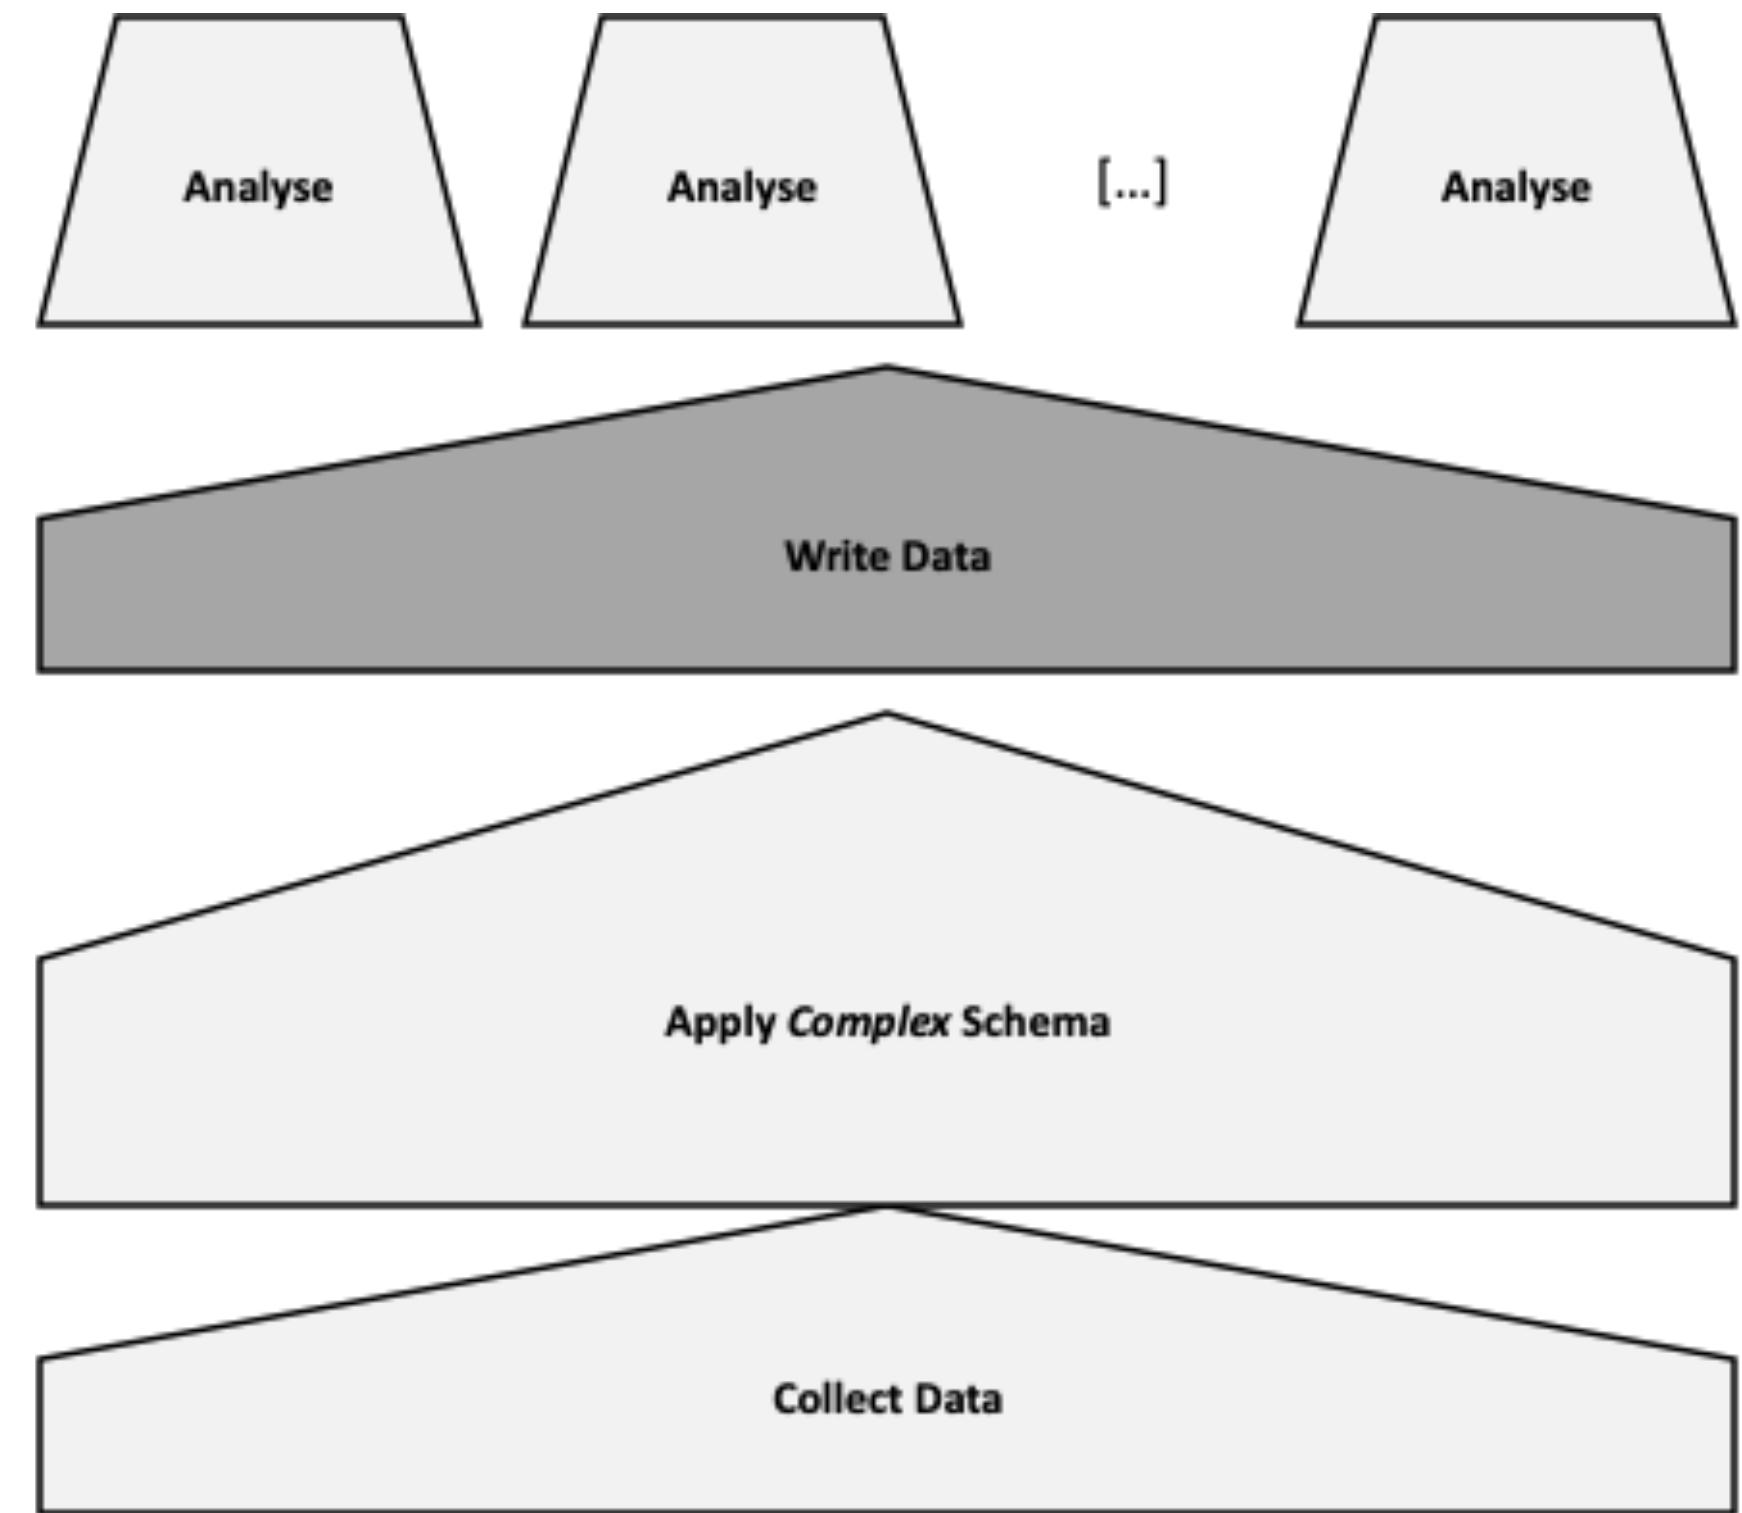
### Leverage data as it is captured



Analyze data *after* it's been processed and landed in a warehouse or mart

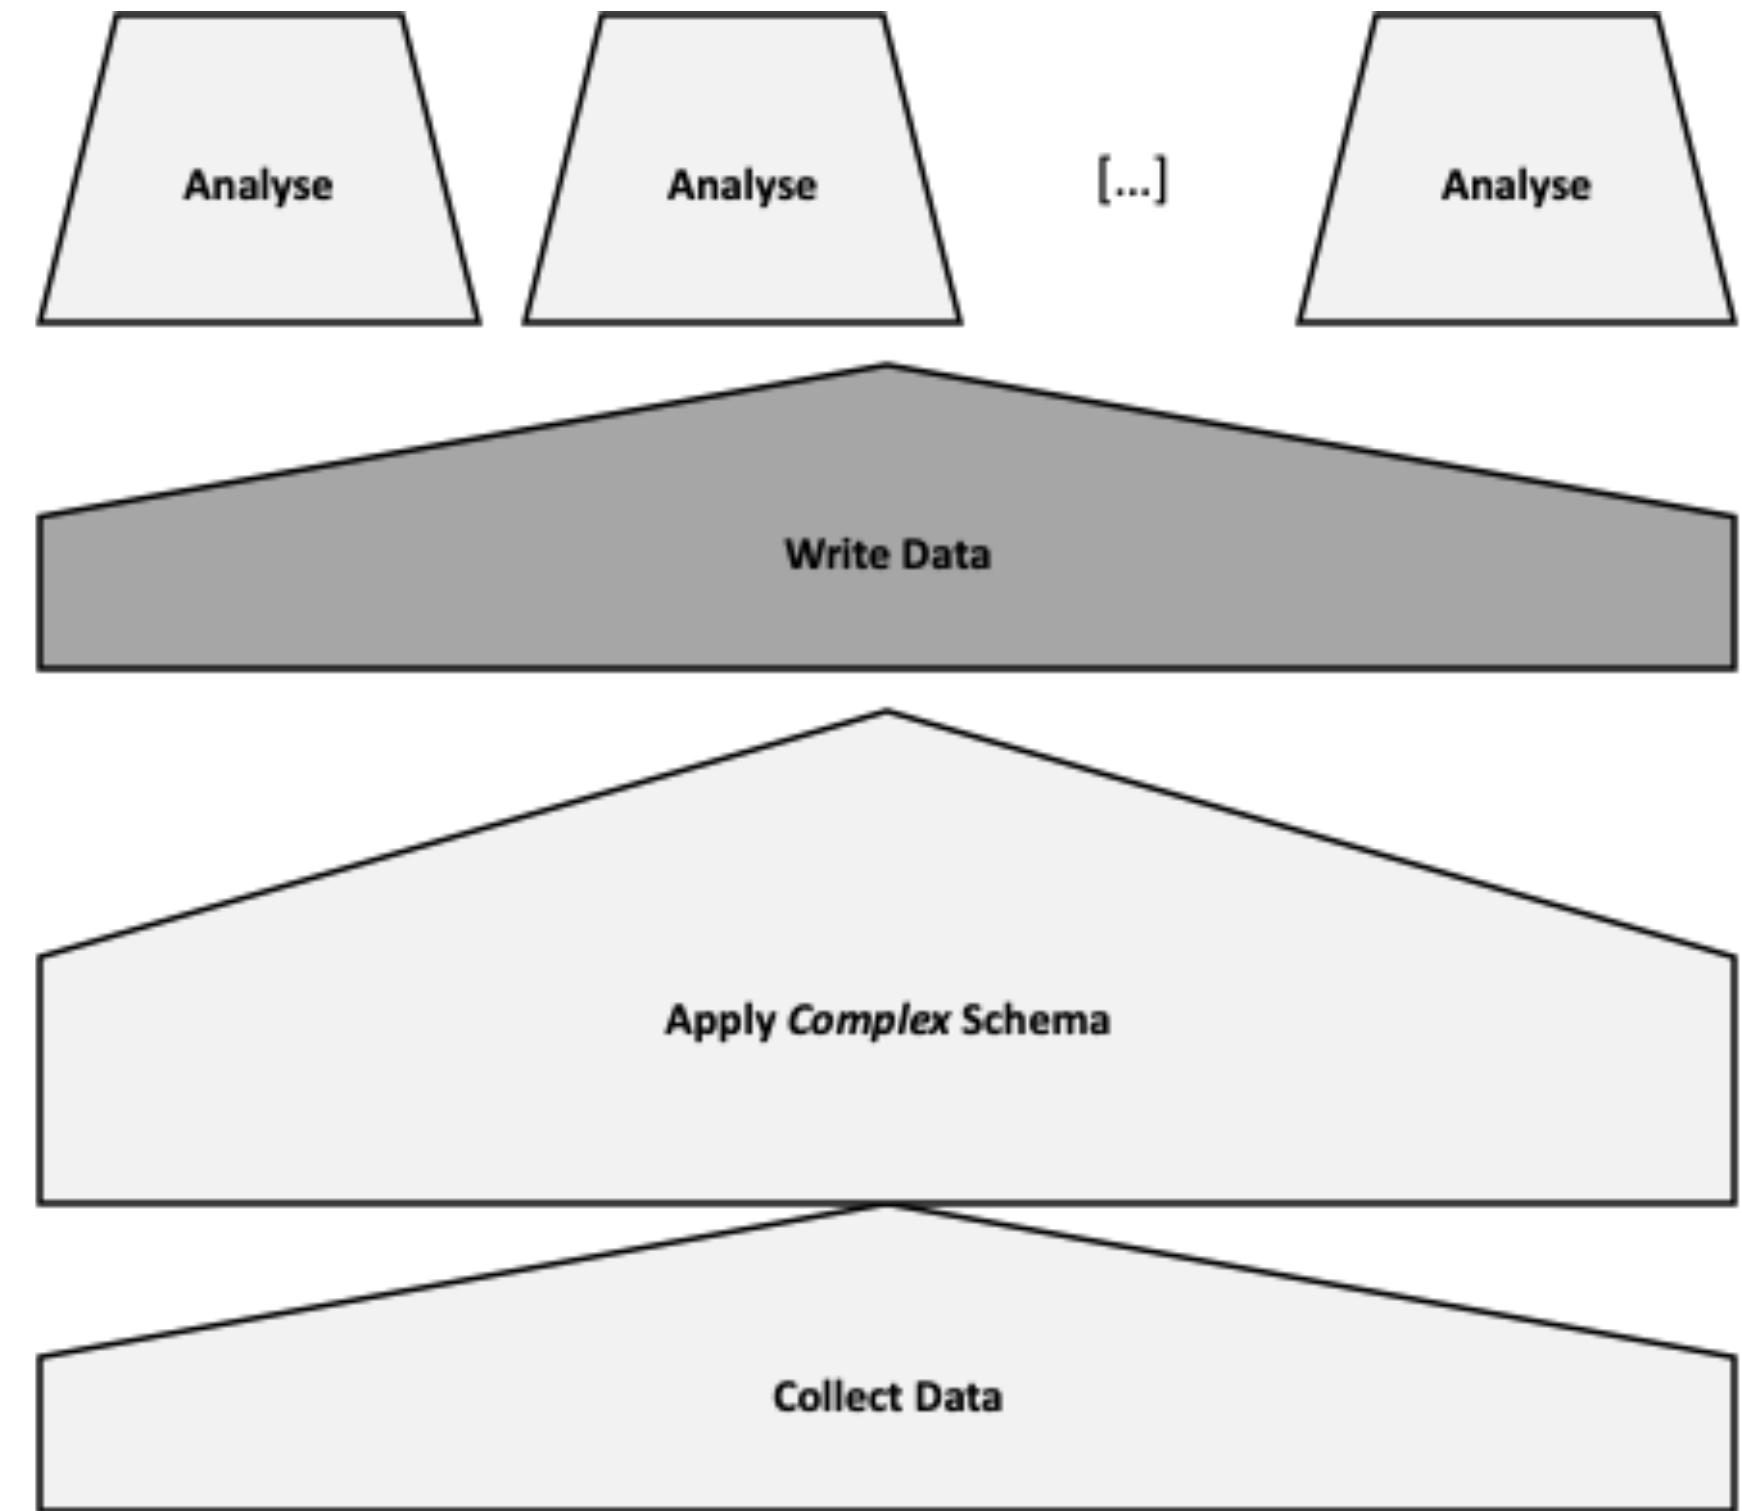
Analyze data *in motion* as it's generated, in real-time

## From Schema on Write



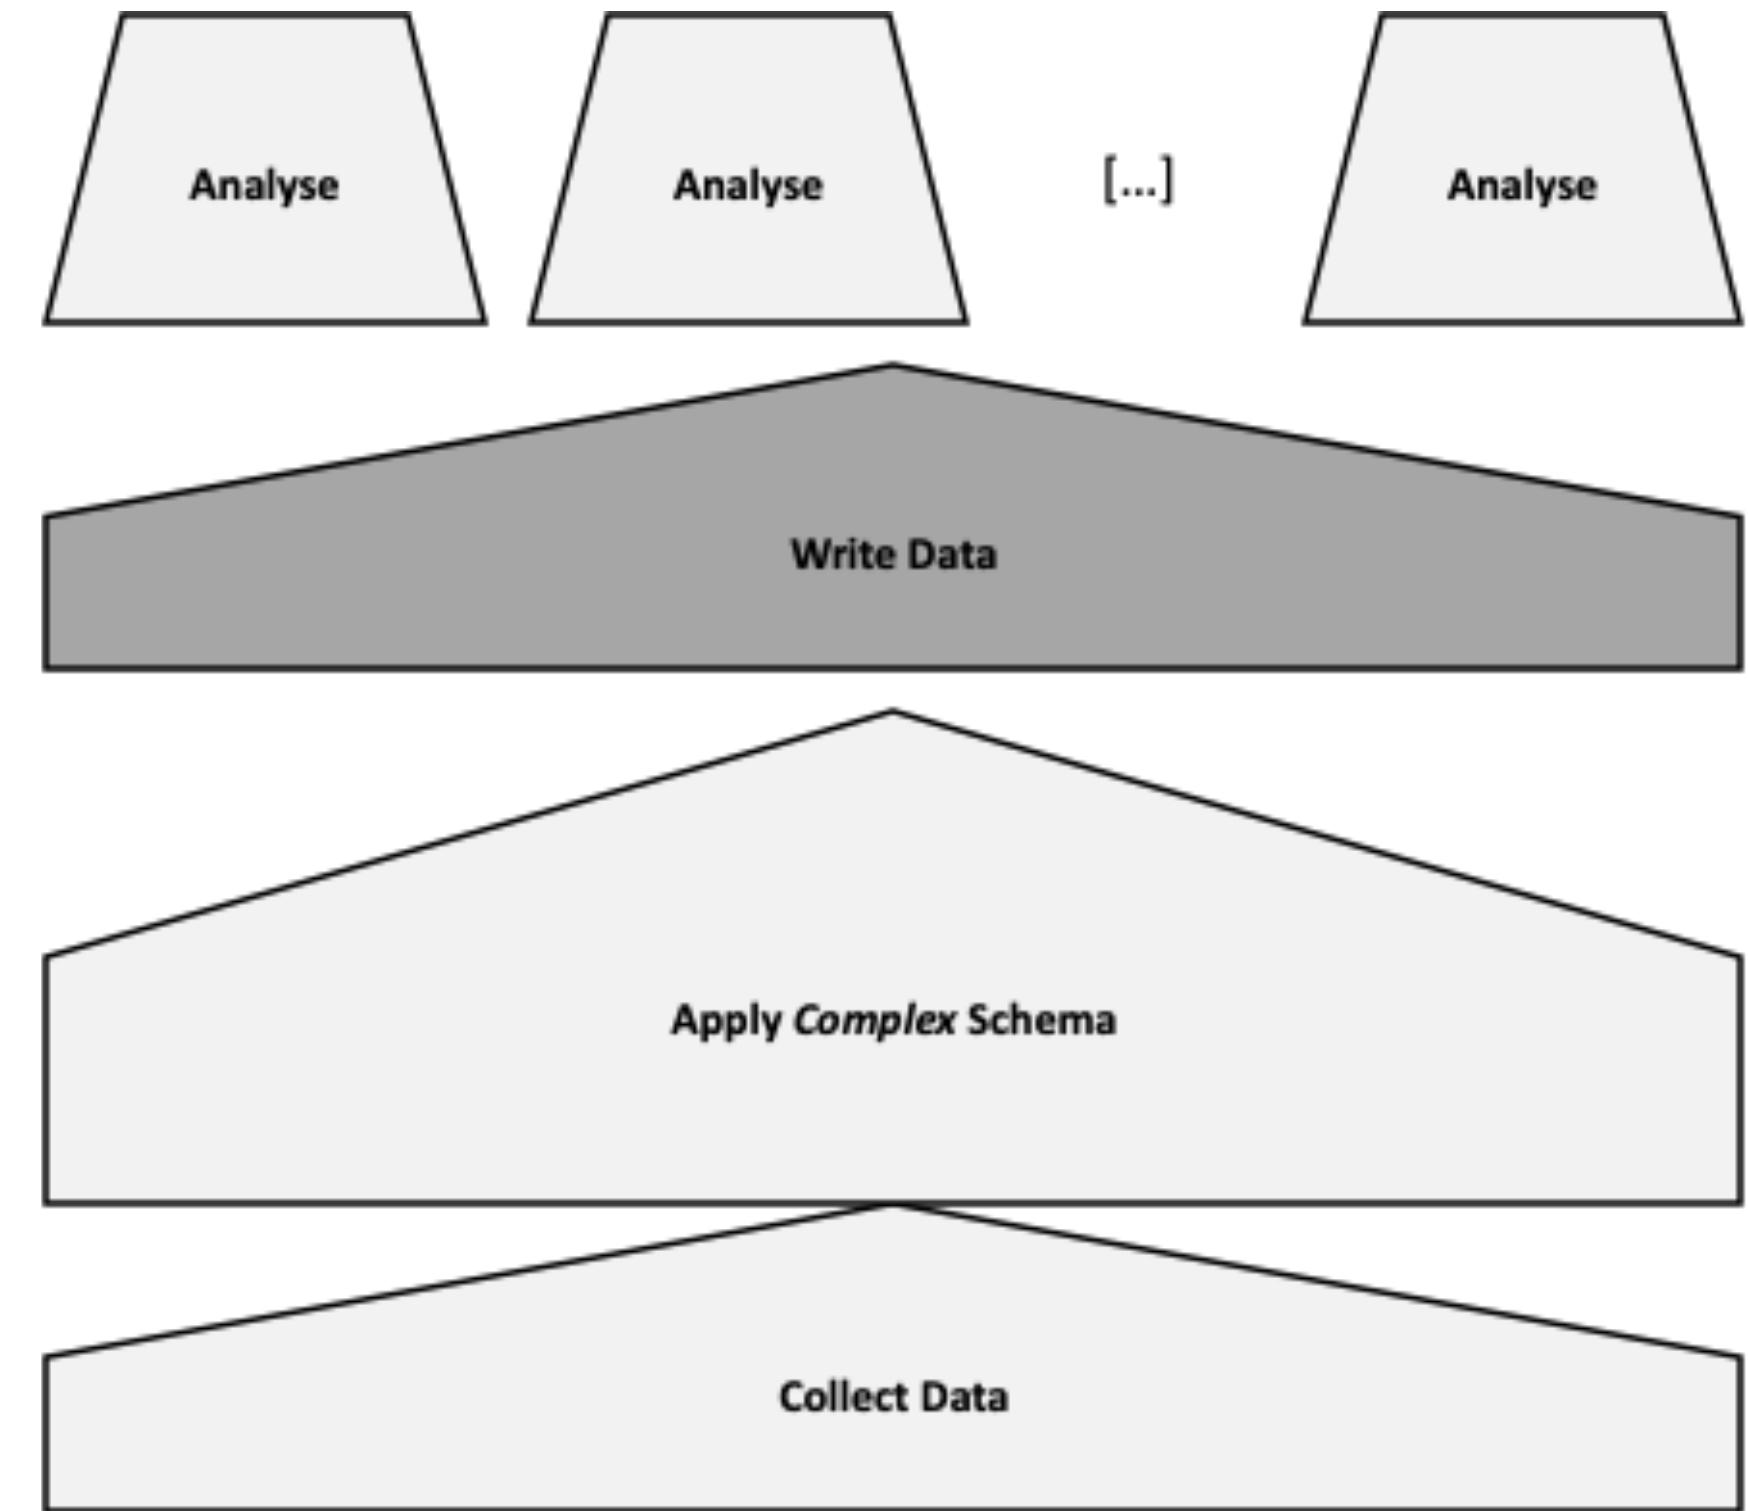
## From Schema on Write

- Focus on the modelling a schema that can accommodate all needs

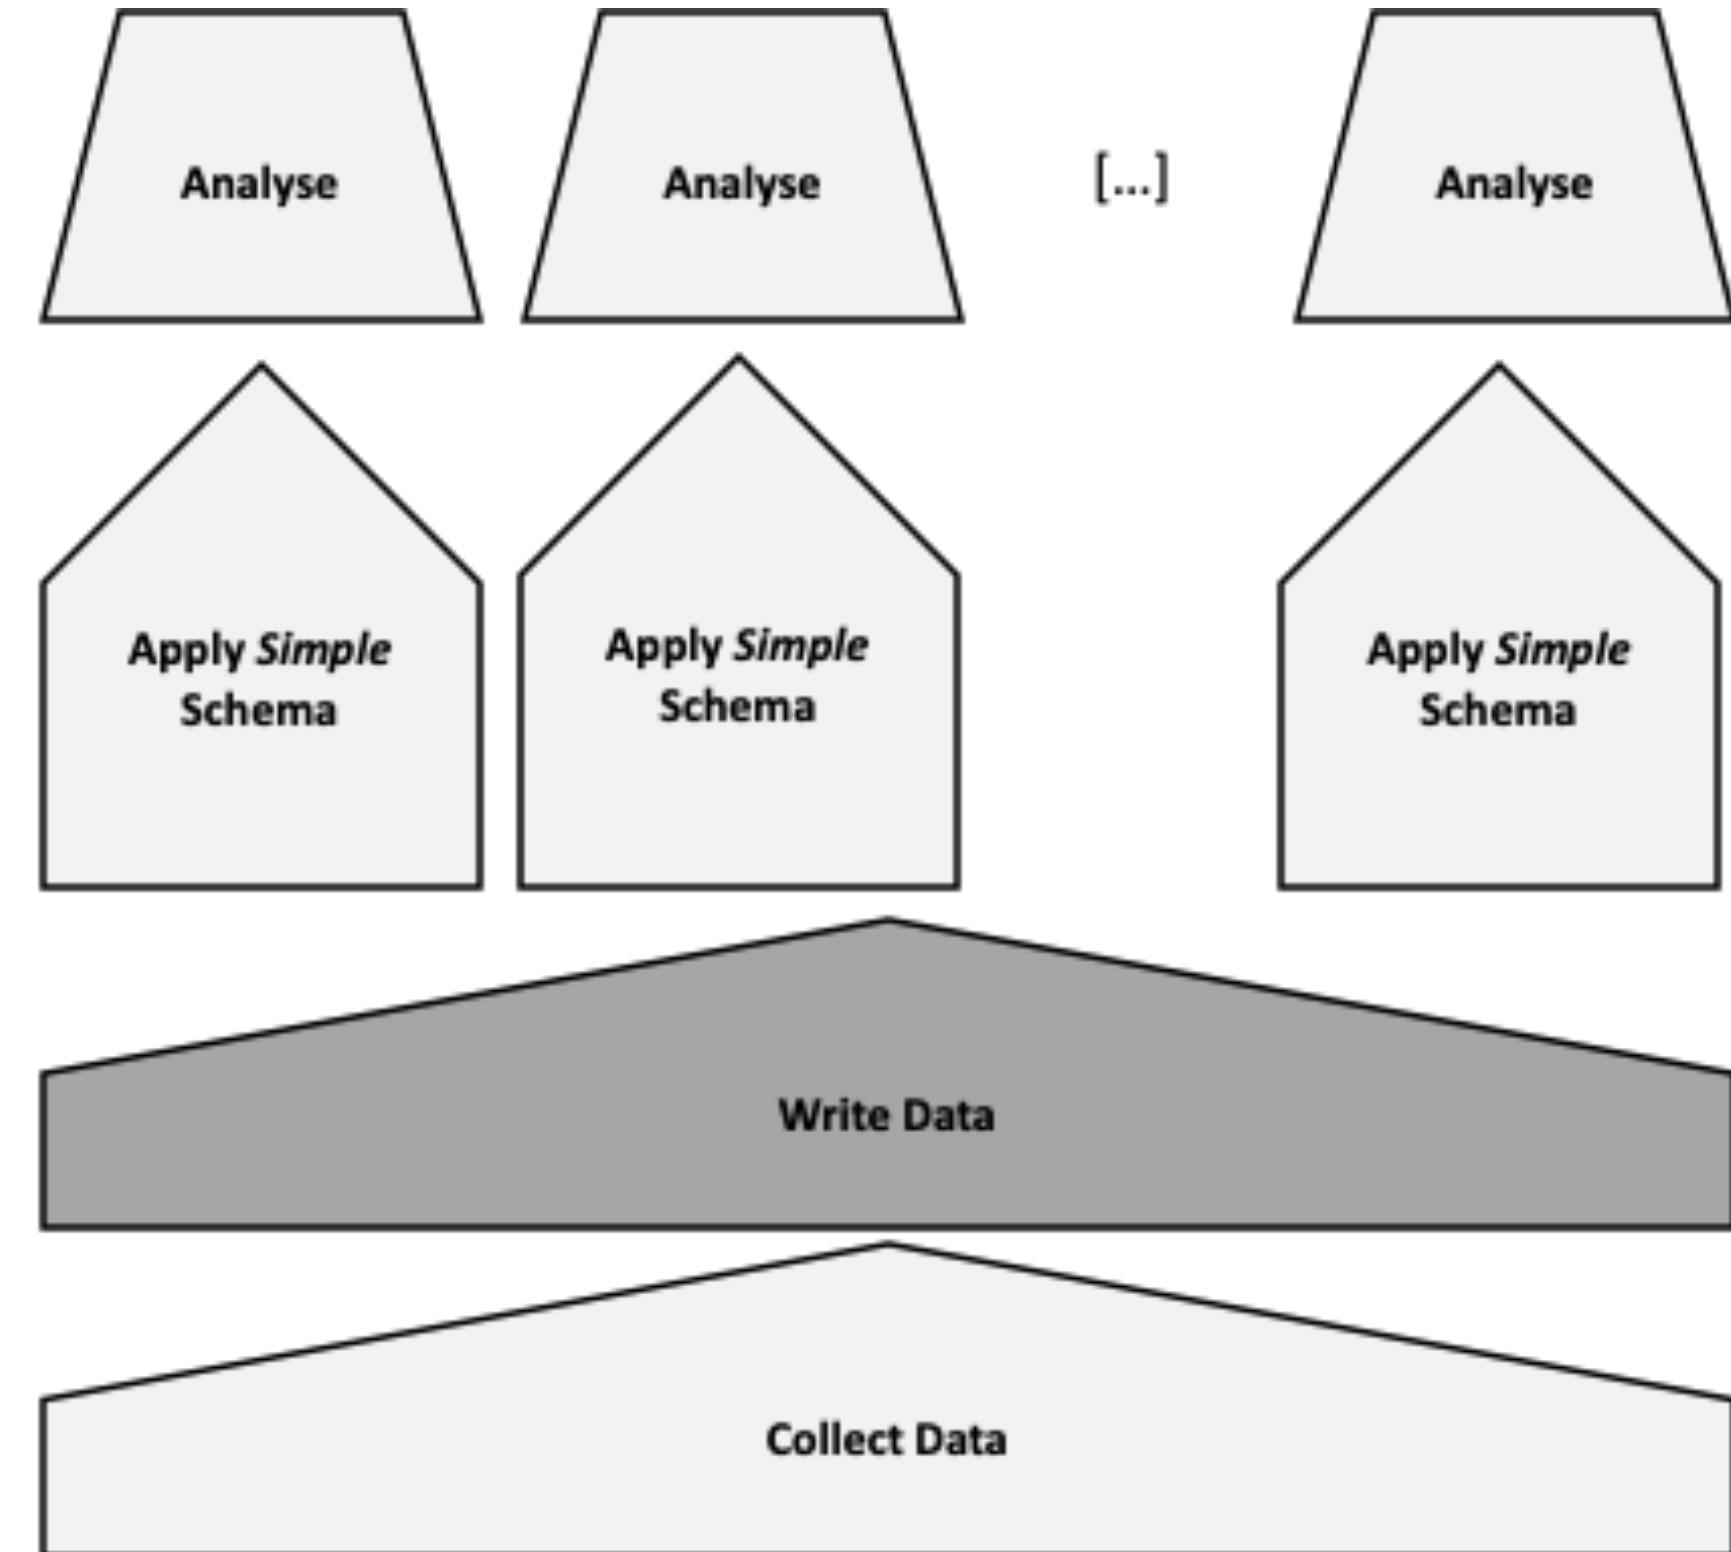


## From Schema on Write

- Focus on the modelling a schema that can accommodate all needs
- Bad impact on those analysis that were not envisioned

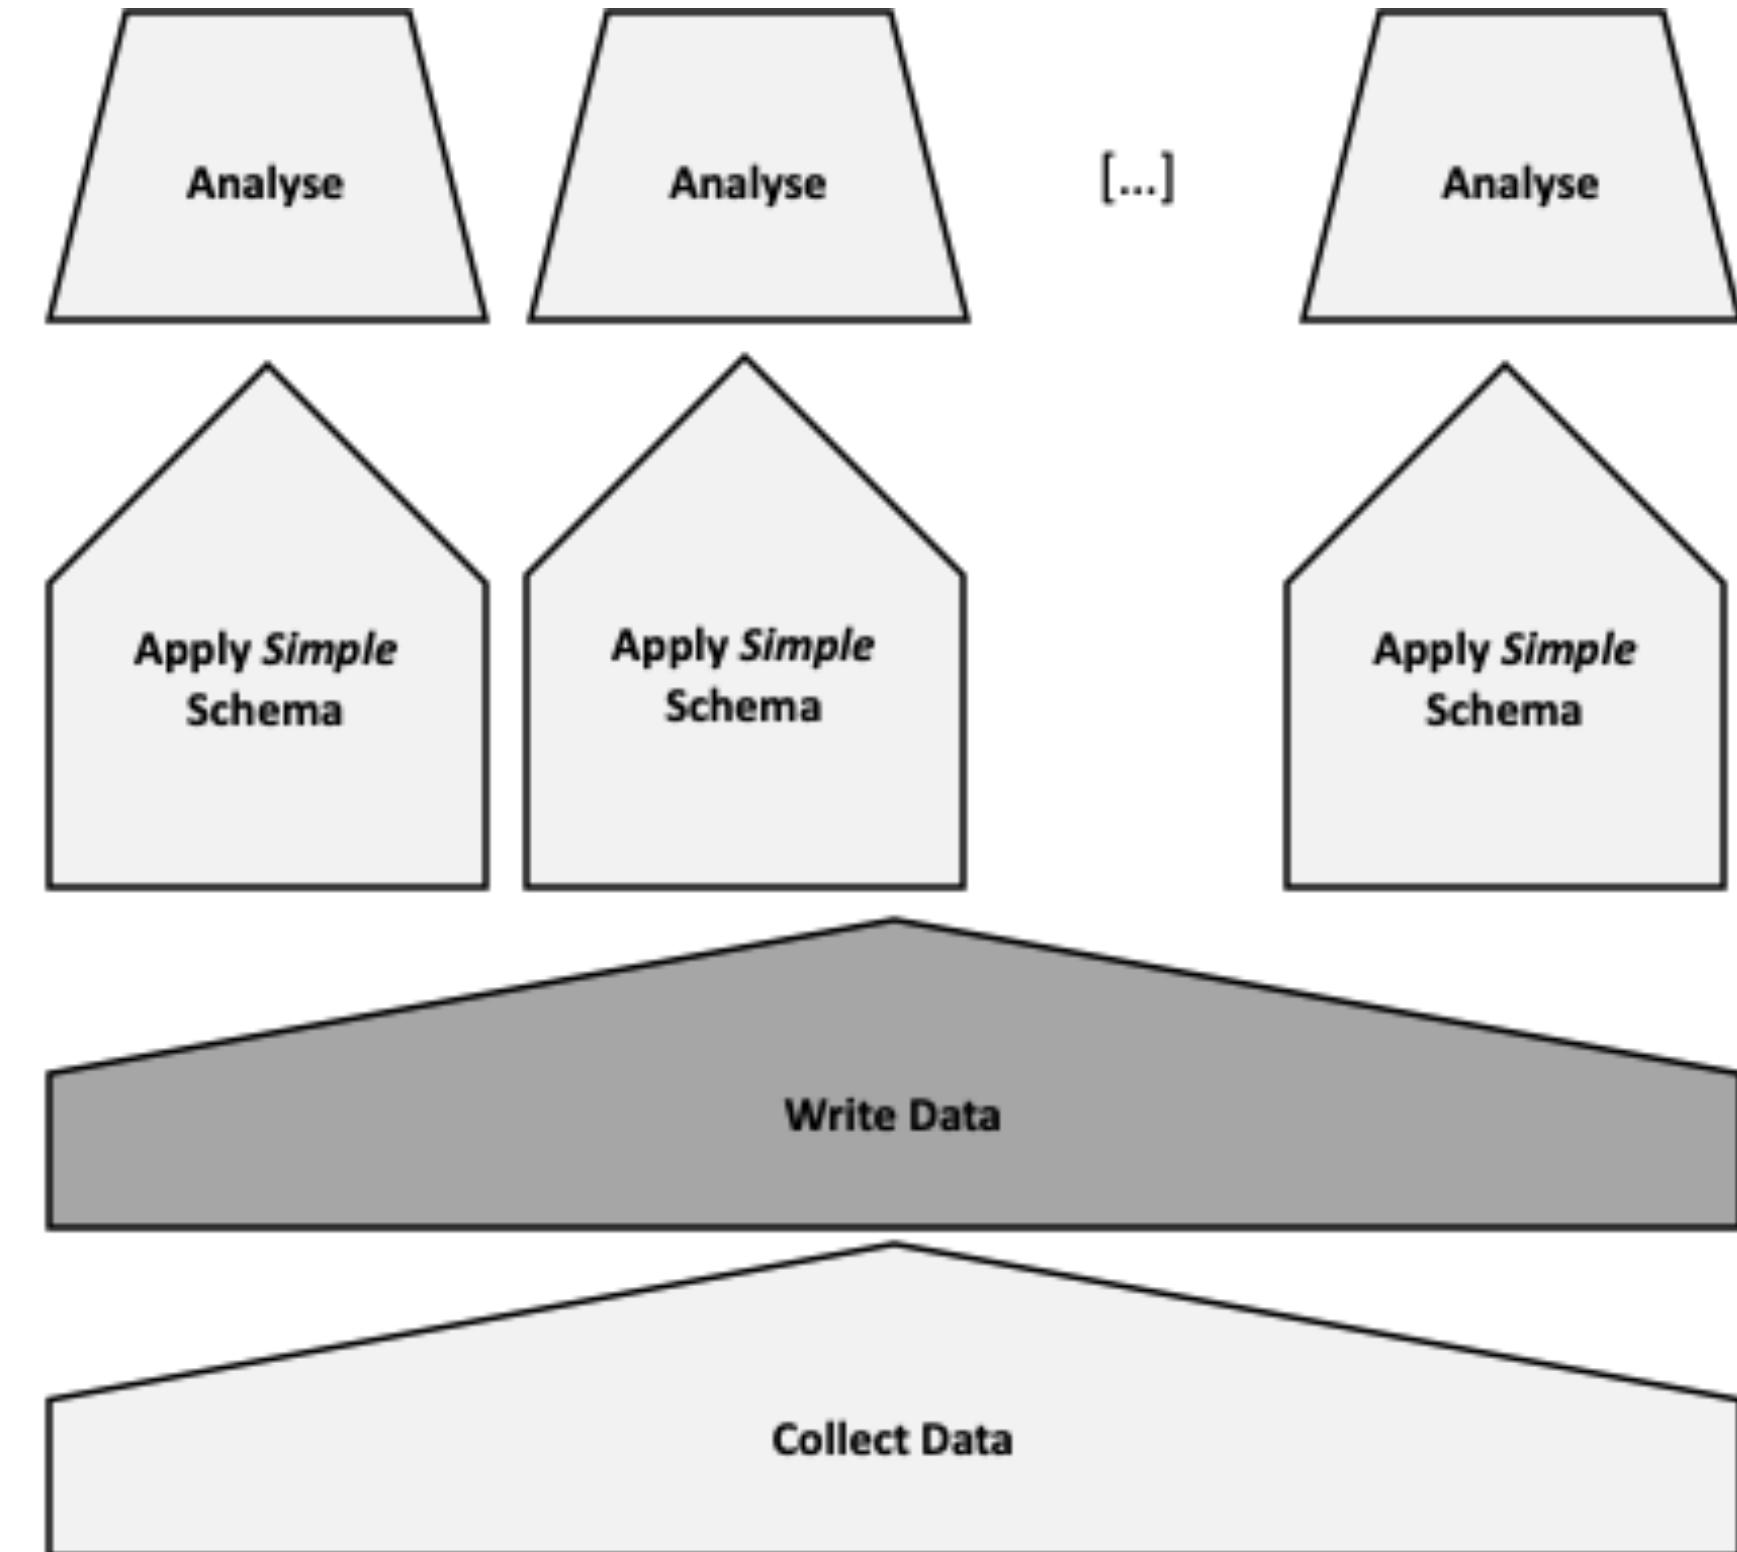


## To Schema on Read



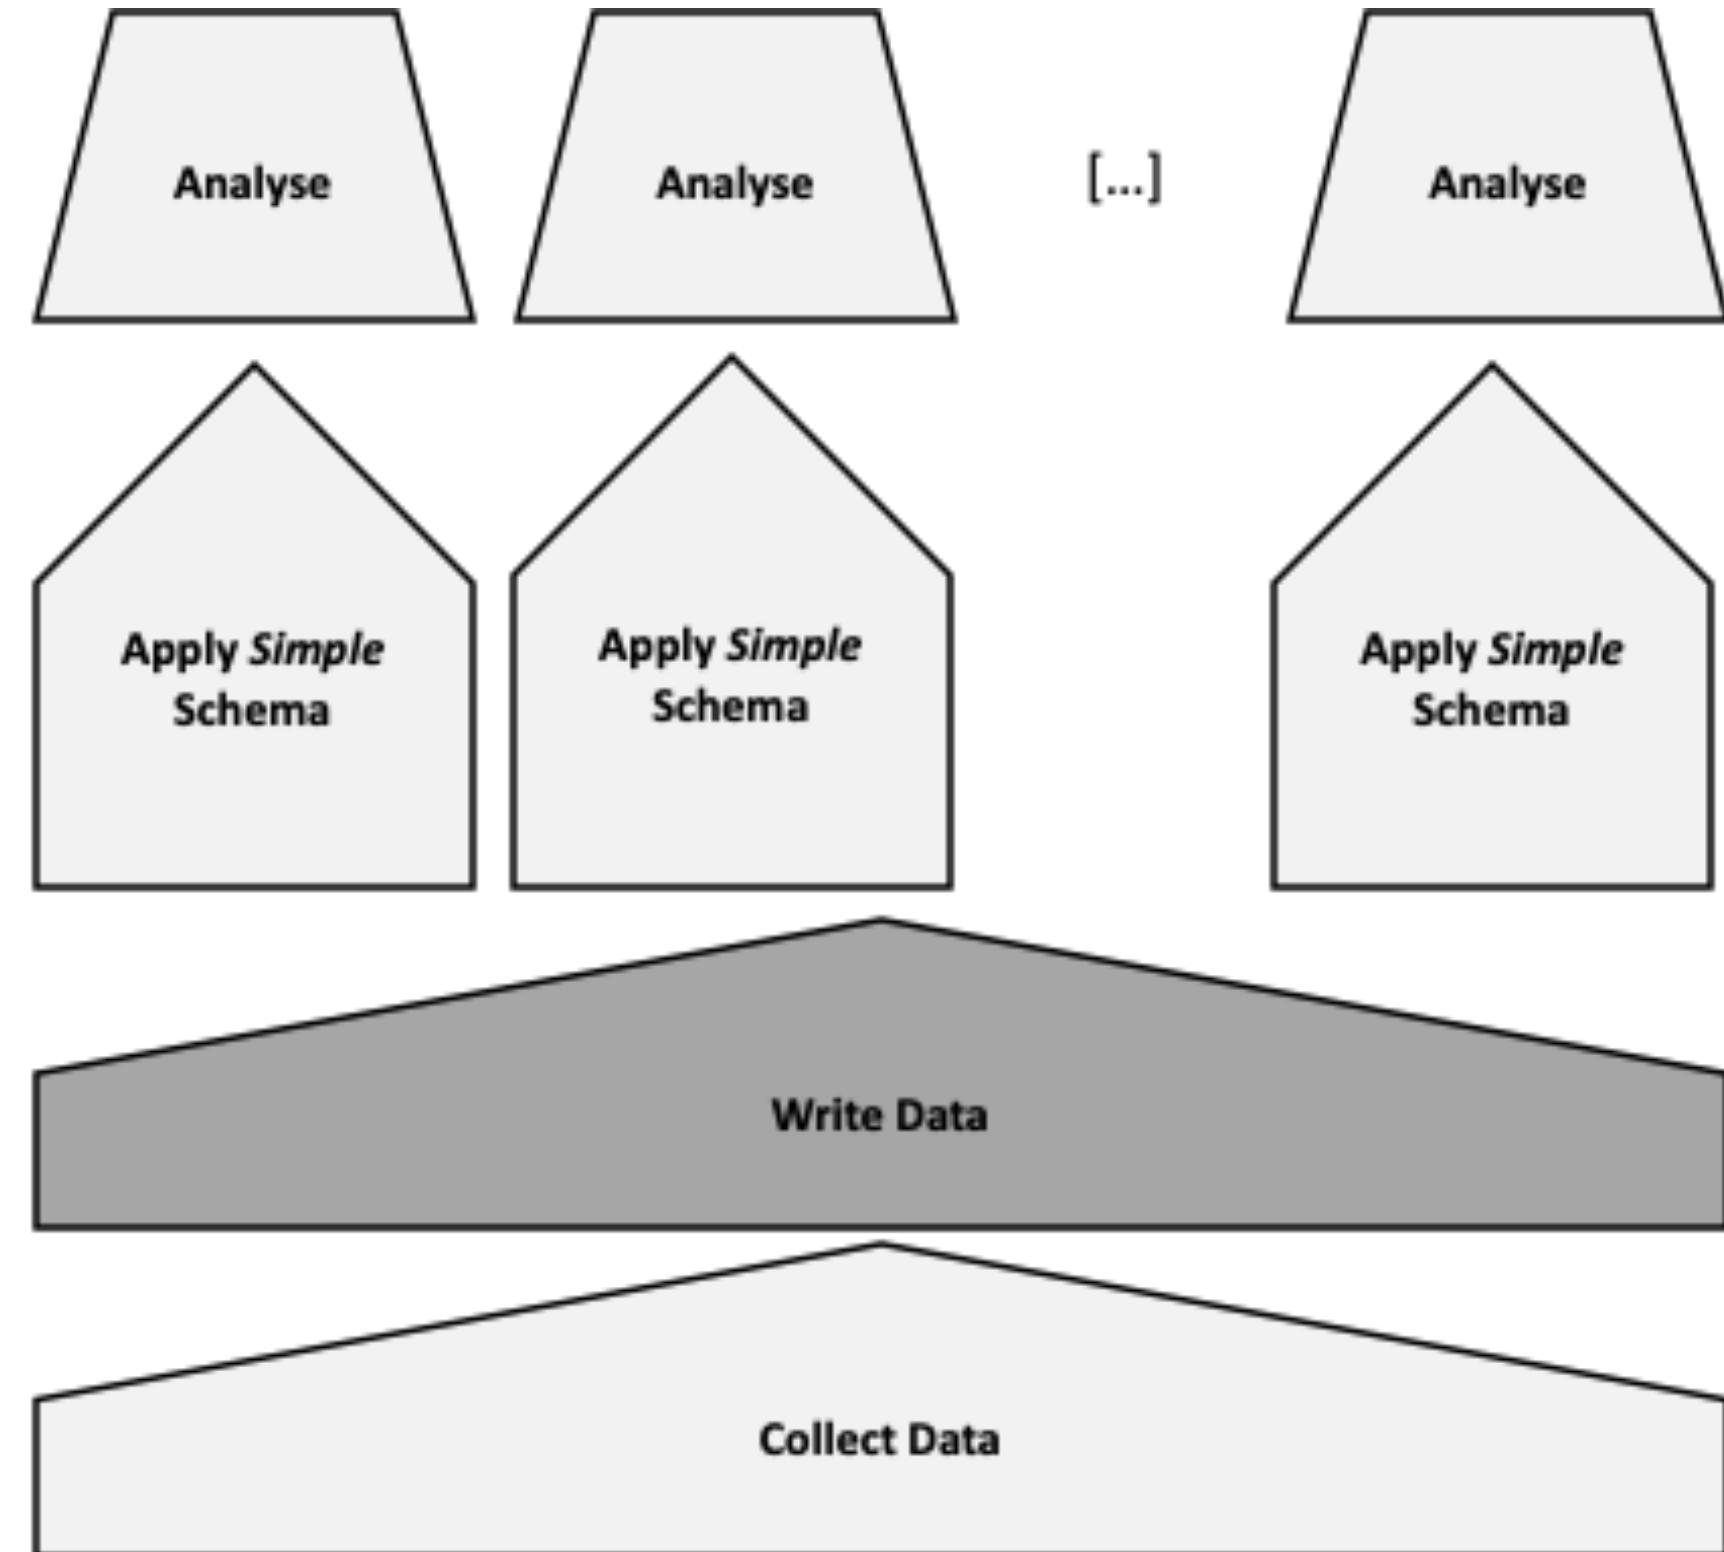
## To Schema on Read

- Load data first, ask question later



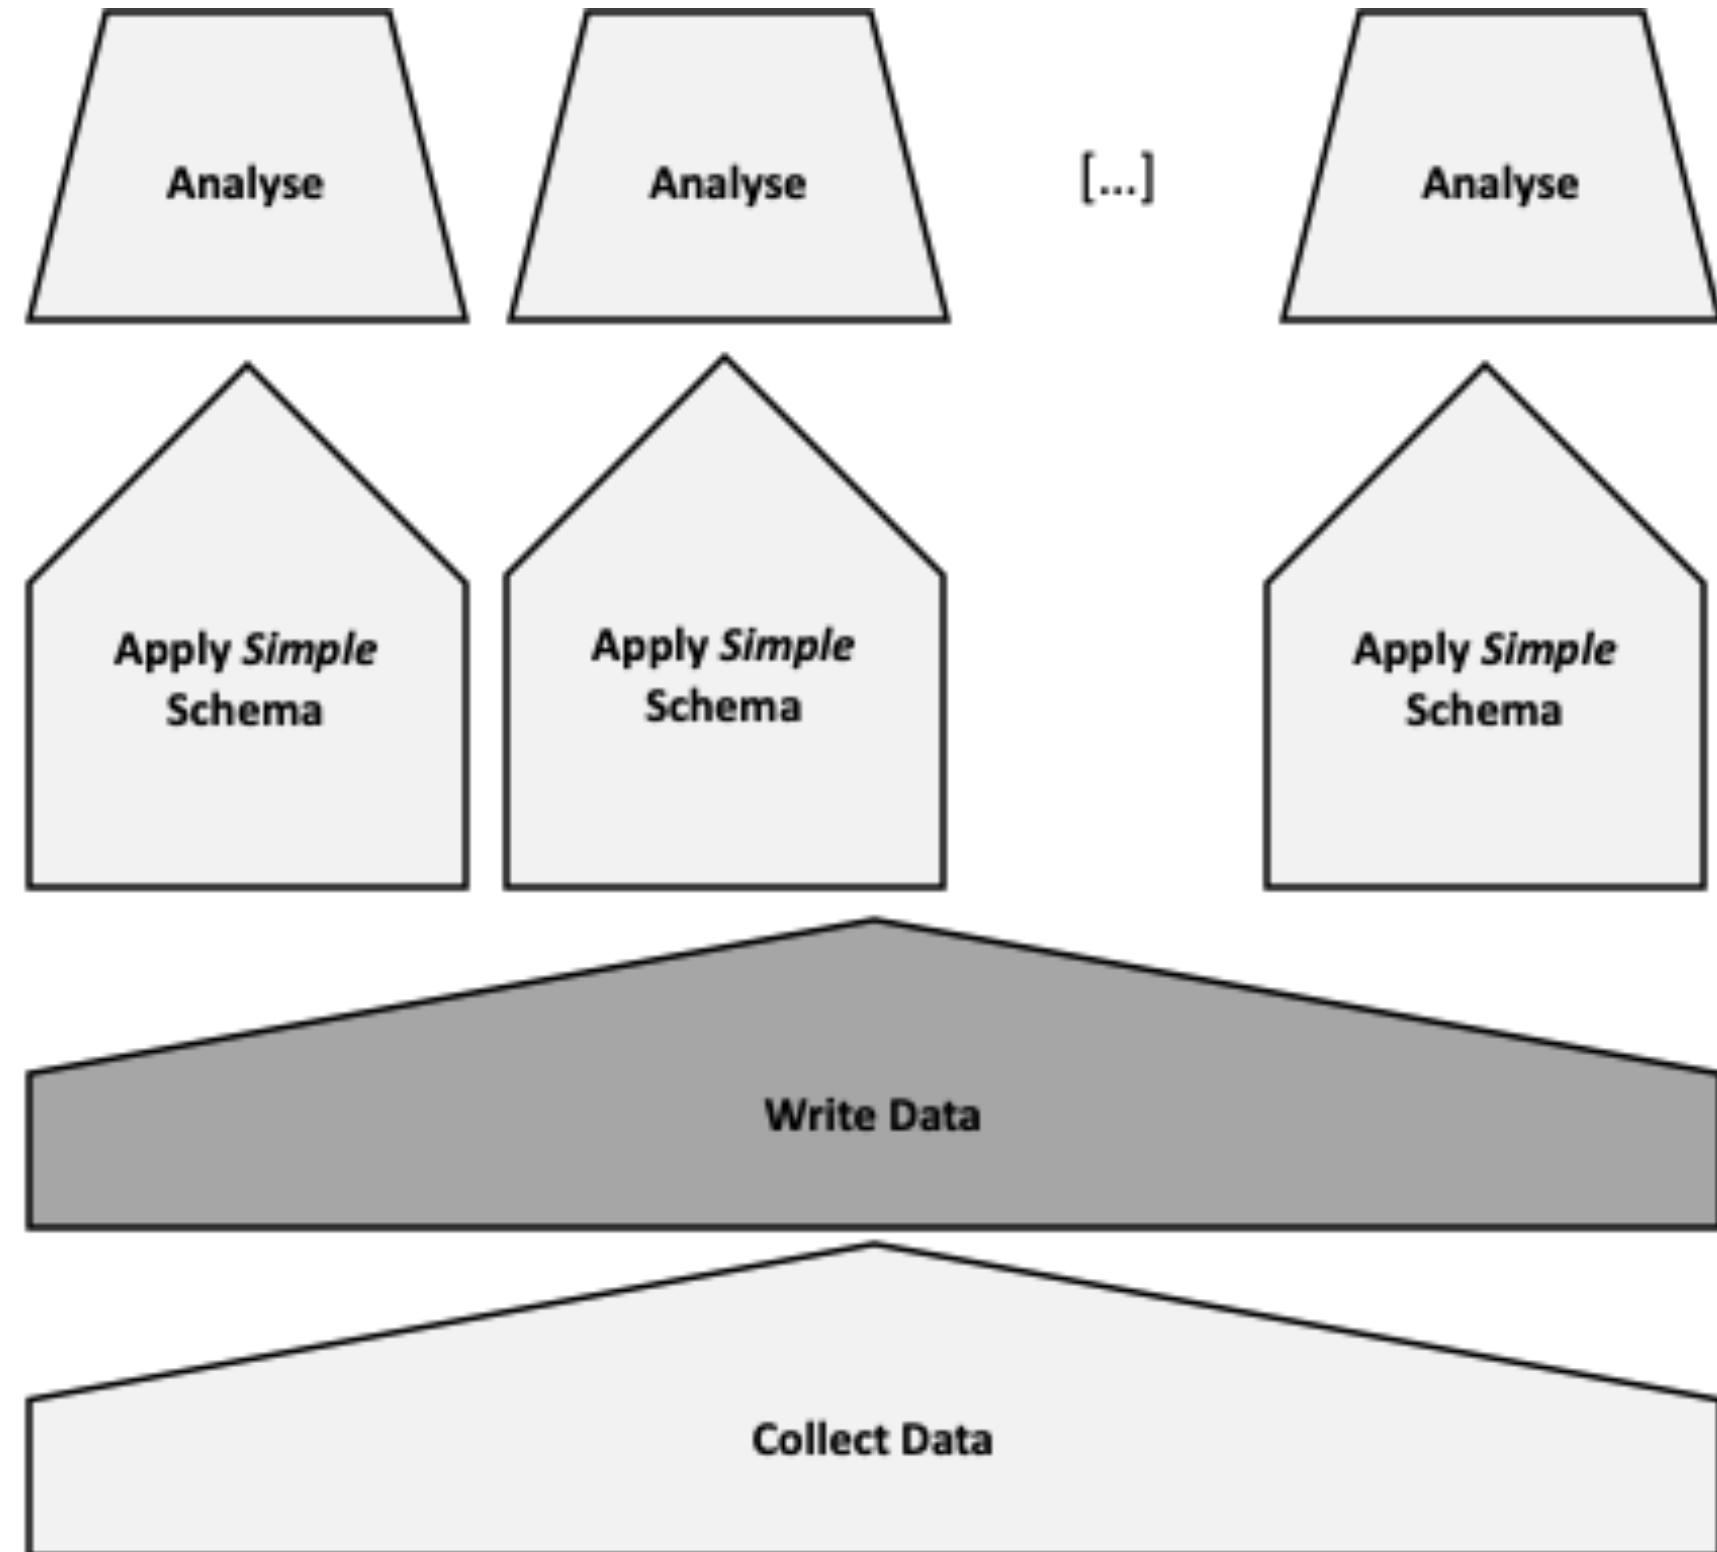
## To Schema on Read

- Load data first, ask question later
- All data are kept, the minimal schema need for an analysis is applied when needed



## To Schema on Read

- Load data first, ask question later
- All data are kept, the minimal schema need for an analysis is applied when needed
- New analyses can be introduced in any point in time



# NoSQL

# NoSQL

- **Queryability:** need for specialized query operations that are not well supported by the relational model

# NoSQL

- **Queryability:** need for specialized query operations that are not well supported by the relational model
- **Schemaless:** desire for a more dynamic and expressive data model than relational

# NoSQL

- **Queryability:** need for specialized query operations that are not well supported by the relational model
- **Schemaless:** desire for a more dynamic and expressive data model than relational
- **Object-Relational Mismatch:** translation between the objects in the application code and the database model

Document Database	Graph Databases
 	 
Wide Column Stores	Key-Value Databases
 	 
	     

## Kinds of NoSQL (2/4)

NoSQL solutions fall into four major areas:

## Kinds of NoSQL (2/4)

NoSQL solutions fall into four major areas:

- **Key-Value Store**

## Kinds of NoSQL (2/4)

NoSQL solutions fall into four major areas:

- **Key-Value Store**
  - A key that refers to a payload (actual content / data)

## Kinds of NoSQL (2/4)

NoSQL solutions fall into four major areas:

- **Key-Value Store**
  - A key that refers to a payload (actual content / data)
  - Examples: MemcacheDB, Azure Table Storage, Redis, HDFS

## Kinds of NoSQL (2/4)

NoSQL solutions fall into four major areas:

- **Key-Value Store**
  - A key that refers to a payload (actual content / data)
  - Examples: MemcacheDB, Azure Table Storage, Redis, HDFS
- **Column Store**

# Kinds of NoSQL (2/4)

NoSQL solutions fall into four major areas:

- **Key-Value Store**

- A key that refers to a payload (actual content / data)
- Examples: MemcacheDB, Azure Table Storage, Redis, HDFS

- **Column Store**

- Column data is saved together, as opposed to row data

# Kinds of NoSQL (2/4)

NoSQL solutions fall into four major areas:

- **Key-Value Store**

- A key that refers to a payload (actual content / data)
- Examples: MemcacheDB, Azure Table Storage, Redis, HDFS

- **Column Store**

- Column data is saved together, as opposed to row data
- Super useful for data analytics

## Kinds of NoSQL (2/4)

NoSQL solutions fall into four major areas:

- **Key-Value Store**

- A key that refers to a payload (actual content / data)
- Examples: MemcacheDB, Azure Table Storage, Redis, HDFS

- **Column Store**

- Column data is saved together, as opposed to row data
- Super useful for data analytics
- Examples: Hadoop, Cassandra, Hypertable

# Kinds of NoSQL (4/4)

## Kinds of NoSQL (4/4)

- **Document / XML / Object Store**

## Kinds of NoSQL (4/4)

- **Document / XML / Object Store**
  - Key (and possibly other indexes) point at a serialized object

## Kinds of NoSQL (4/4)

- **Document / XML / Object Store**
  - Key (and possibly other indexes) point at a serialized object
  - DB can operate against values in document

## Kinds of NoSQL (4/4)

- **Document / XML / Object Store**
  - Key (and possibly other indexes) point at a serialized object
  - DB can operate against values in document
  - Examples: MongoDB, CouchDB, RavenDB

# Kinds of NoSQL (4/4)

- **Document / XML / Object Store**
  - Key (and possibly other indexes) point at a serialized object
  - DB can operate against values in document
  - Examples: MongoDB, CouchDB, RavenDB
- **Graph Store**

# Kinds of NoSQL (4/4)

- **Document / XML / Object Store**
  - Key (and possibly other indexes) point at a serialized object
  - DB can operate against values in document
  - Examples: MongoDB, CouchDB, RavenDB
- **Graph Store**
  - Nodes are stored independently, and the relationship between nodes (edges) are stored with data

# Kinds of NoSQL (4/4)

- **Document / XML / Object Store**
  - Key (and possibly other indexes) point at a serialized object
  - DB can operate against values in document
  - Examples: MongoDB, CouchDB, RavenDB
- **Graph Store**
  - Nodes are stored independently, and the relationship between nodes (edges) are stored with data
  - Examples: AllegroGraph, Neo4j

## ACID vs. BASE trade-off

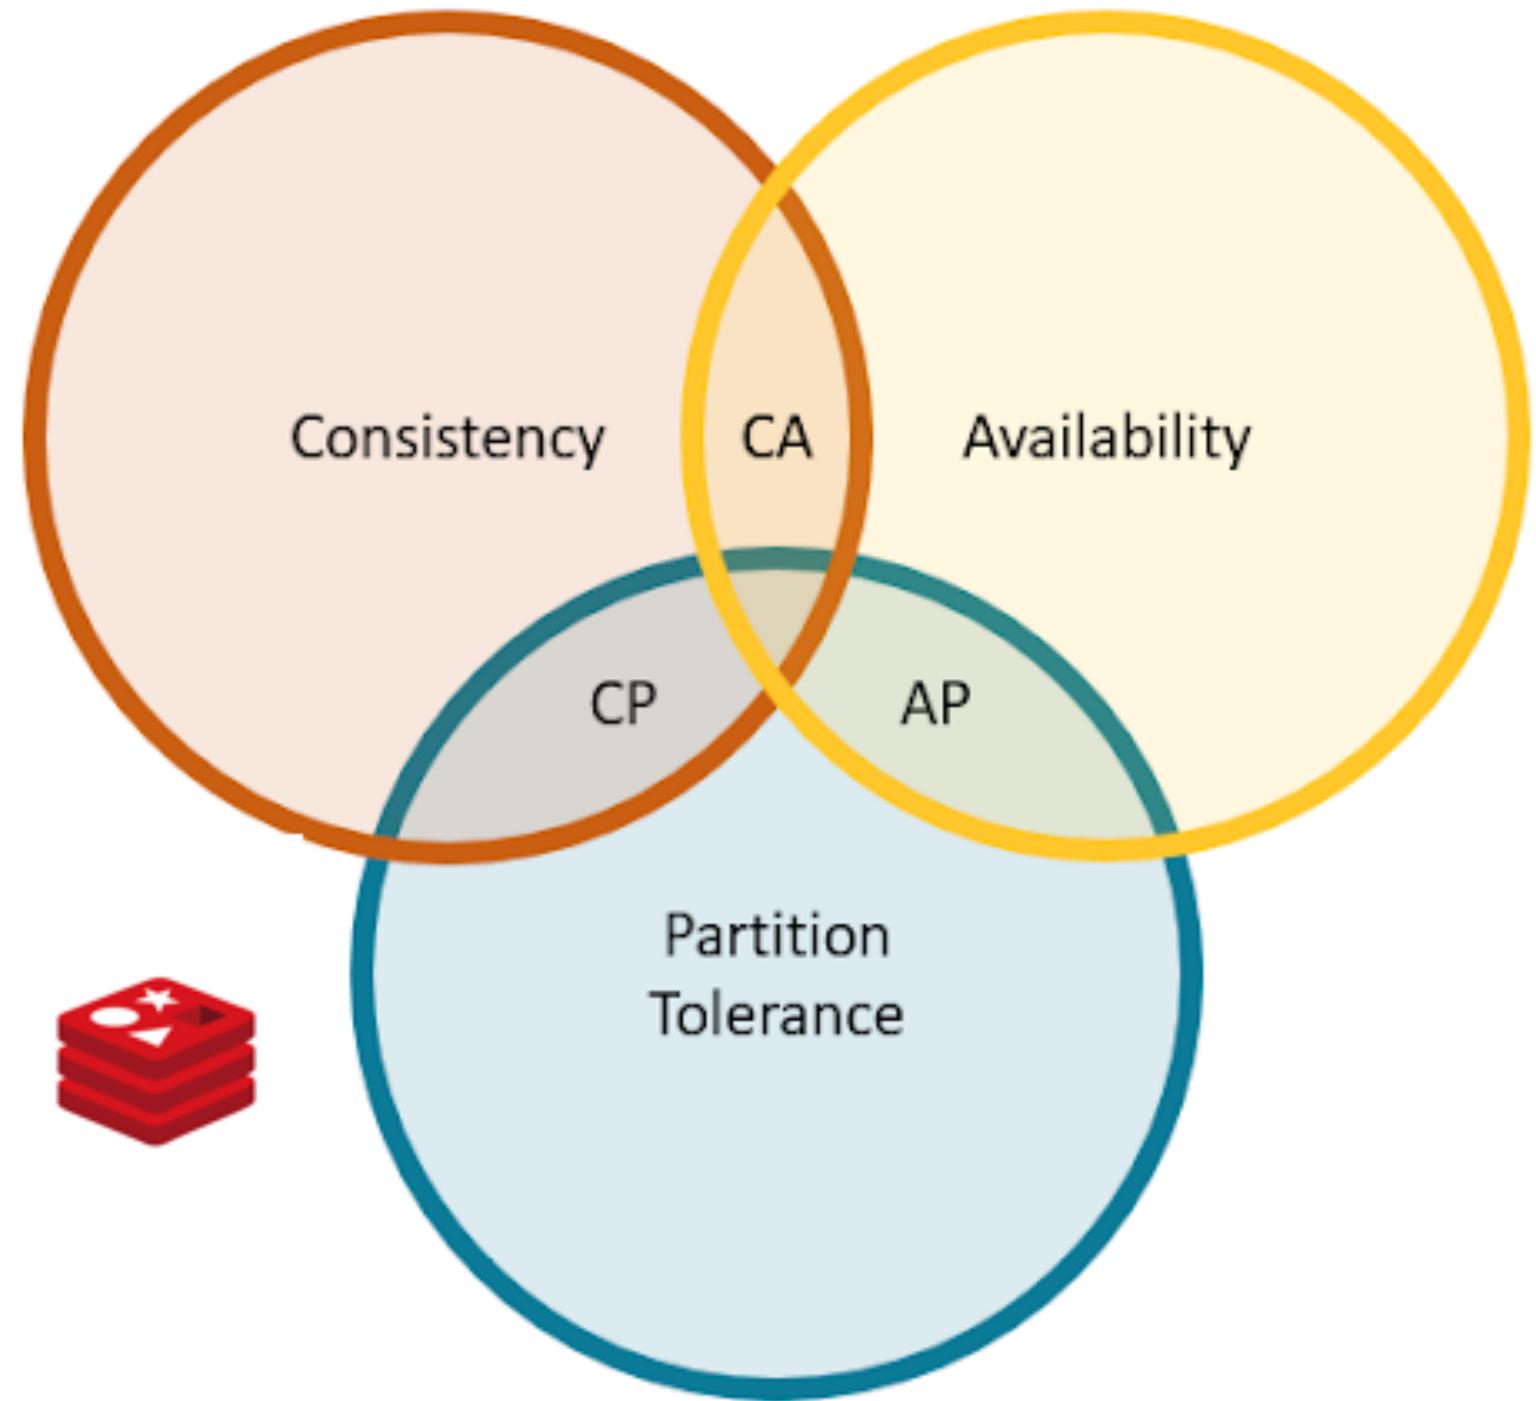
**No general answer** to whether your application needs an ACID versus BASE consistency model.

Given **BASE**'s loose consistency, developers **need to** be more knowledgeable and **rigorous** about **consistent** data if they choose a BASE store for their application.

Planning around **BASE** limitations can sometimes be a major **disadvantage** when compared to the simplicity of ACID transactions.

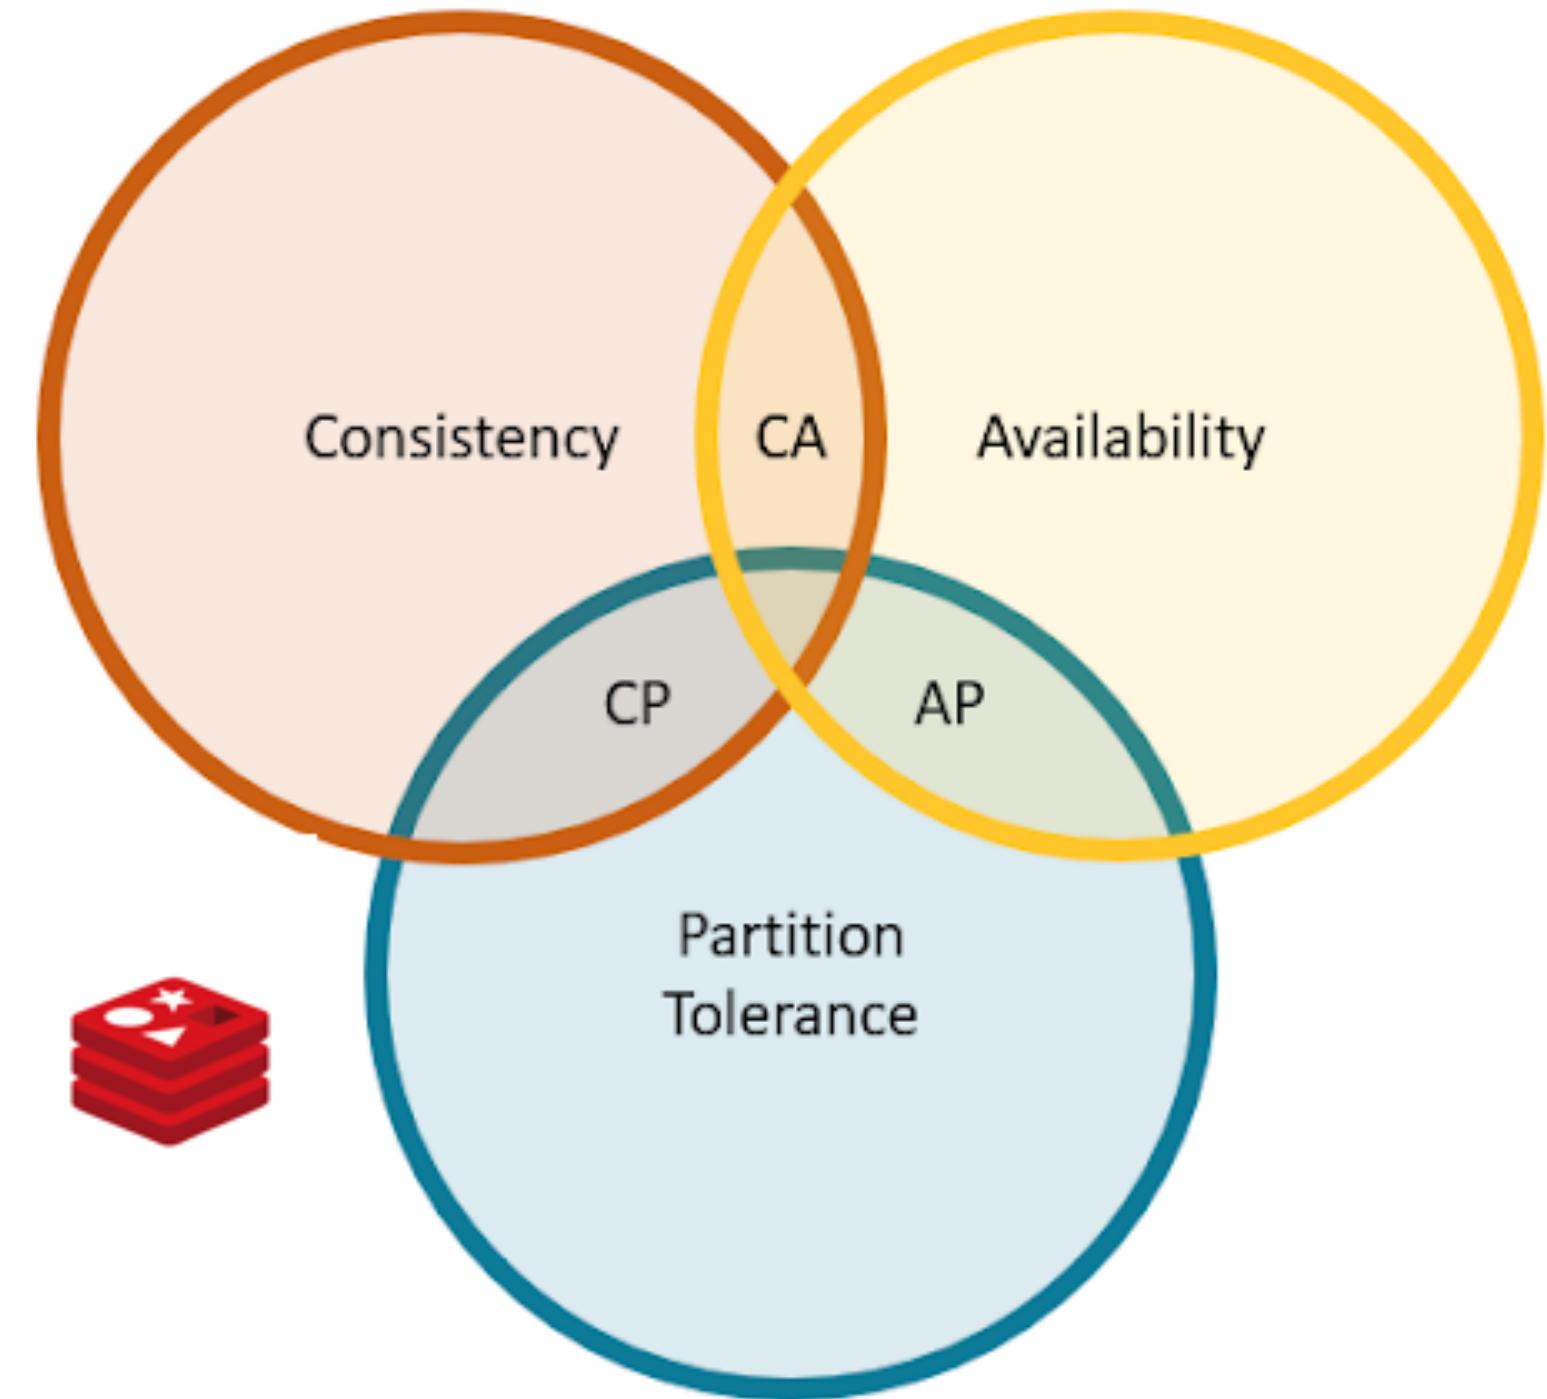
A fully **ACID** database is the perfect fit for use cases where data **reliability** and **consistency** are essential.

# Redis



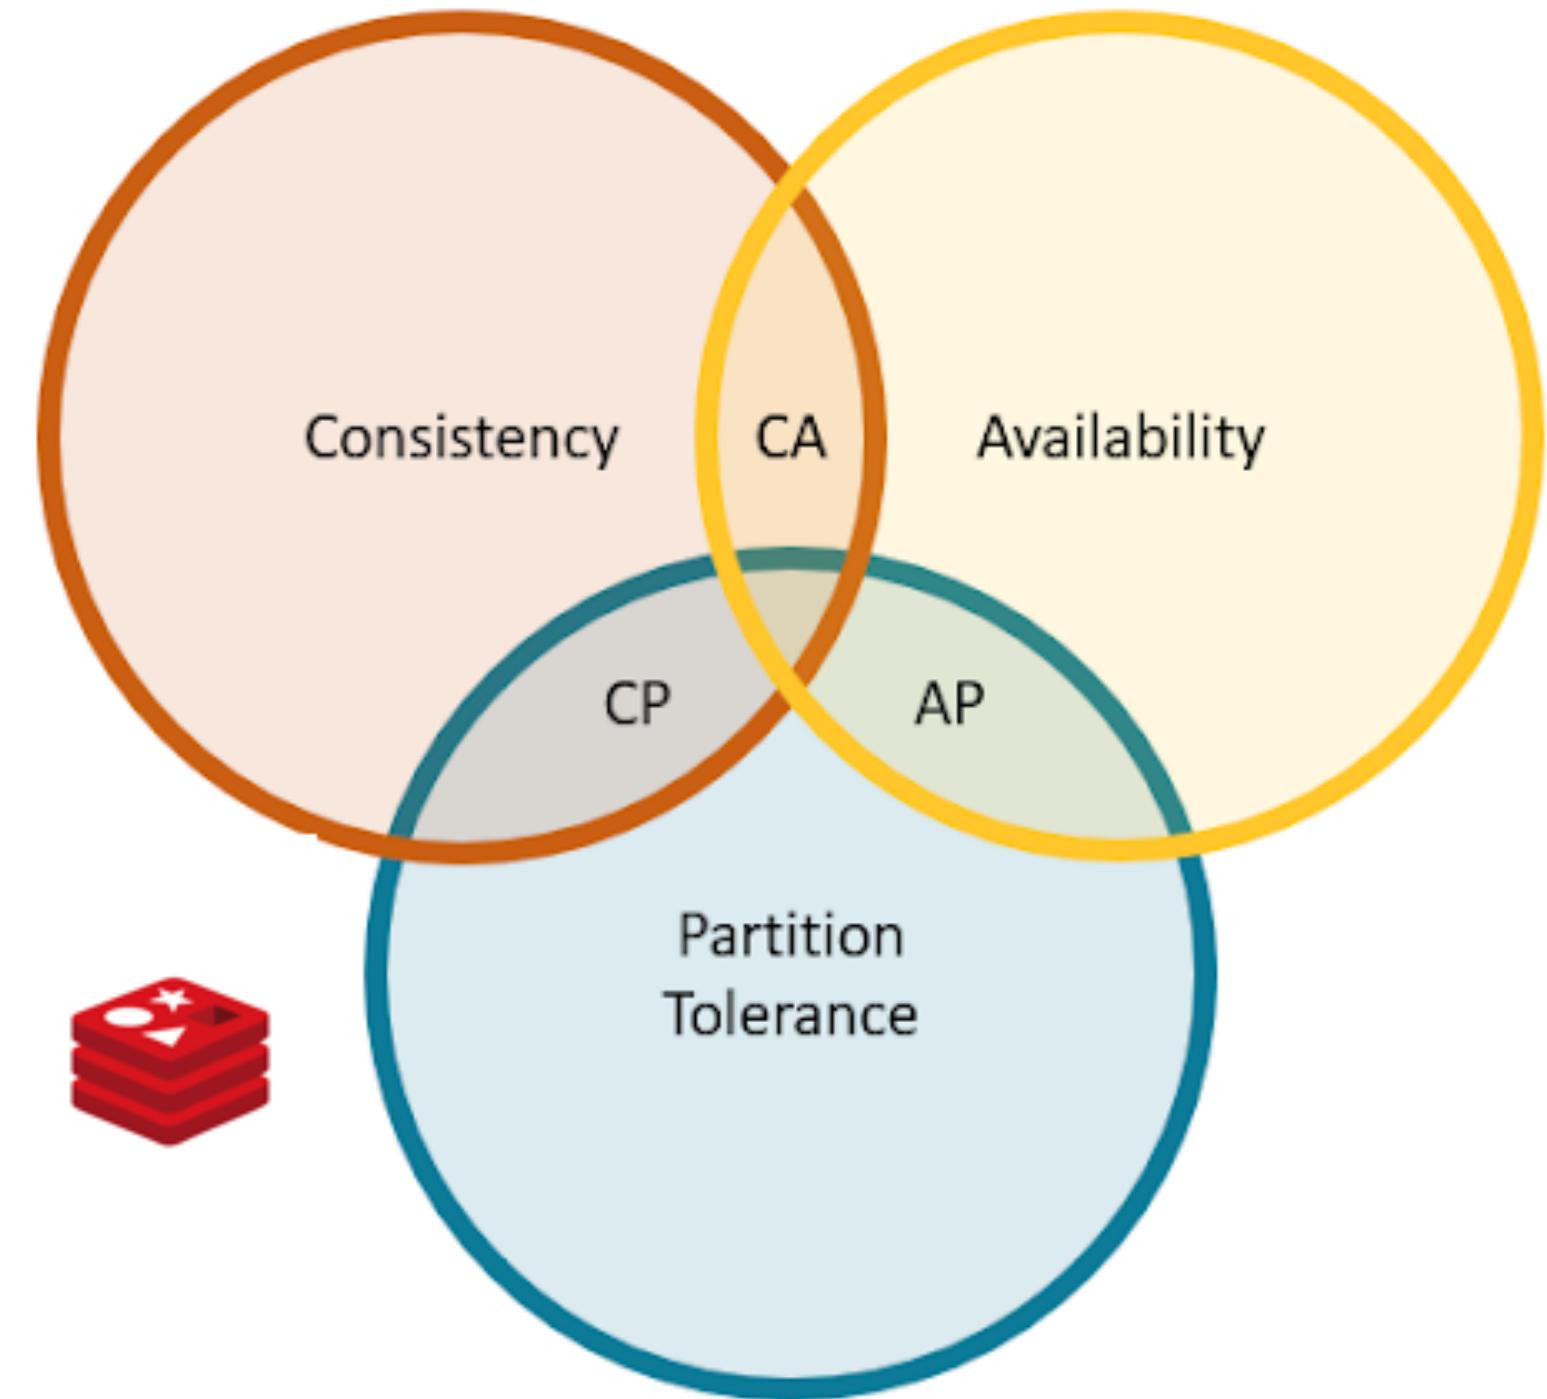
# Redis

- is an In-Memory [[Key-Value Store]]



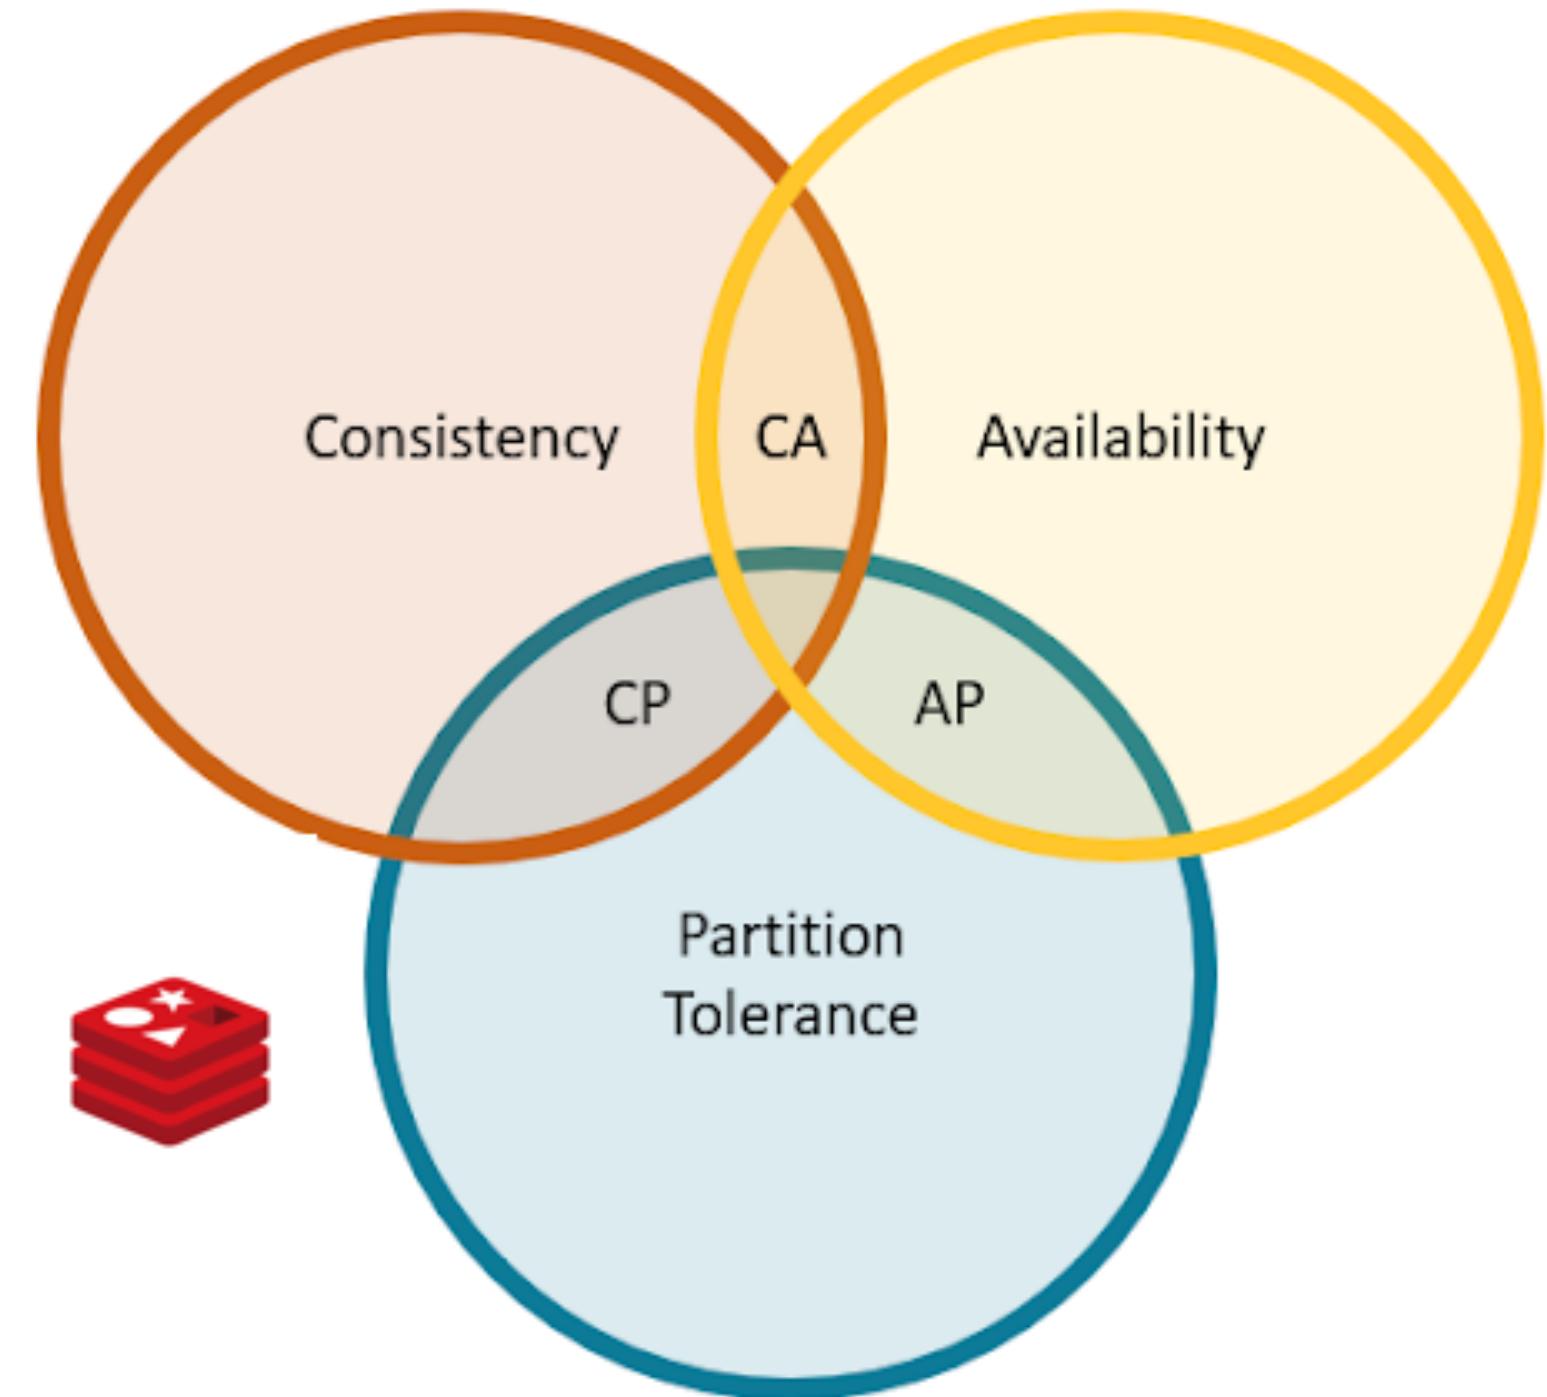
# Redis

- is an In-Memory [[Key-Value Store]]
- Strong consistency (**C**)



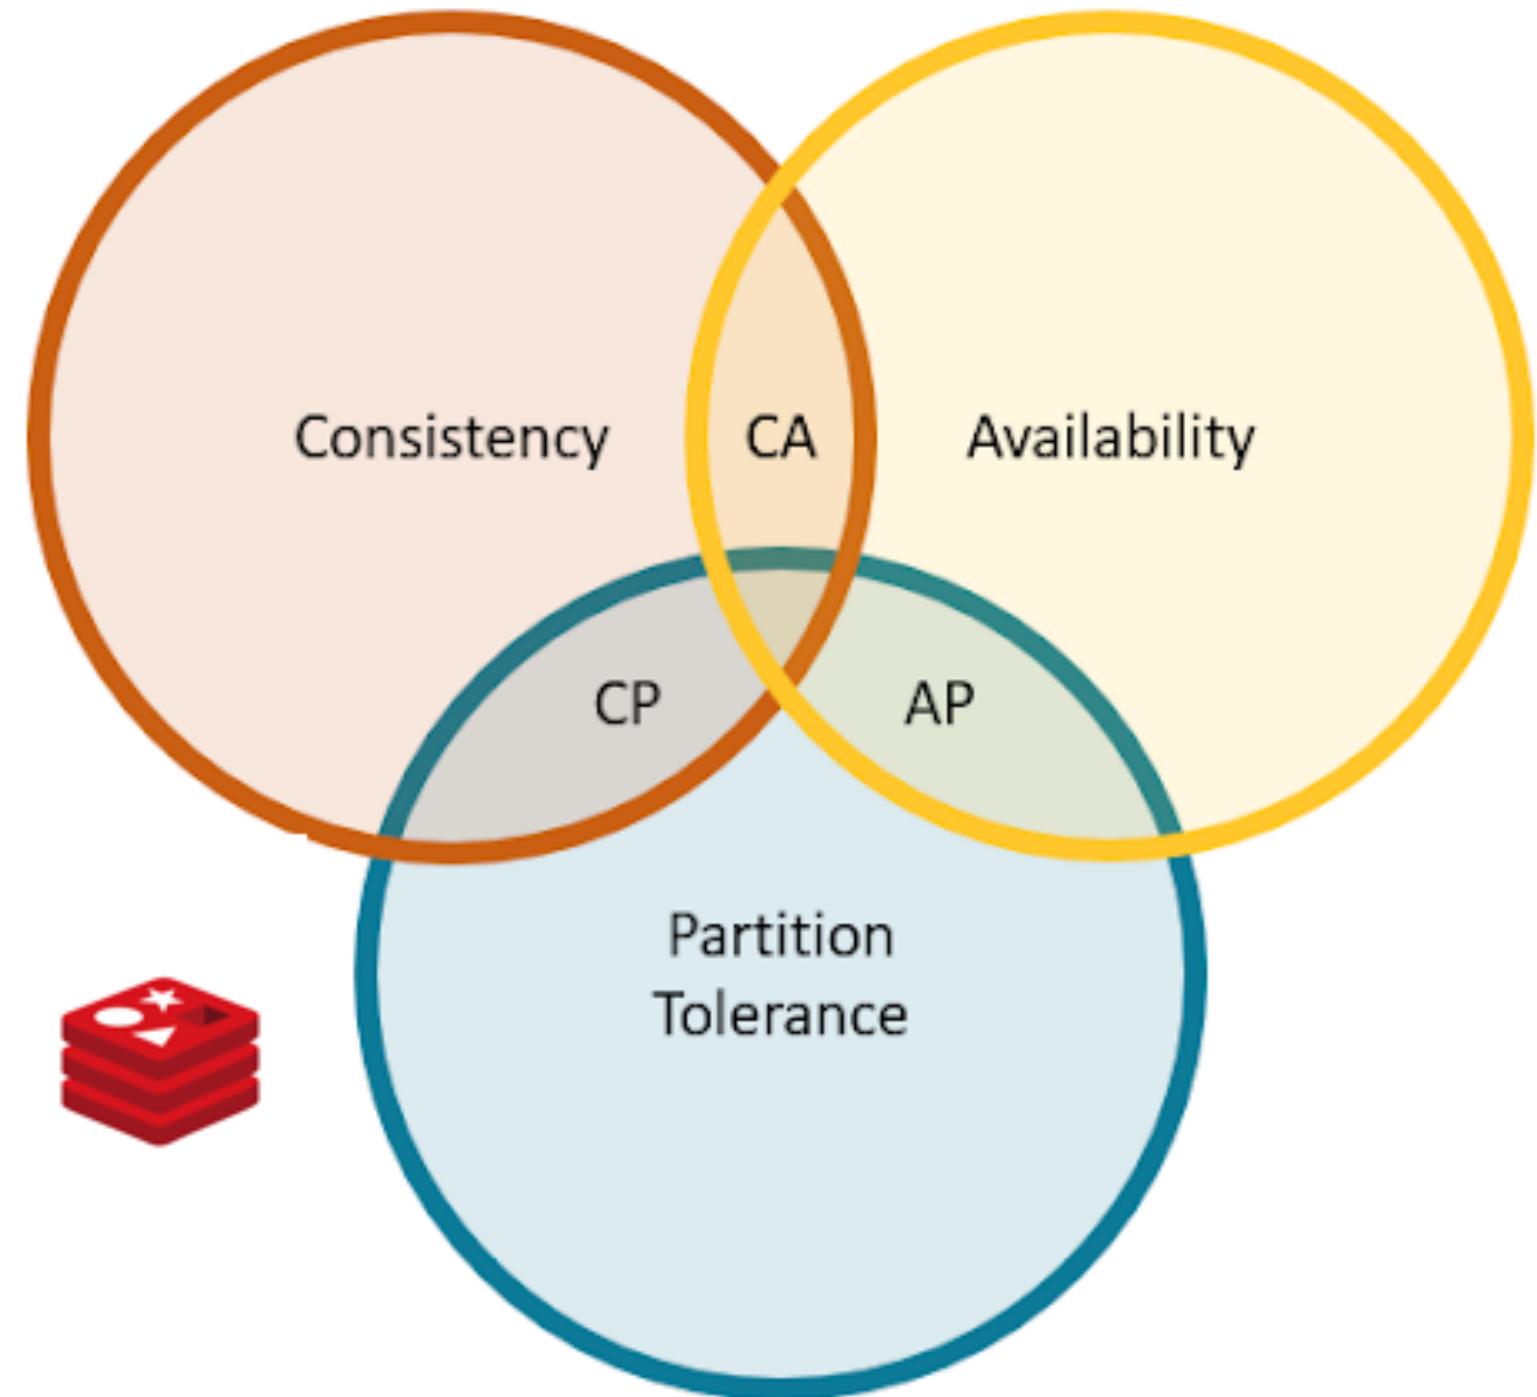
# Redis

- is an In-Memory [[Key-Value Store]]
- Strong consistency (**C**)
- *Tuneably available* (**A**)



# Redis

- is an In-Memory [[Key-Value Store]]
- Strong consistency (**C**)
- *Tuneably* available (**A**)
- Horizontal Scalable (**P**)



# What to remember

# What to remember

- Data Model: Hash Table

# What to remember

- Data Model: Hash Table
- Data Types: ASCII and Collections

# What to remember

- Data Model: Hash Table
- Data Types: ASCII and Collections
- Operations: Lookups and Iterations

# What to remember

- Data Model: Hash Table
- Data Types: ASCII and Collections
- Operations: Lookups and Iterations
- Replication: Master-Slaves(read only)

# What to remember

- Data Model: Hash Table
- Data Types: ASCII and Collections
- Operations: Lookups and Iterations
- Replication: Master-Slaves(read only)
- Partitioning: Client/Proxy/Query Router

# What to remember

- Data Model: Hash Table
- Data Types: ASCII and Collections
- Operations: Lookups and Iterations
- Replication: Master-Slaves(read only)
- Partitioning: Client/Proxy/Query Router
- Architecture

# What to remember

- Data Model: Hash Table
- Data Types: ASCII and Collections
- Operations: Lookups and Iterations
- Replication: Master-Slaves(read only)
- Partitioning: Client/Proxy/Query Router
- Architecture
  - Standalone

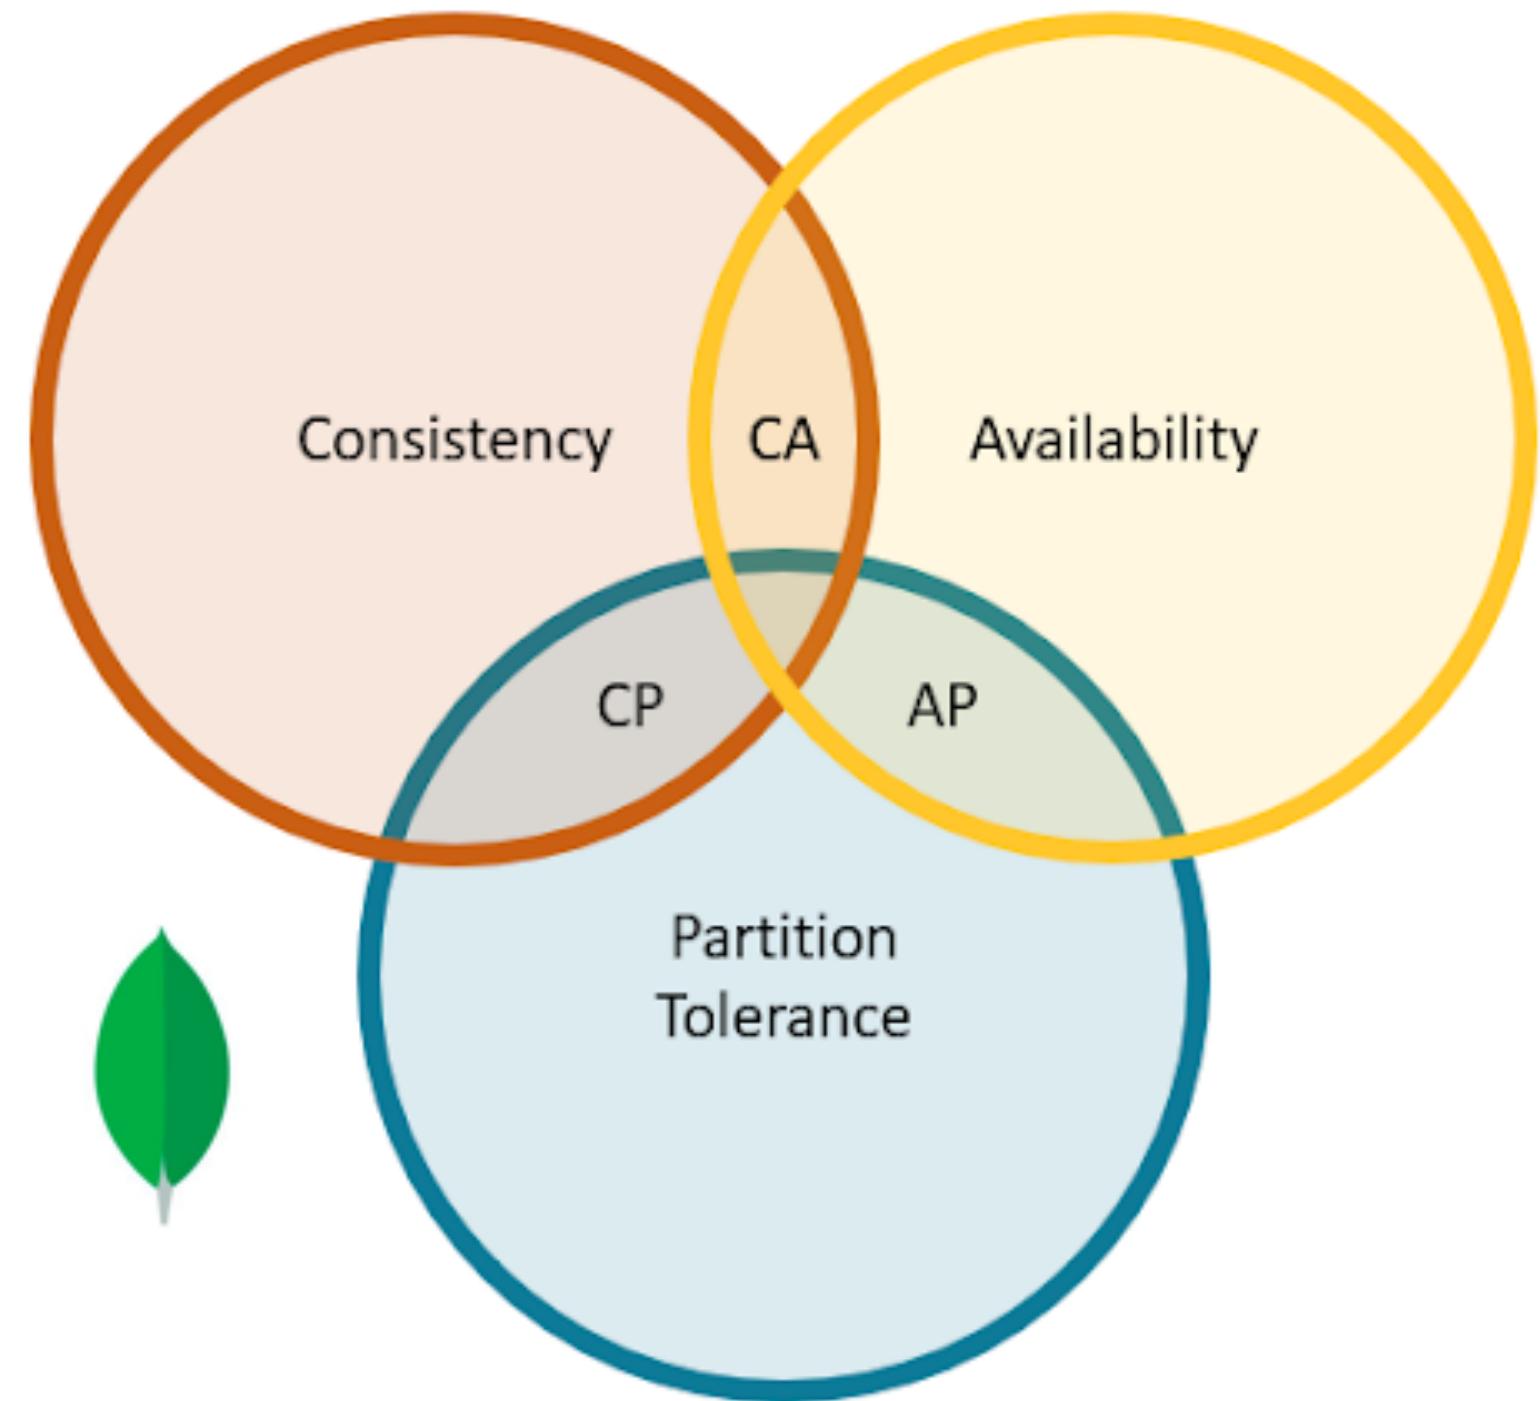
# What to remember

- Data Model: Hash Table
- Data Types: ASCII and Collections
- Operations: Lookups and Iterations
- Replication: Master-Slaves(read only)
- Partitioning: Client/Proxy/Query Router
- Architecture
  - Standalone
  - Sentinel

# What to remember

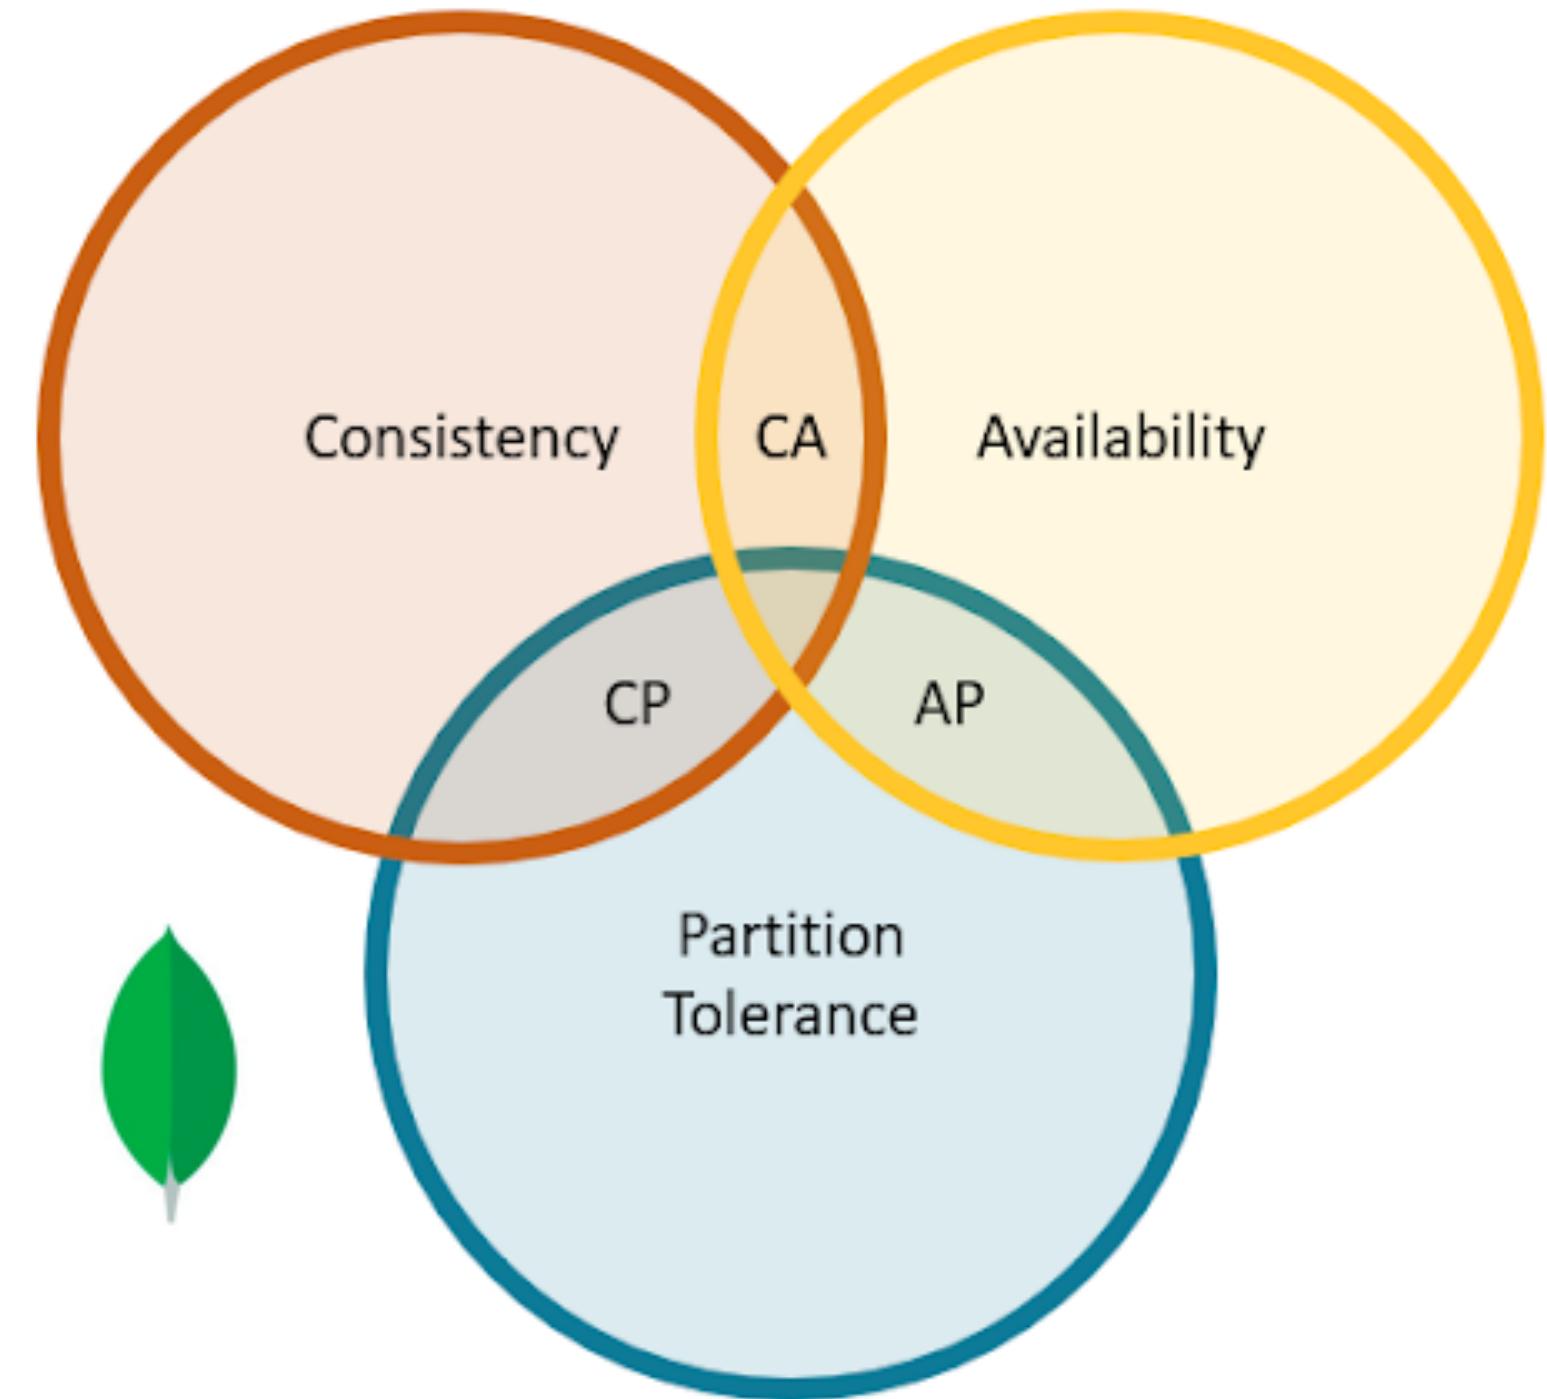
- Data Model: Hash Table
- Data Types: ASCII and Collections
- Operations: Lookups and Iterations
- Replication: Master-Slaves(read only)
- Partitioning: Client/Proxy/Query Router
- Architecture
  - Standalone
  - Sentinel
  - Cluster

# MongoDB



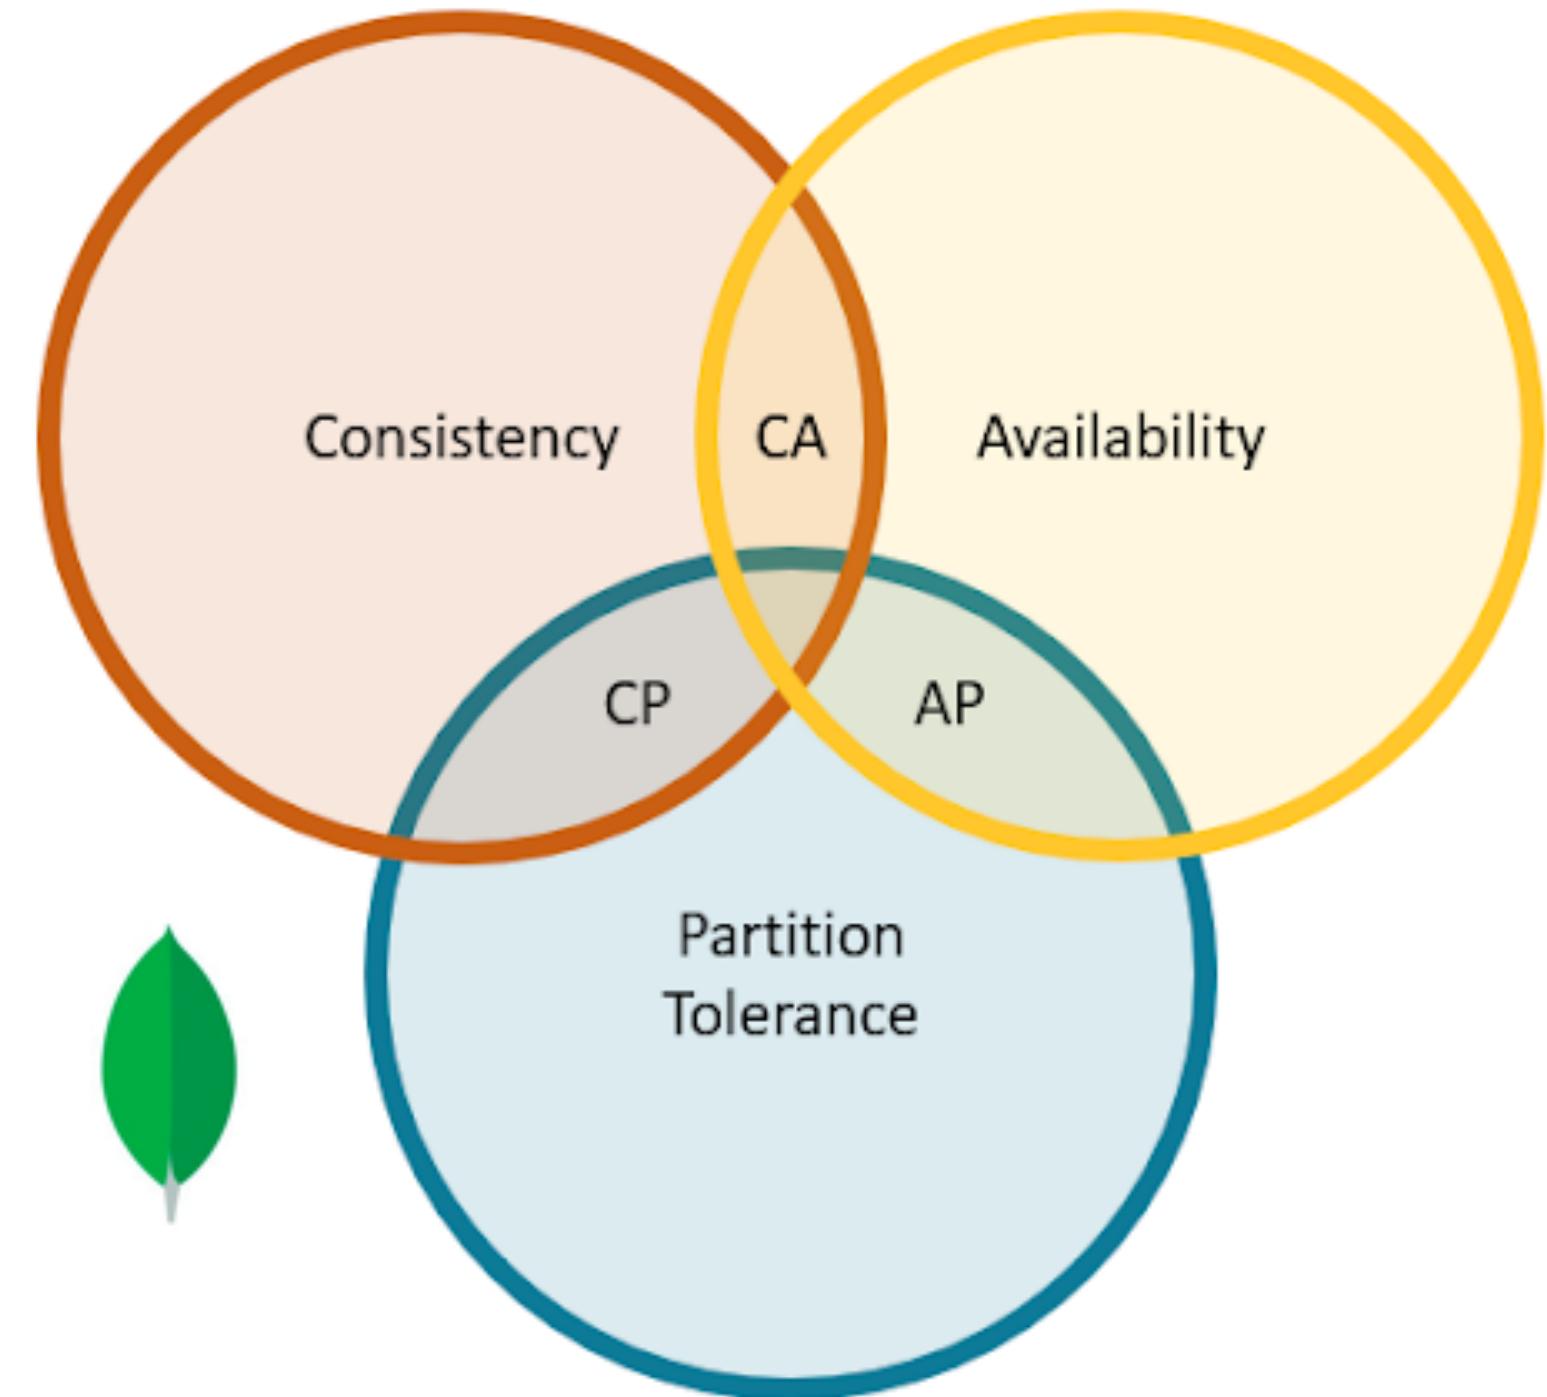
# MongoDB

- An In-Memory [[Document Databases]]



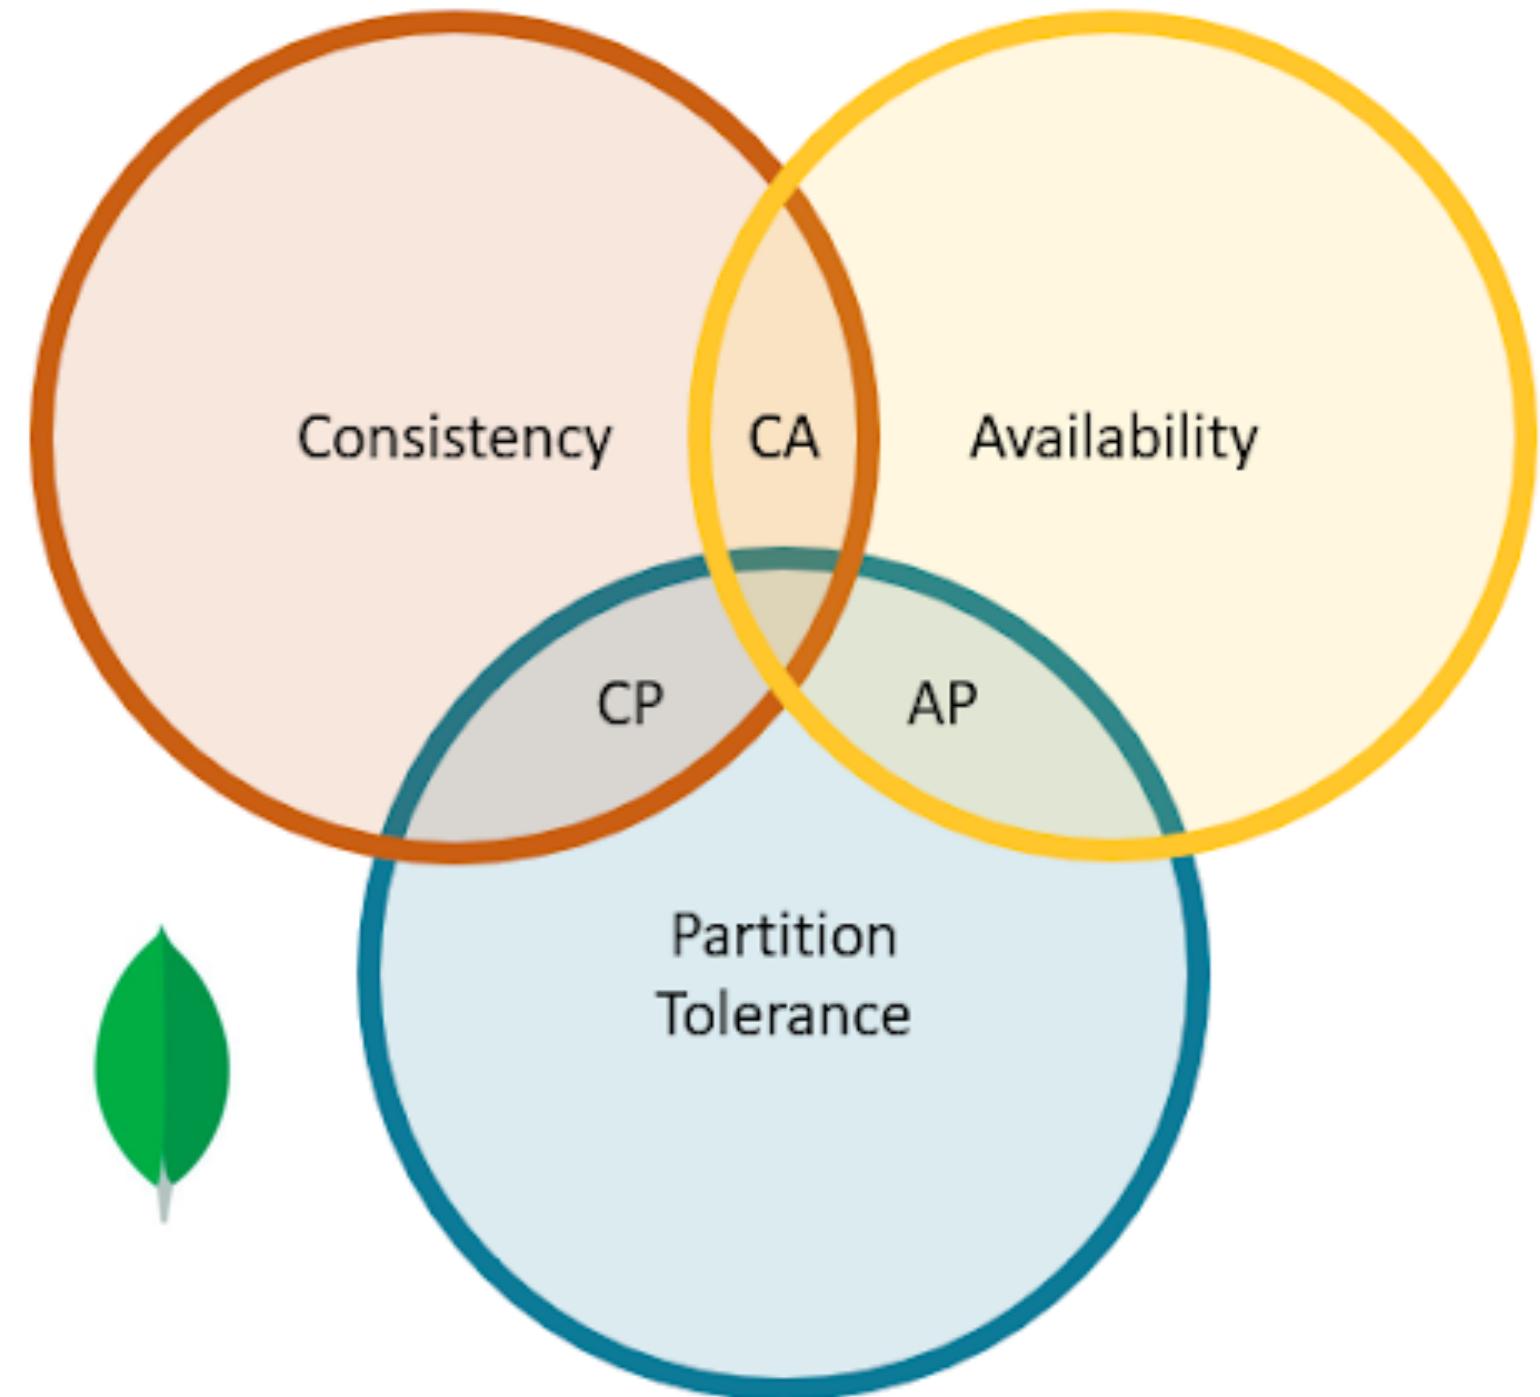
# MongoDB

- An In-Memory [[Document Databases]]
- Strong consistency (**C**)



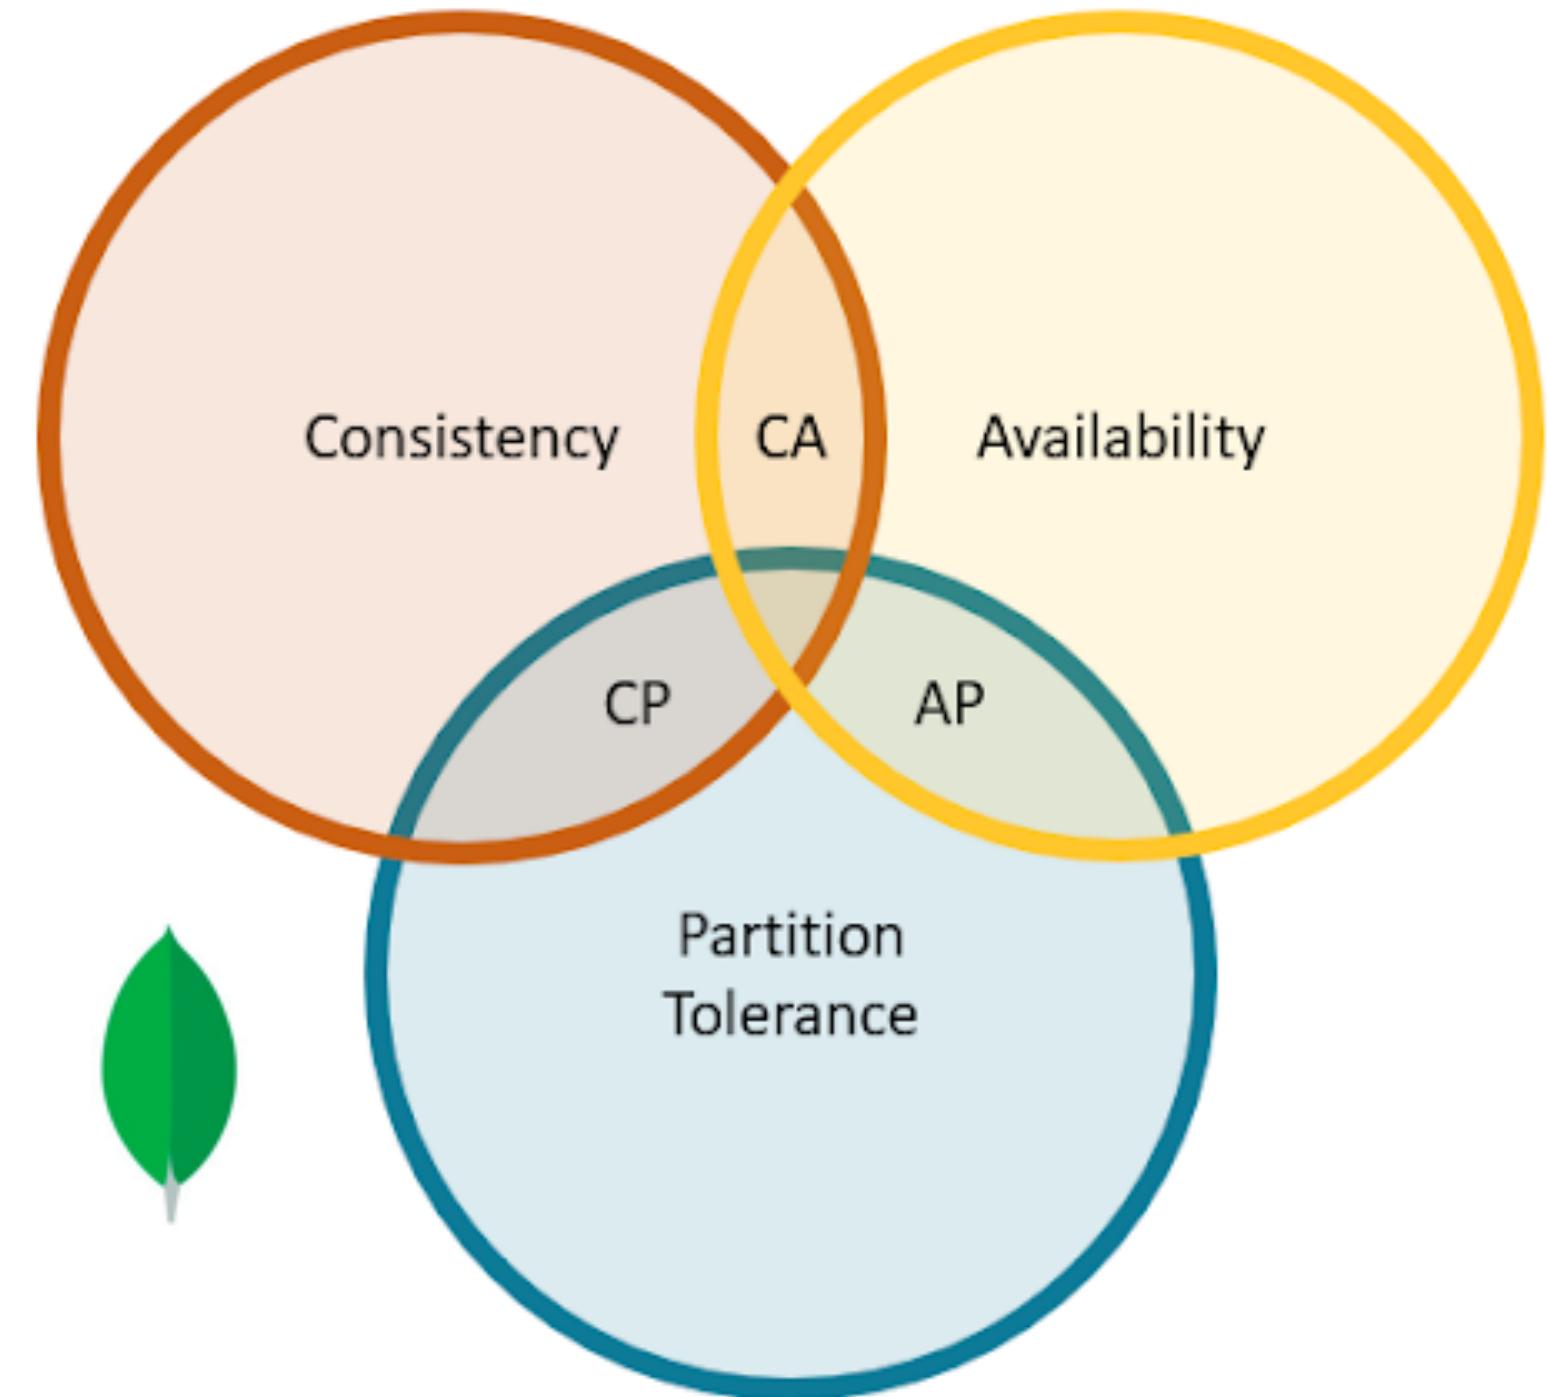
# MongoDB

- An In-Memory [[Document Databases]]
- Strong consistency (**C**)
- *Tuneably* available (**A**)



# MongoDB

- An In-Memory [[Document Databases]]
- Strong consistency (**C**)
- *Tuneably* available (**A**)
- Horizontal Scalable (**P**)



# What to remember

# What to remember

- Data Model: JSON

# What to remember

- Data Model: JSON
- Storage: BSON

# What to remember

- Data Model: JSON
- Storage: BSON
- Operations: iterations

# What to remember

- Data Model: JSON
- Storage: BSON
- Operations: iterations
- Replication: Replica Sets

# What to remember

- Data Model: JSON
- Storage: BSON
- Operations: iterations
- Replication: Replica Sets
- Partitioning: Sharding

# What to remember

- Data Model: JSON
- Storage: BSON
- Operations: iterations
- Replication: Replica Sets
- Partitioning: Sharding
- Architecture

# What to remember

- Data Model: JSON
- Storage: BSON
- Operations: iterations
- Replication: Replica Sets
- Partitioning: Sharding
- Architecture
  - Mongod

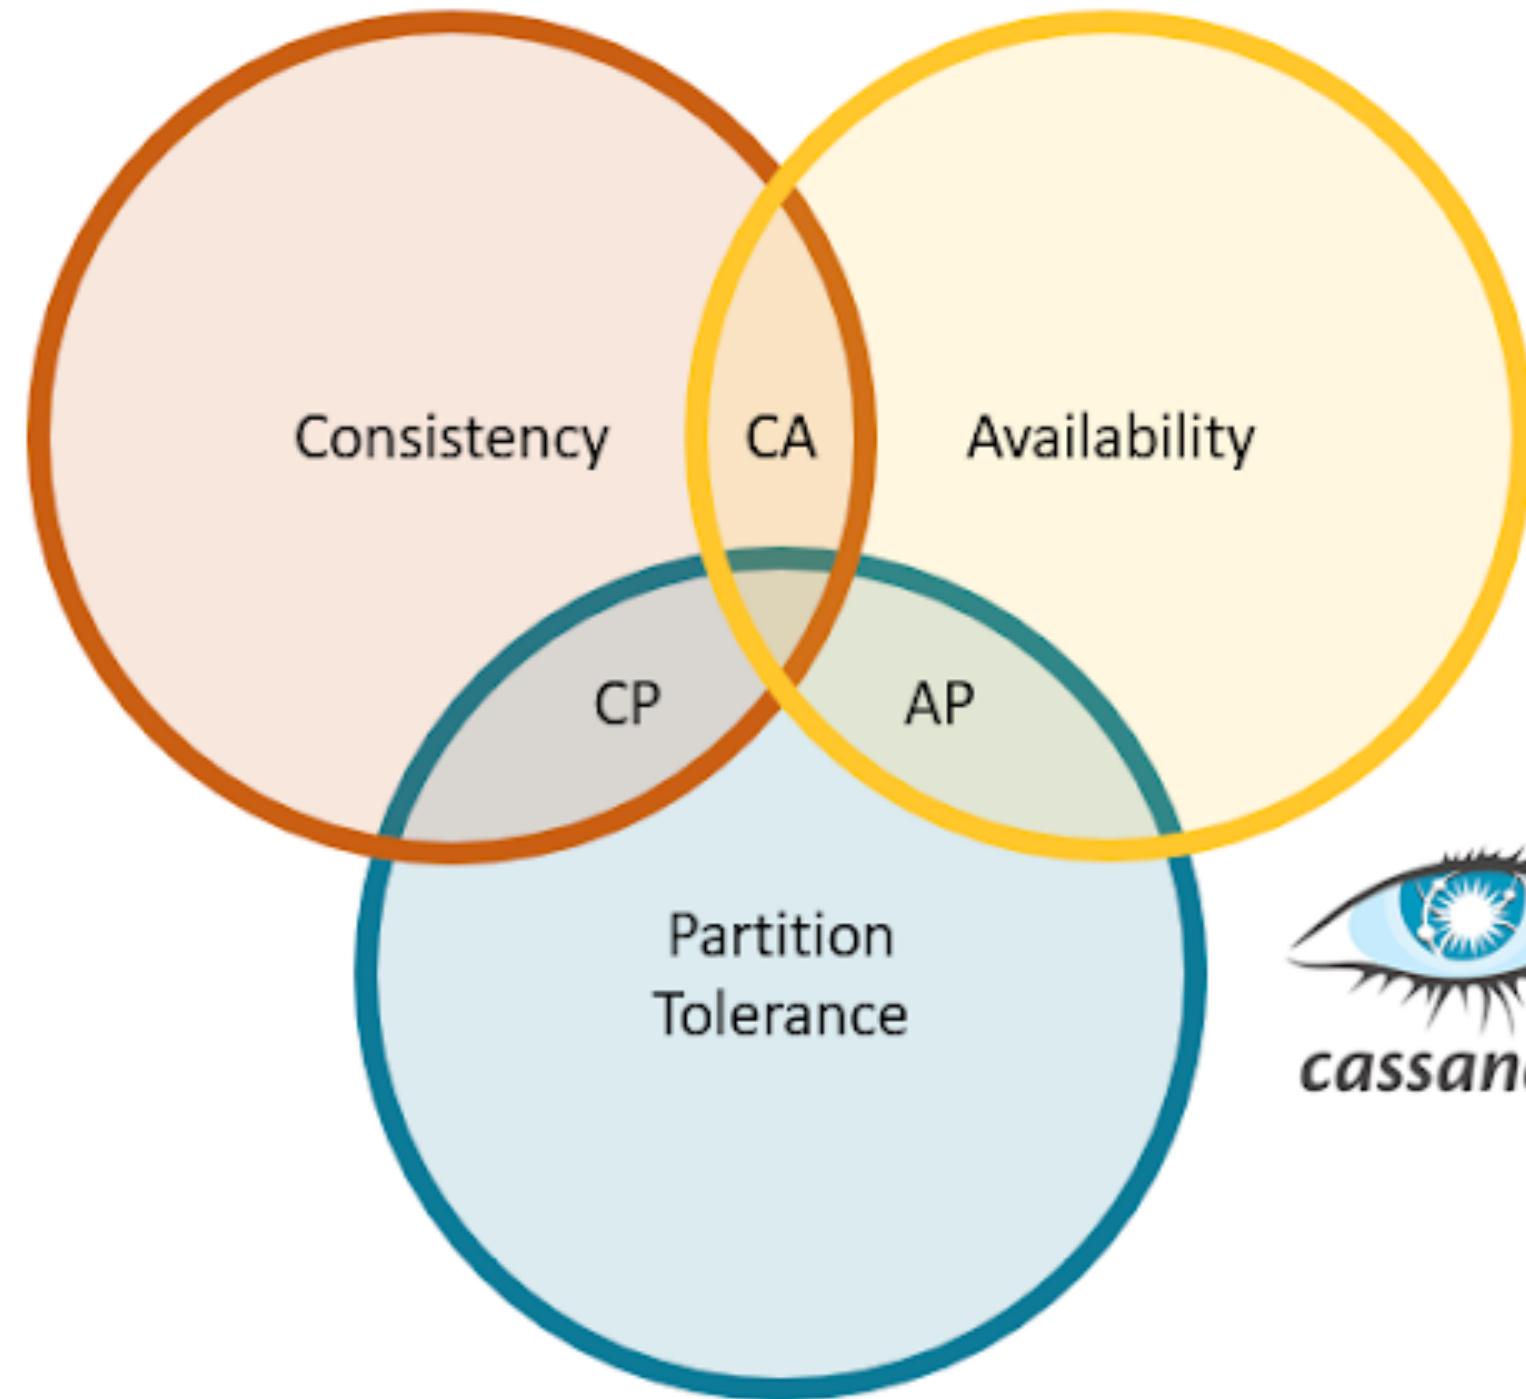
# What to remember

- Data Model: JSON
- Storage: BSON
- Operations: iterations
- Replication: Replica Sets
- Partitioning: Sharding
- Architecture
  - Mongod
  - Mongos

# What to remember

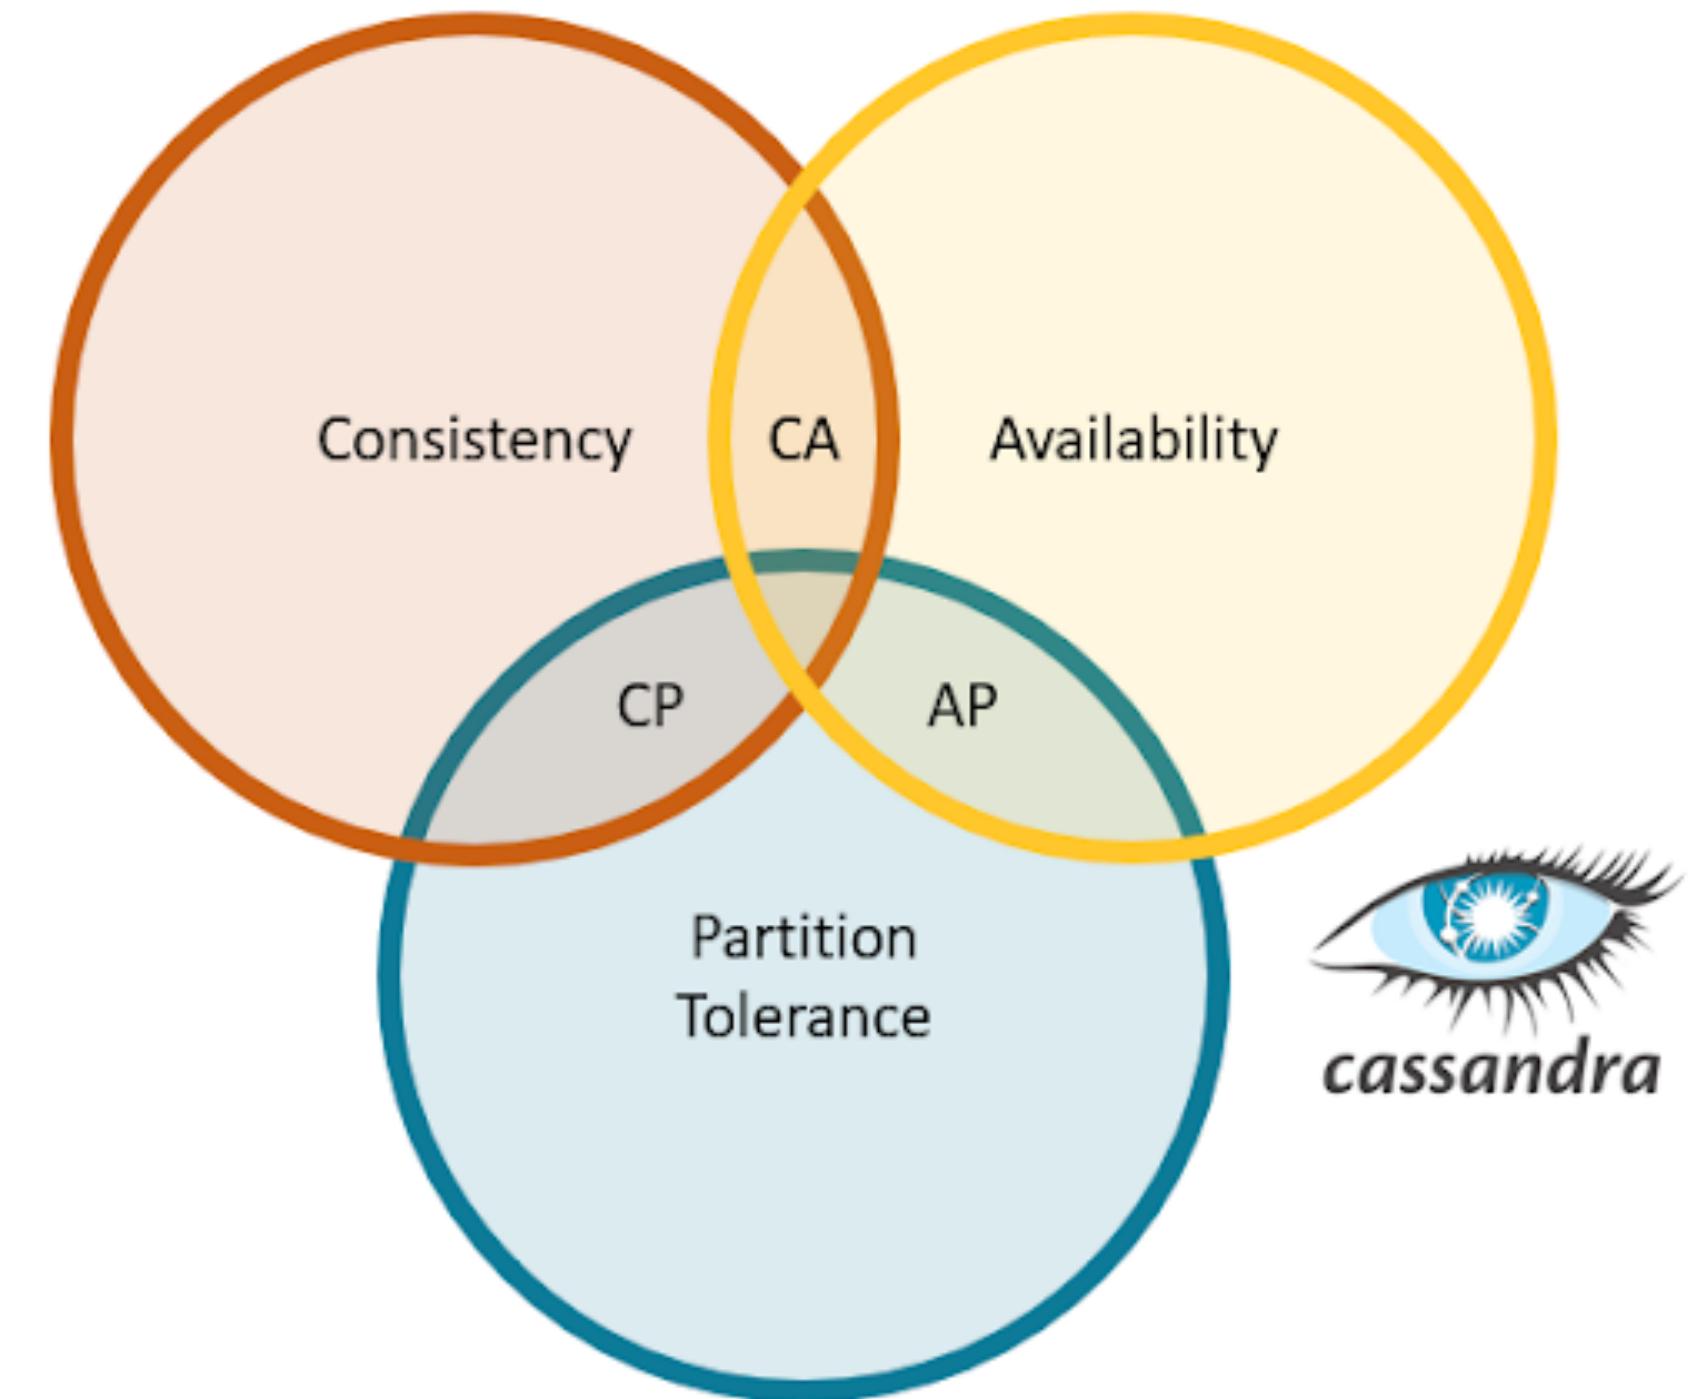
- Data Model: JSON
- Storage: BSON
- Operations: iterations
- Replication: Replica Sets
- Partitioning: Sharding
- Architecture
  - Mongod
  - Mongos
  - Config Server

# Cassandra



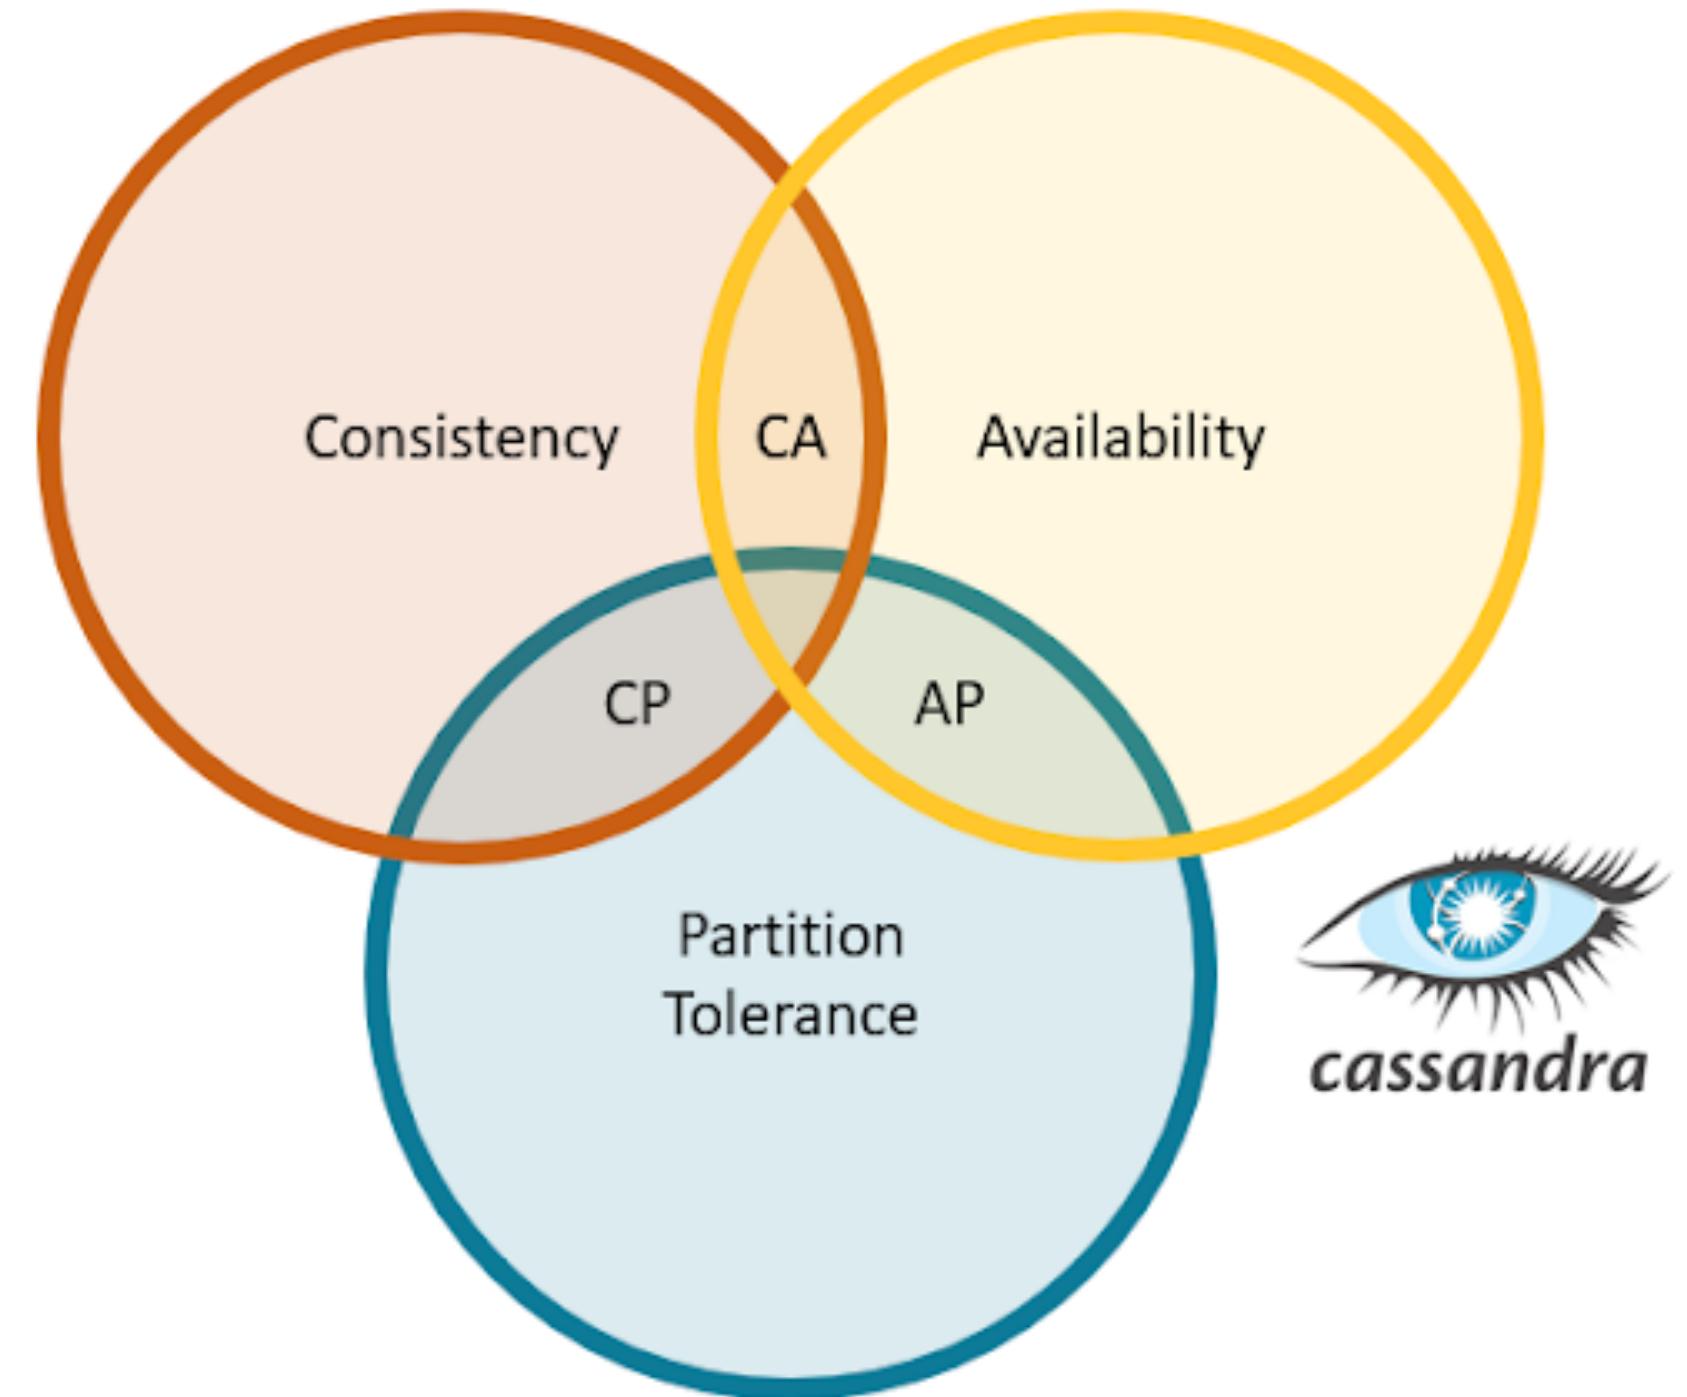
# Cassandra

- A Wide [[Column Oriented Database]]



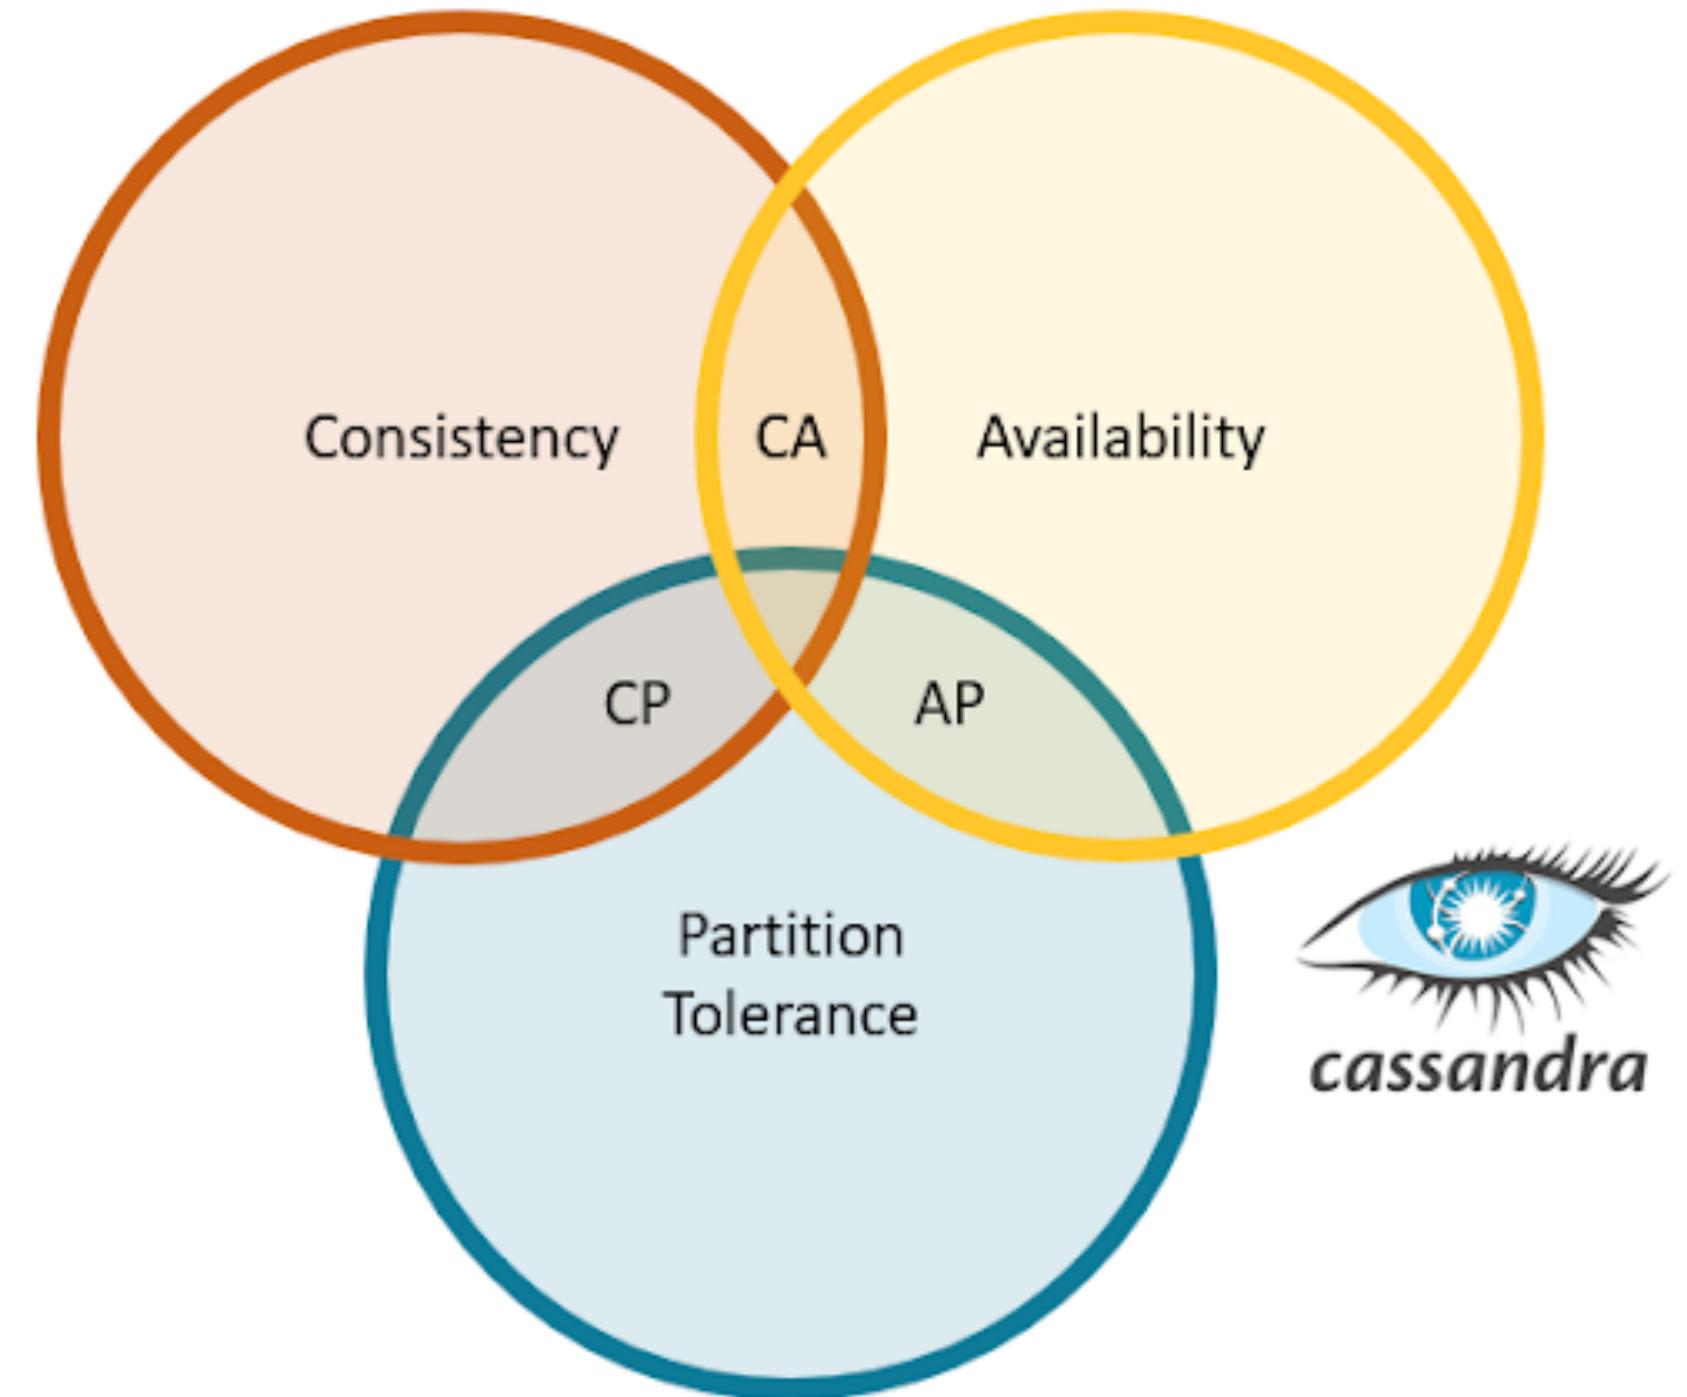
# Cassandra

- A Wide [[Column Oriented Database]]
- *tuneably consistent (C)*



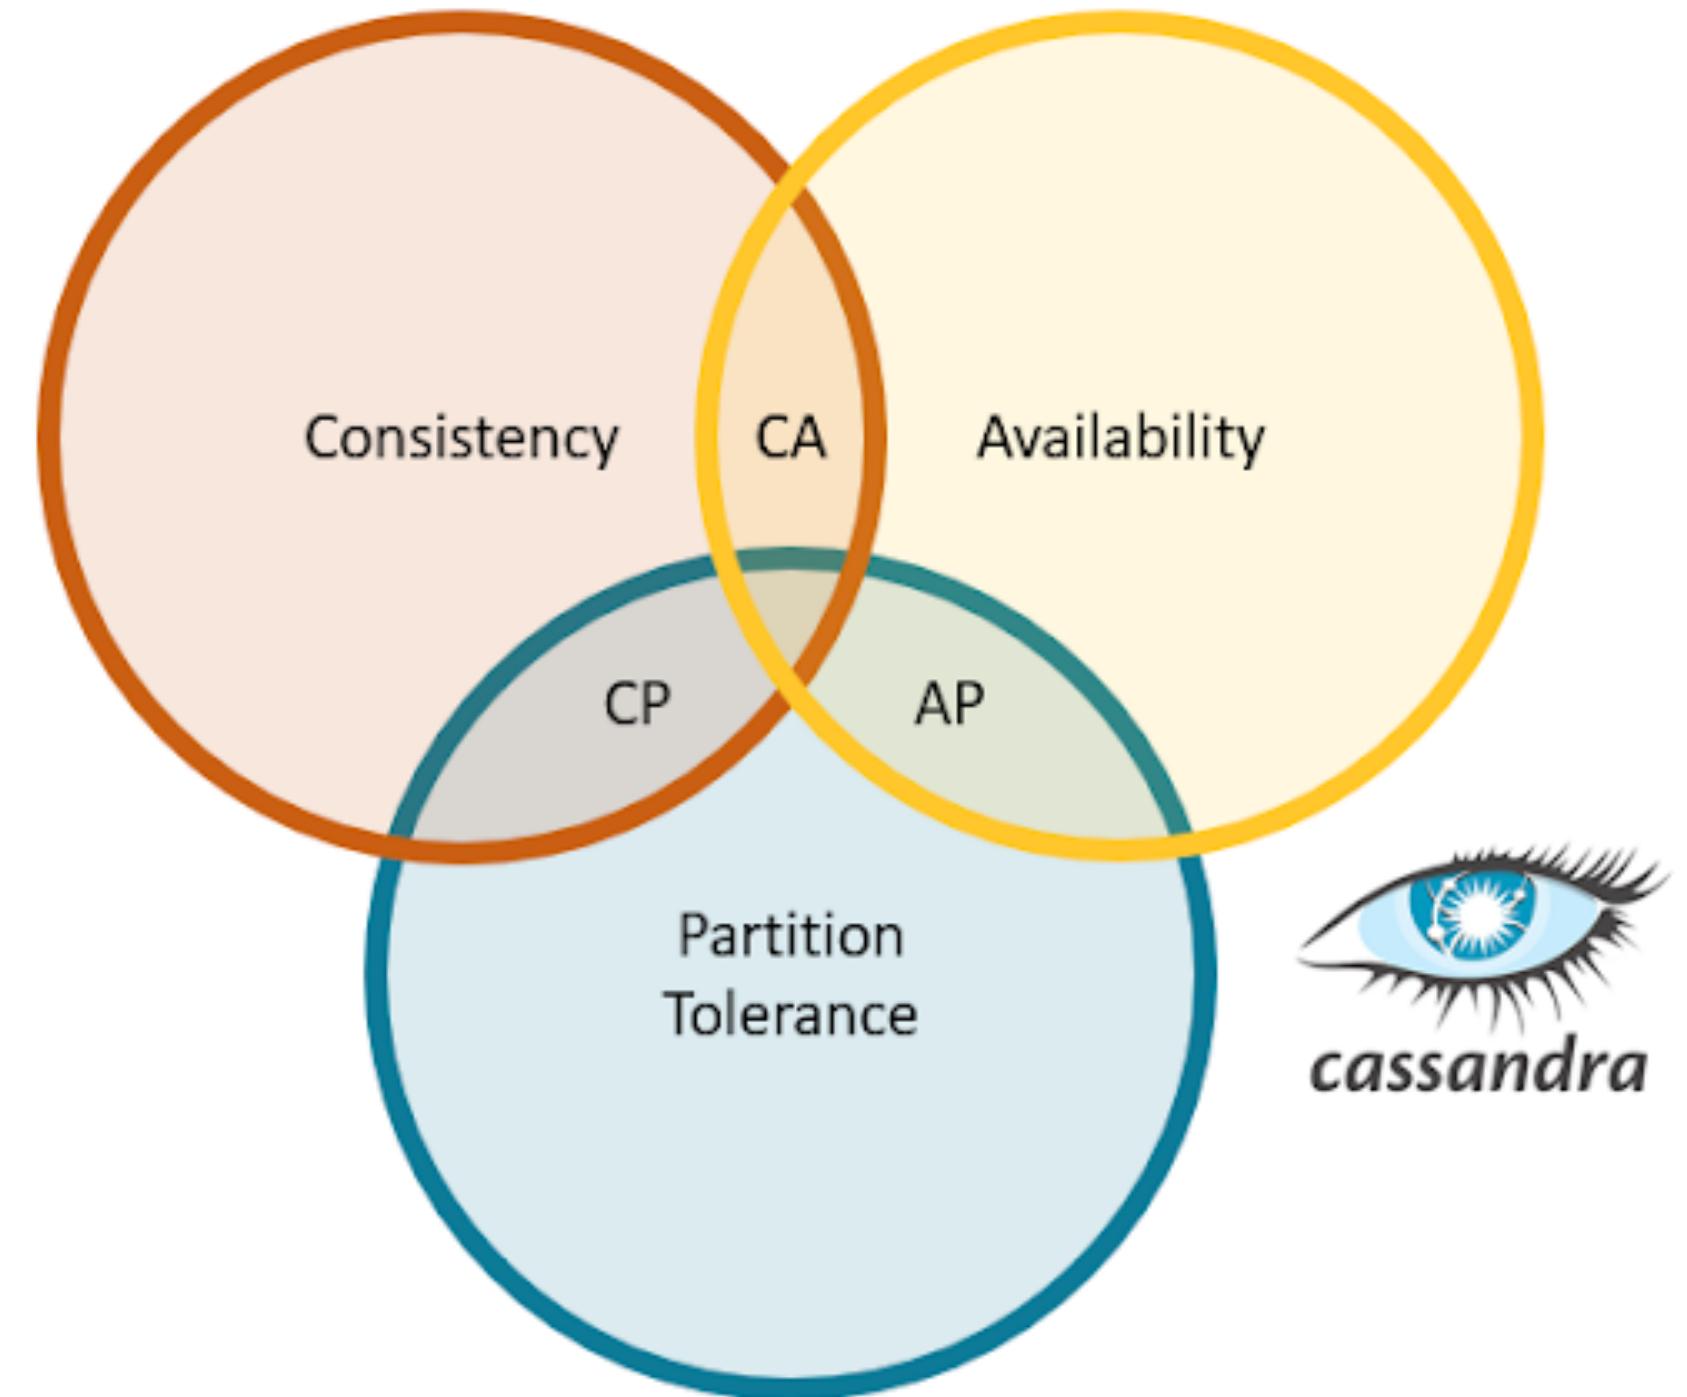
# Cassandra

- A Wide [[Column Oriented Database]]
- *tuneably* consistent ( $\epsilon$ )
- very fast in writes



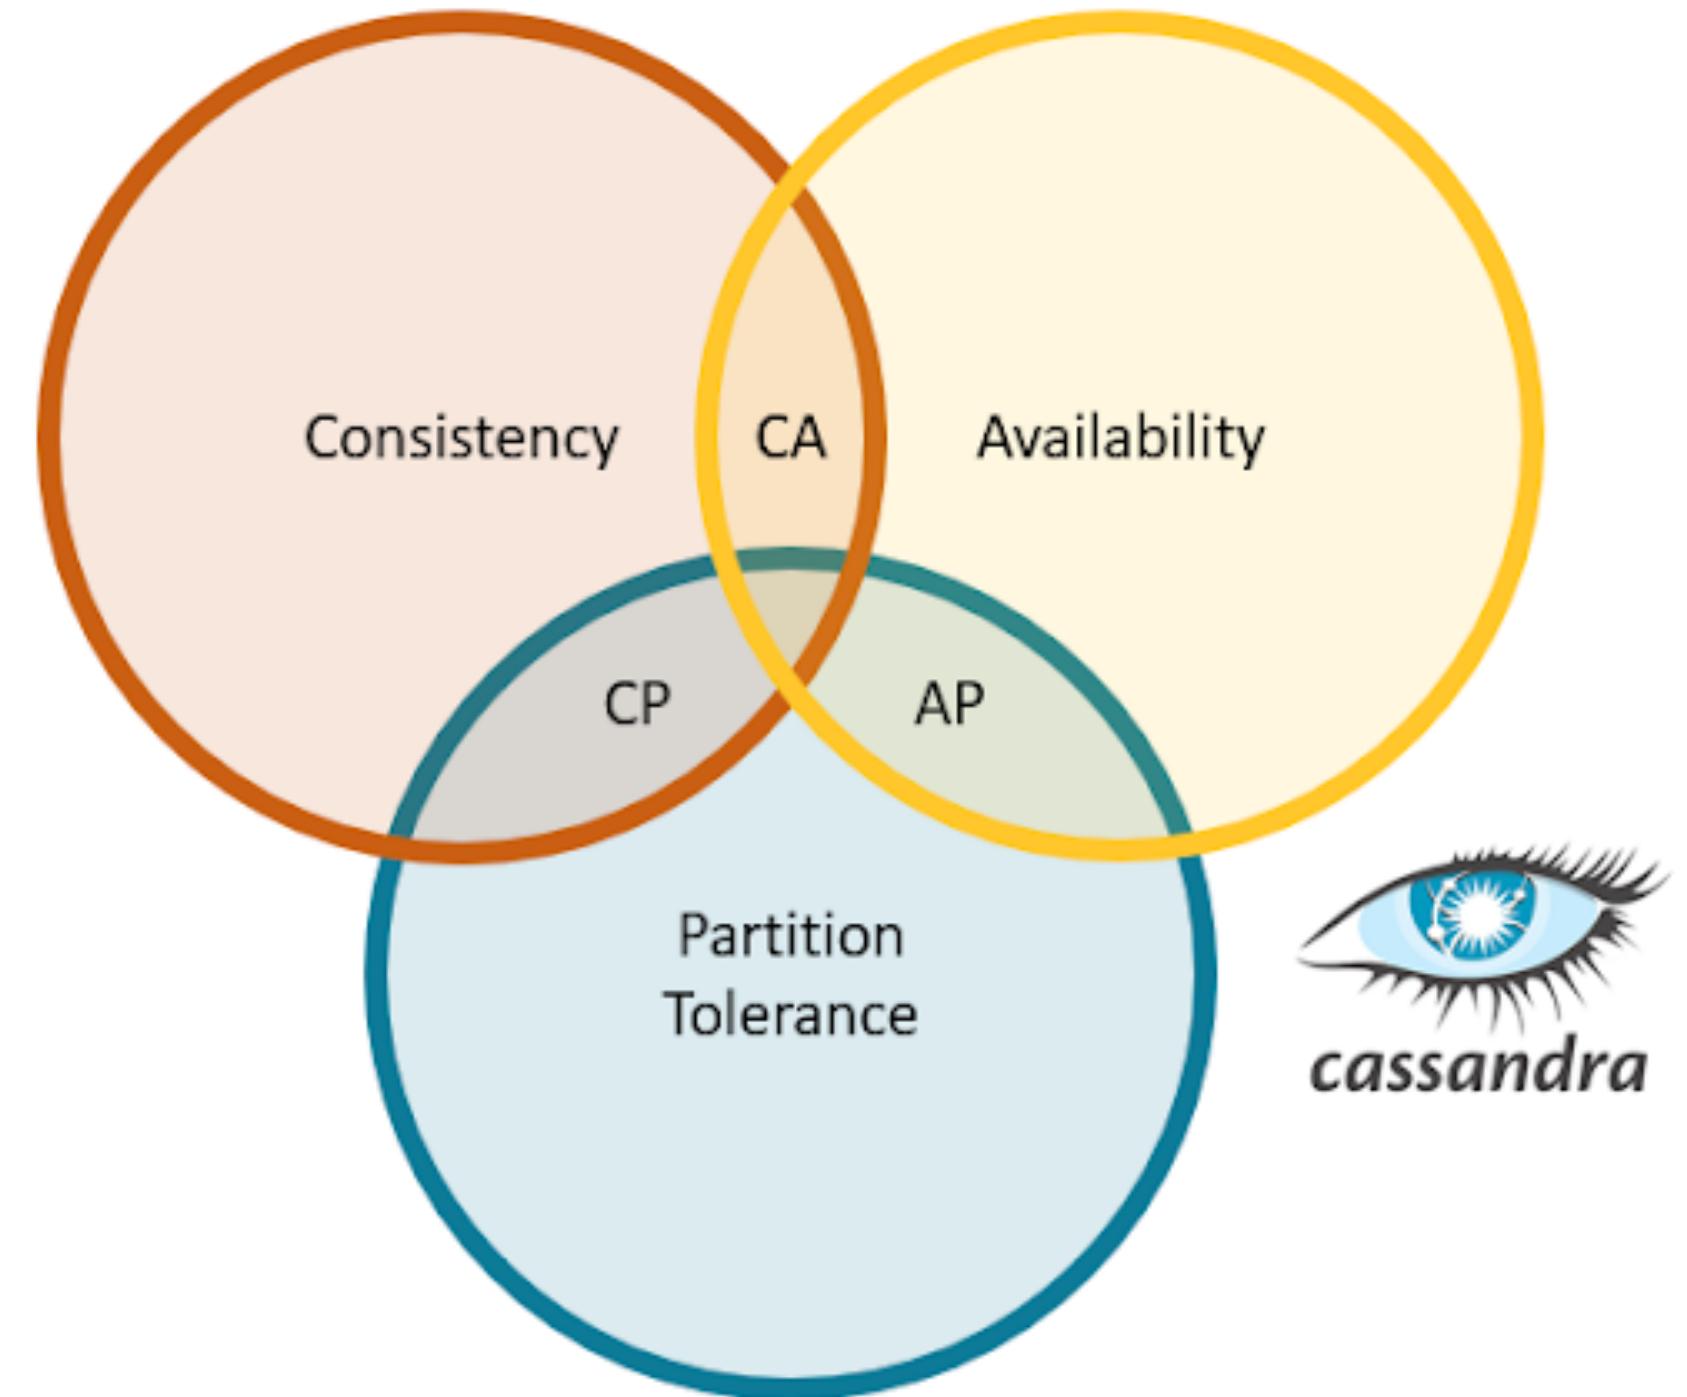
# Cassandra

- A Wide [[Column Oriented Database]]
- *tuneably* consistent ( $\epsilon$ )
- very fast in writes
- highly avaeng



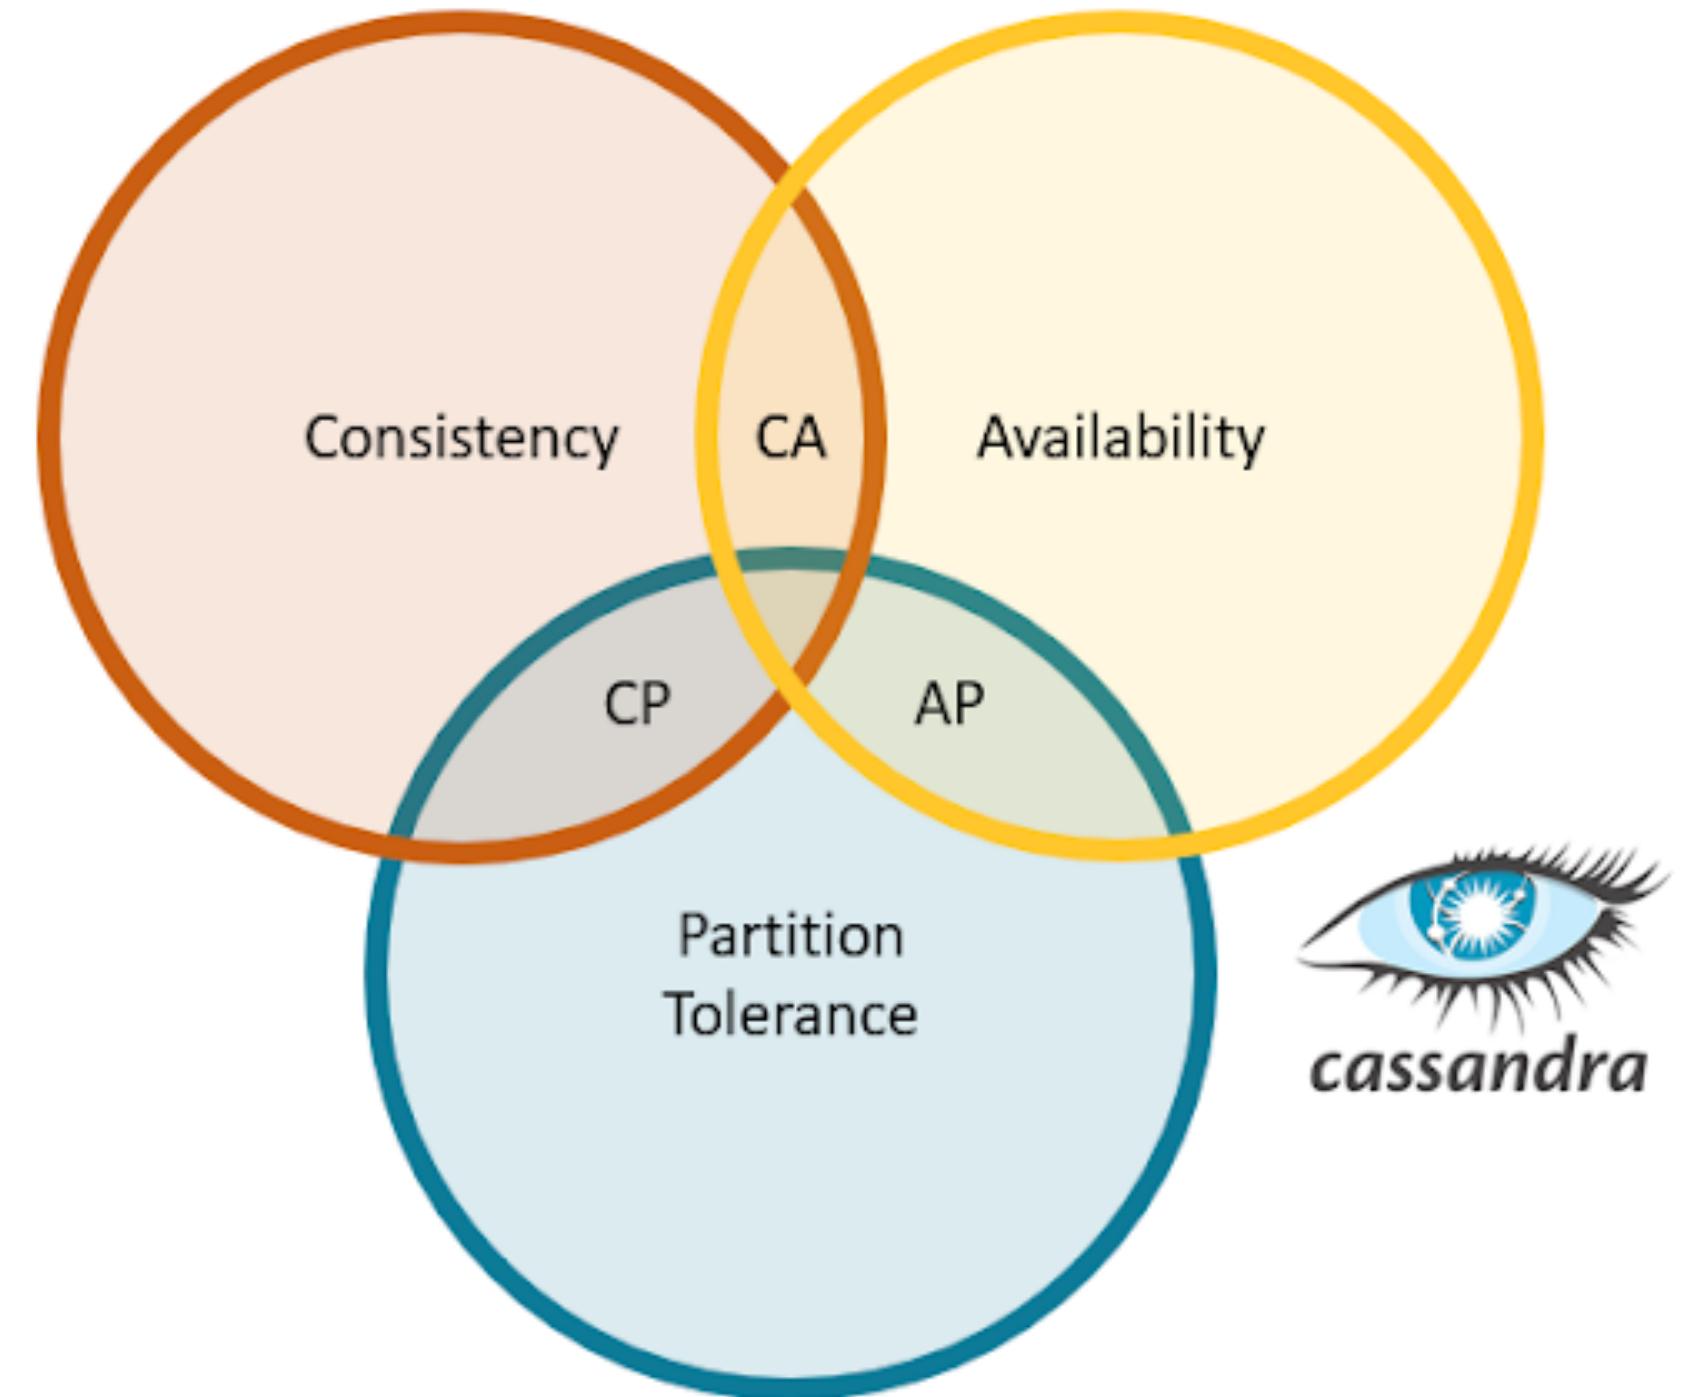
# Cassandra

- A Wide [[Column Oriented Database]]
- *tuneably* consistent ( $\epsilon$ )
- very fast in writes
- highly avaeng
- ailable (A)



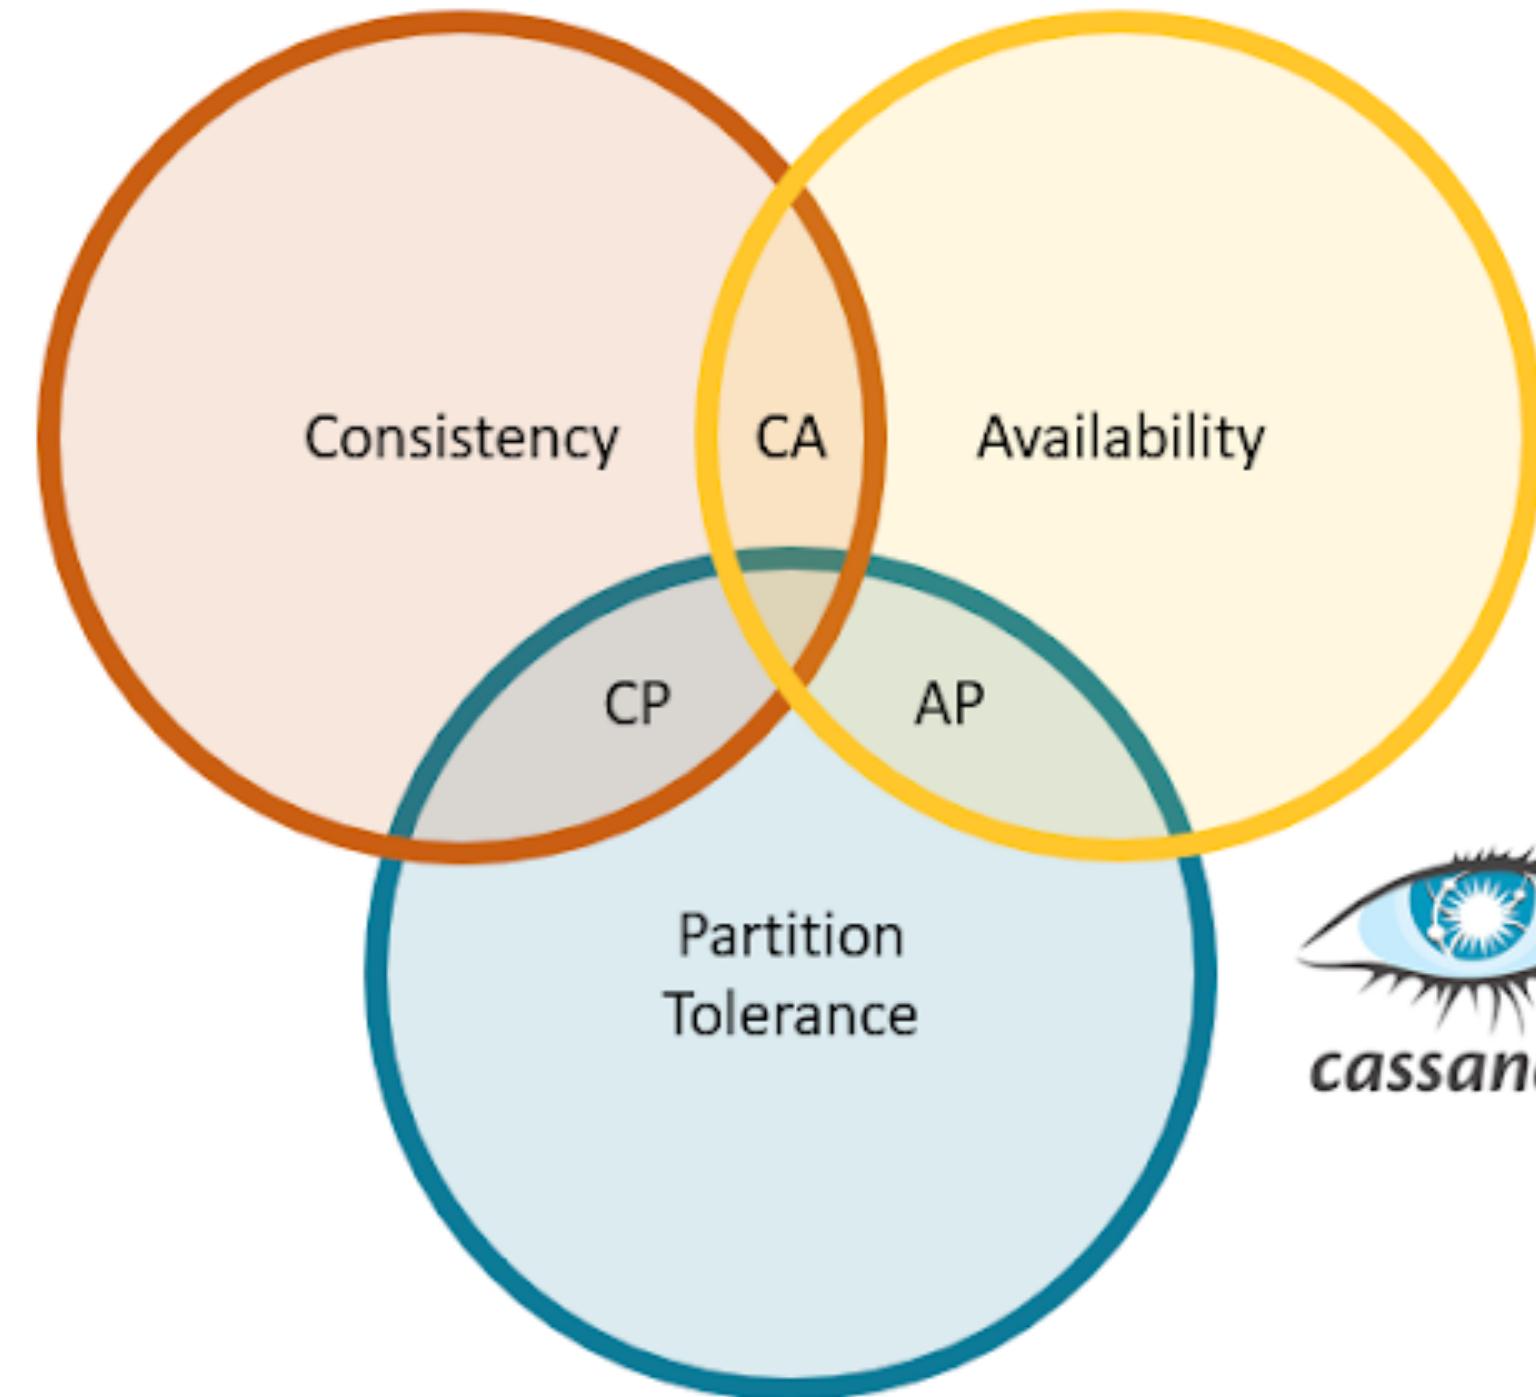
# Cassandra

- A Wide [[Column Oriented Database]]
- *tuneably* consistent (**C**)
- very fast in writes
- highly avaeng
- ailable (**A**)
- fault tolerant (**P**)



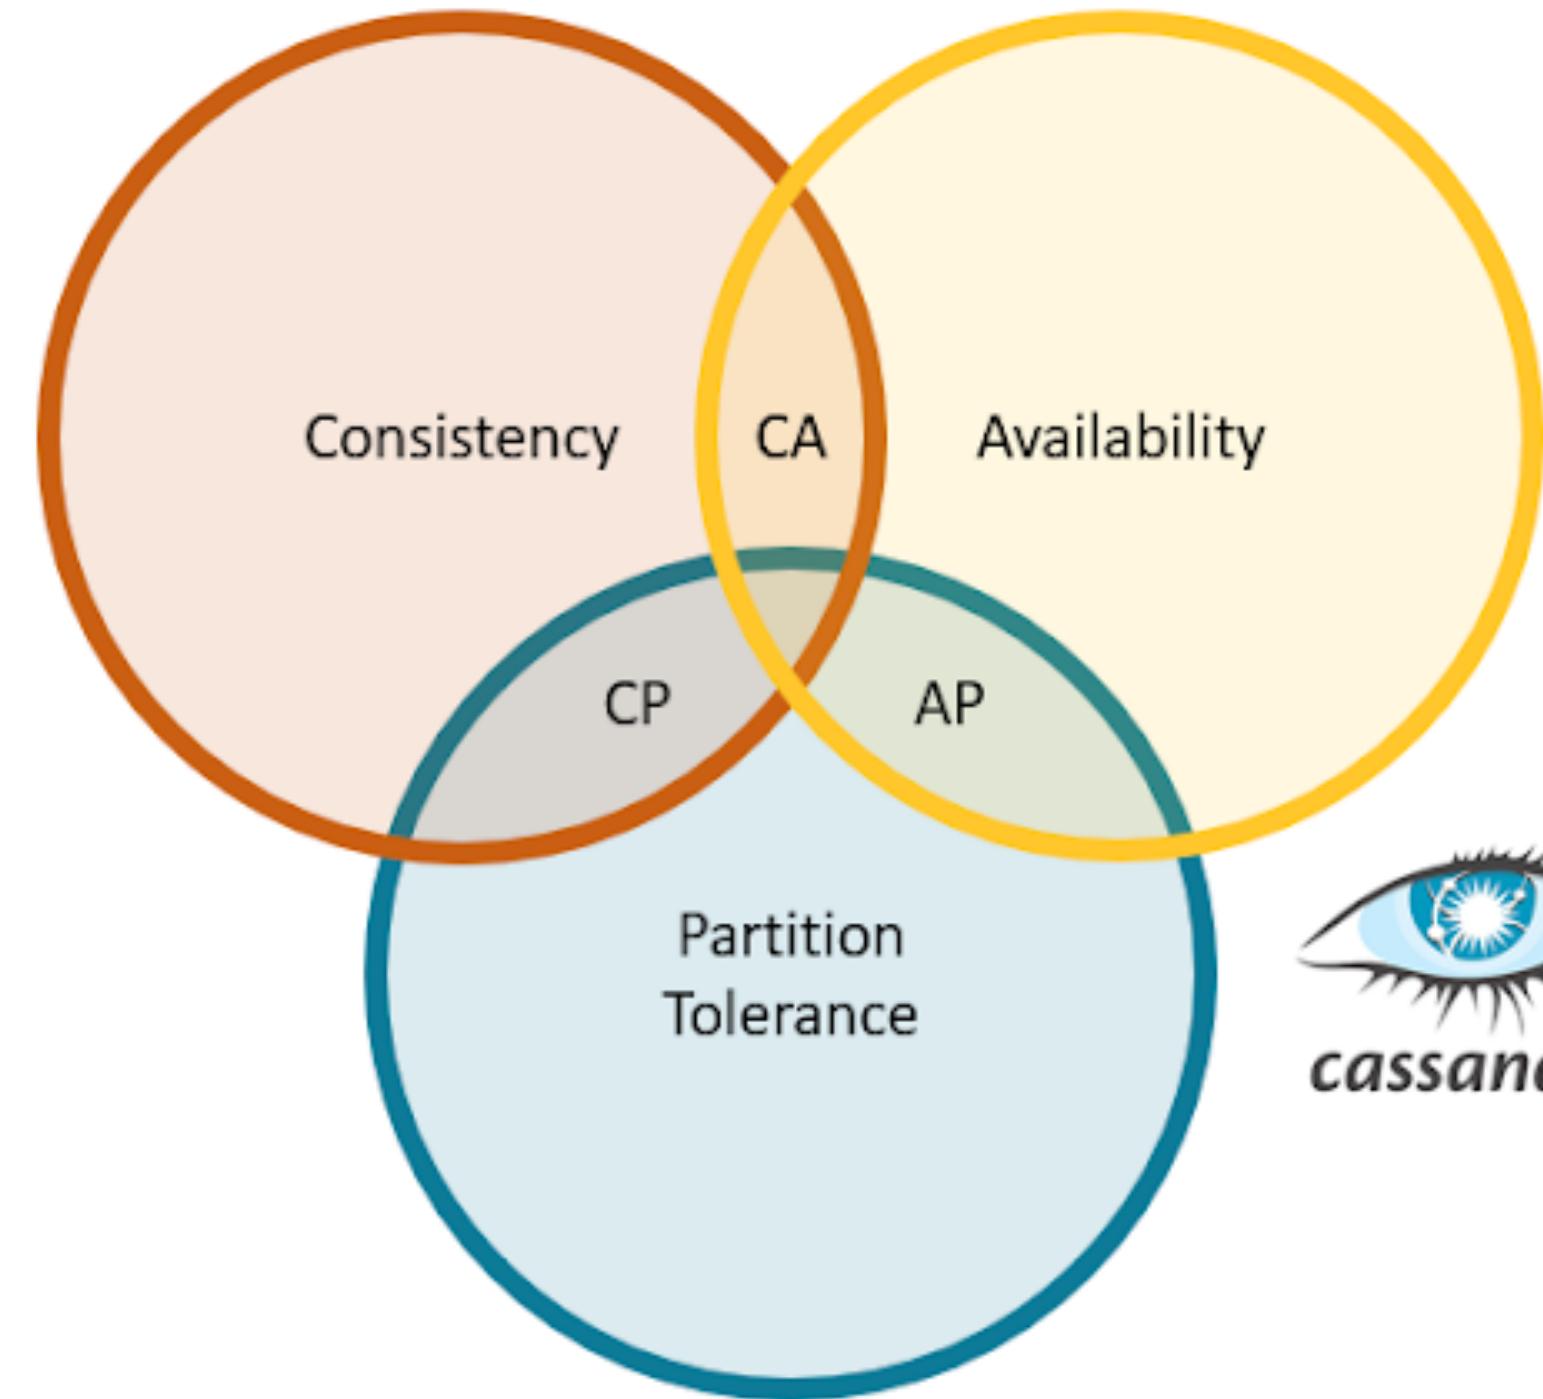
# Cassandra

- A Wide [[Column Oriented Database]]
- *tuneably* consistent (**C**)
- very fast in writes
- highly avaeng
- ailable (**A**)
- fault tolerant (**P**)
- linearly scalable, elastic scalability



## Cassandra

- A Wide [[Column Oriented Database]]
- *tuneably* consistent (**C**)
- very fast in writes
- highly avaeng
- ailable (**A**)
- fault tolerant (**P**)
- linearly scalable, elastic scalability
- Cassandra is very good at writes, okay with reads.



# What to remember

# What to remember

- Data Model: KeySpace, Rows, Column Families

# What to remember

- Data Model: KeySpace, Rows, Column Families
- Storage: Commit Log, SSTables, Bloom Filters

# What to remember

- Data Model: KeySpace, Rows, Column Families
- Storage: Commit Log, SSTables, Bloom Filters
- Operations: ~~SQL~~ (CQL!)

# What to remember

- Data Model: KeySpace, Rows, Column Families
- Storage: Commit Log, SSTables, Bloom Filters
- Operations: ~~SQL~~ (CQL!)
  - Different Consistency Levels for reads and writes

# What to remember

- Data Model: KeySpace, Rows, Column Families
- Storage: Commit Log, SSTables, Bloom Filters
- Operations: ~~SQL~~ (CQL!)
  - Different Consistency Levels for reads and writes
- Replication: Replication factor

# What to remember

- Data Model: KeySpace, Rows, Column Families
- Storage: Commit Log, SSTables, Bloom Filters
- Operations: ~~SQL~~ (CQL!)
  - Different Consistency Levels for reads and writes
- Replication: Replication factor
  - Simple vs Network Topology

# What to remember

- Data Model: KeySpace, Rows, Column Families
- Storage: Commit Log, SSTables, Bloom Filters
- Operations: ~~SQL~~ (CQL!)
  - Different Consistency Levels for reads and writes
- Replication: Replication factor
  - Simple vs Network Topology
- Partitioning: Order Preserving

# What to remember

- Data Model: KeySpace, Rows, Column Families
- Storage: Commit Log, SSTables, Bloom Filters
- Operations: ~~SQL~~ (CQL!)
  - Different Consistency Levels for reads and writes
- Replication: Replication factor
  - Simple vs Network Topology
- Partitioning: Order Preserving
- Architecture: Gossip Protocol

# GraphDBs (Neo4J)

# GraphDBs (Neo4J)

- Back to Centralized Architecture

# GraphDBs (Neo4J)

- Back to Centralized Architecture
  - Relationship First

# GraphDBs (Neo4J)

- Back to Centralized Architecture
  - Relationship First
  - Storage

# GraphDBs (Neo4J)

- Back to Centralized Architecture
  - Relationship First
  - Storage
    - Native vs Non Native

# GraphDBs (Neo4J)

- Back to Centralized Architecture
  - Relationship First
  - Storage
    - Native vs Non Native
  - Graph Pattern Matching

# GraphDBs (Neo4J)

- Back to Centralized Architecture
  - Relationship First
  - Storage
    - Native vs Non Native
  - Graph Pattern Matching
    - Union

# GraphDBs (Neo4J)

- Back to Centralized Architecture
  - Relationship First
  - Storage
    - Native vs Non Native
  - Graph Pattern Matching
    - Union
    - Projection

# GraphDBs (Neo4J)

- Back to Centralized Architecture
  - Relationship First
  - Storage
    - Native vs Non Native
  - Graph Pattern Matching
    - Union
    - Projection
    - Different

# GraphDBs (Neo4J)

- Back to Centralized Architecture
  - Relationship First
  - Storage
    - Native vs Non Native
  - Graph Pattern Matching
    - Union
    - Projection
    - Different
    - Optional

# GraphDBs (Neo4J)

- Back to Centralized Architecture
  - Relationship First
  - Storage
    - Native vs Non Native
  - Graph Pattern Matching
    - Union
    - Projection
    - Different
    - Optional
    - Filter

# GraphDBs (Neo4J)

- Back to Centralized Architecture
  - Relationship First
  - Storage
    - Native vs Non Native
  - Graph Pattern Matching
    - Union
    - Projection
    - Different
    - Optional
    - Filter
  - Navigational Queries (Paths)

# Questions Template



- What are the three levels of data modeling

- What are the three levels of data modeling
  - A) Conceptual, Logical, Physical

- What are the three levels of data modeling
  - A) Conceptual, Logical, Physical
  - B) Perceptual, Latent, Physical

- What are the three levels of data modeling
  - A) Conceptual, Logical, Physical
  - B) Perceptual, Latent, Physical
  - C) Coherent, Logical, Phenomenal



- What are the three levels of data modeling

- What are the three levels of data modeling
  - A) **Conceptual, Logical, Physical**

- What are the three levels of data modeling
  - A) **Conceptual, Logical, Physical**
  - B) Perceptual, Latent, Physical

- What are the three levels of data modeling
  - A) **Conceptual, Logical, Physical**
  - B) Perceptual, Latent, Physical
  - C) Coherent, Logical, Phenomenal



- What is a fact table

- What is a fact table
  - A) a table that contains the descriptive attributes used by BI applications for filtering and grouping the facts

- What is a fact table
  - A) a table that contains the descriptive attributes used by BI applications for filtering and grouping the facts
  - B) a table with a one-to-one relationship to a measurement event as described by the fact table's grain.

- What is a fact table
  - A) a table that contains the descriptive attributes used by BI applications for filtering and grouping the facts
  - B) a table with a one-to-one relationship to a measurement event as described by the fact table's grain.
  - C) a table that contains the numeric measures produced by an operational measurement event in the real world.



- What is a fact table

- What is a fact table
  - A) a table that contains the descriptive attributes used by BI applications for filtering and grouping the facts

- What is a fact table
  - A) a table that contains the descriptive attributes used by BI applications for filtering and grouping the facts
  - B) a table with a one-to-one relationship to a measurement event as described by the fact table's grain.

- What is a fact table
  - A) a table that contains the descriptive attributes used by BI applications for filtering and grouping the facts
  - B) a table with a one-to-one relationship to a measurement event as described by the fact table's grain.
  - C) a table that contains the numeric measures produced by an operational measurement event in the real world.

What of the following are the valid ADT to represent graphs

- A) Edge List
- B) Adjacency List
- C) Adjacency Incident
- D) Edge Matrix

What of the following are the motivations behind the NoSQL movement

- need for greater scalability than relational databases can't offer
- need for specialized query operations
- need for a more dynamic and expressive data model
- the Object-Relational Mismatch

What of the following are the motivations behind the NoSQL movement

- need for greater scalability than relational databases can't offer
- need for specialized query operations
- need for a more dynamic and expressive data model \*\*
- \*\*the Object-Relational Mismatch

What of the following are the valid ADT to represent graphs

- A) **Edge List**
- B) **Adjacency List**
- C) Adjacency Incident
- D) Edge Matrix