

# Data Engineering

LTAT.02.007

Ass Prof. Riccardo Tommasini

Assistants: **Fabiano Spiga, Mohamed Ragab, Hassan Eldeeb**



[https://courses.cs.ut.ee/2020/  
dataeng](https://courses.cs.ut.ee/2020/dataeng)

Forum

Moodle



# Column Oriented Database

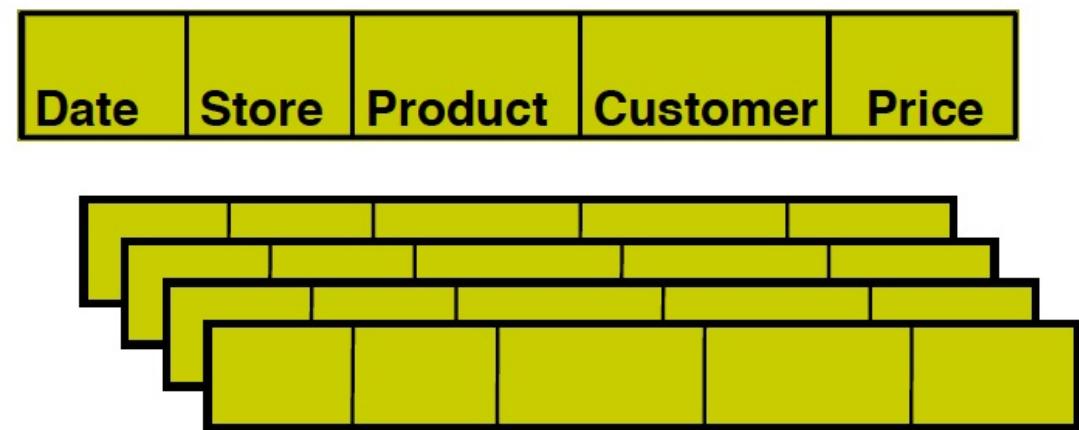
The approach to store and process data by column instead of row has its origin in analytics and business intelligence

Column-stores operating in a **shared-nothing** massively parallel processing architecture can be used to build high-performance applications.

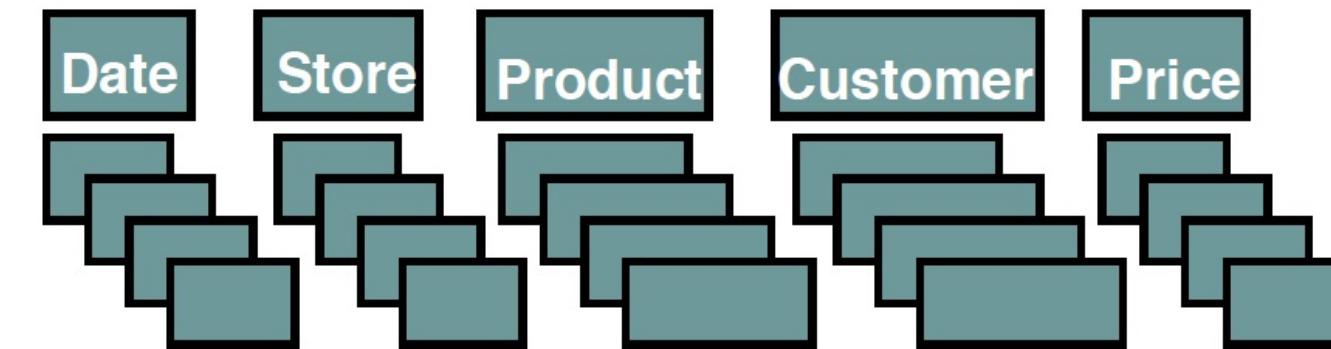
The class of column-oriented stores, which sees in Google's BigTable it's first member, is seen less puristic, also subsuming datastores that integrate column- and row-orientation.

# Column storage

**row-store**



**column-store**



+ easy to add/modify a record

- might read in unnecessary data

+ only need to read in relevant data

- tuple writes require multiple accesses

=> *suitable for read-mostly, read-intensive, large data repositories*

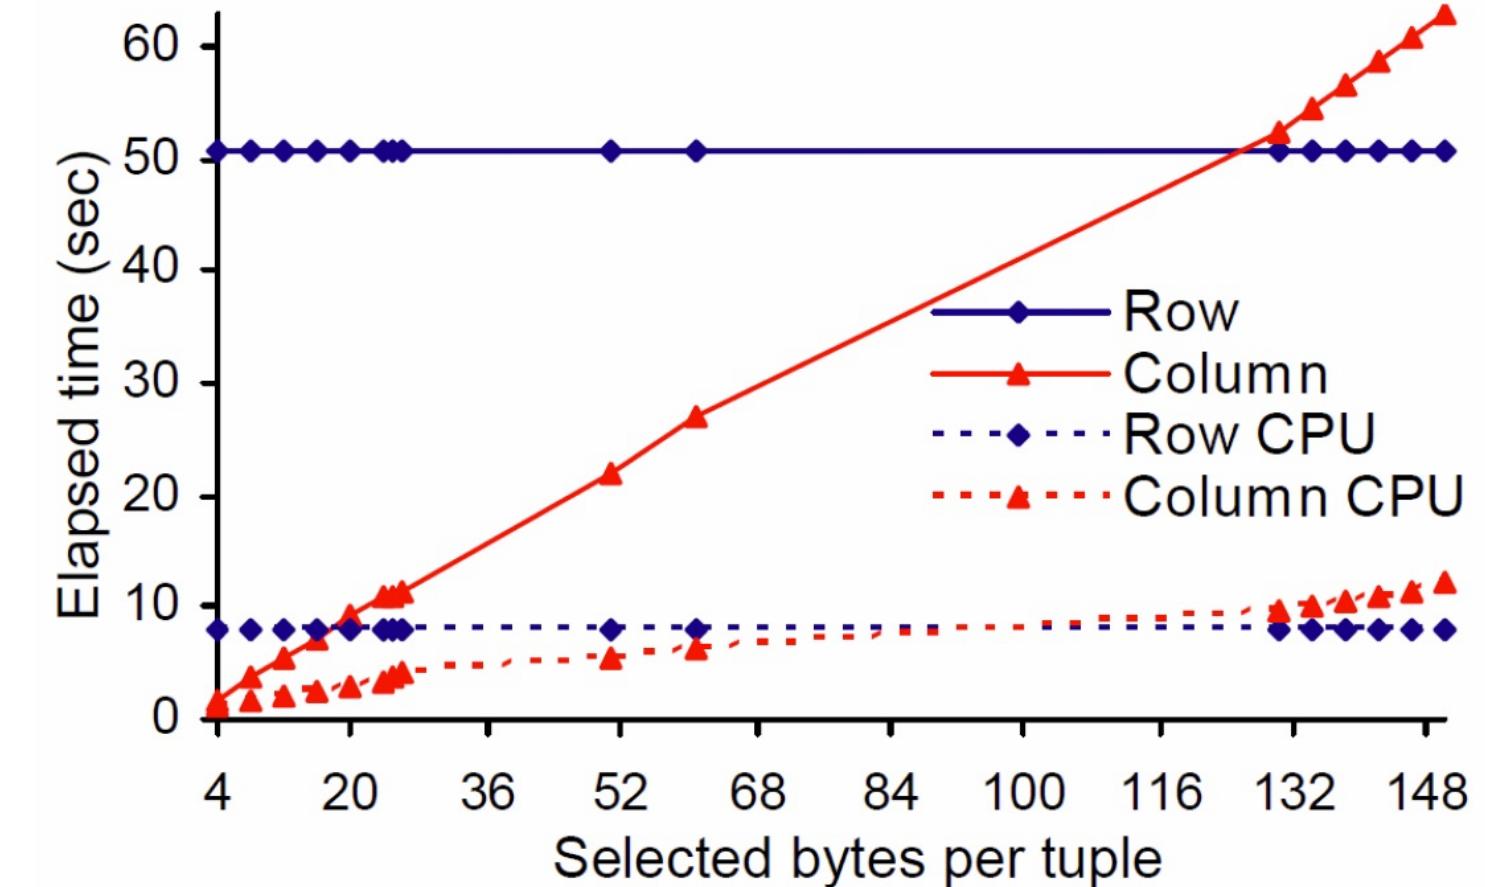
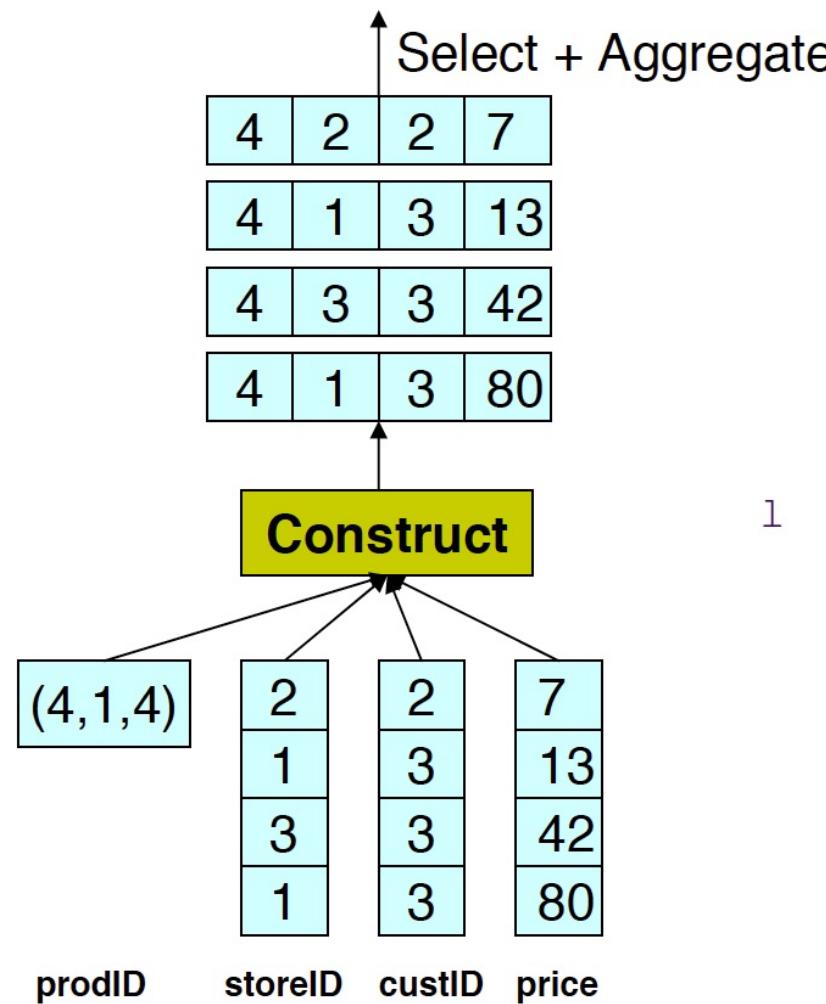
## Pros and Cons

- Data compression
- Improved Bandwidth Utilization
- Improved Code Pipelining
- Improved cache locality
- Increased Disk Seek<sup>70</sup> Time
- Increased cost of Inserts
- Requires disk prefetching
- Adds tuple reconstruction costs

---

<sup>70</sup> Seek time is the time taken for a hard disk controller to locate a specific piece of stored data

# Tuple Reconstruction



source

# Compression

- Increased column-store opportunities
  - Higher data value locality in column stores
  - Can use extra space to store multiple copies of data in different sort orders
  - Techniques such as run length encoding far more useful

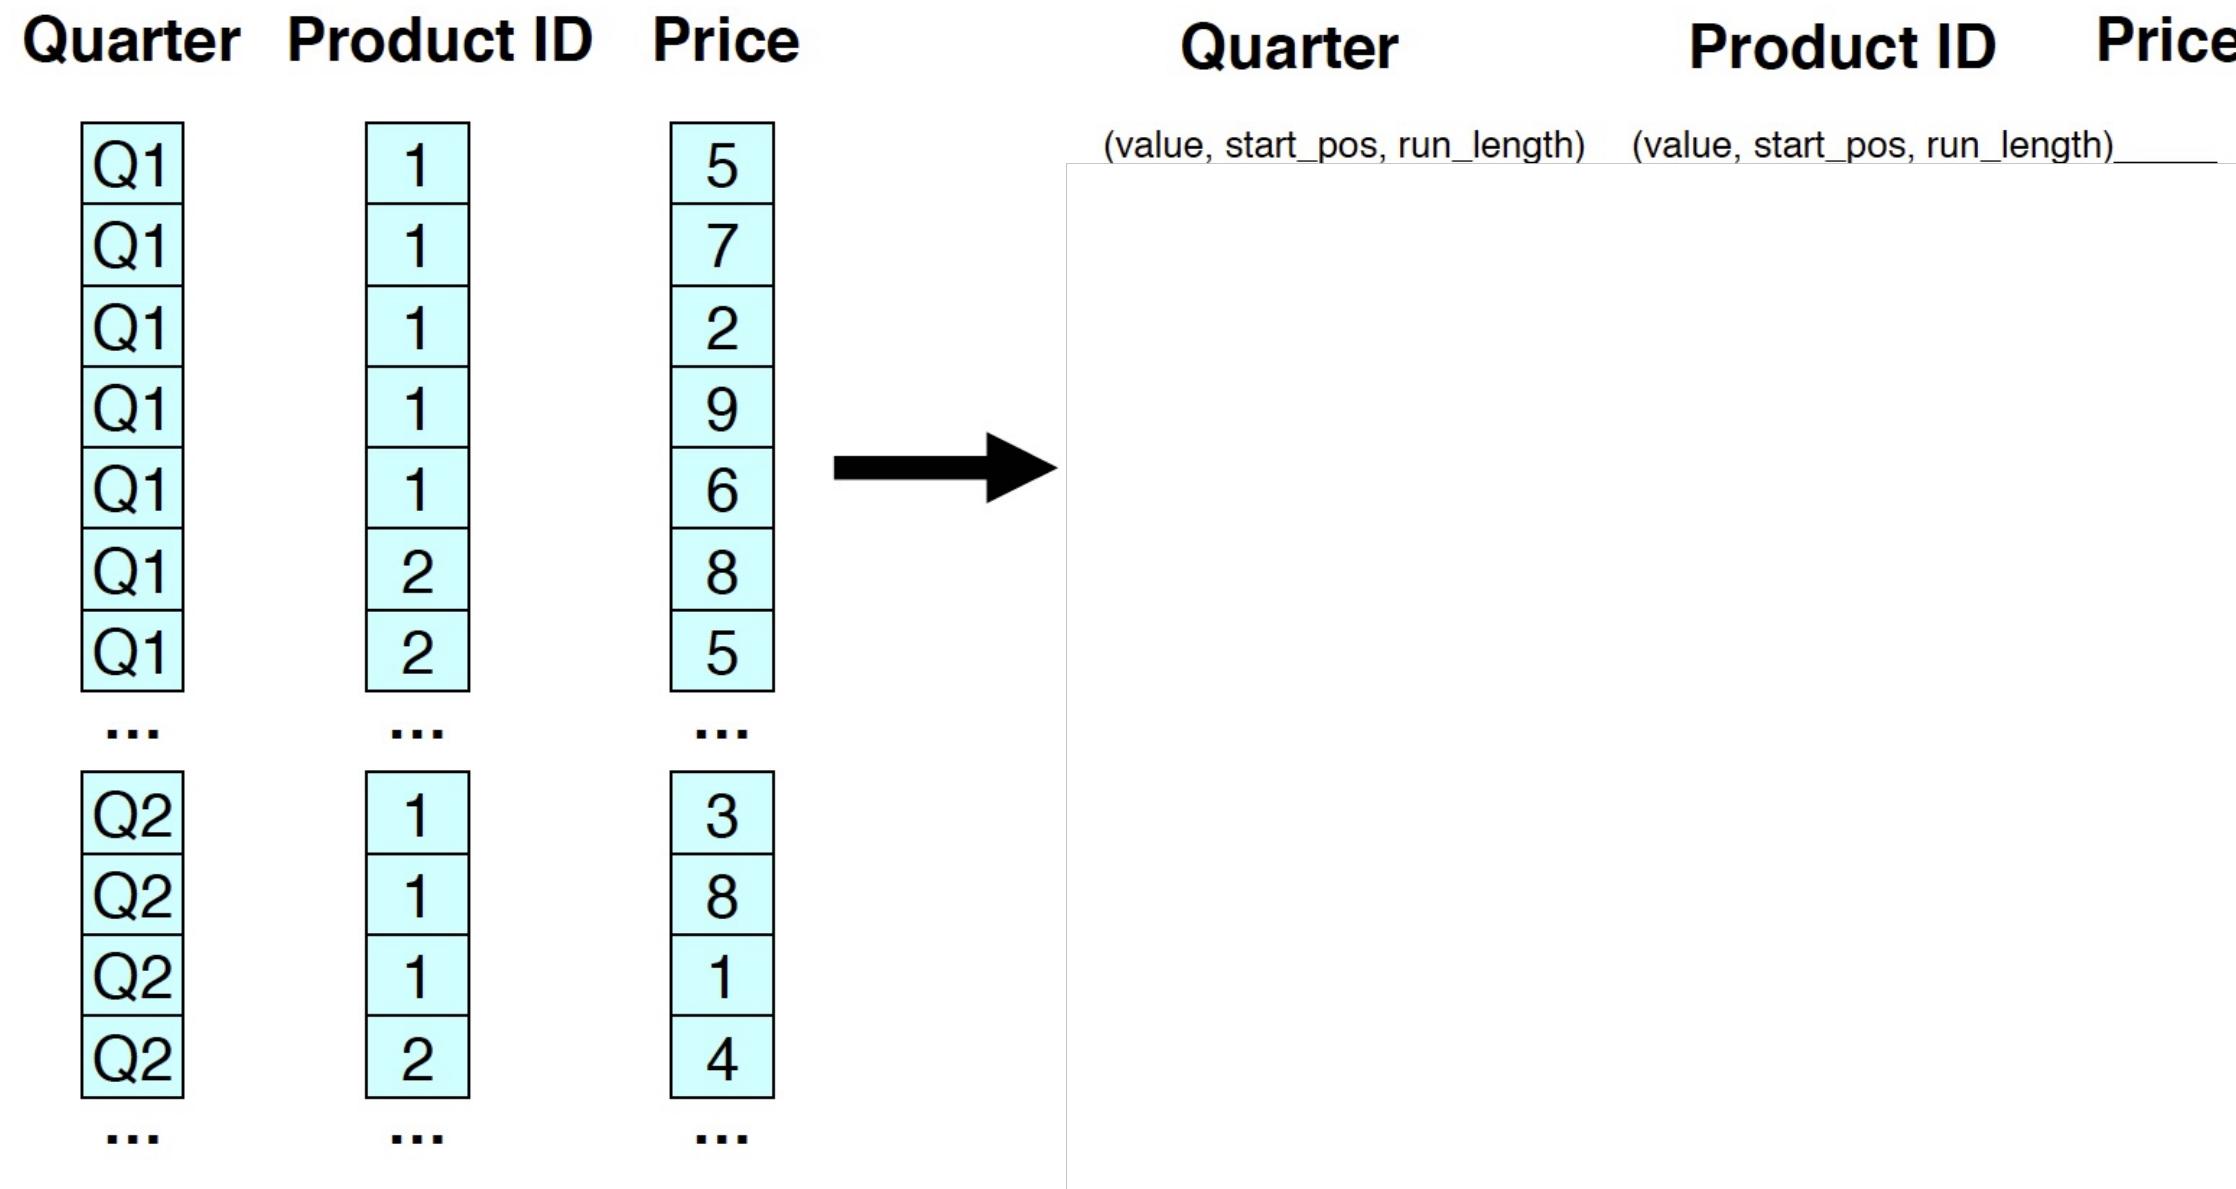
**Extra Available**  **Paper Summary**

# Example (String): Run-Length Encoding

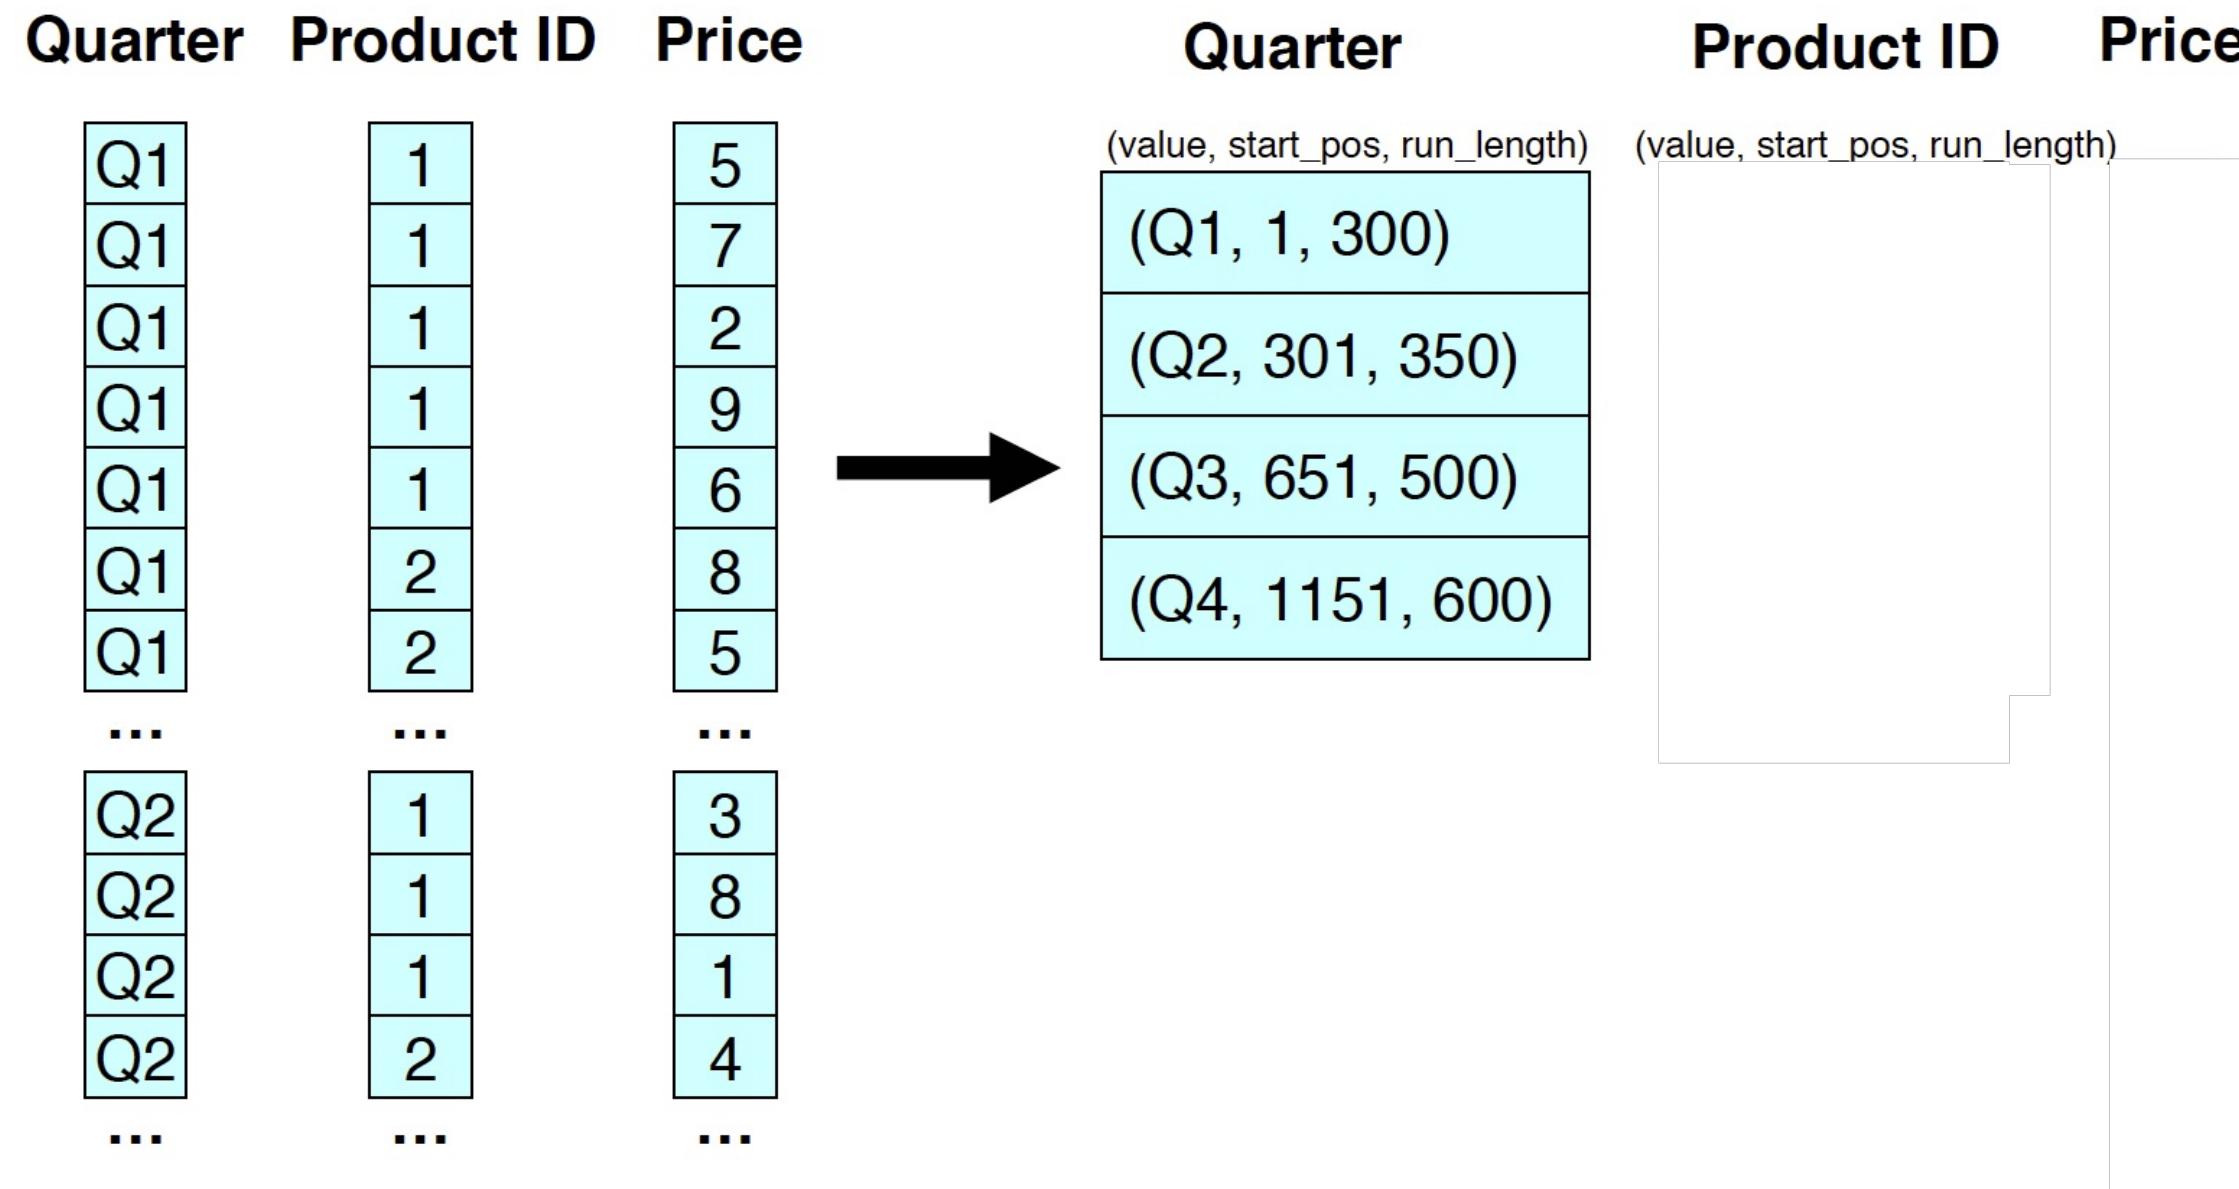
aaaabbbbbbbbbbccccccddddeeeeeedddd

4a13b7c7d5e5f

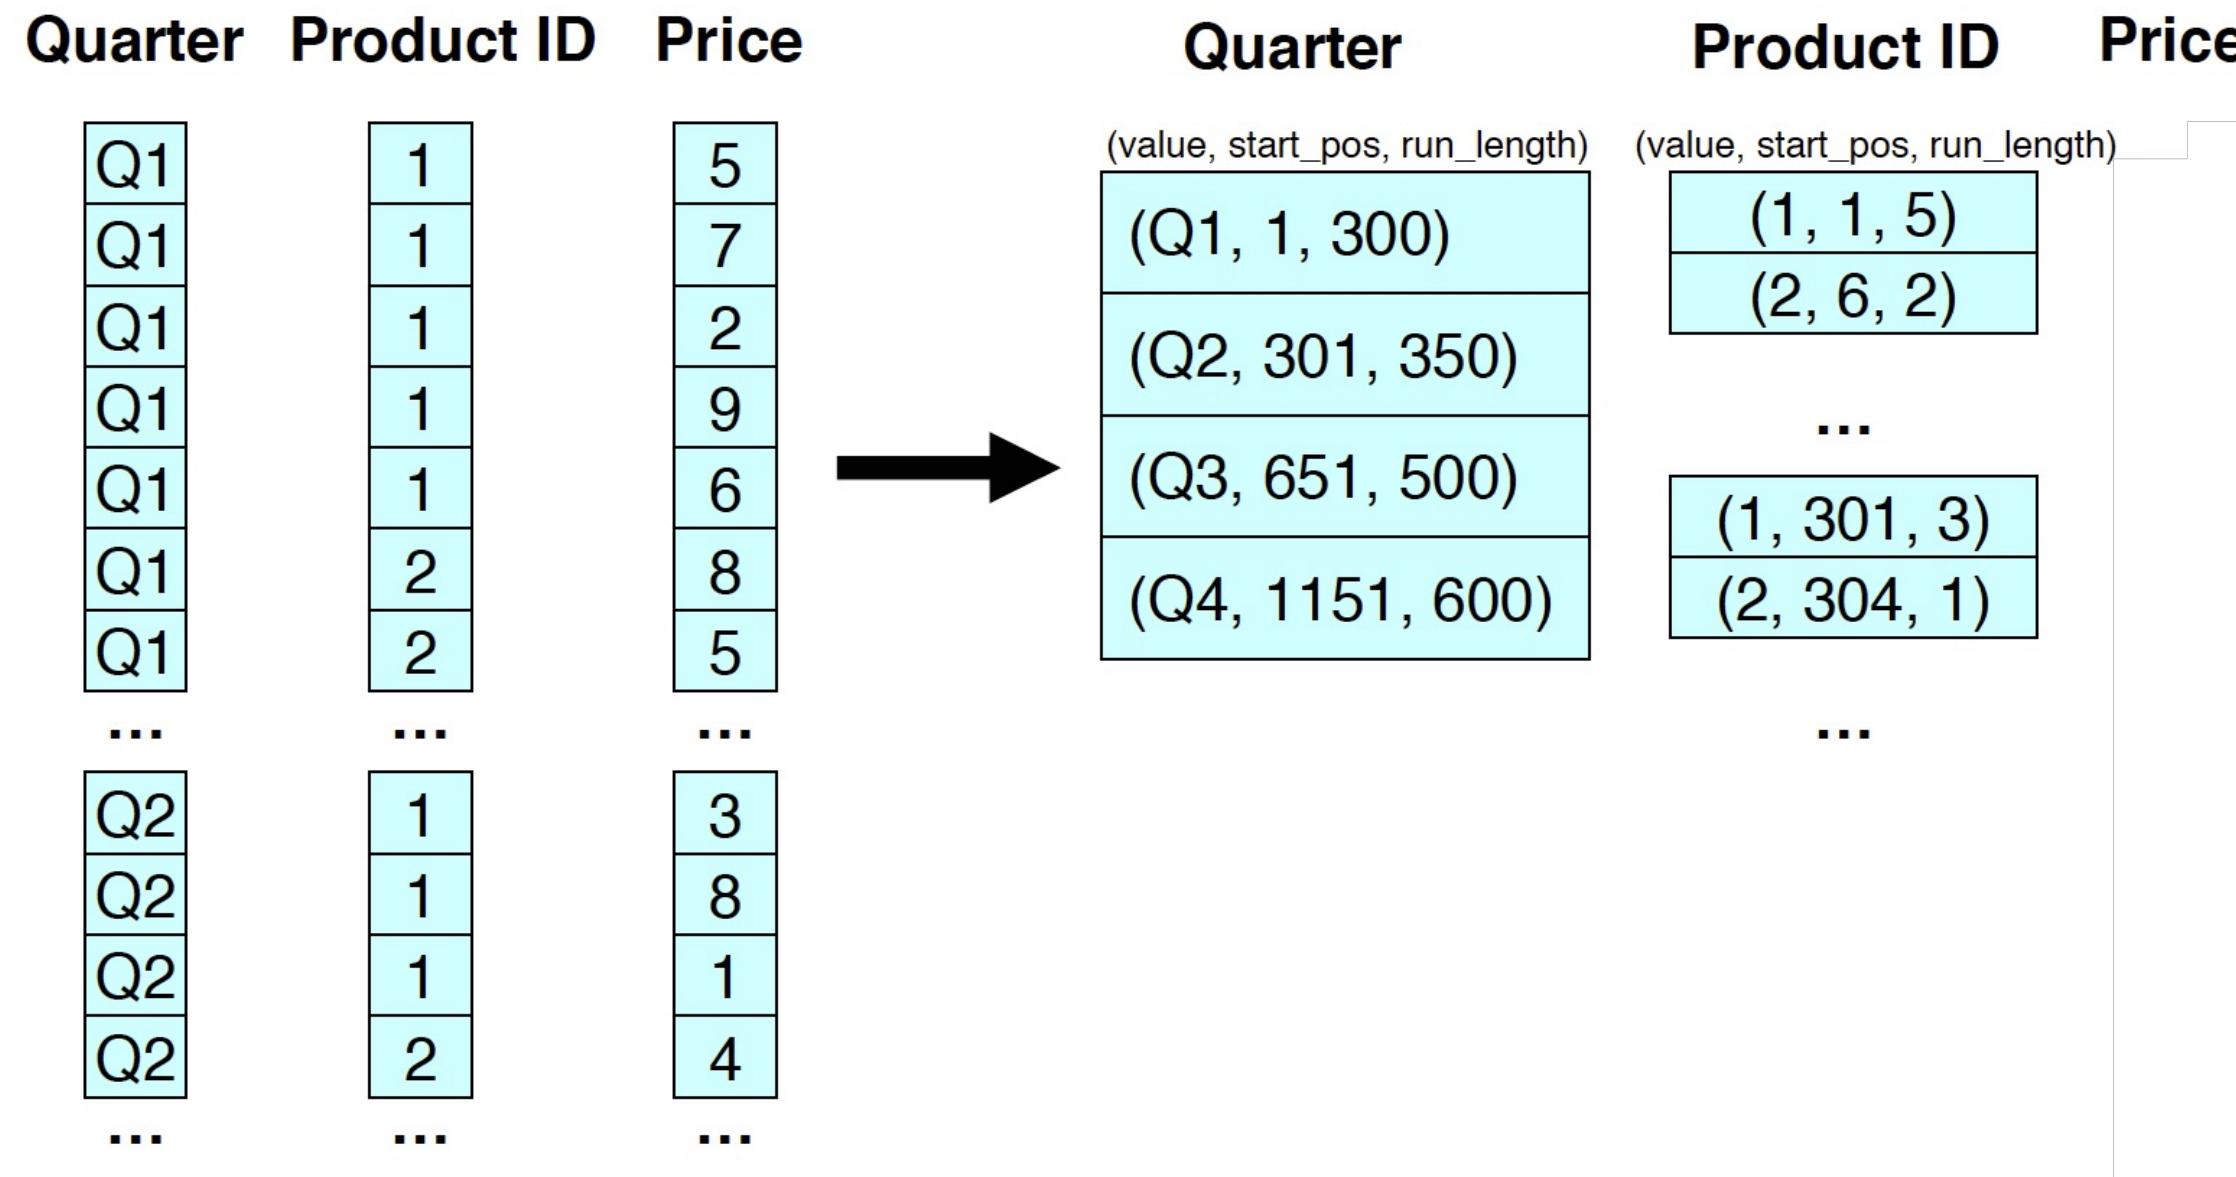
# Example (DB): Run-Length Encoding



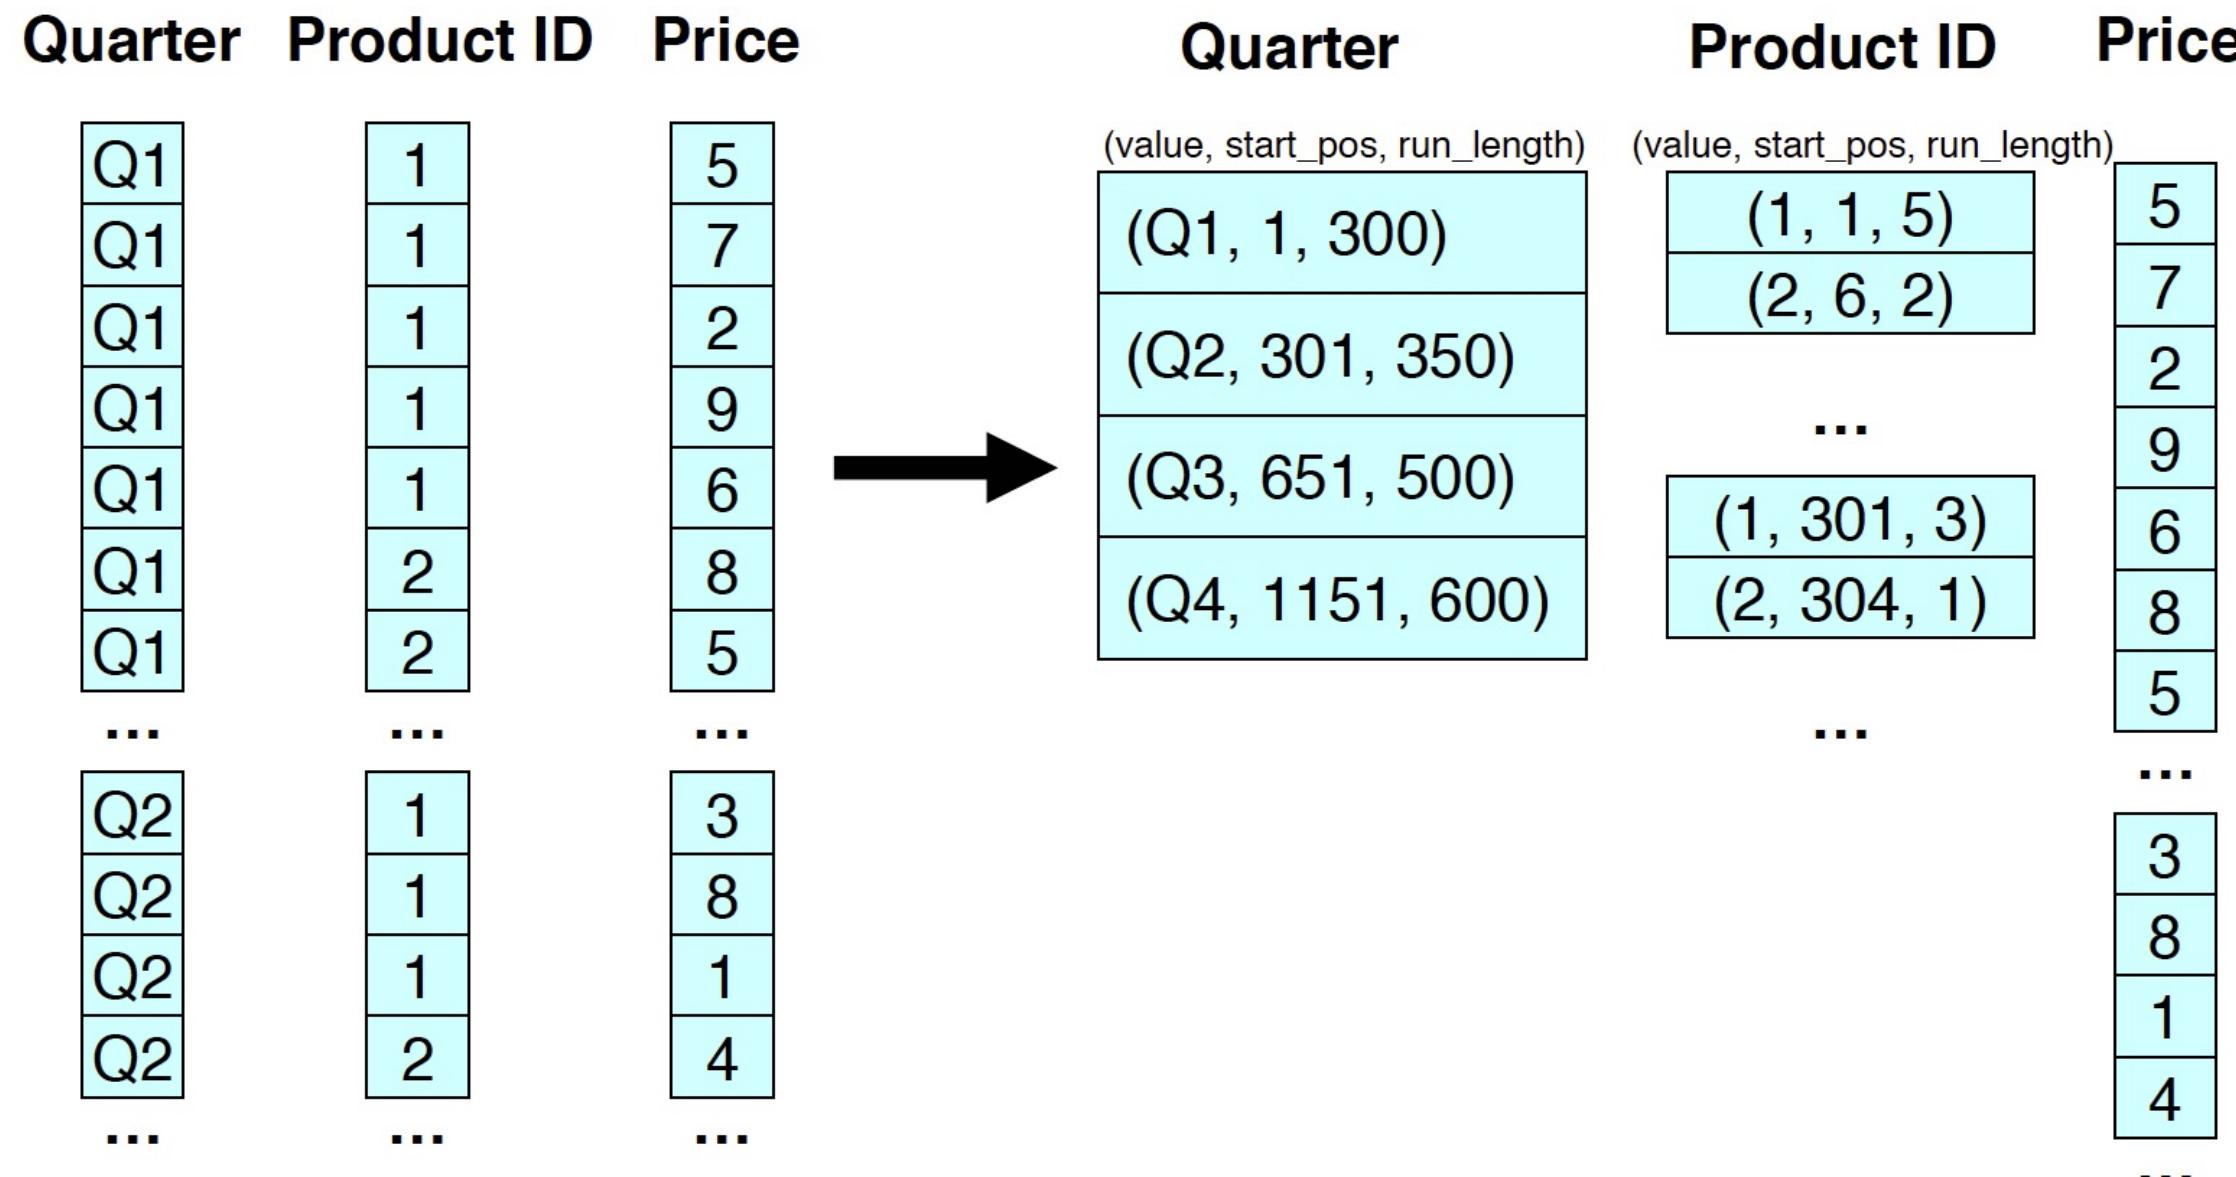
# Compression: Run-Length Encoding



# Compression: Run-Length Encoding



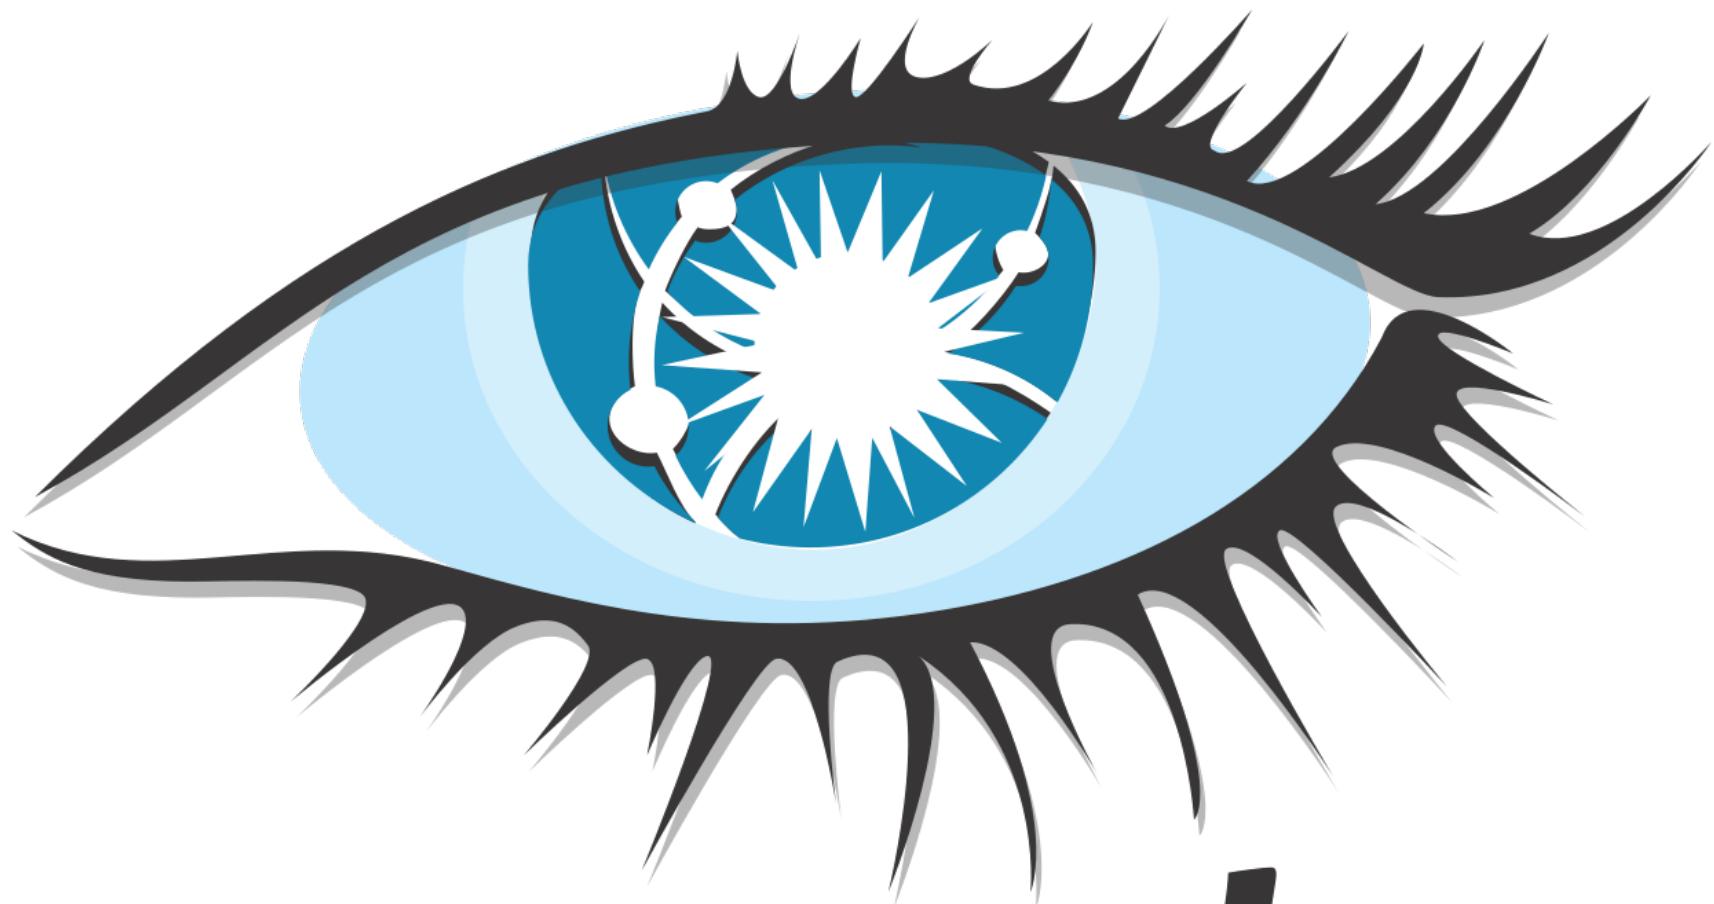
# Compression: Run-Length Encoding



# List of Databases

- **Cassandra**
- Vertica
- SybaseIQ
- C-Store
- BigTable/HBASE
- MonetDB
- LucidDB

# Cassandra



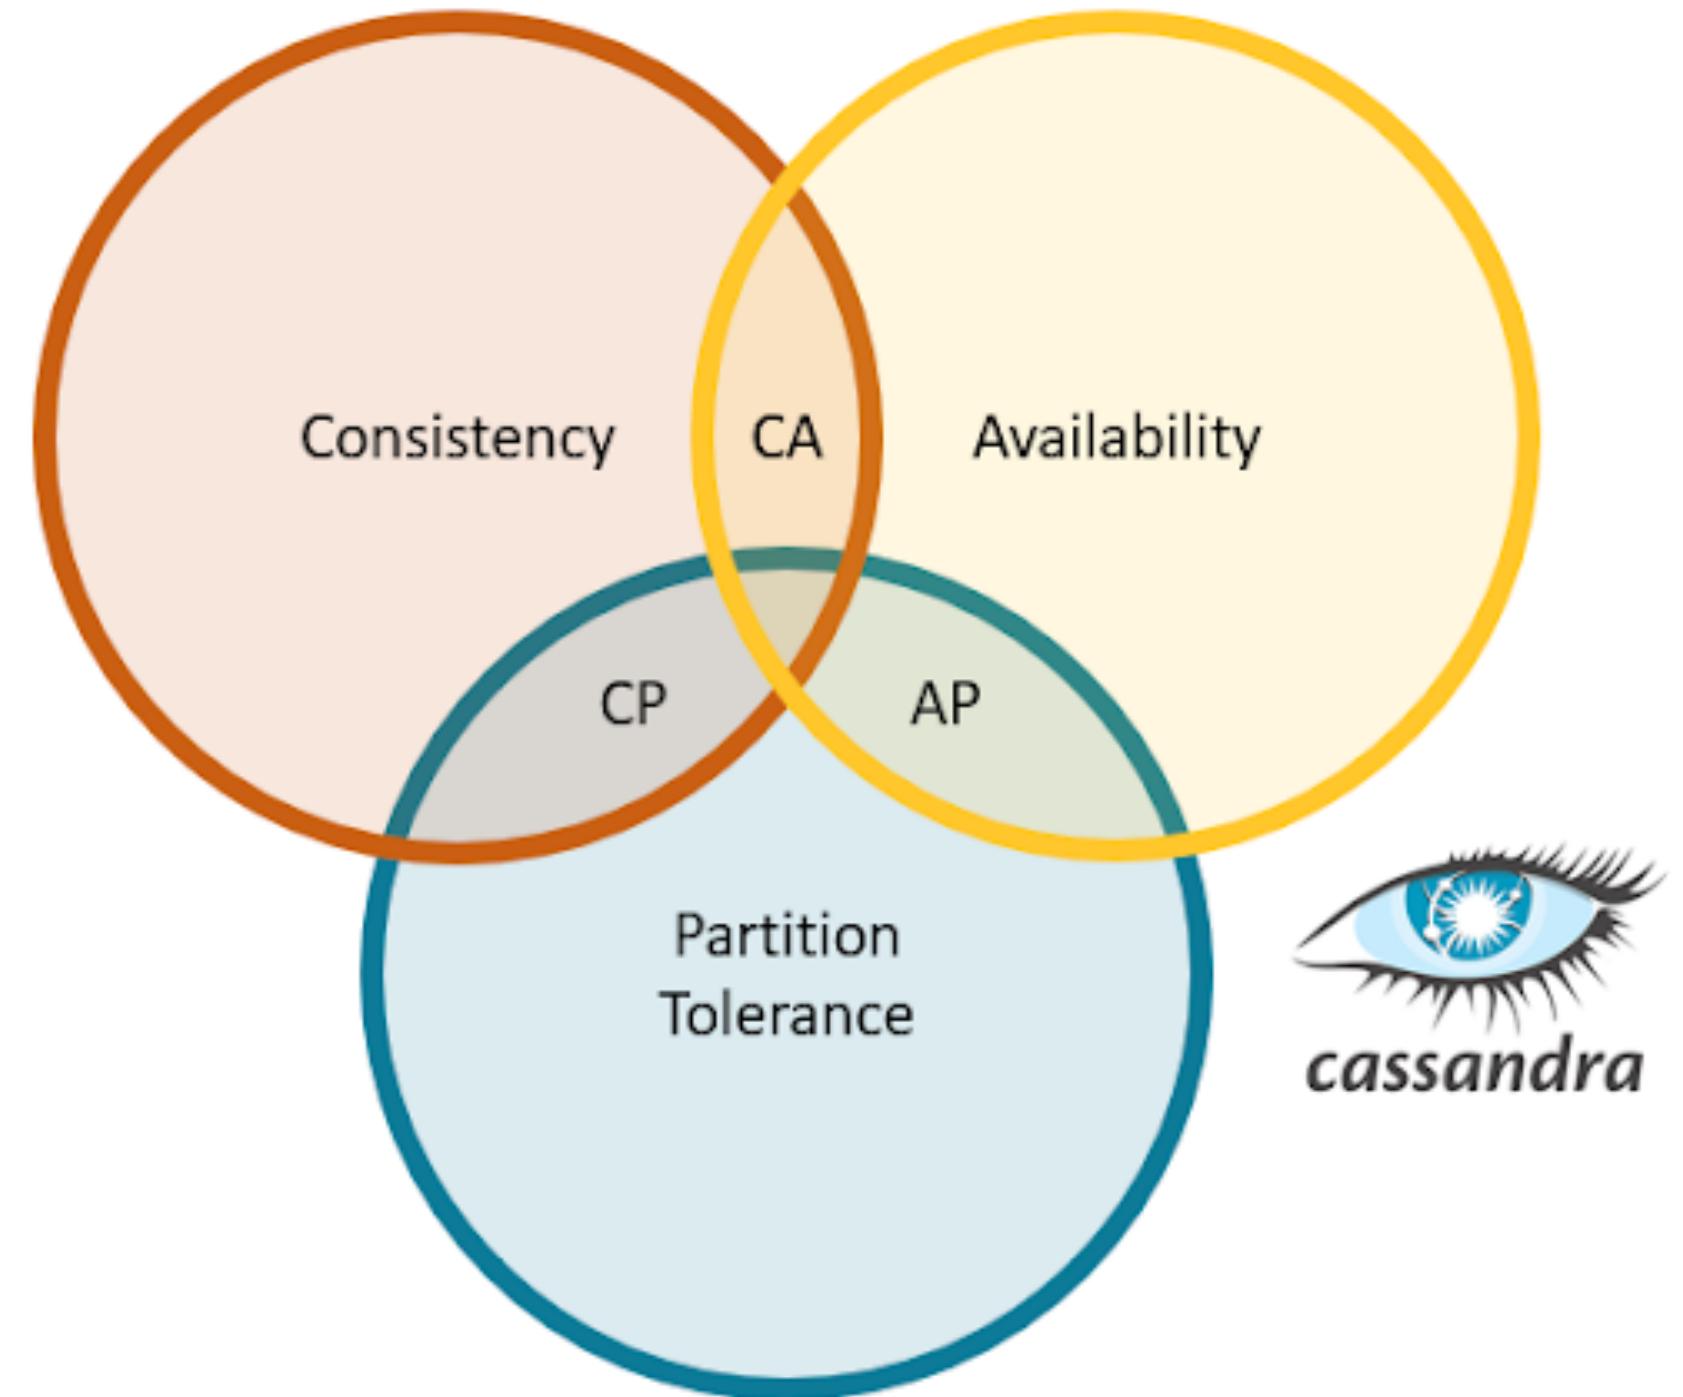
# History of Cassandra

Originally designed at Facebook

Open-sourced and now within Apache foundation

# What Cassandra is

- tuneably consistent (**C**)
- very fast in writes
- highly available (**A**)
- fault tolerant (**P**)
- linearly scalable, elastic scalability
- Cassandra is very good at writes, okay with reads.



# What Cassandra is not

- Cassandra is not a replacement for Relational Databases
- Tables should **not** have multiple access paths
- Cassandra does not support aggregates, if you need to do a lot of them, think another database.
- Updates and deletes are implemented as special cases of writes and that has consequences that are not immediately obvious.

# Comparison with RDBMS

Property	Cassandra	RDBMS
Core Architecture	Masterless (no single point of failure)	Master-slave (single points of failure)
High Availability	Always-on continuous availability	General replication with master-slave
Data Model	Dynamic; structured and unstructured data	Legacy RDBMS; Structured data
Scalability Model	Big data/Linear scale performance	Oracle RAC or Exadata
Multi-Data Center Support	Multi-directional, multi-cloud availability	Nothing specific
Enterprise Search	Integrated search on Cassandra data.	Handled via Oracle search
In-Memory Database Option	Built-in in-memory option	Columnar in-memory option

Property	Cassandra	RDBMS
Joining	Doesn't support joining	Supports joining
Referential Integrity	Cassandra has no concept of referential integrity across tables. No cascading deletes.	Supports foreign keys in a table to reference the primary key of a another table. Supports cascading delete.
Normalization	Tables contain duplicate denormalize data.	Tables are normalized to avoid redundancy.

# Use Cases

The use-case leading to the initial design and development of Cassandra was the so entitled Inbox Search problem at Facebook.

-  [Purchases, test scores](#)
- Storing time series data (as long as you do your own aggregates).
  - Storing health tracker data.
  - Weather service history.
  -  [User Activity](#)
- Internet of things status and event history.
-  [IOT for cars and trucks](#)
- Email envelopes—not the contents.

# When to consider Cassandra

- you need really fast writes
- you need durability
- you have lots of data (> GBs) and (>=) three servers
- your app is evolving
  - startup mode, fluid data structure
- loose domain data
  - “points of interest”
- your programmers can handle
  - complexity
  - consistency model
  - change
  - visibility tools
- your operations can deal
  - hardware considerations
  - data transport
  - JMX monitoring

## Advantages

A general-purpose framework for high concurrency & load conditioning

Decomposes applications into stages separated by queues

Adopt a structured approach to event-driven concurrency

# Data Model

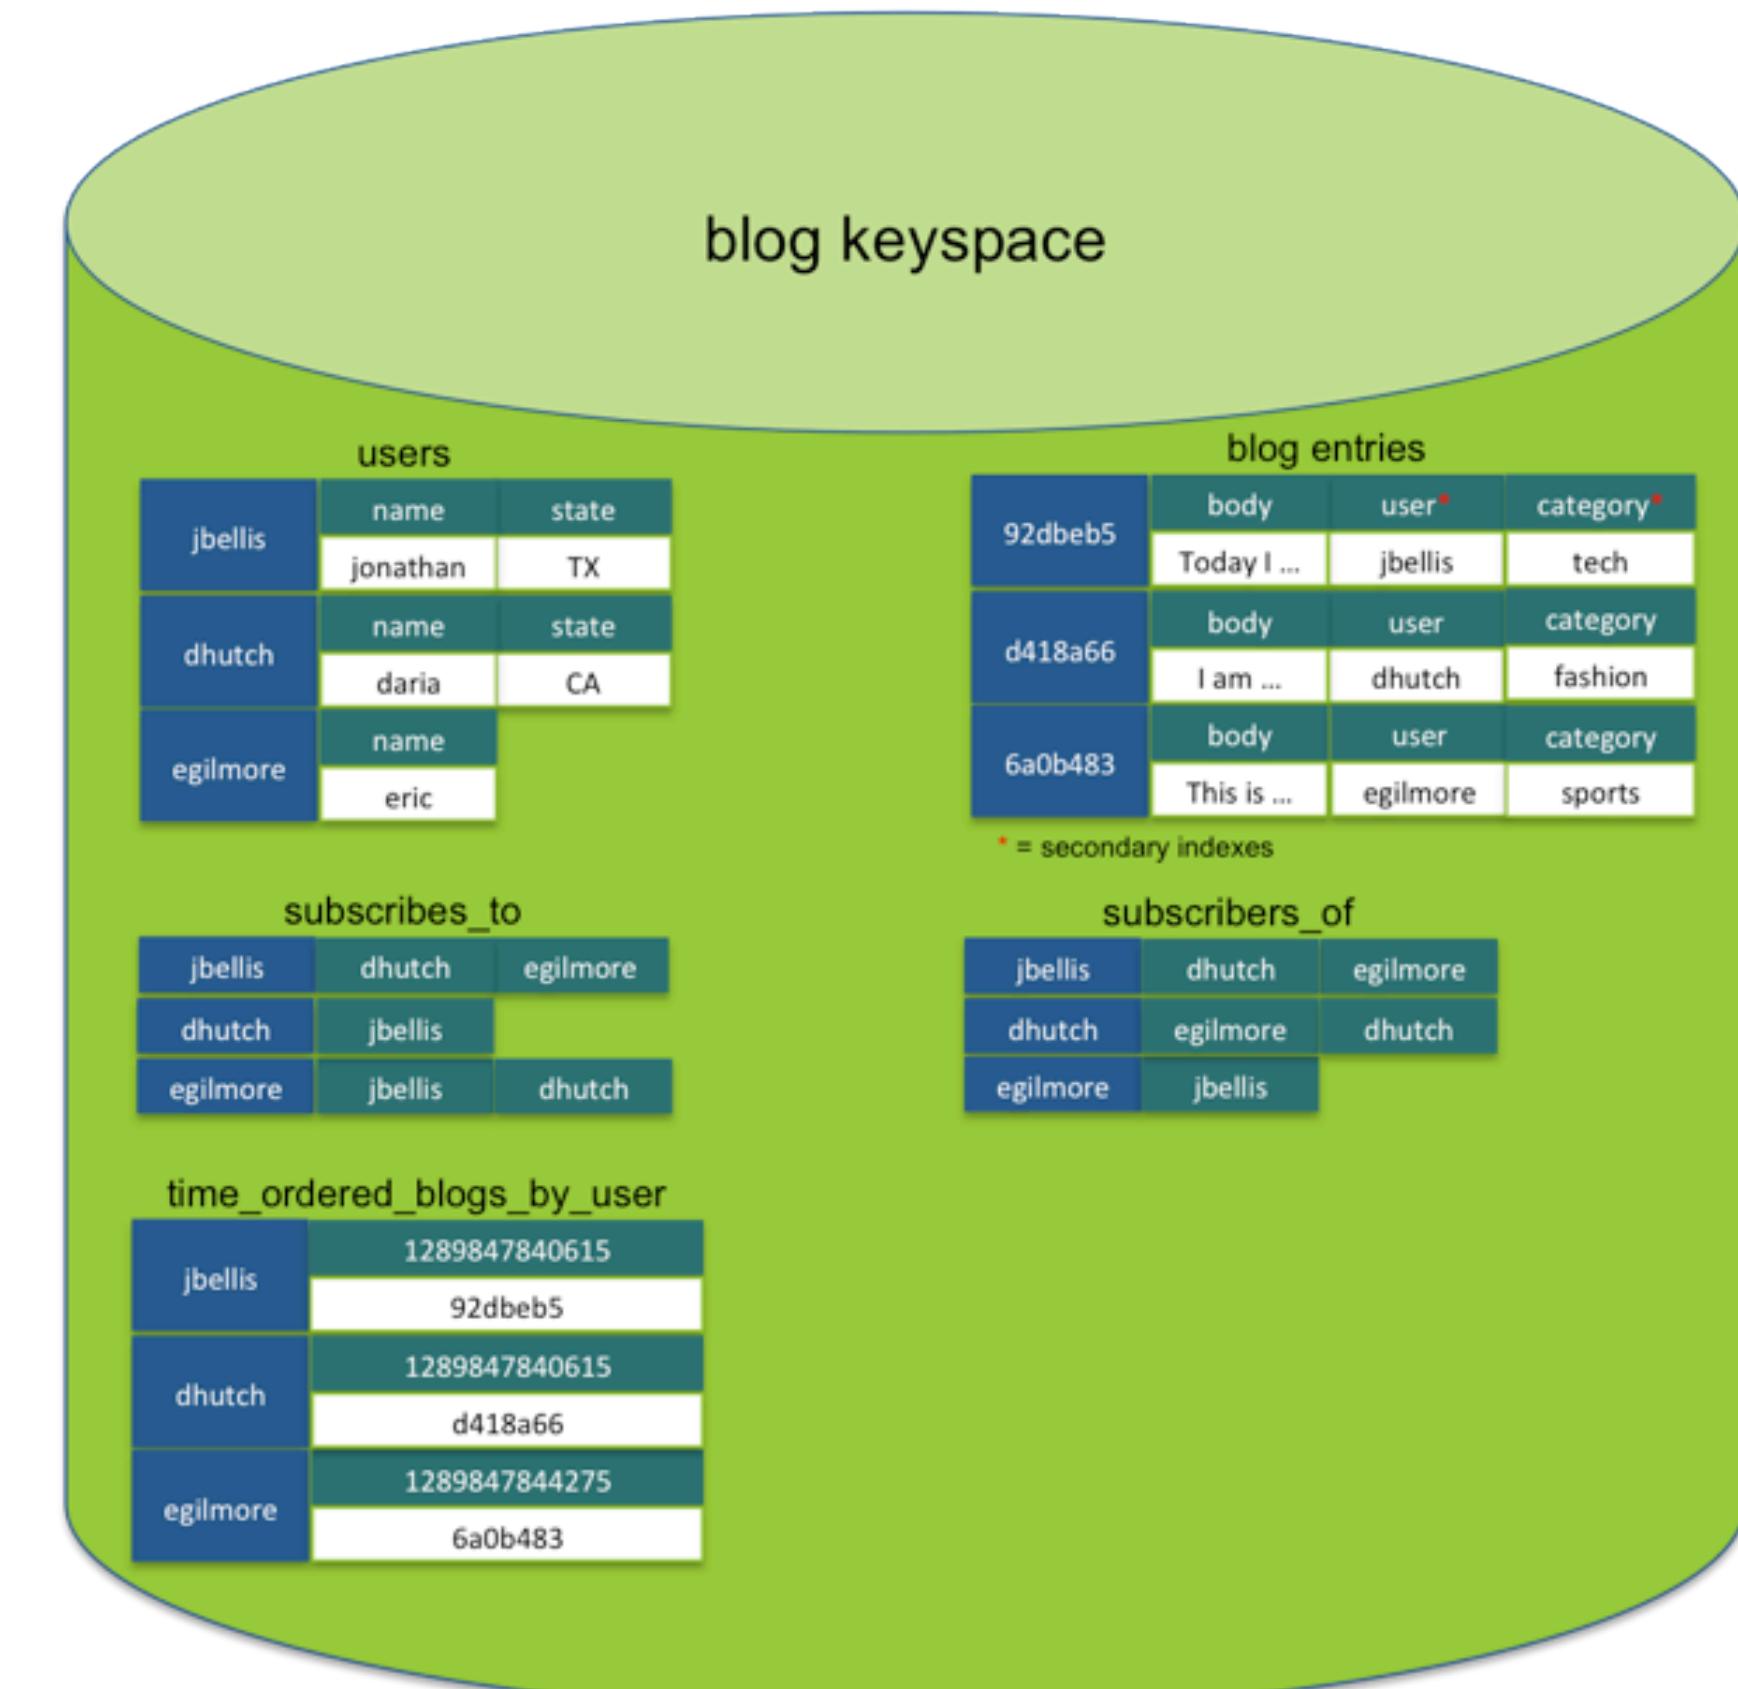
**RDBMSs:** domain-based model  
-> what answers do I have?

**Cassandra:** query-based model  
-> what questions do I have?

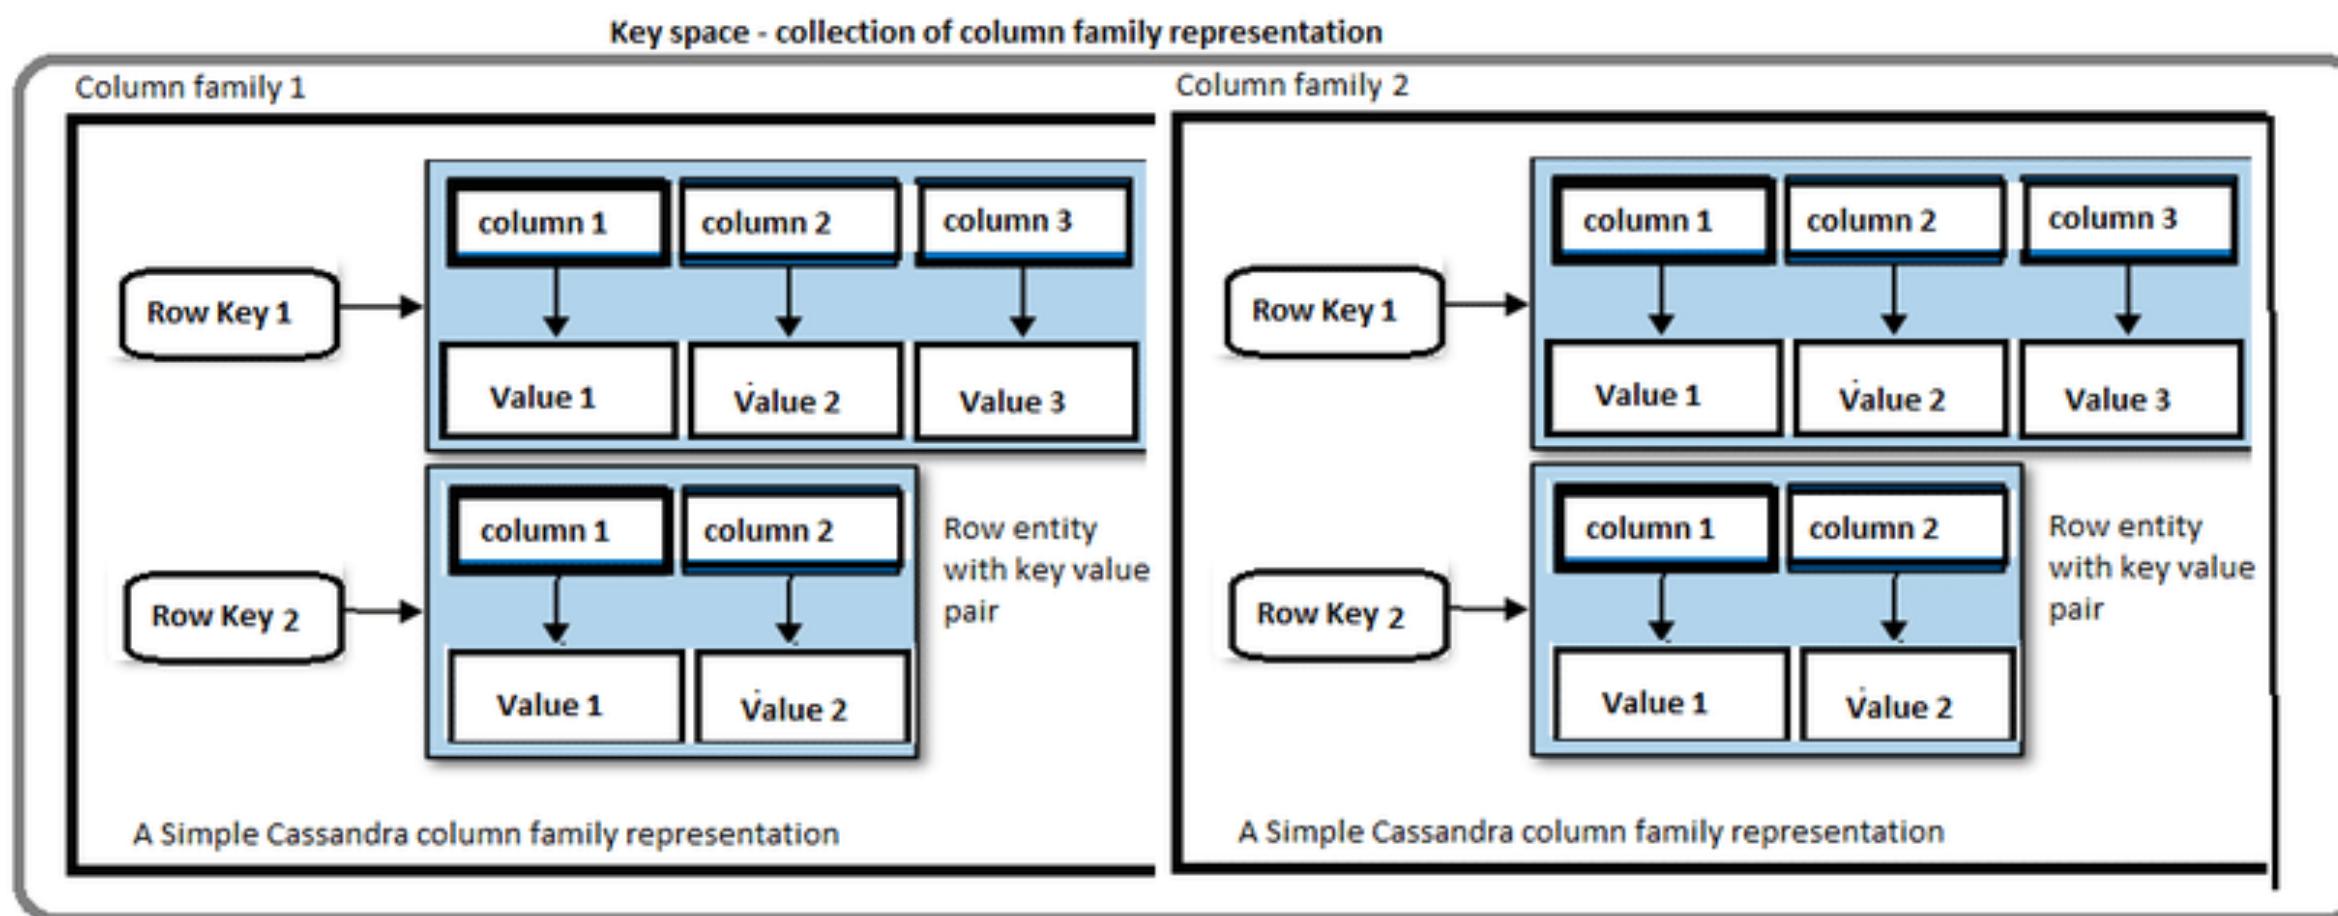
Cassandra does **not** support a full relational data model.

Instead, it provides clients with a simple data model that supports **dynamic control** over data layout and formats.

An instance of Cassandra typically consists of one distributed multidimensional map indexed by key which contains one or more **column families** that, in turn, **rows**

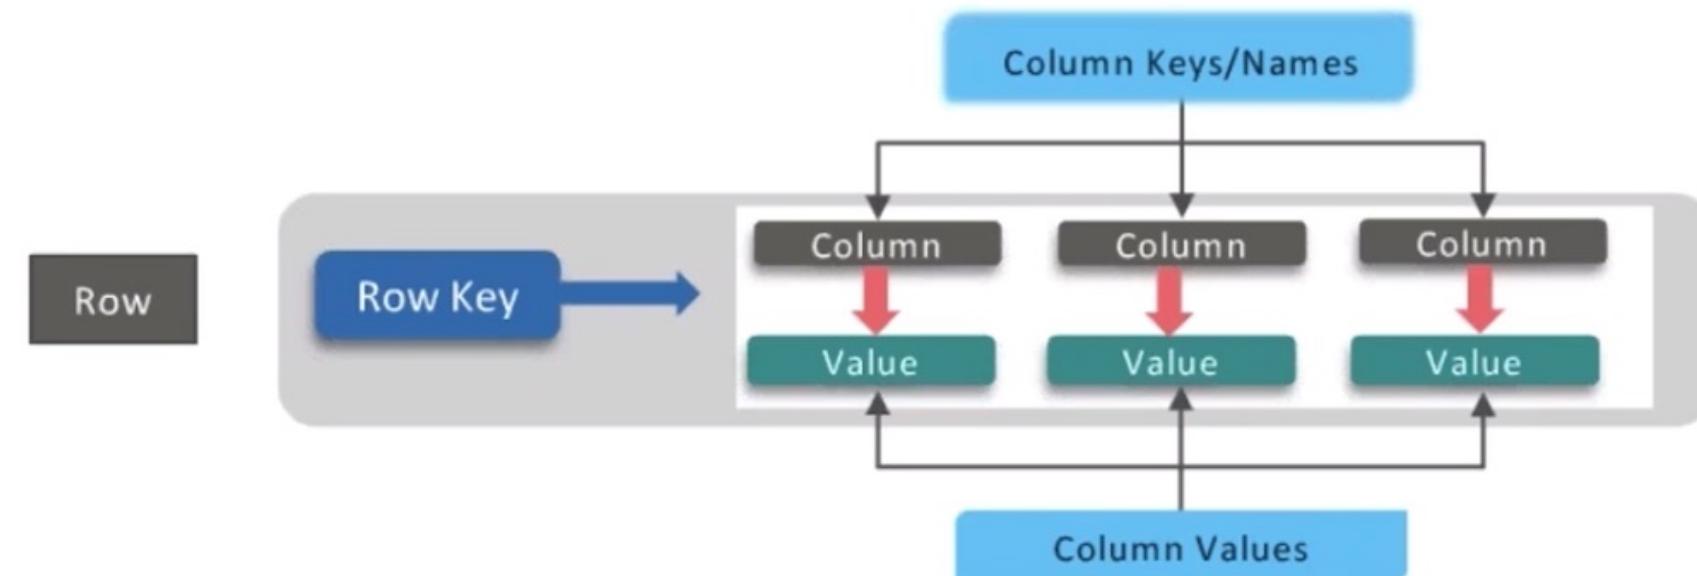


- **Rows** are identified by a string-key
- **Column Families** corresponds to tables in RDBMS but may be unstructured. A column family consists of
  - **(Simple) Columns Families** have a name and store a number of values per row which are identified by a timestamp
  - **Super Columns Families** have a name and an arbitrary number of columns associated with them



# Keyspace

- Key space is typically one per application
- Keys are similar to those of databases
- Some settings are configurable only per keyspace
- Each Row must have a key



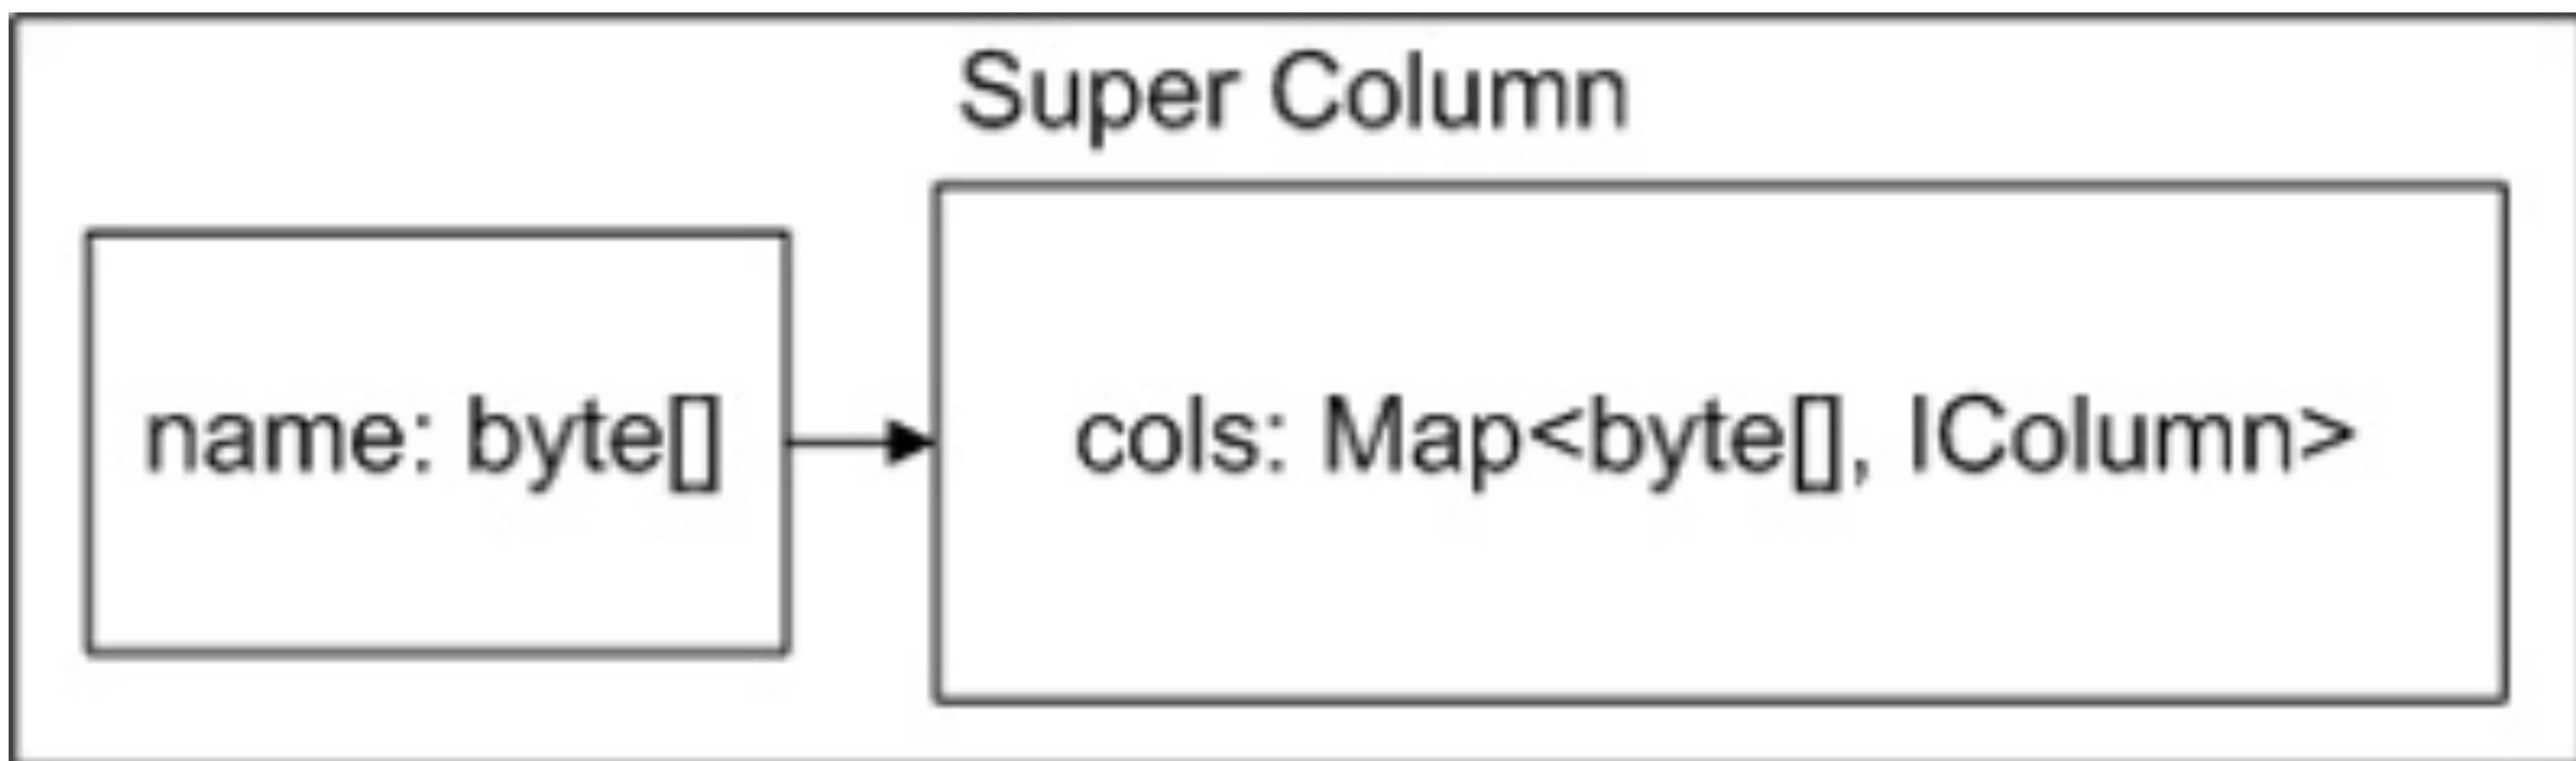
# Columns Families

A Column consists of three parts

- name
  - byte[]
  - determines sort order
  - used in queries
  - indexed
- value
  - byte[]
  - you don't query on column values
- timestamp
  - long (clock)
  - last write wins conflict resolution

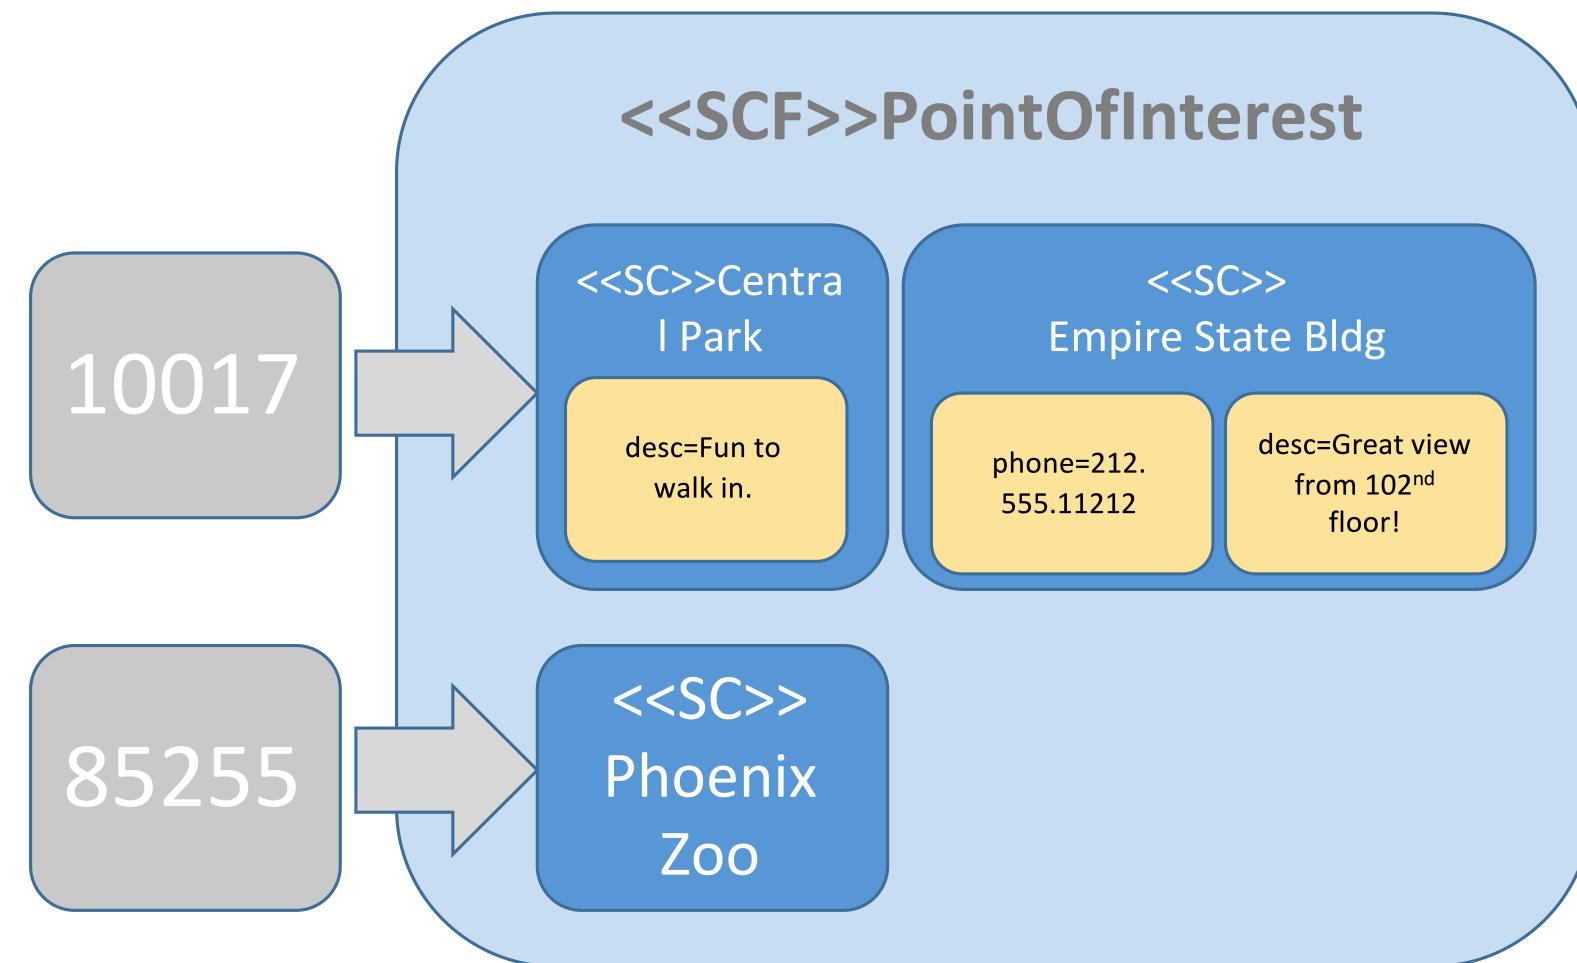
# Super Column Families

- Super columns group columns under a common name
- sub-column names in a Super Column Family are **not** indexed
  - top level columns (Super Column Family Name) are **always** indexed
- often used for **denormalizing** data from standard Column Families



## Example

# super column family



## Example (Json Notation)

```
PointOfInterest { //Supercolumn Family
    key:85255 {
        Phoenixzoo { phone: 480-555-5555, //column
                     desc: They have animals here //column },
        Spring Training {
            phone: 623-333-3333, //column
            desc: Fun for baseball fans. //column
        }
    } //end phoenix,

    key: 10019 {
        Central Park //super column
        { desc: Walk around. It's pretty. // missing phone column } ,
        Empire State Building { phone: 212-777-7777,
                               desc: Great view from 102nd floor. }
    } //end nyc
}
```

# Architecture

Cassandra is required to be incrementally scalable.

Therefore machines can join and leave a cluster (or they may crash).

Data have to be **partitioned** and **distributed** among the nodes of a cluster in a fashion that allows *repartitioning* and *redistribution*.

# Partitioning

- Data of a Cassandra table get partitioned and distributed among the nodes by a consistent **order-preserving** hashing function.
- The order preservation property of the hash function is important to support **range scans** over the data of a table.
- Cassandra performs a **deterministic** load balancing
  - it measures and analyzes the load information of servers and moves nodes on the consistent hash ring to get the data and processing load balanced.

# Replication

- Data get replicated to a number of nodes which can be defined as a **replication factor** per Cassandra instance.
- Replication is managed by a **coordinator node** for the particular **key** being modified.
- The coordinator node for any key is the **first node on the consistent hash ring** that is visited when walking from the key's position on the ring in **clockwise** direction.

## ~~Replication Strategies~~ (Used to be)

- **Rack Unaware:** the non-coordinator replicas are chosen by picking  $N-1$  successors of the coordinator on the ring
- **Rack Aware** and **Datacenter Aware** rely on Zookeeper for leader election.
  - the elected leader is in charge of maintaining the invariant that no node is responsible for more than  $N-1$  ranges in the ring

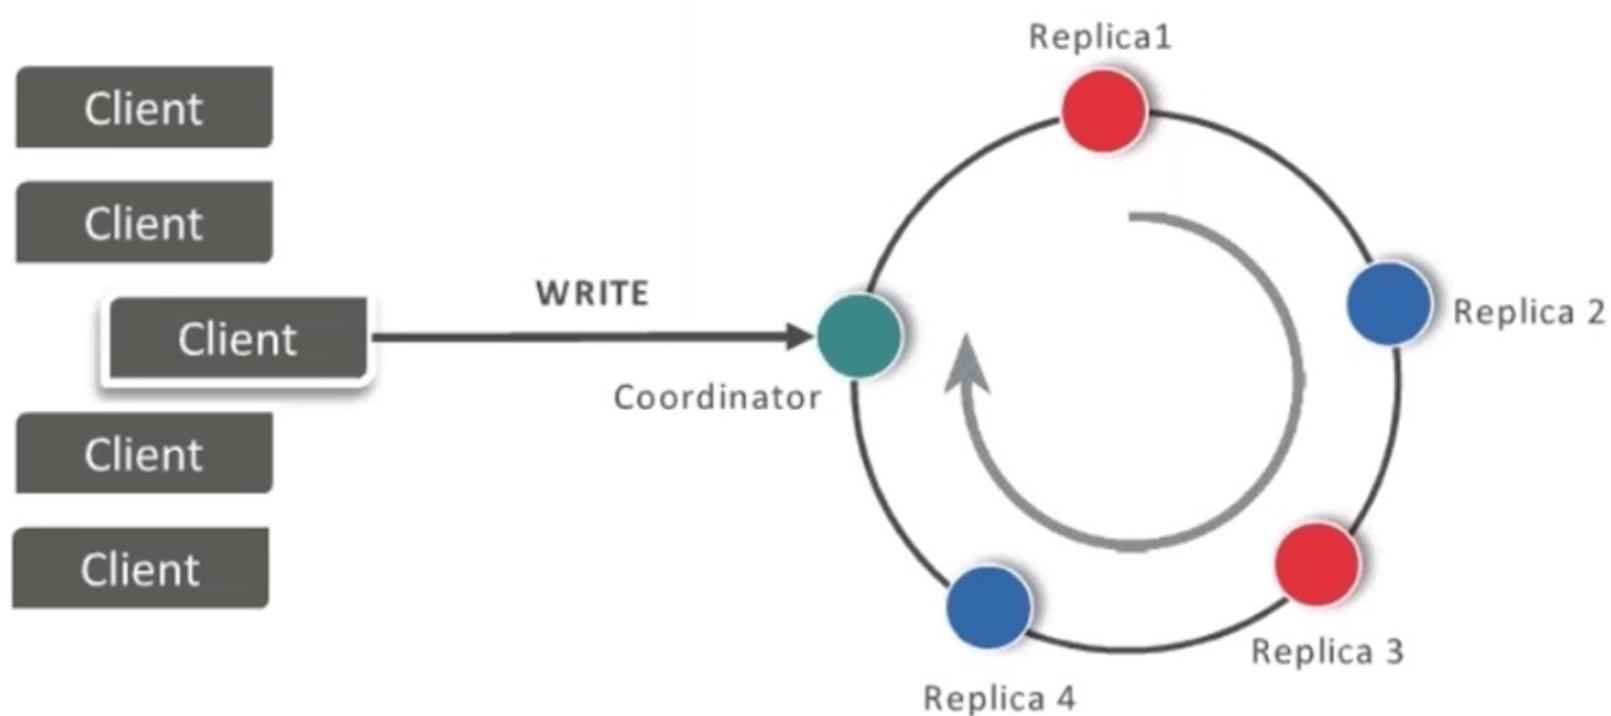
## Replica placement strategies Today<sup>81</sup>

---

<sup>81</sup> [docs](#)

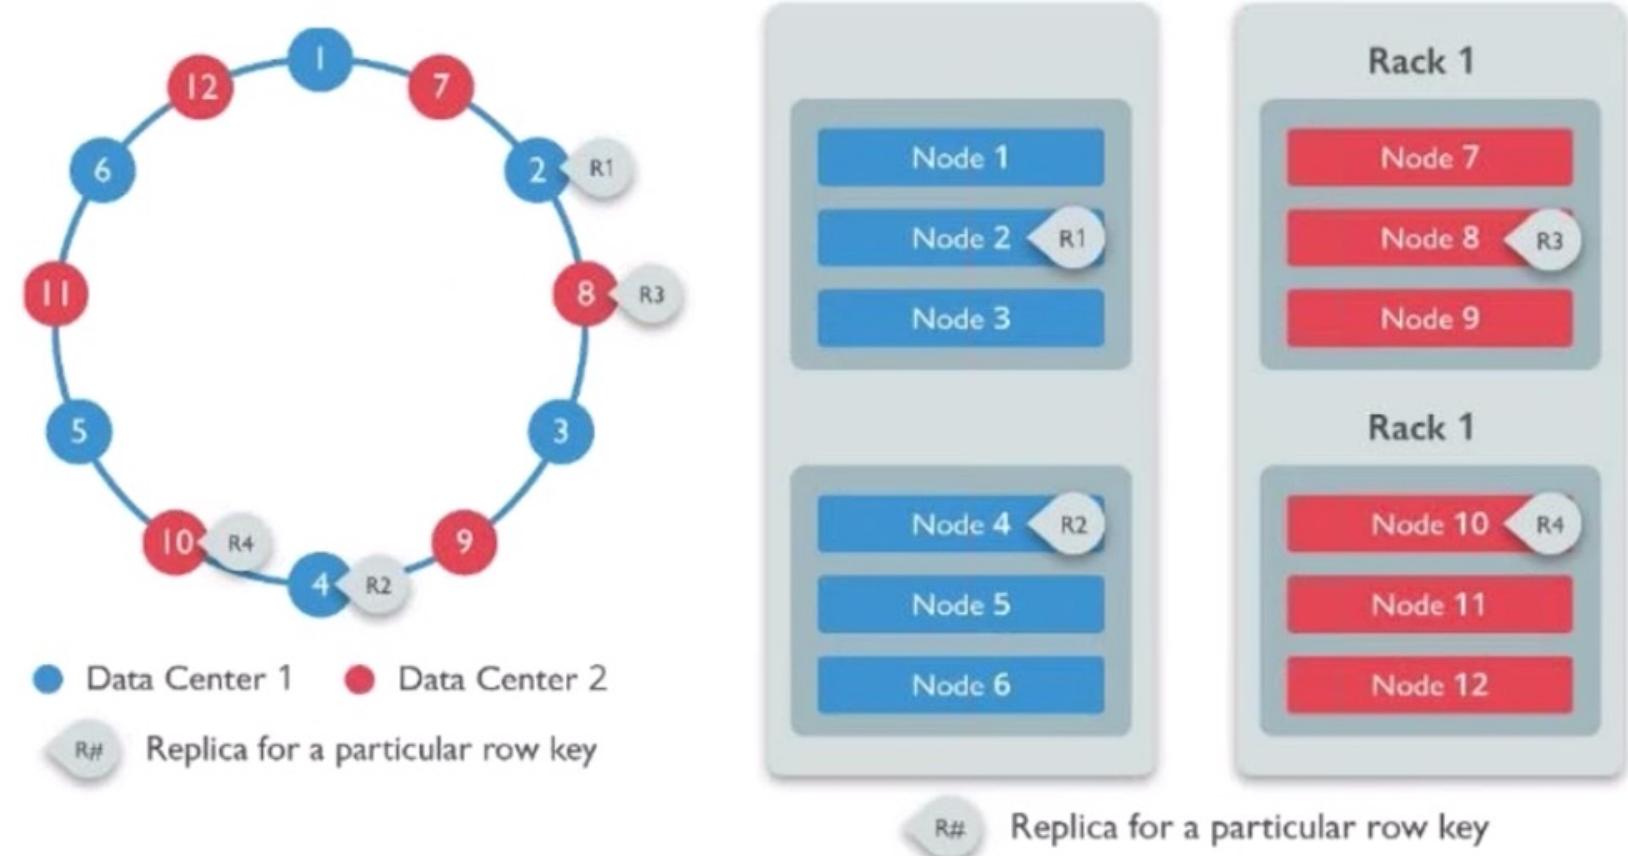
## Simple Strategy

- Allows a single integer *replication\_factor* to be defined
- Single datacenter
- Clockwise placement to the next node(s)
- all nodes are treated equally, ignoring any configured data centers or racks.



## Network Topology Strategy

- Multiple datacenters
- Allows a single integer *replication\_factor* to be defined per data center
- Attempts to choose replicas within a data center from different racks as specified by the Snitch<sup>82</sup>
- Supports local (reads) queries



<sup>82</sup> Snitch teaches Cassandra about your network topology to route requests efficiently.

# Partitioner Smack-Down

## Random Preserving

- system will use MD5(key) to distribute data across nodes
- even distribution of keys from one Column Family across ranges/nodes

## Order Preserving

- key distribution determined by token
- lexicographical ordering
- required for range queries
- can specify the token for this node to use

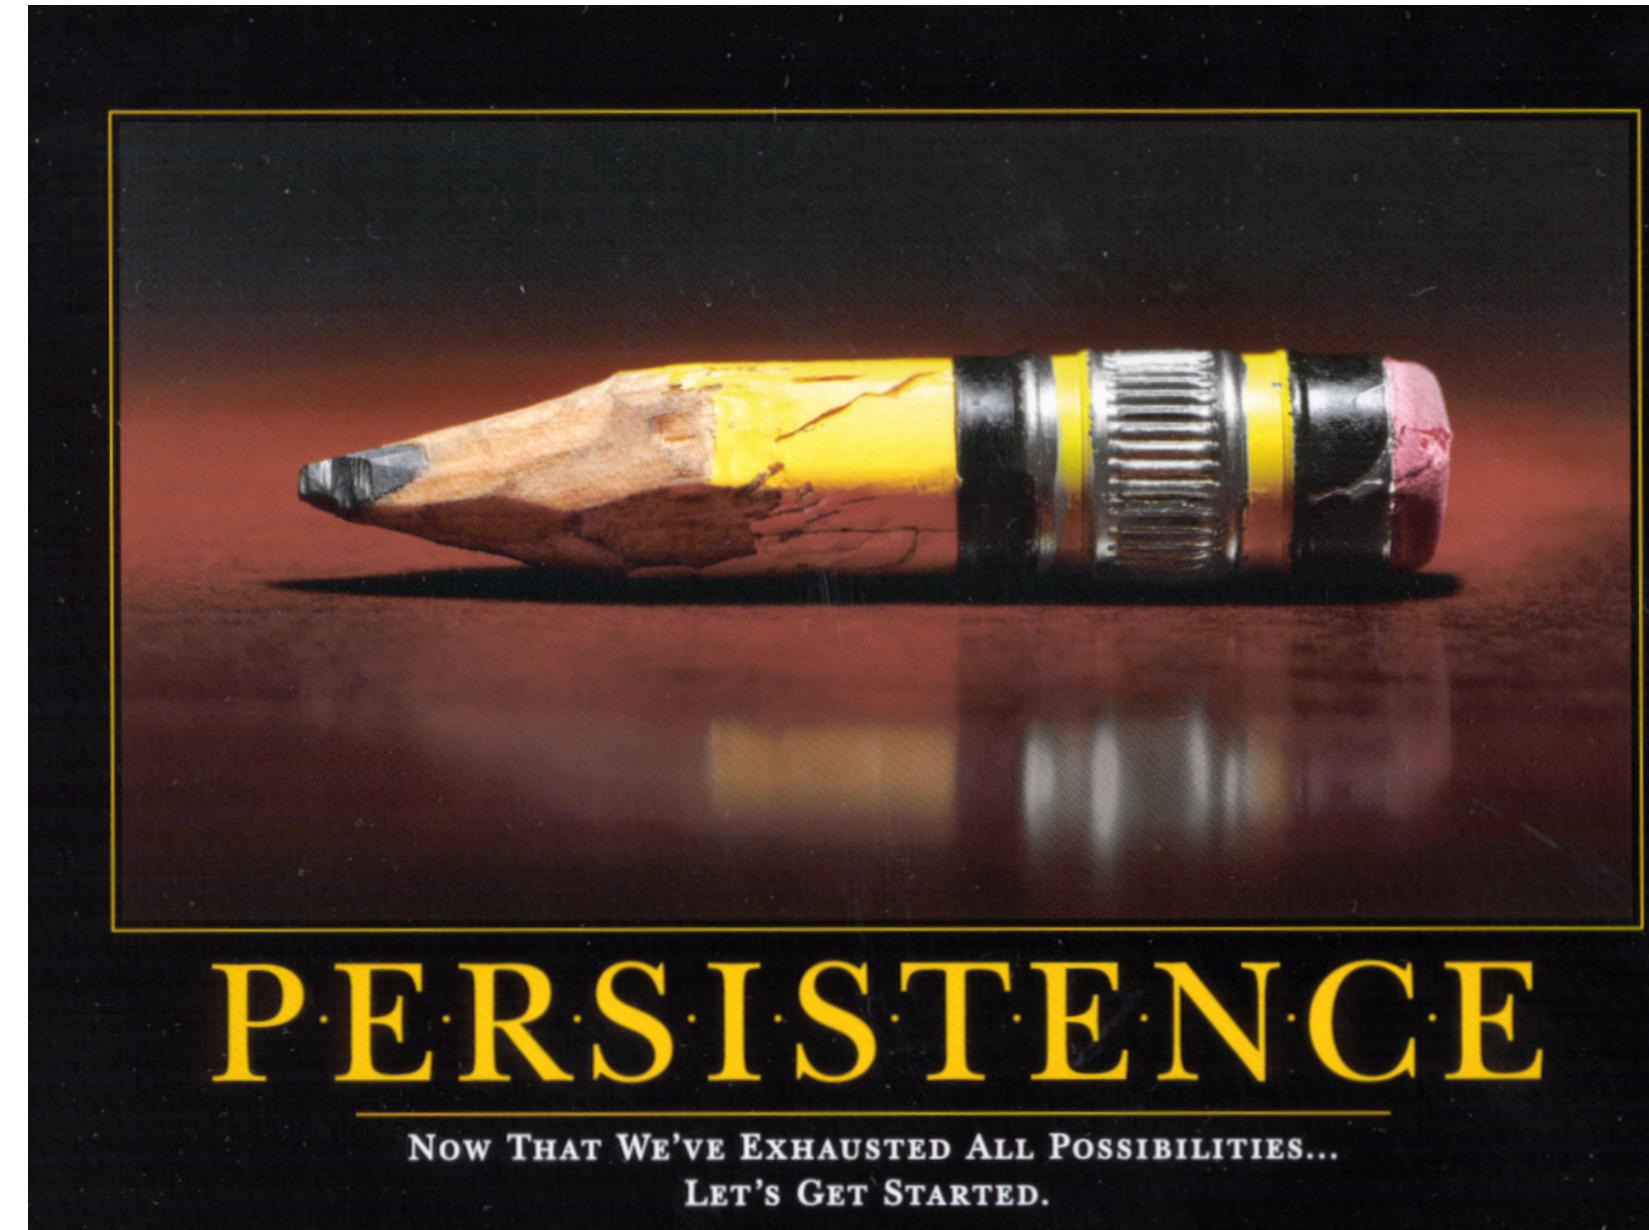
## Persistence

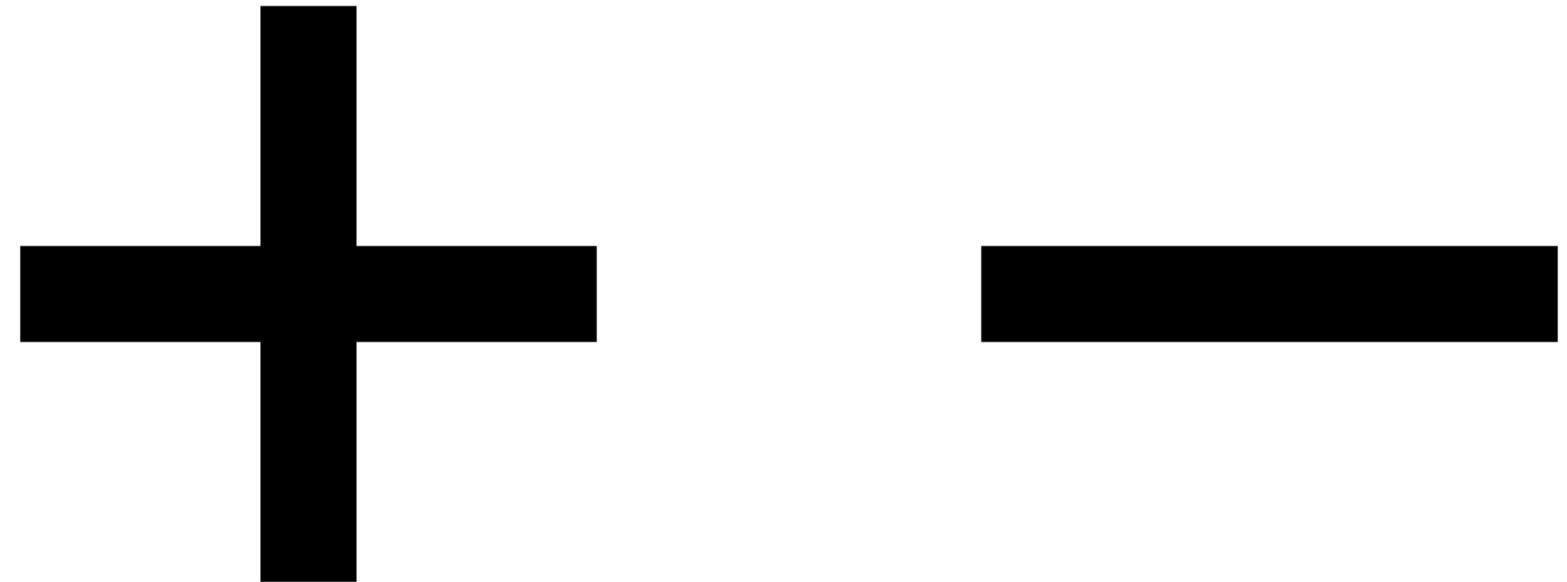
Cassandra provides **durability guarantees** in the presence of node failures and network partitions by relaxing the quorum requirements

The Cassandra system relies on the **local file system** for data persistence.

The data is **represented** on disk using a format that lends itself to **efficient** data **retrieval**.

Typical write operation involves a write into a **commit** log for durability and recoverability and an update into an in-memory data structure.





# Operations



# Writes

- Need to be lock-free and fast (no reads or disk seeks)
- A Client sends write to one front-end node in Cassandra cluster (Coordinator)
- Coordinator sends it to all replica nodes responsible for that key
- A write is atomic at the partition-level, meaning inserting columns in a row is treated as one write operation.

## Hinted Handoff

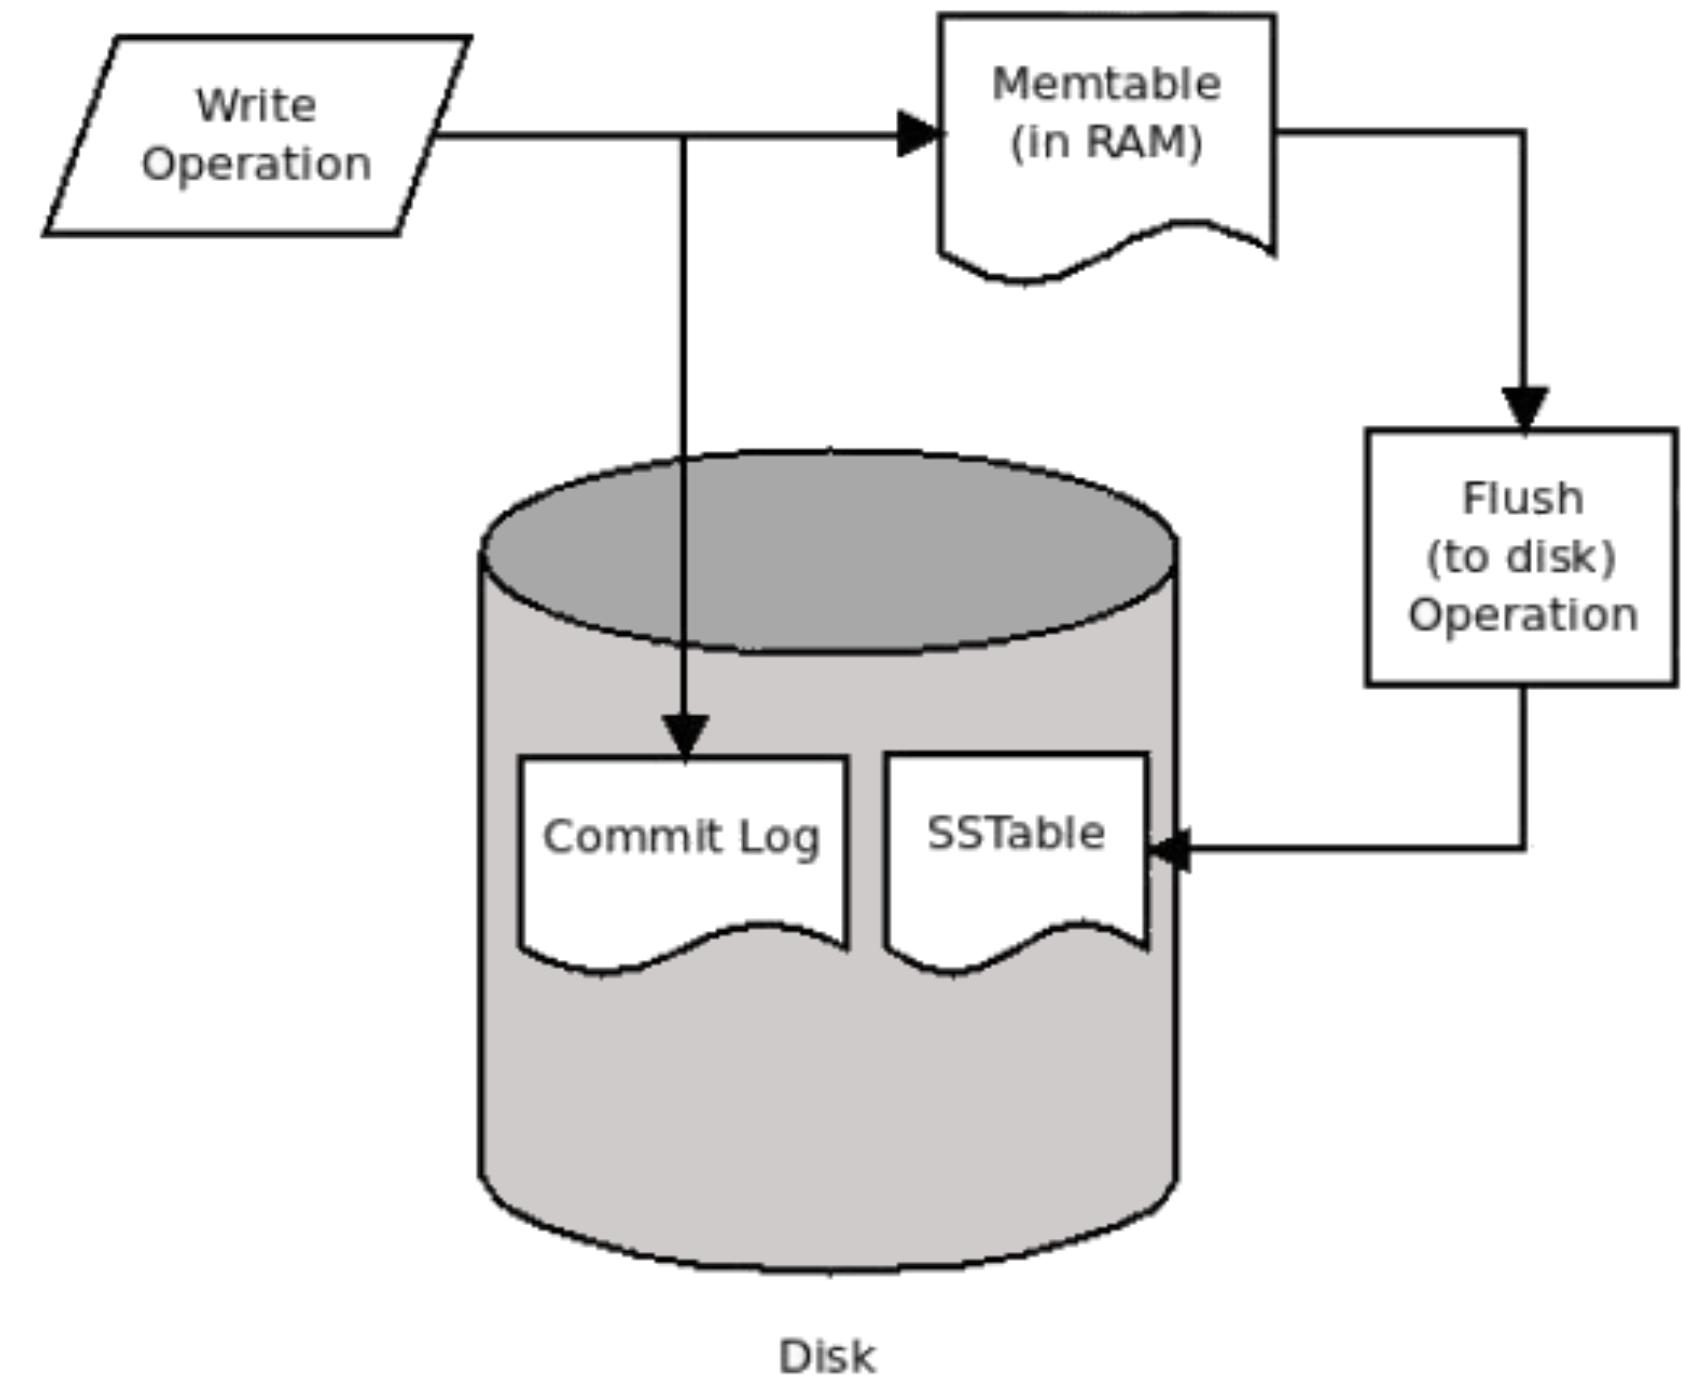
If any replica is down, the coordinator writes to all other replicas, and keeps the write until down replica comes back up.

When all replicas are down, the Coordinator (front end) buffers writes (for up to an hour).

## Writing Flow

1. Cassandra logs it in disk commit log (disk)
2. Adds *values* to appropriate *memtables*<sup>83</sup>
3. When memtable is full or old, flush to disk using a Sorted String Table

[source](#)



---

<sup>83</sup> In-memory representation of multiple key-value pairs

# Consistency levels for a write operations<sup>87</sup>

- ANY: any node (may not be replica)
- ONE: at least one replica
- QUORUM: quorum across all replicas in all datacenters
- LOCAL-QUORUM: in coordinator's datacenter
- EACH-QUORUM: quorum in every datacenter
- ALL: all replicas all datacenters

---

<sup>87</sup> [detailed discussion](#)

# Write Consistency

Level	Description
ZERO	Good luck with that
ANY	1 replica (hints count)
ONE	1 replica. read repair in bkgnd
QUORUM (DCQ for RackAware)	$(N / 2) + 1$
ALL	$N = \text{replication factor}$

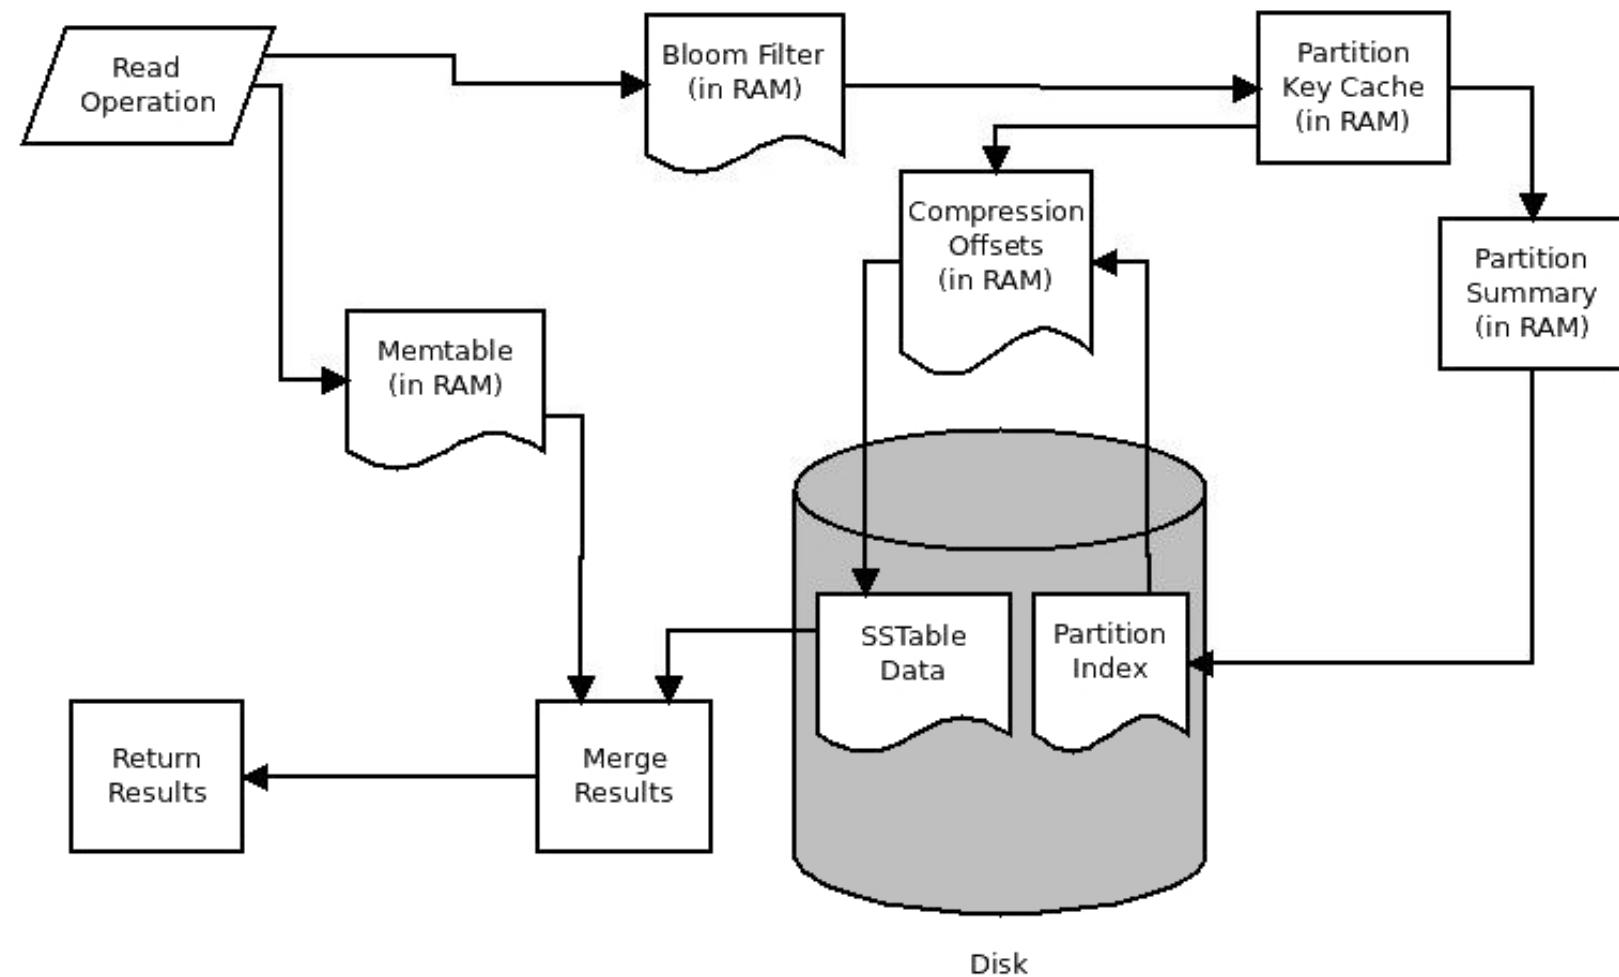
# Reads

- Coordinator can contact closest replica (e.g., in same rack)
- Coordinator also fetches from multiple replicas
  - check consistency in the background,
  - Makes read slower than writes (but still fast)
  - initiating a **read-repair** if any two values are different using gossip

## Reading Flow

1. Check row cache, if enabled
2. Checks partition key cache, if enabled
3. Check the memtable
4. Fetches the data from the SSTable on disk
5. If Row cache is enabled the data is added to the row cache

[source](#)



# Read Consistency

Level	Description
ZERO	Ummm...
ANY	Try ONE instead
ONE	1 replica
QUORUM (DCQ for RackAware)	Return most recent TS after $(N / 2) + 1$ report
ALL	$N = \text{replication factor}$

## Read-Repair<sup>84</sup>

Read-repair is a **lazy** mechanism in Cassandra that ensures that the data you request from the database is accurate and consistent.

For every read request, the coordinator node requests to all the nodes having the data requested by the client. All nodes return the data which client requested for.

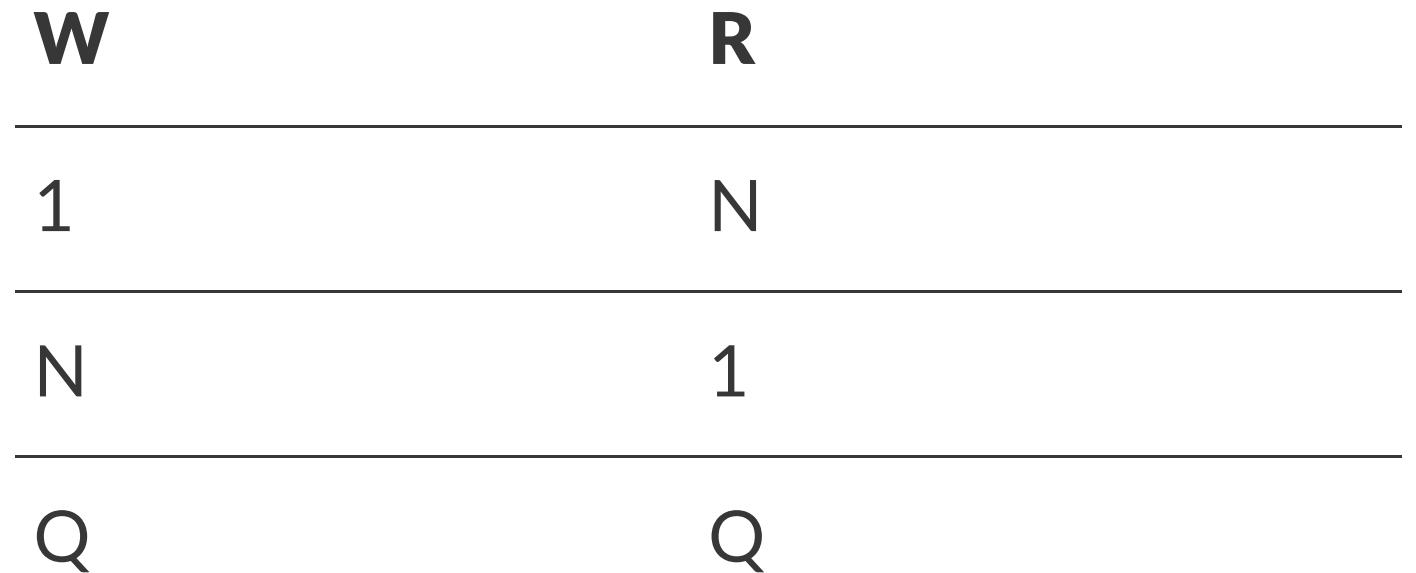
The most recent data is sent to the client and asynchronously, the coordinator identifies any replicas that return obsolete data and issues a read-repair request to each of these replicas to update their data based on the latest data.

---

<sup>84</sup> [source](#)

# Consistency Level: Quorum

- N is replication factor
- R is read replica count,
- W is write replica count
- Quorum  $Q = N/2 + 1$
- If  $W+R > N$  and  $W > N/2$ , you have consistency
- Allowed:



# Consistency Level: Explained

Reads

- Wait for R replicas (R specified by clients)
- In background check for consistency of remaining N-R replicas

Writes

- **Block** until quorum is reached
- **Async**: Write to any node

# Deletes

- Delete: don't delete item right away
  - add a tombstone to the log
  - Compaction will remove tombstone and delete item



# Digression Time

# The Data Structure That Powers Your Database<sup>85</sup>

---

<sup>85</sup> Chapter 3 - Designing Data Intensive Applications

# Log

A log is an append-only sequence of records. It doesn't have to be human-readable;

log-structured storage segments are typically a sequence of key-value pairs.

These pairs appear in the order that they were written, and values later in the log take precedence over values for the same key earlier in the log.

*Commit Log*

Offset	0	1	2	3	4	5	6	7	8
key	$k_1$	$k_2$	$k_1$	$k_3$	$k_4$	$k_5$	$k_5$	$k_2$	$k_6$
Value	$v_2$	$v_3$	$v_4$	$v_5$	$v_6$	$v_7$	$v_8$	$v_9$	$v_{10}$

## Sorted String Table (SSTable)

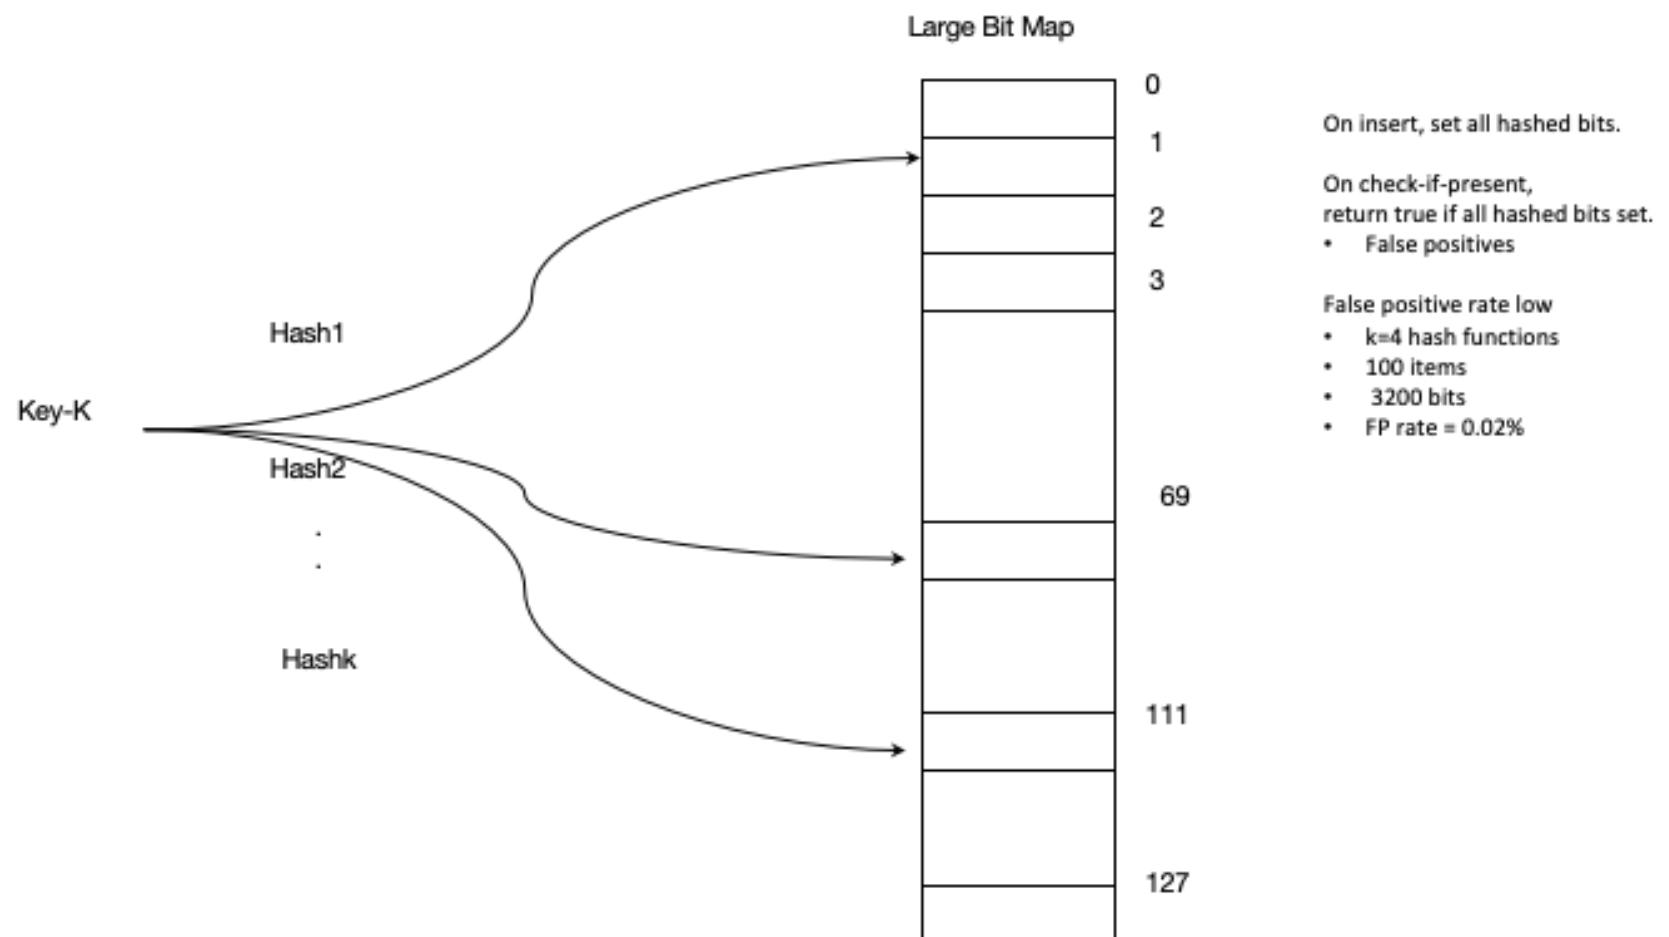
Make a simple change to logs: sequence of key-value pairs is sorted by key.

Merging segments is simple and efficient, even if the files are bigger than the available memory (mergesort algorithm).

In order to find a particular key in the file, you no longer just need a spare index of the offsets

# Bloom Filter

- Compact way of representing a set of items
- Checking for existence in set is cheap
- Some probability of false positives: an item not in set may check true as being in set
- Never false negatives



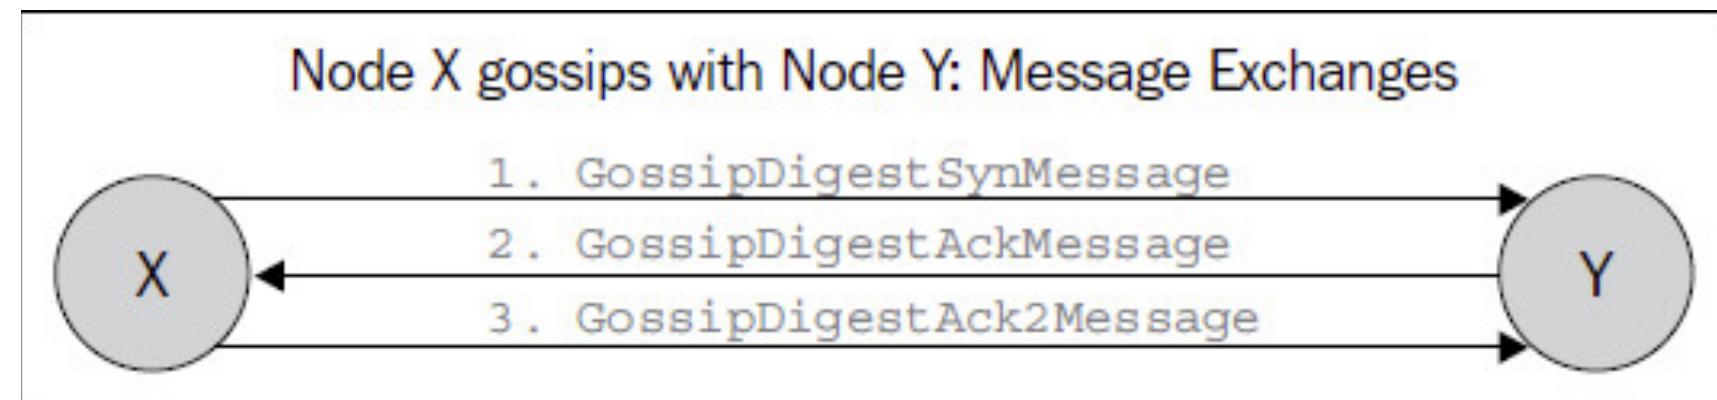
</Digression Time>

# Cluster Membership

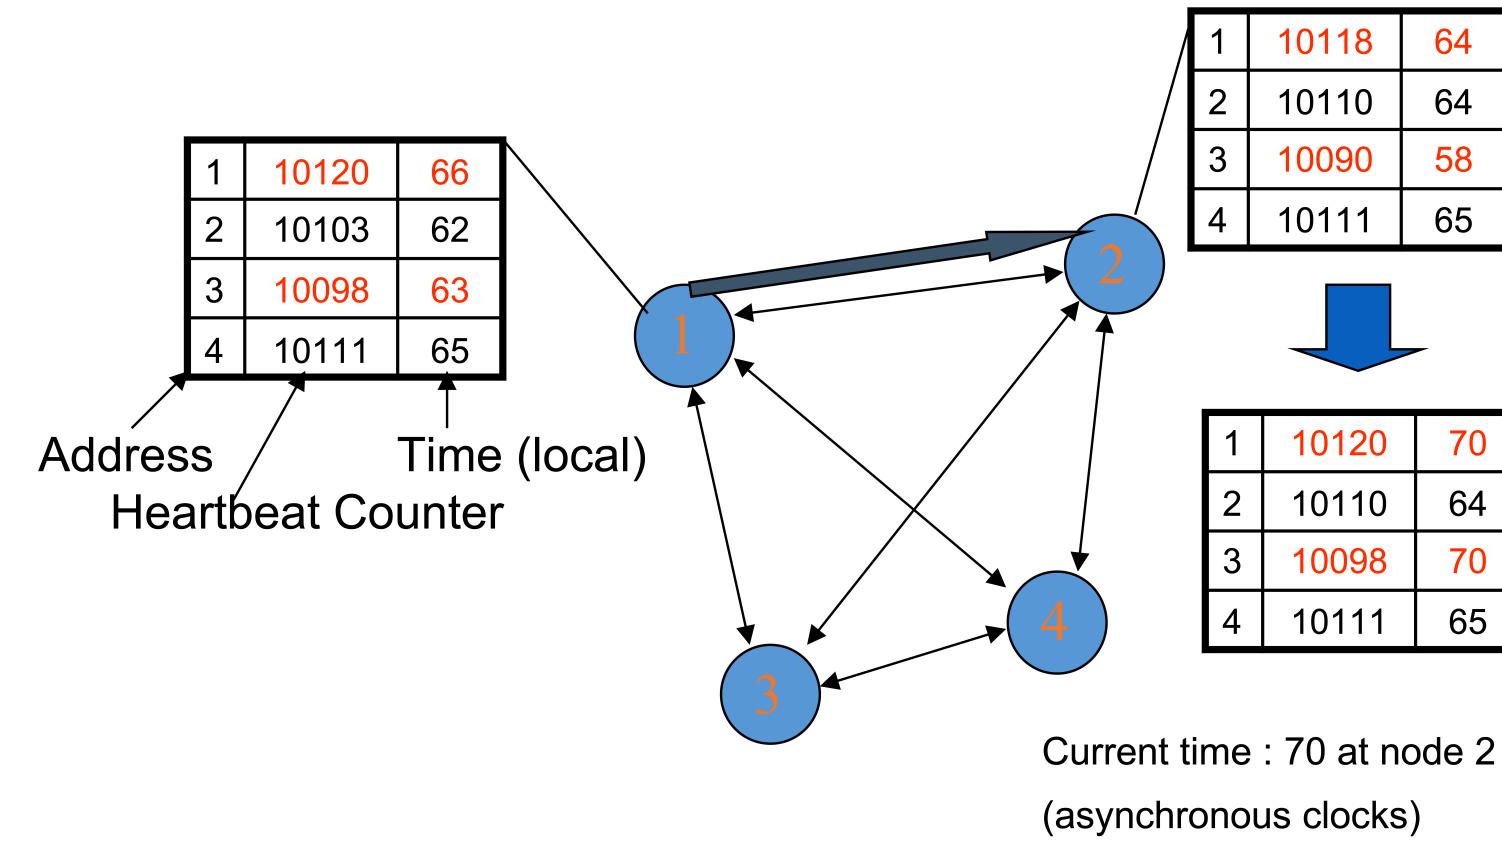
- Any server in cluster could be the coordinator
- So every server needs to maintain a list of all the other servers that are currently in the server
- List needs to be updated automatically as servers join, leave, and fail

# Gossip Protocol

- Each node picks its discussants (up to 3)
- Having three messages for each round of gossip adds a degree of *anti-entropy* .
- This process allows obtaining "**convergence**" of data shared between the two interacting nodes much faster.
- Always a constant amount of network traffic (except for gossip storms)



# Gossip Protocol in practice



- regulates cluster membership
- Nodes periodically gossip their membership list
- On receipt, the local membership list is updated

# Cluster Membership, contd.

- Suspicion mechanisms
- Accrual detector: FD outputs a value (PHI) representing suspicion
- Apps set an appropriate threshold
- $\text{PHI} = 5 \Rightarrow 10\text{-}15 \text{ sec detection time}$
- PHI calculation for a member
  - Inter-arrival times for gossip messages
  - $\text{PHI}(t) = -\log(\text{CDF or Probability}(t_{\text{now}} - t_{\text{last}})) / \log 10$
  - PHI basically determines the detection timeout, but is sensitive to actual inter-arrival time variations for gossiped heartbeats

# Queries

Values in Cassandra are addressed by the triple (row-key, column-key, timestamp) with column- key as

- column-family:column (for simple columns contained in the column family)
- column-family: supercolumn:column (for columns subsumed under a supercolumn).

what about... CQL?

SELECT WHERE  
ORDER BY  
JOIN ON  
GROUP

# SELECT WHERE

Column Family: USER

Key: UserID

Columns: username, email, birth date, city, state

How to support this query?

```
SELECT * FROM User WHERE city = 'Scottsdale'
```

Create a new columns family called **UserCity**:

Column Family: USERCITY

Key: city

Columns: IDs of the users in that city.

Also uses the Valueless Column pattern

## SELECT WHERE pt 2

- Use an aggregate key
  - **state:city: { user1, user2}**

Get rows between **AZ:** & **AZ;** for all Arizona users

Get rows between **AZ:Scottsdale** & **AZ:Scottsdale1** for all Scottsdale users

# ORDER BY

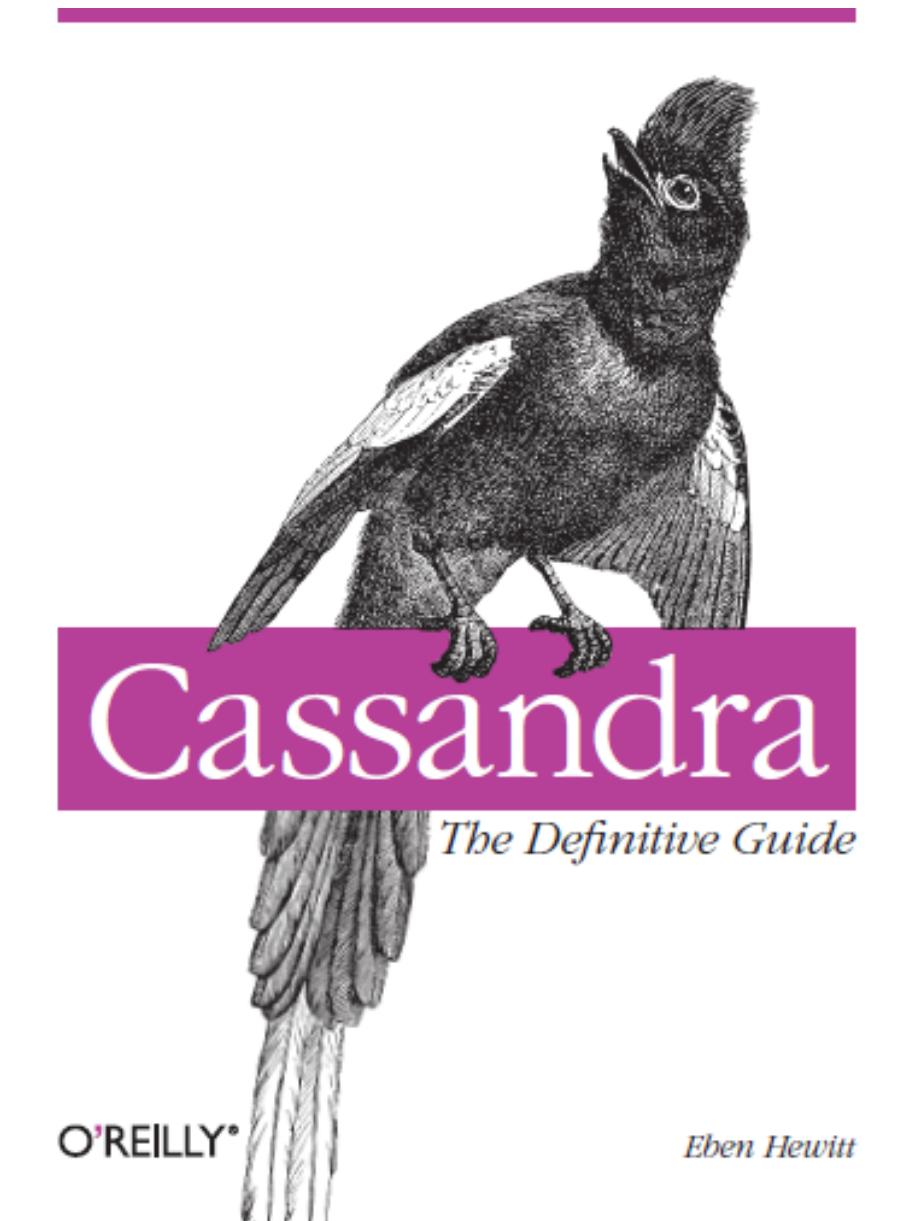
## Columns

are sorted according to  
`CompareWith` or  
`CompareSubcolumnsWith`

## Rows

- are *sorted* by key, regardless of partitioner
- are *placed* according to their Partitioner:
  - Random: MD5 of key
  - Order-Preserving: actual key

# References



## Extra (10)

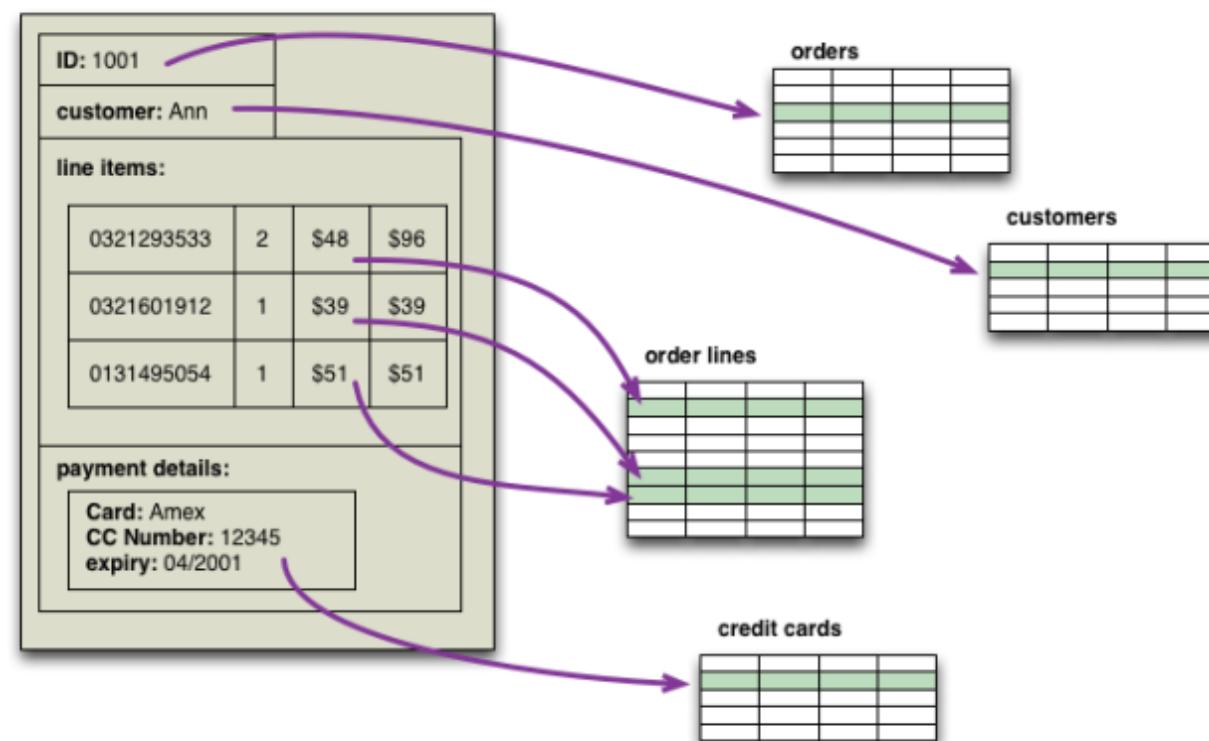
Prepare Cassandra Practice

hints

# Document Databases

# Why document-based?

- Handles Schema Changes Well (easy development)
- Solves Impedance Mismatch problem
- Rise of JSON
- python module: simplejson



# What is a document?

```
{  
  "business_id": "rncjoVoEFUJGUoC1JgnUA",  
  "full_address": "8466 W Peoria AvenSte 6nPeoria, AZ 85345",  
  "open": true,  
  "categories": ["Accountants", "Professional Services", "Tax Services"],  
  "city": "Peoria",  
  "review_count": 3,  
  "name": "Peoria Income Tax Service",  
  "neighborhoods": [],  
  "longitude": -112.241596,  
  "state": "AZ",  
  "stars": 5.0,  
  "latitude": 33.58186700000003,  
  "type": "business":  
}
```

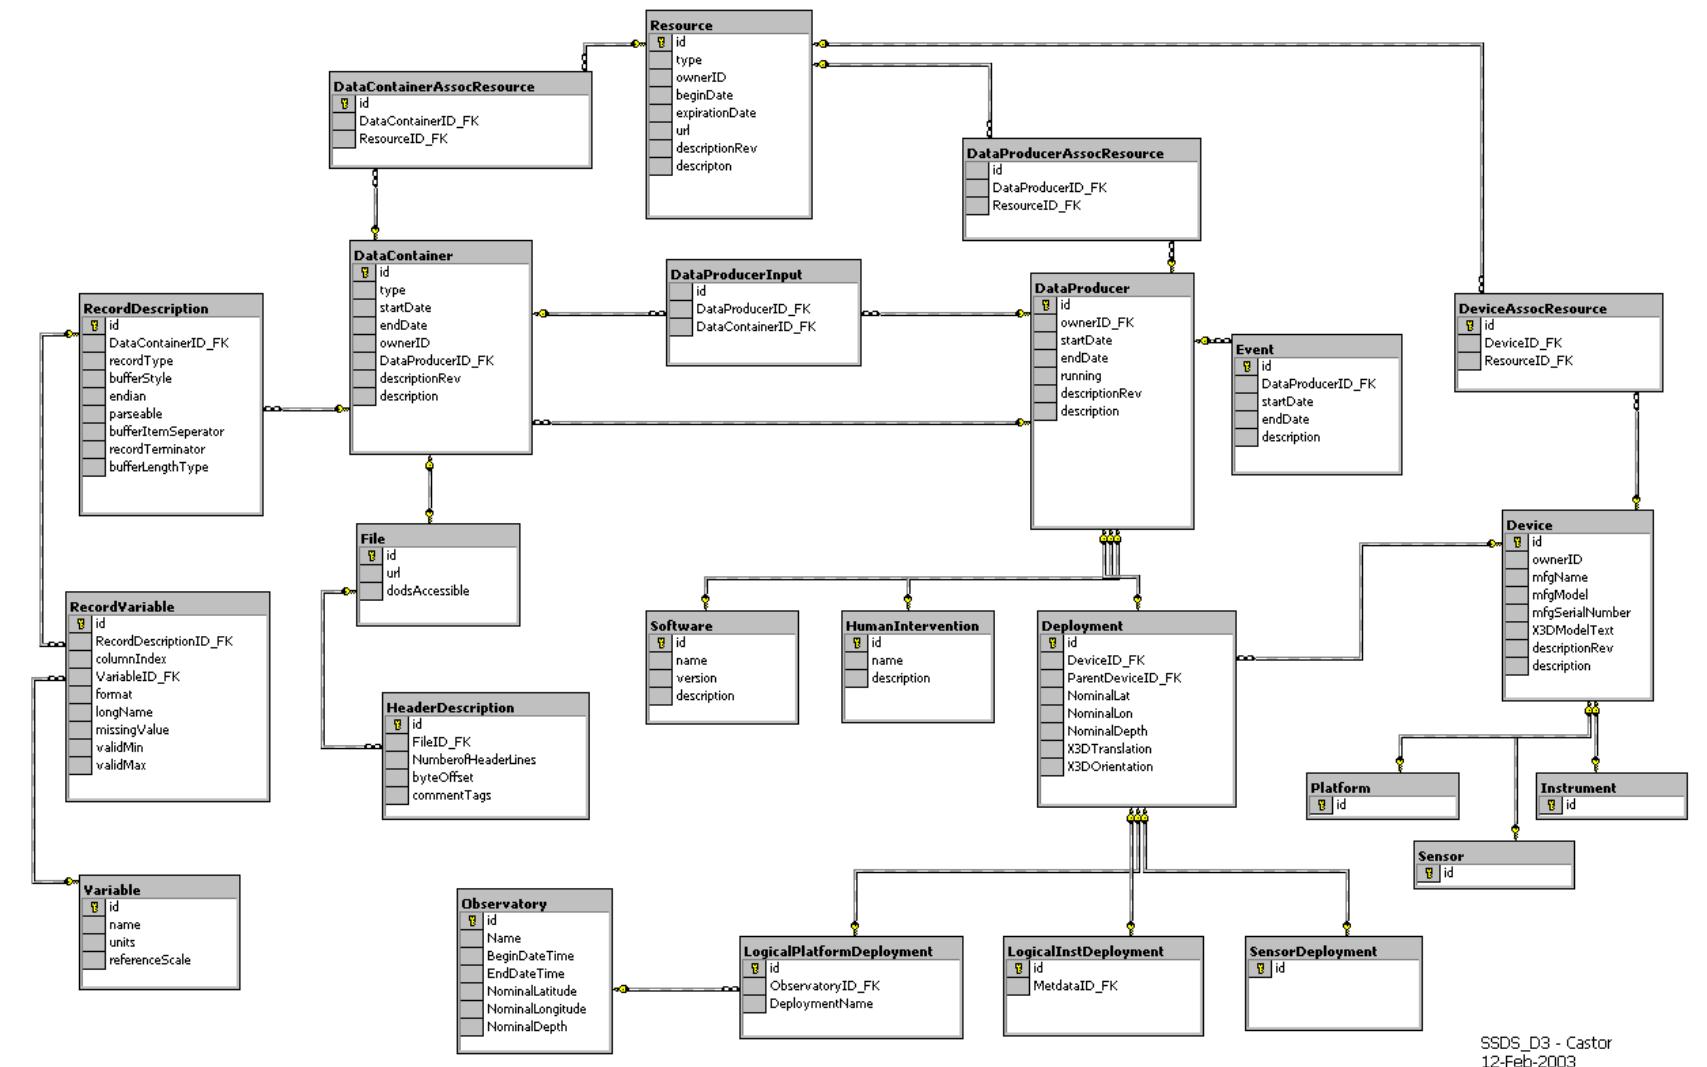
# JSON Format

- Data is in name / value pairs
- A name/value pair consists of a field name followed by a colon, followed by a value:
  - Example: "name": "R2-D2"
- Data is separated by commas
  - Example: "name": "R2-D2", race : "Droid"
- Curly braces hold objects
  - Example: {"name": "R2-D2", race : "Droid", affiliation: "rebels"}
- An array is stored in brackets []
  - Example [ {"name": "R2-D2", race : "Droid", affiliation: "rebels"}, {"name": "Yoda", affiliation: "rebels"} ]

# Designing NoSQL Data Structures

- NoSQL data structures driven by application design.
  - Need to take into account necessary CRUD operations
- To embed or not to imbed. That is the question!
  - Rule of thumb is to imbed whenever possible.
  - No modeling standards or CASE tools!

# Relational to Document



```
{
  "title" : "MongoDB",
  "contributors": [
    { "name" : "Eliot Horowitz",
      "email" : "eliot@10gen.com" },
    { "name" : "Dwight Merriman"
      "email" : "dwight@10gen.com" } ],
  "model" : {
    "relational" : false,
    "awesome" : true
  }
}
```

## A normalized structure

```
{  
  "_id"      : "First Post",  
  "author"   : "Rick",  
  "text"     : "This is my first post."  
  
}  
  
{  
  "_id" : ObjectId(...),  
  "post_id" : "First Post",  
  "author" : "Bob",  
  "text"  : "Nice Post!"  
}
```

# A (denormalized) embedded structure

```
{  
  "_id" : "First Post",  
  "comments" : [  
    { "author" : "Bob",  
      "text" : "Nice Post!"},  
    { "author" : "Tom",  
      "text" : "Dislike!"}],  
  "comment_count" : 2  
}
```

# A polymorphic structure

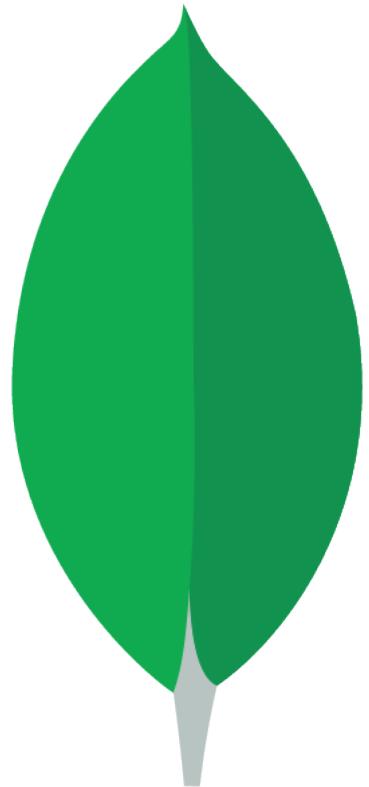
- When all the documents in a collection are similarly, but not identically structured.
- Enables simpler schema migration.
- no more of custom \_field\_1
- Better mapping of object - oriented inheritance and polymorphism.

```
{  
  "_id" : 1,  
  "title": "Welcome",  
  "url": "/",  
  "type": "page",  
  "content": "Welcome to my wonderful wiki."  
}  
  
{  
  "_id": 3,  
  "title": "Cool Photo",  
  "url": "/photo.jpg",  
  "type": "photo",  
  "content": Binary(...)  
}
```

# List of Systems

- **MongoDB**
- CouchDB
- OrientDB

# MongoDB



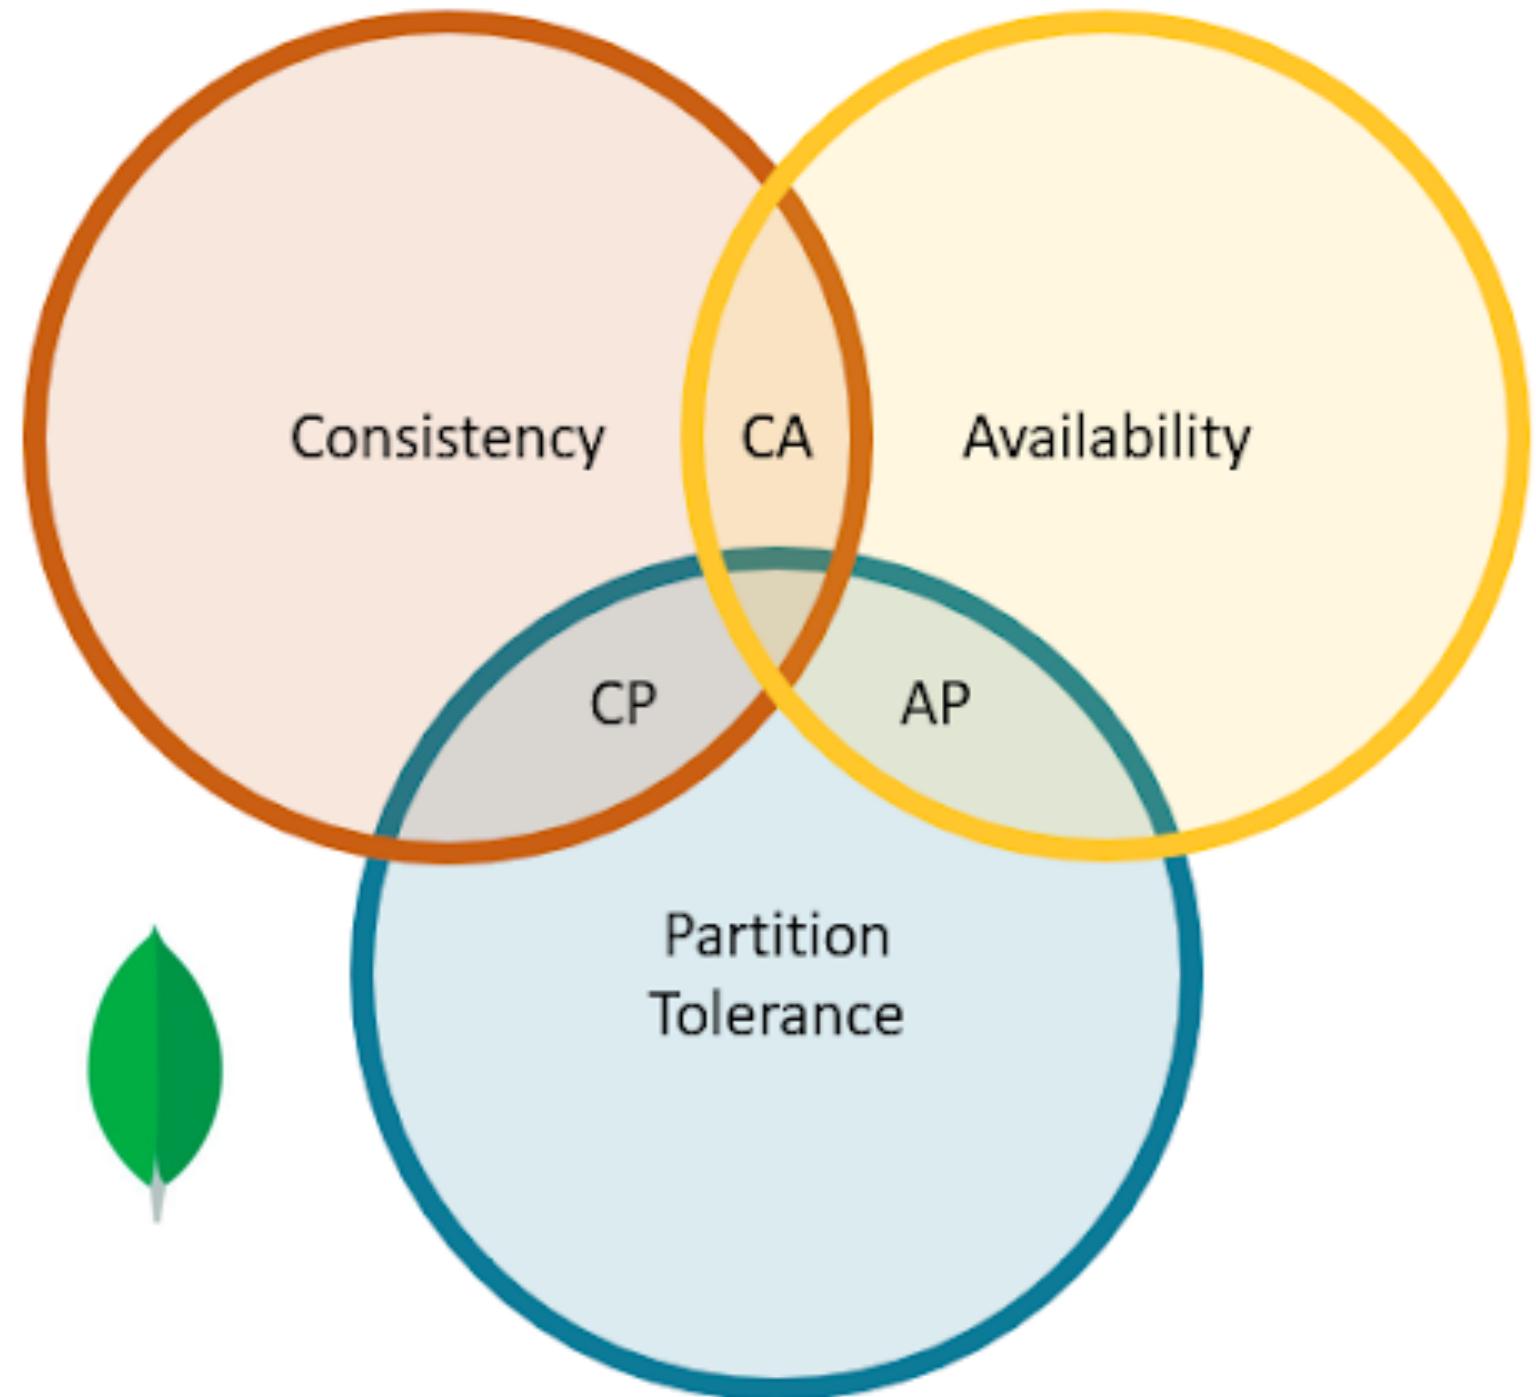
# mongoDB®

# History and Motivation

- An open source and document-oriented database.
- Data is stored in JSON-like documents.
- Designed with both scalability and developer agility.
- Dynamic schemas.
- Automatic data sharding

# What MongoDB is :

- In-Memory
- Strong consistency (**C**)
- *Tuneably* available (**A**)
- Horizontal Scalable (**P**)



# What MongoDB is not

- Always Available<sup>91</sup>
- No Schemas
- No transactions
- No joins
- Max document size of 16MB<sup>92</sup>

---

<sup>91</sup> Larger documents handled with GridFS

<sup>92</sup> there will always be downtime when (i) the new leader is getting elected or (ii) the client driver disconnects from the leader

# Use Cases

-  Capture **game** events, scaling to meet high-write workloads.
-  Financial Services: Risk Analytics & Reporting, Trade Repository
-  **BOSCH** manufacturing, automotive, retail, and energy
-  **ThermoFisher** SCIENTIFIC fast-changing sensor data captured from multiple devices and experiments

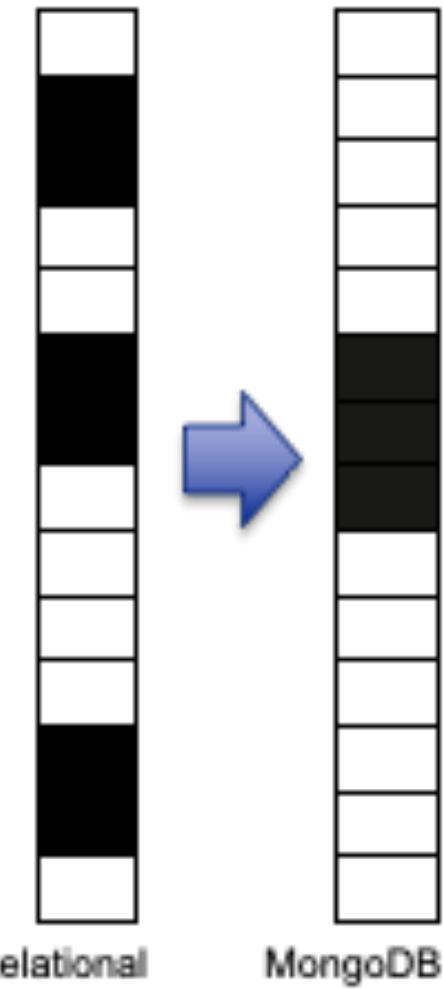
# When to consider MongoDB

- When you need high availability of data
- when you need fast and instant data recovery
- when do not want to sustain schema migration costs

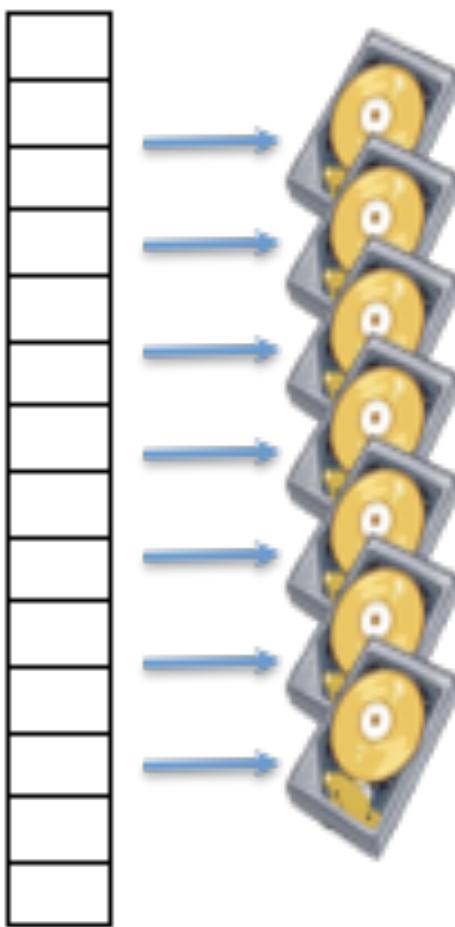
# Advantages

- Full featured indexes
- Sophisticated query language
- Easy mapping to object oriented code
- Native language drivers in all popular language
- Simple to setup and manage
- Operates at in-memory speed wherever possible
- Auto-sharding built in
- Dynamically add / remove capacity with no downtime

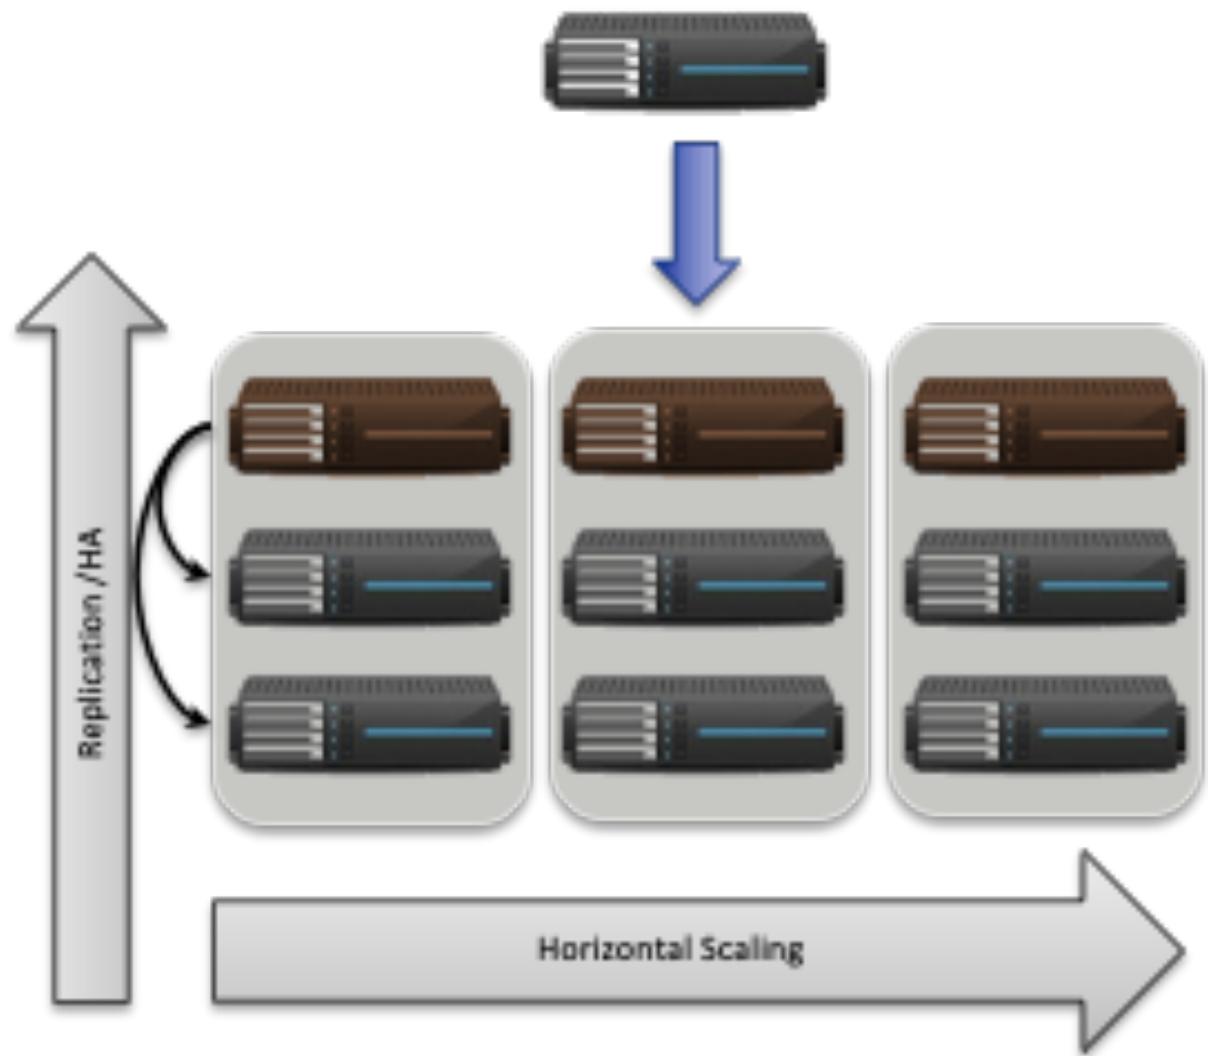
Better data locality



In-Memory  
Caching



Distributed Architecture



# Terminology: SQL vs MongoDB

## SQL Terms/Concepts

database

table

row

column

index

table joins (e.g. select queries)

Primary keys

Aggregation (e.g. group by)

## MongoDB Terms/Concepts

database

collection

document

field

index

embedded documents and linking

\_id field is always the primary key

aggregation pipeline

# Data Model

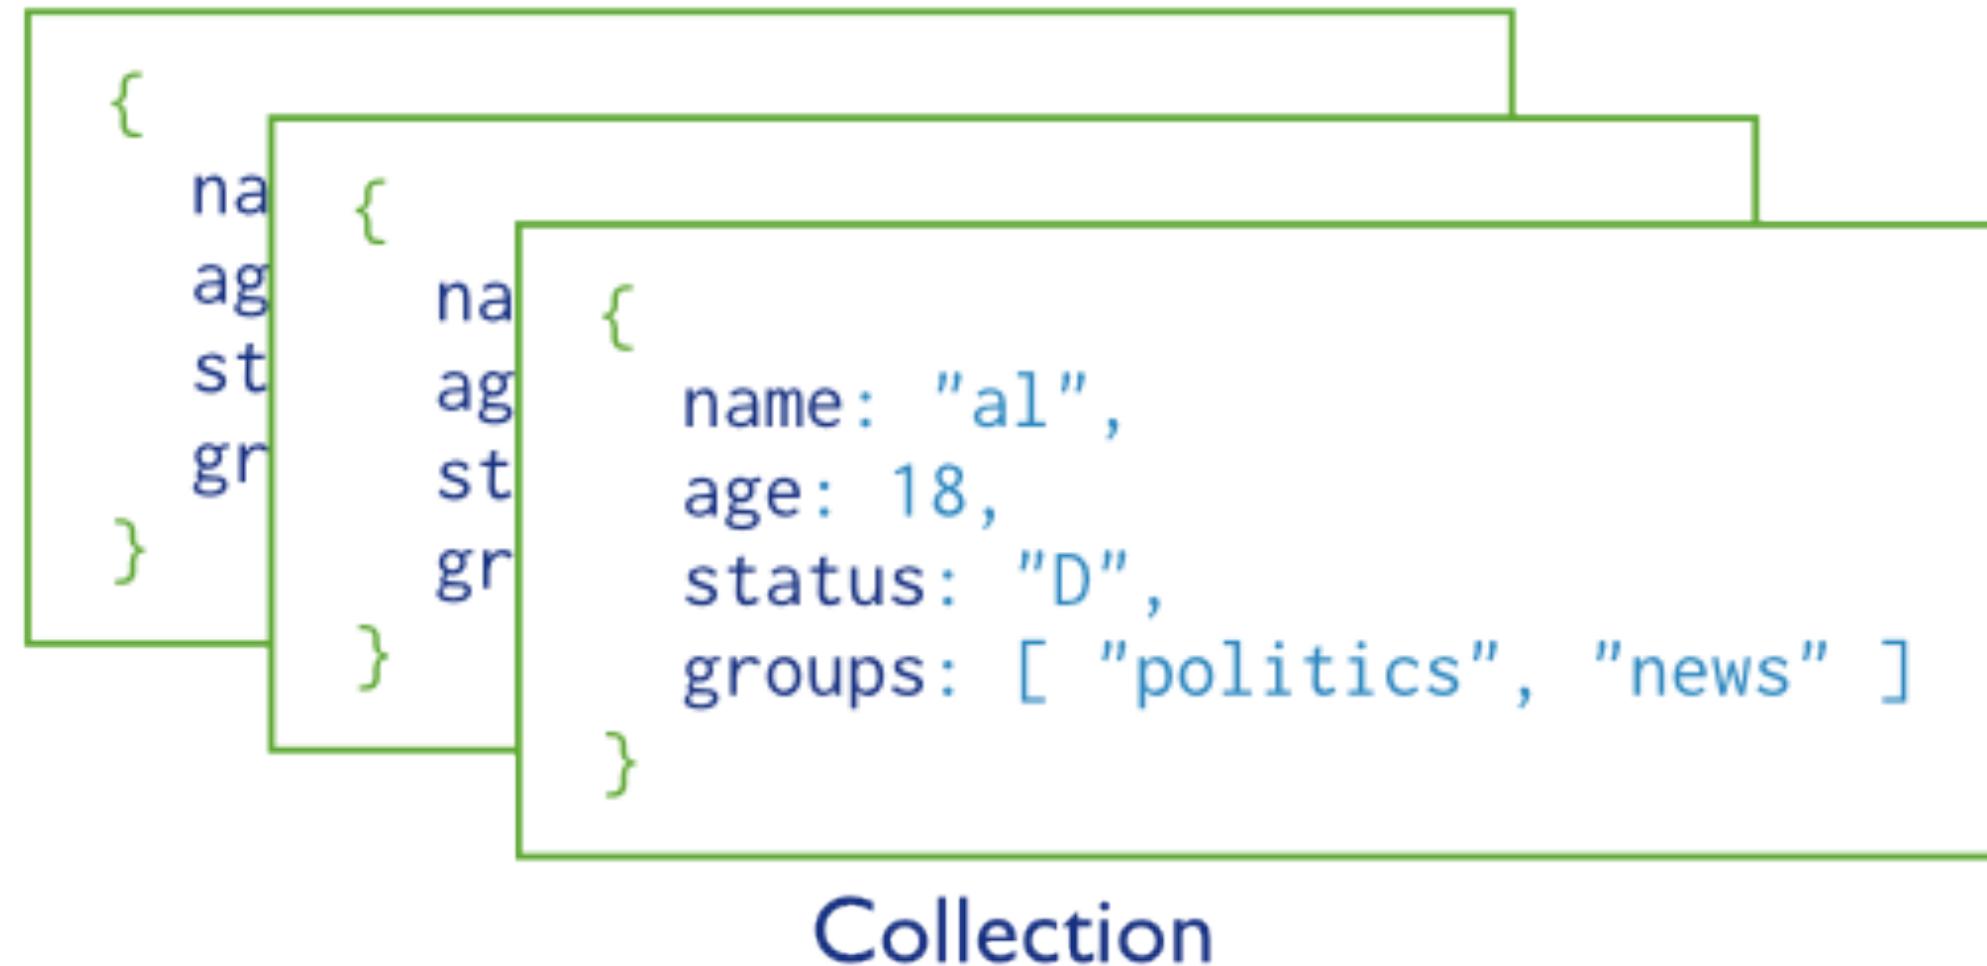
## Structure of a JSON-document:

```
{  
    name: "sue",  
    age: 26,  
    status: "A",  
    groups: [ "news", "sports" ]  
}  
  
← field:value  
← field:value  
← field:value  
← field:value
```

The value of field:

- Native data types
- Arrays
- Other documents

# Collections of Documents



Rule: Every document must have an `_id`.

# Embedded documents:

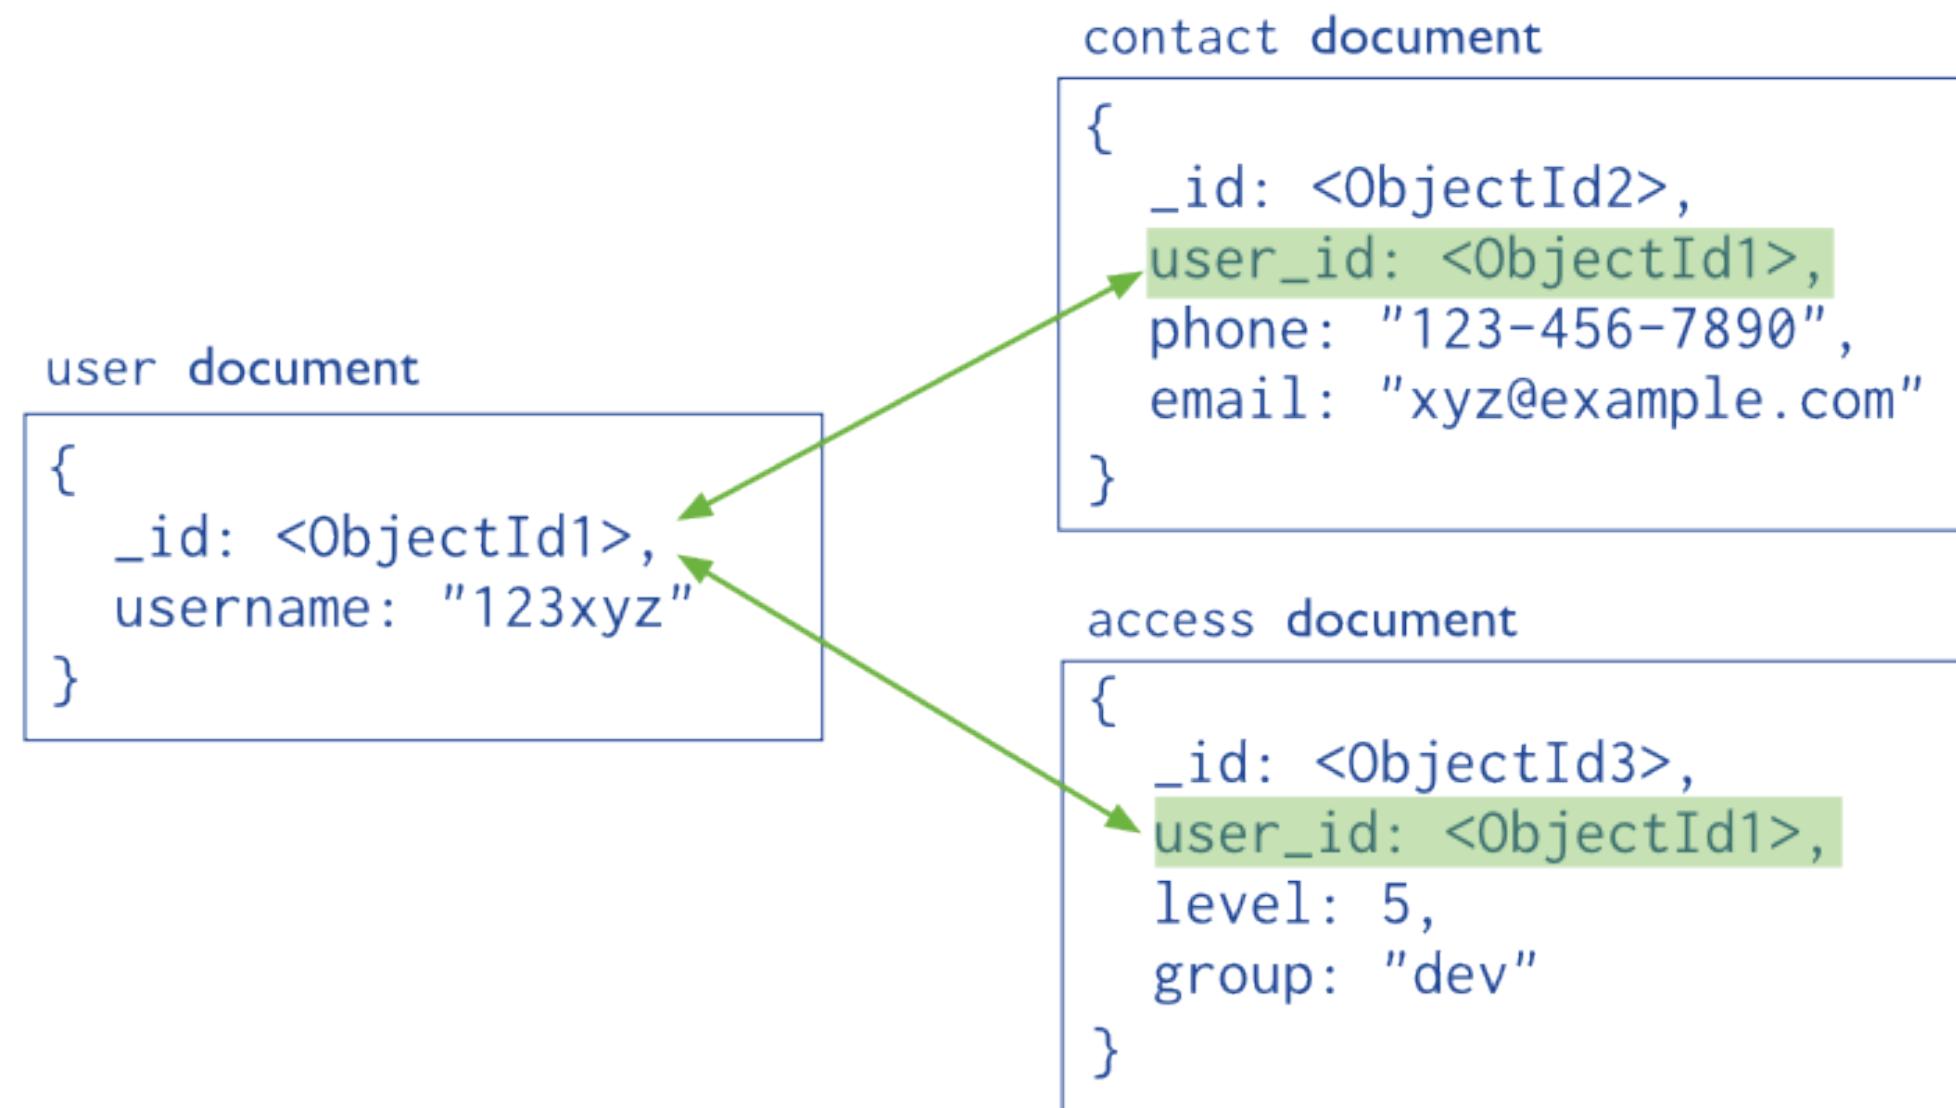
```
{  
  _id: <ObjectId1>,  
  username: "123xyz",  
  contact: {  
    phone: "123-456-7890",  
    email: "xyz@example.com"  
  },  
  access: {  
    level: 5,  
    group: "dev"  
  }  
}
```



Embedded sub-document

Embedded sub-document

# Reference documents



## Storage: BSON Format

- Binary-encoded serialization of JSON-like documents optimized for space and speed
- BSON types are a superset of JSON types<sup>94</sup>
- Zero or more key/value pairs are stored as a single entity<sup>93</sup>
- Large entities are prefixed with a length field to facilitate scanning

### BSON:

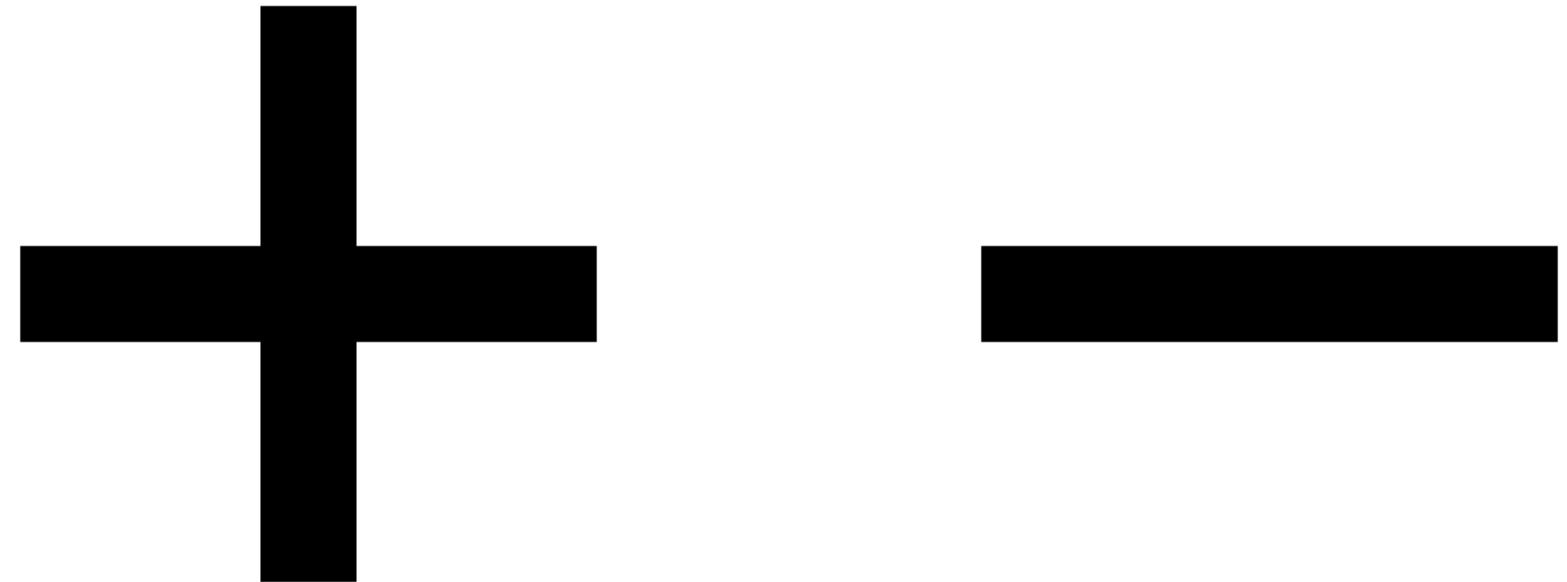
```
\x16\x00\x00\x00  
 \x02  
  name\x00  
  \x06\x00\x00\x00Devang\x00  
  \x00
```

```
// total document size  
// 0x02 = type String  
// field name  
// field value
```

---

<sup>94</sup> JSON does not have a date or a byte array type, for example

<sup>93</sup> Each entry consists of a field name, a data type, and a value



# Operations



# Create

Create a database

```
use database_name
```

Create a collection

```
db.createCollection(name, options)
```

# Insert

```
db.<collection_name>.insert({"name": "nguyen", "age": 24, "gender": "male"})
```

```
db.employee.insert({  
    name: "sally",  
    salary: 15000,  
    designation: "MTS",  
    teams: ["cluster-management"]  
})`
```

# Read

```
db.<collection_name>.find().pretty()  
  
db.employee #collection  
  .find( {  
    salary: {$gt:18000}, #condition  
    {name:1} #projection  
  })  
  .sort({salary:1}) #modifier
```

# Update

```
db.employee #collection
.update(
  {salary:{$gt:18000}}, #Update Criteria
  {$set: {designation: "Manager"}}, #Update Action
  {multi: true} #Update Option
)
```

# Delete

```
db.employee.remove(  
    {salary:{$lt:10000}}, #Remove Criteria  
)
```

# Aggregates

SQL-like aggregation functionality

Pipeline documents from a collection pass through an aggregation pipeline

Expressions produce output documents based on calculations performed on input documents

Example:

```
db.parts.aggregate(  
  {$group : {_id: type, totalquantity :  
    { $sum: quantity}  
  }})
```

## Save

```
db.employee.save(  
  { _id:ObjectId('string_id'),  
    "name": "ben",  
    "age": 23,  
    "gender":  
    "male"  
)
```

## Drop

- Drop a database
- Drop it:  
db.dropDatabase()
- Drop a collection:

db.<collection\_name>.drop()

# Mapping to SQL

## SQL Statement

SELECT \* FROM table

## MongoDB commands

db.collection.find()

SELECT \* FROM table WHERE artist = 'Nirvana' db.collection.find({Artist:"Nirvana"})

SELECT\* FROM table ORDER BY Title

db.collection.find().sort>Title:1

DISTINCT

.distinct()

GROUP BY

.group()

>=, <

\$gte, \$lt

# Comparison Operators

Name	Description
\$eq	Matches value that are equal to a specified value
\$gt, \$gte	Matches values that are greater than (or equal to) a specified value
\$lt, \$lte	Matches values less than or (equal to) a specified value
\$ne	Matches values that are not equal to a specified value
\$in	Matches any of the values specified in an array
\$nin	Matches none of the values specified in an array
\$or	Joins query clauses with a logical OR returns all
\$and	Join query clauses with a logical AND
\$not	Inverts the effect of a query expression
\$nor	Join query clauses with a logical NOR
\$exists	Matches documents that have a specified field

[source](#)

# Indexes

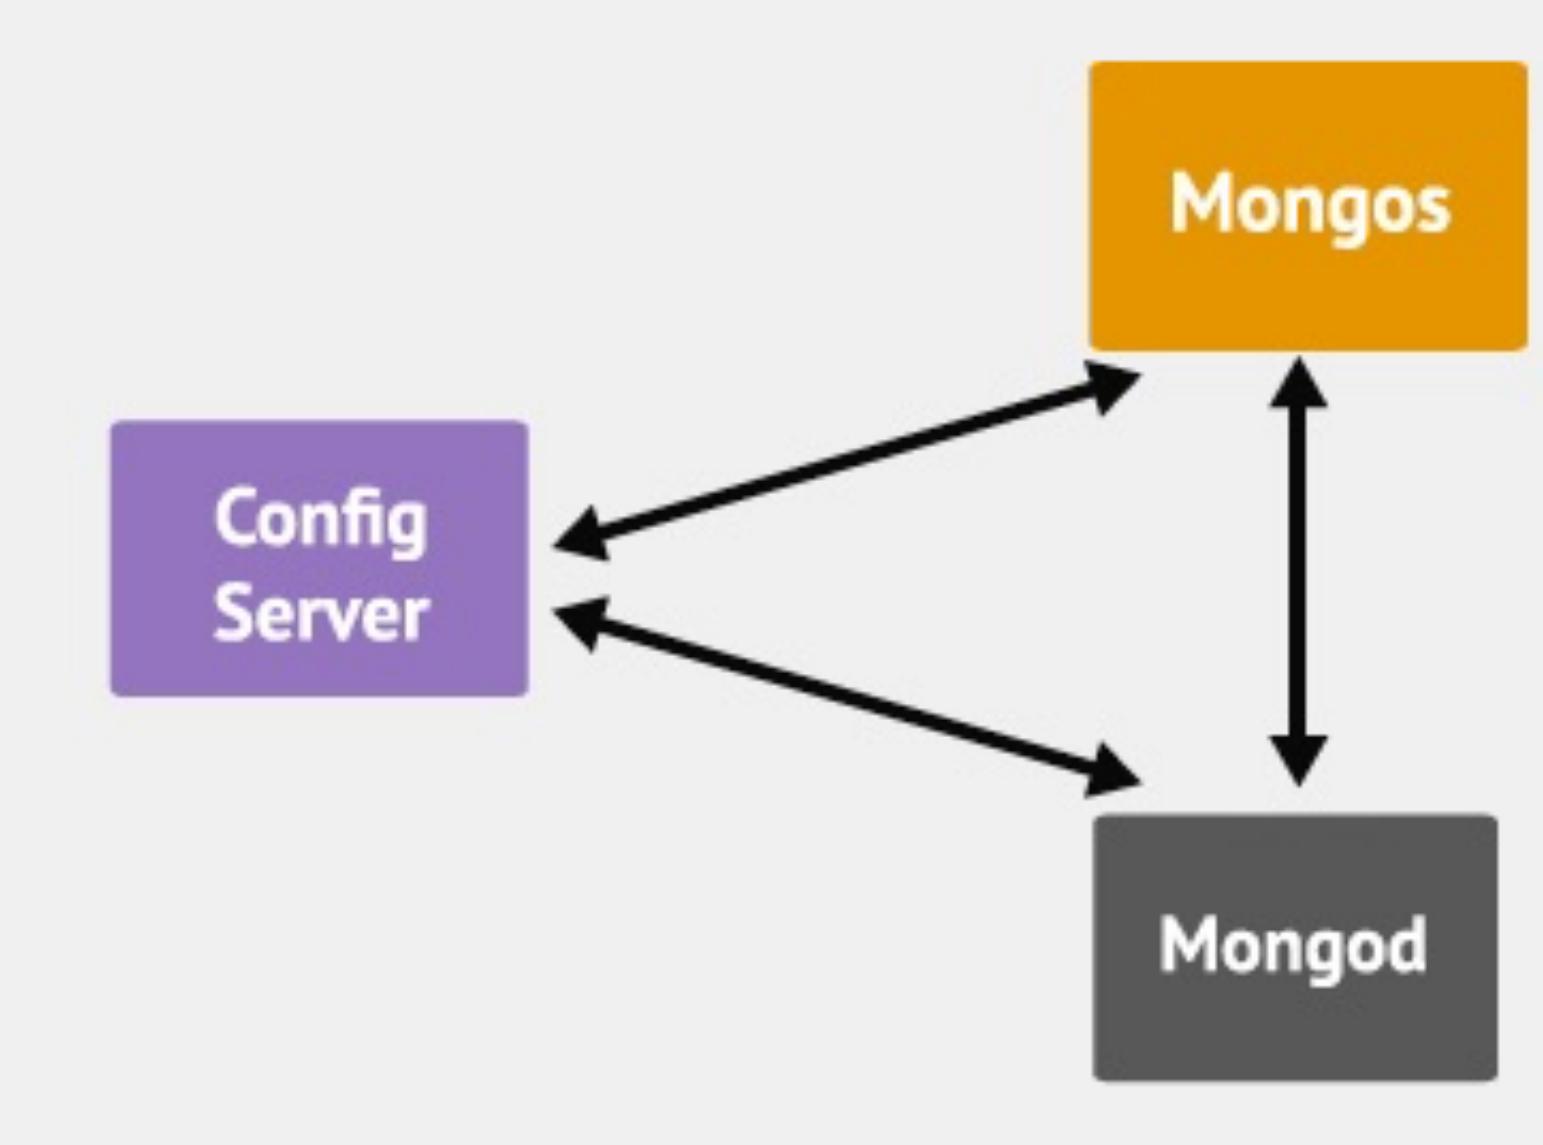
- B+ tree indexes
- An index is automatically created on the `_id` field (the primary key)
- Users can create other indexes to improve query performance or to enforce Unique values for a particular field
- Supports single field index as well as Compound index
- Like SQL order of the fields in a compound index matters
- If you index a field that holds an array value, MongoDB creates separate index entries for every element of the array

# Sparse Indexes

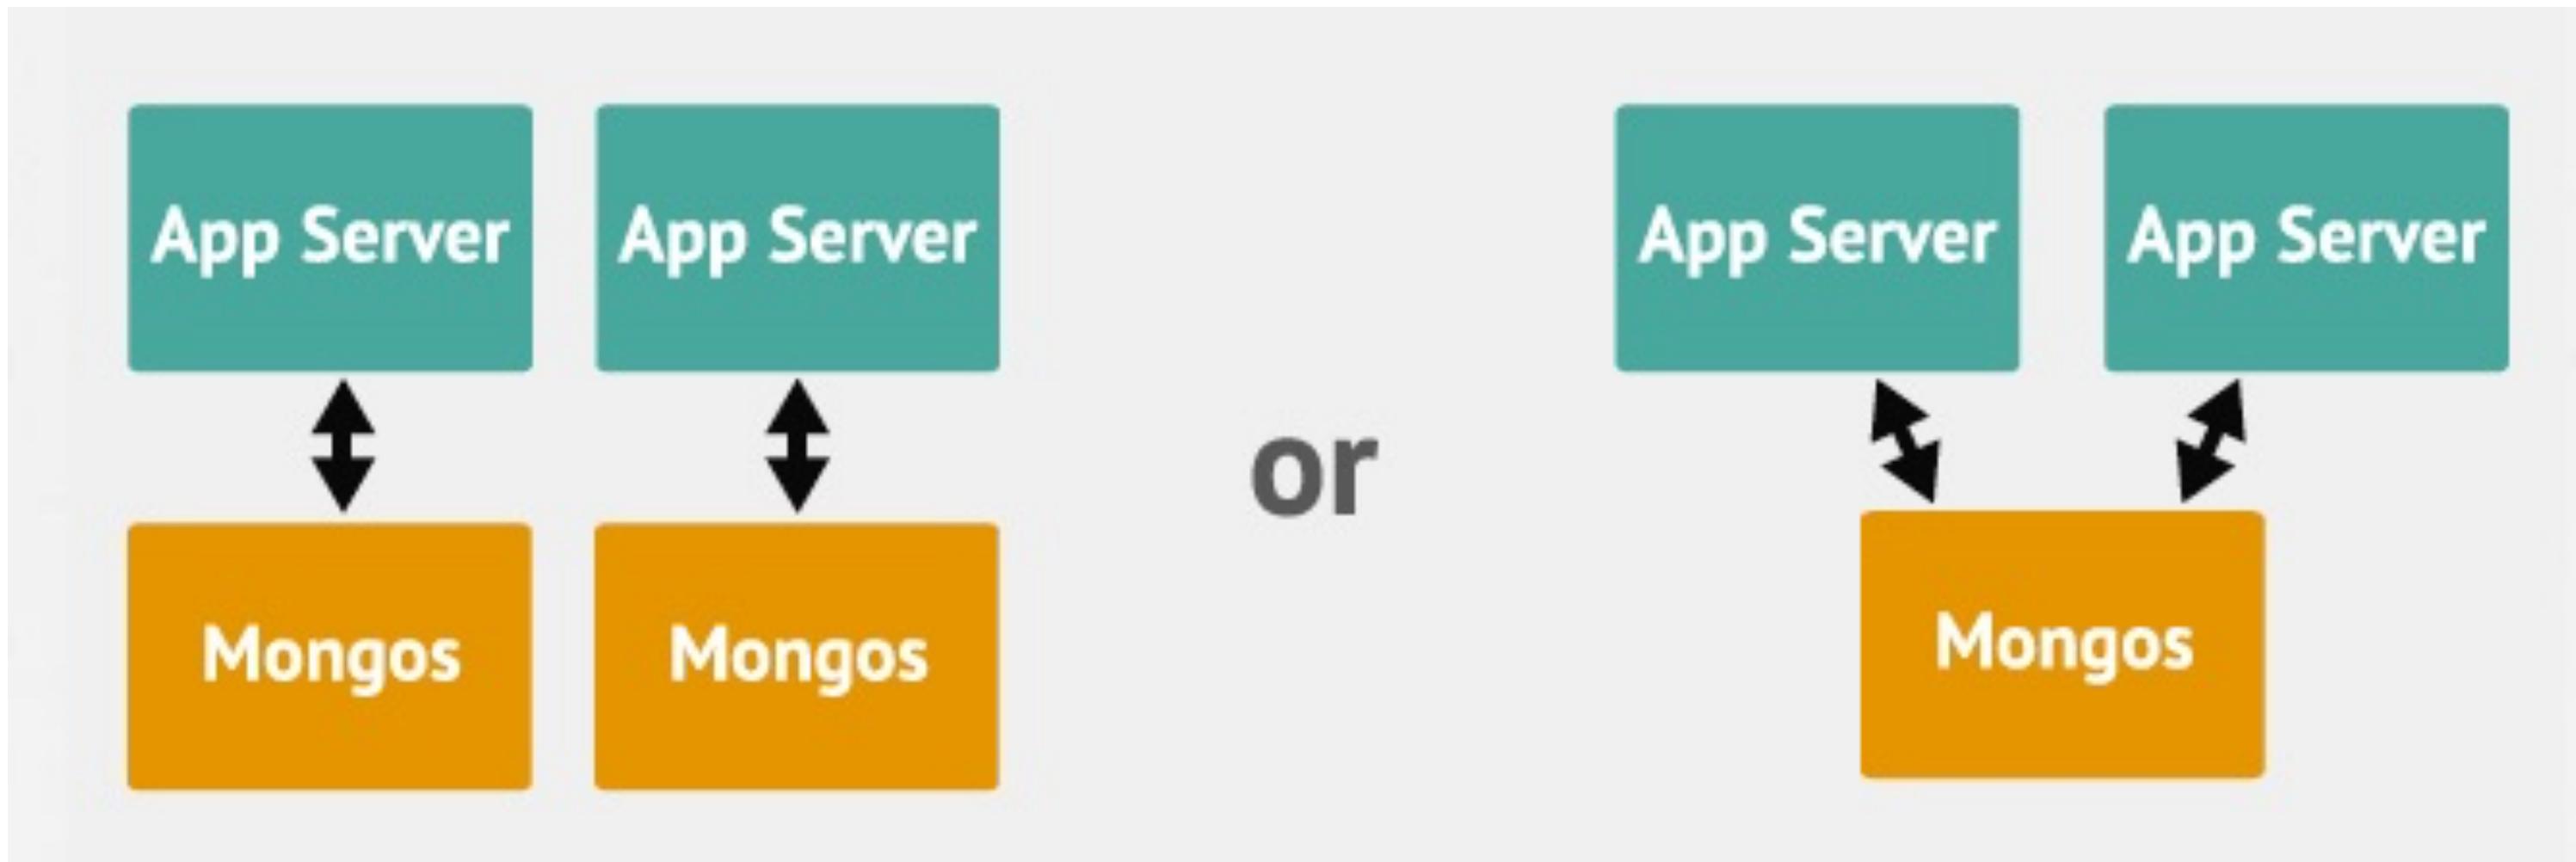
- Sparse\_- property of an index ensures that the index only contain entries for documents that have the indexed field. (so ignore records that do not have the field defined)
- If an index is both unique and sparse – then the system will reject records that have a duplicate key value but allow records that do not have the indexed field defined

# Architecture

- Mongod – Database instance
- Mongos - Sharding processes:
  - Like a database router processes all requests
  - Decides how many and which *mongod* should receive the query
  - No local data
  - Collects the results, and sends it back to the client.
- Config Server
  - Stores cluster chunk ranges and locations
  - Can have only 1 or 3 (production must have 3)



# Mongod and Mongos



# Client

- Mongo – an interactive shell ( a client)
- Fully functional JavaScript environment for use with a MongoDB
- You can have one mongos for the whole system no matter how many mongods you have
- OR you can have one local mongos for every client if you wanted to minimize network latency.

# Replication

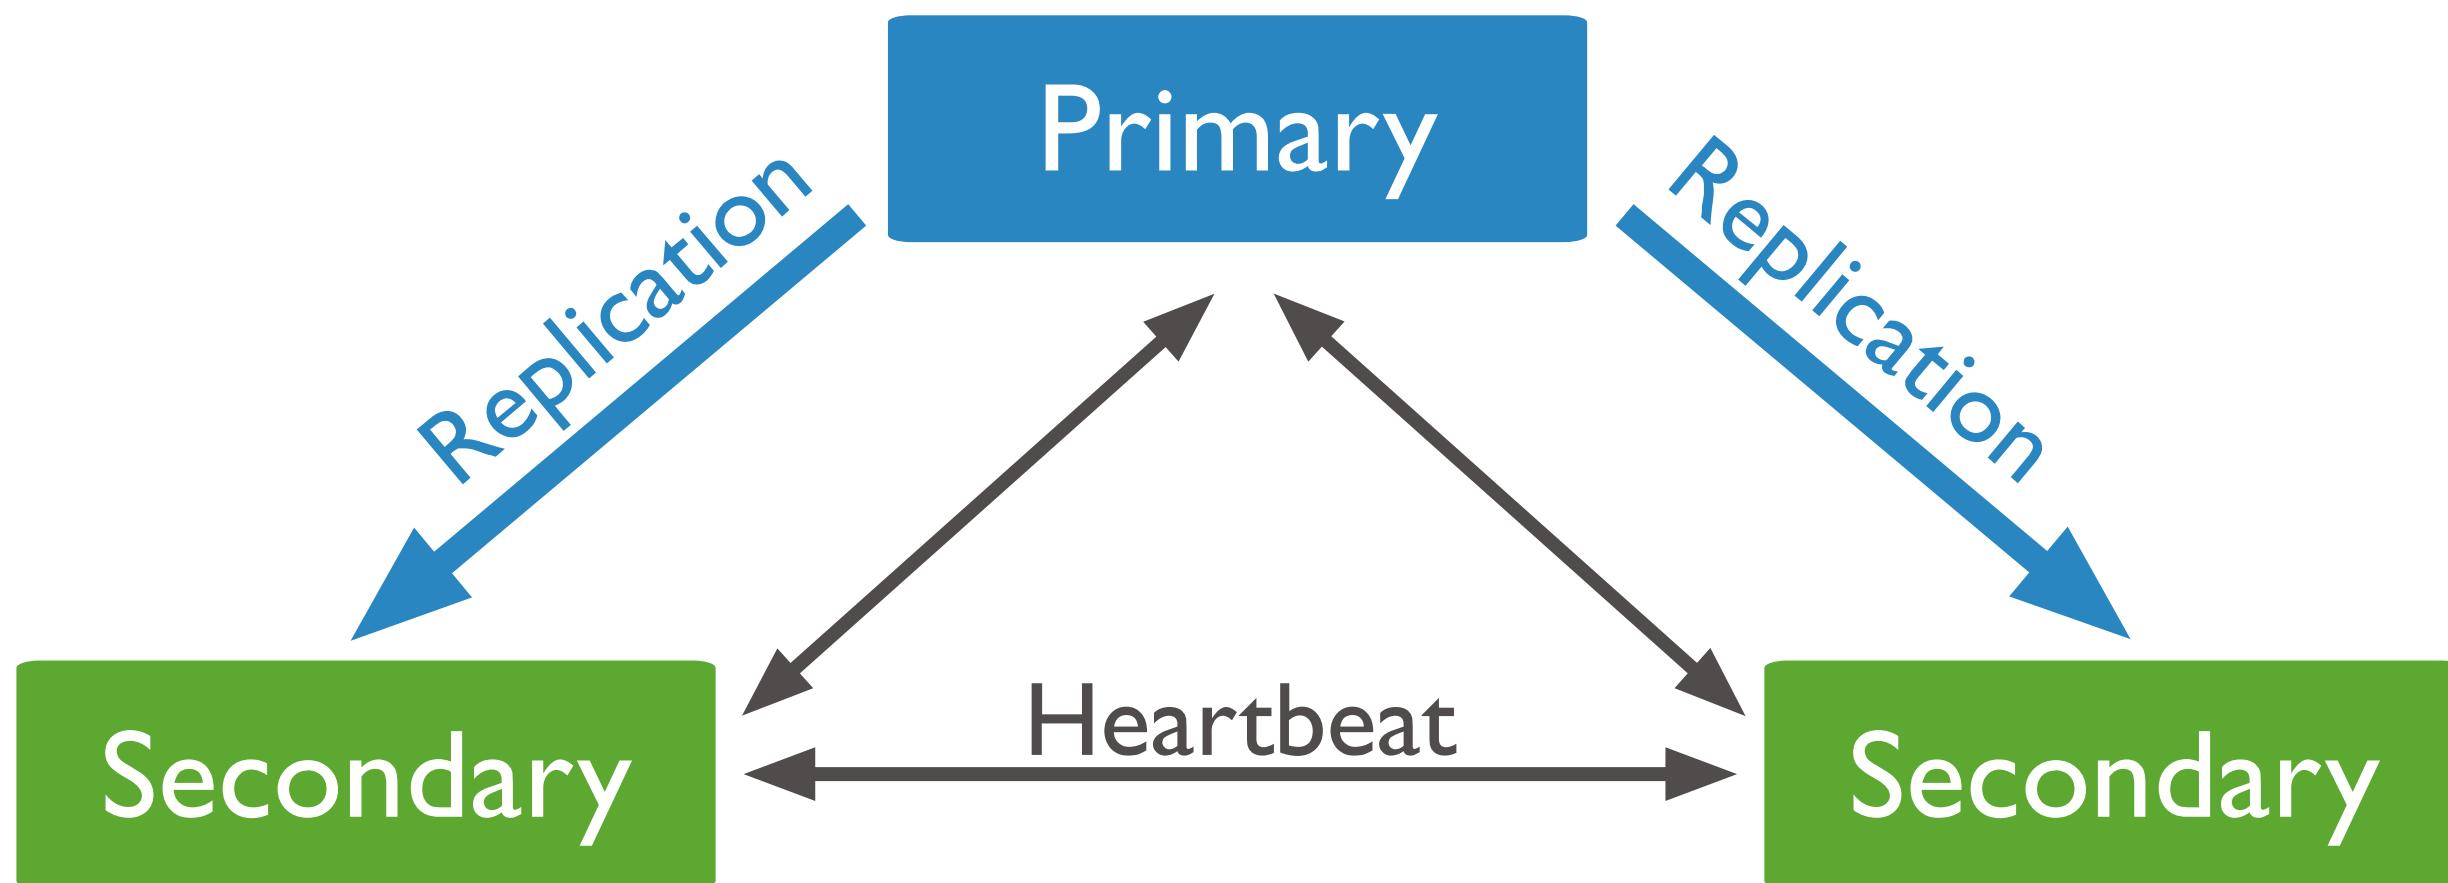
For redundancy MongoDB provides asynchronous replication.

Only one database node is in charge of write operations at any given time (called primary server/node).

Read operations may go to this same server for strong consistency semantics or to any of its replica peers if eventual consistency is sufficient.

# Master Slave Replication

Consists of two servers out of one which takes the role of a master handling write requests and replicating those operations to the second server, the slave.



# Replica Sets

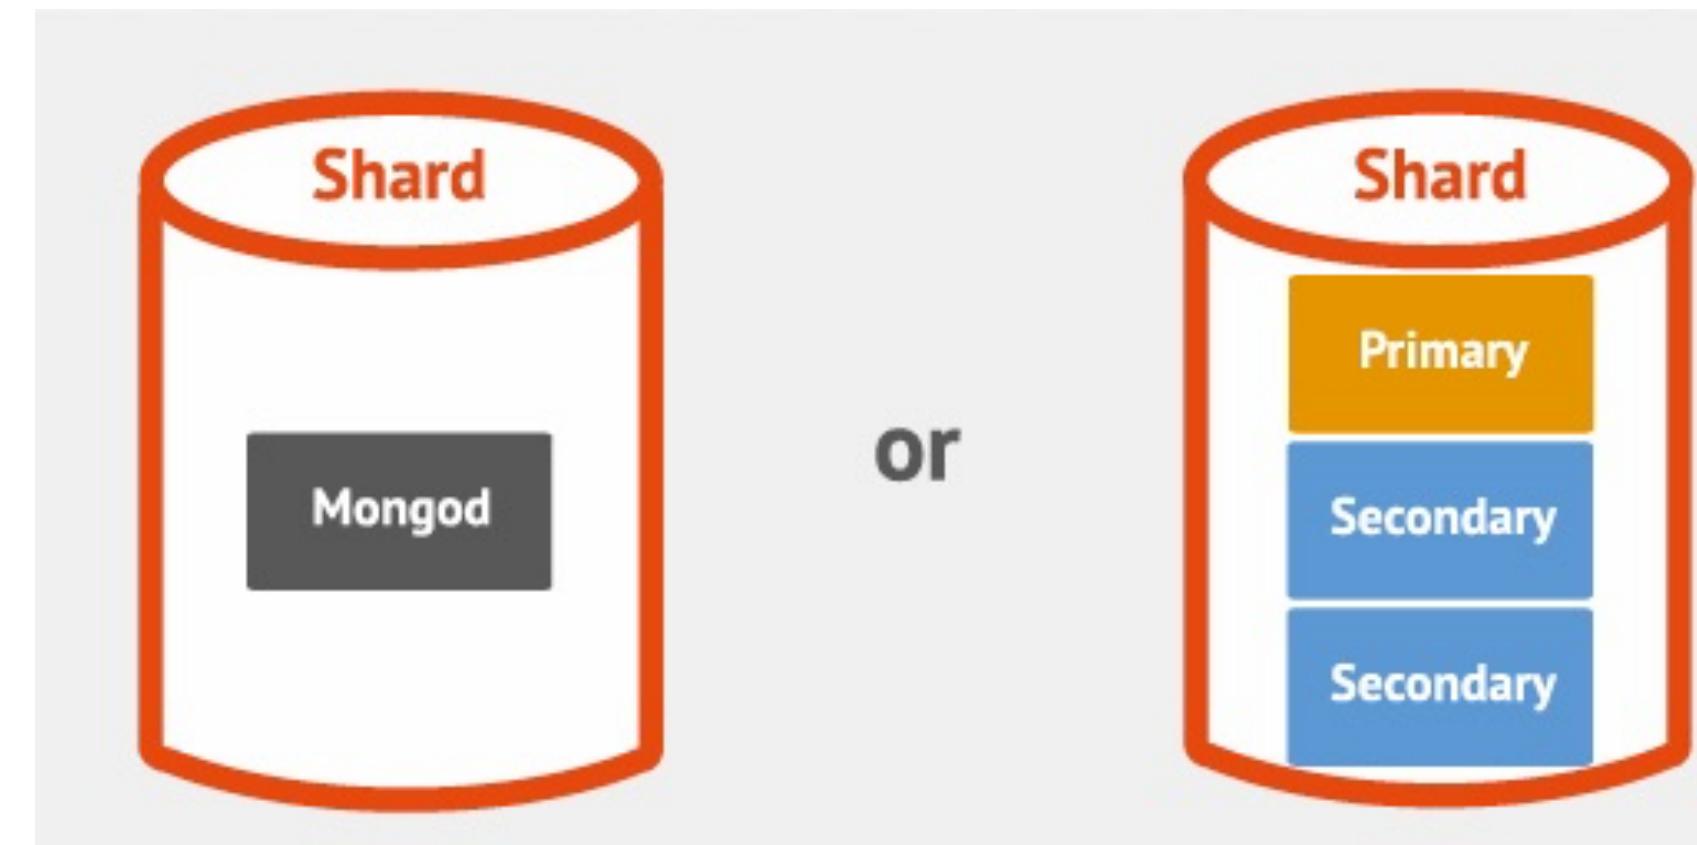
Consists of groups of MongoDB nodes that work together to provide automated failover.

# Partitioning

- called Sharding in MongoDB
- User defines shard key for partitioning
- Shard key defines range of data
- Key space is like points on a line
- Range is a segment of that line

## What is a Shard?

- Shard is a node of the cluster
- Shard can be a single mongod or a replica set
- Default max chunk size: 64mb
- MongoDB automatically splits & migrates chunks when max reached



# Auto-sharding

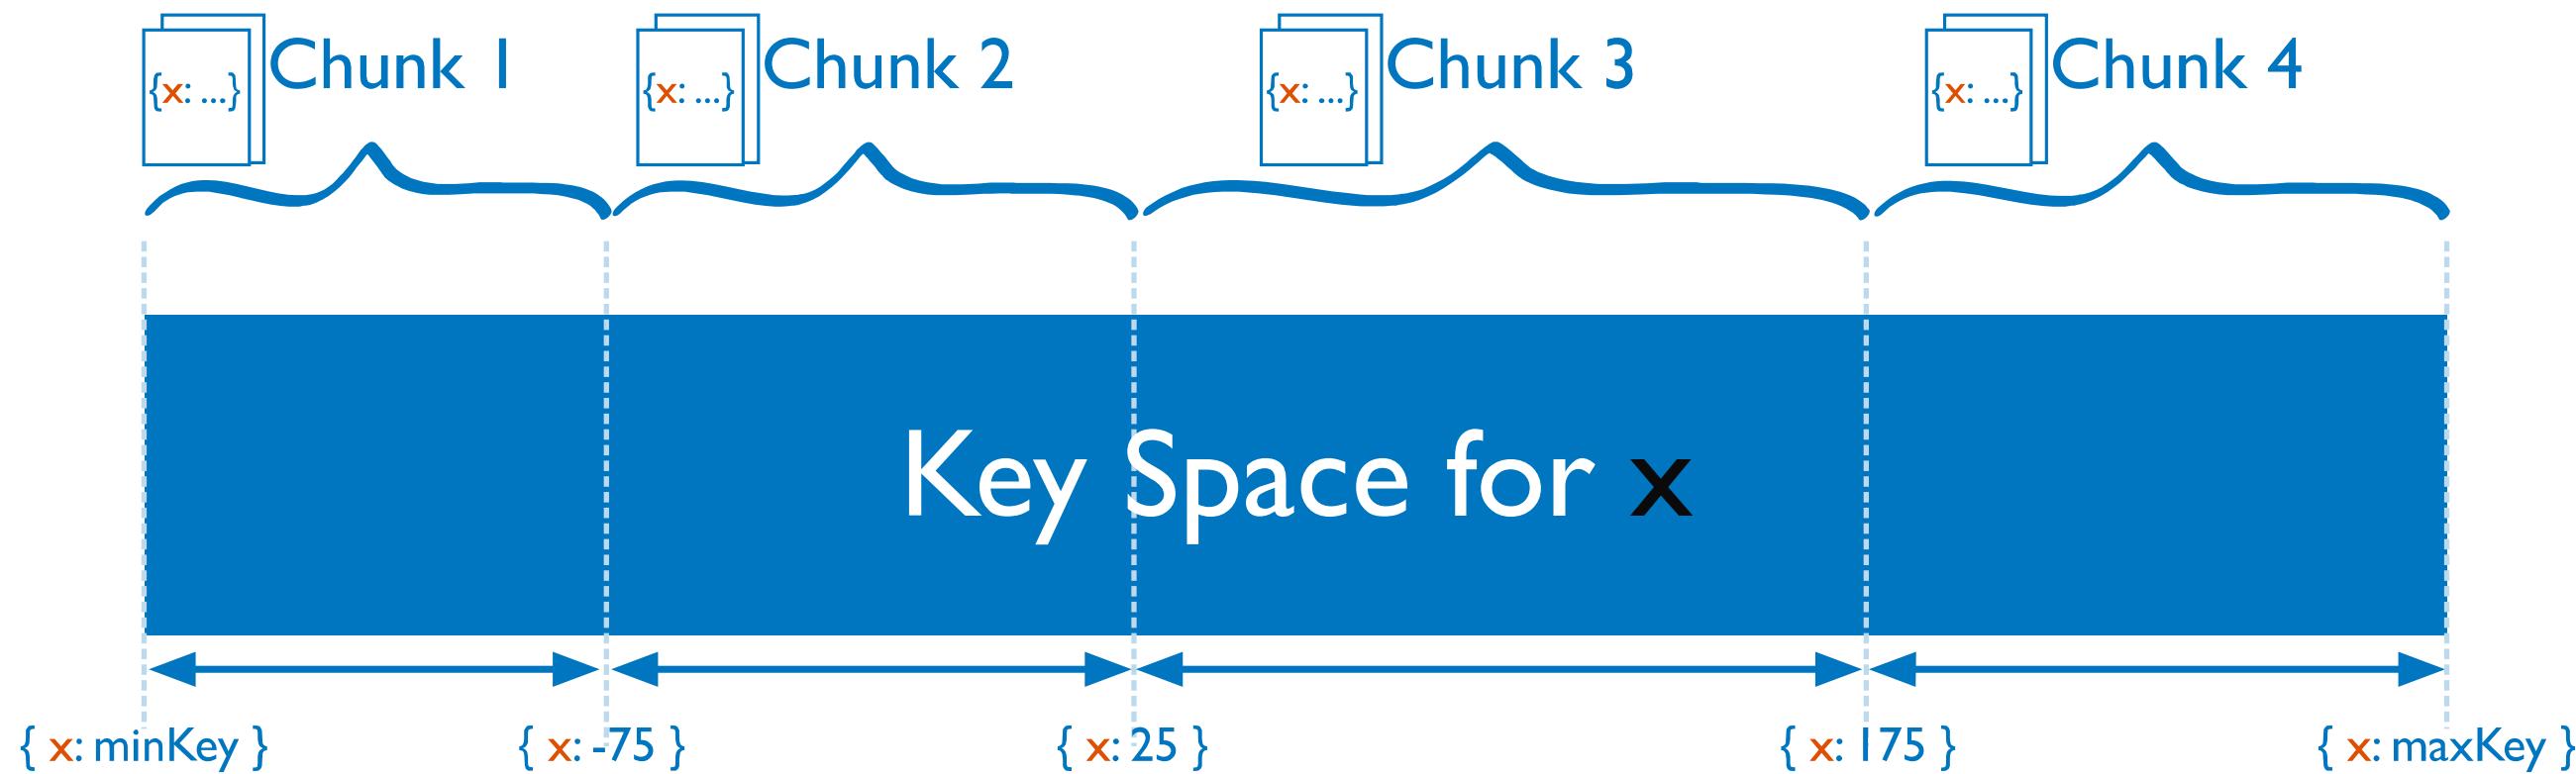
- Minimal effort required
  - Enable Sharding for a database
  - Shard collection within database
  - Decide Sharding Strategy

# MongoDB Sharding Strategies

- Ranged
- Hashed
- Tag-aware

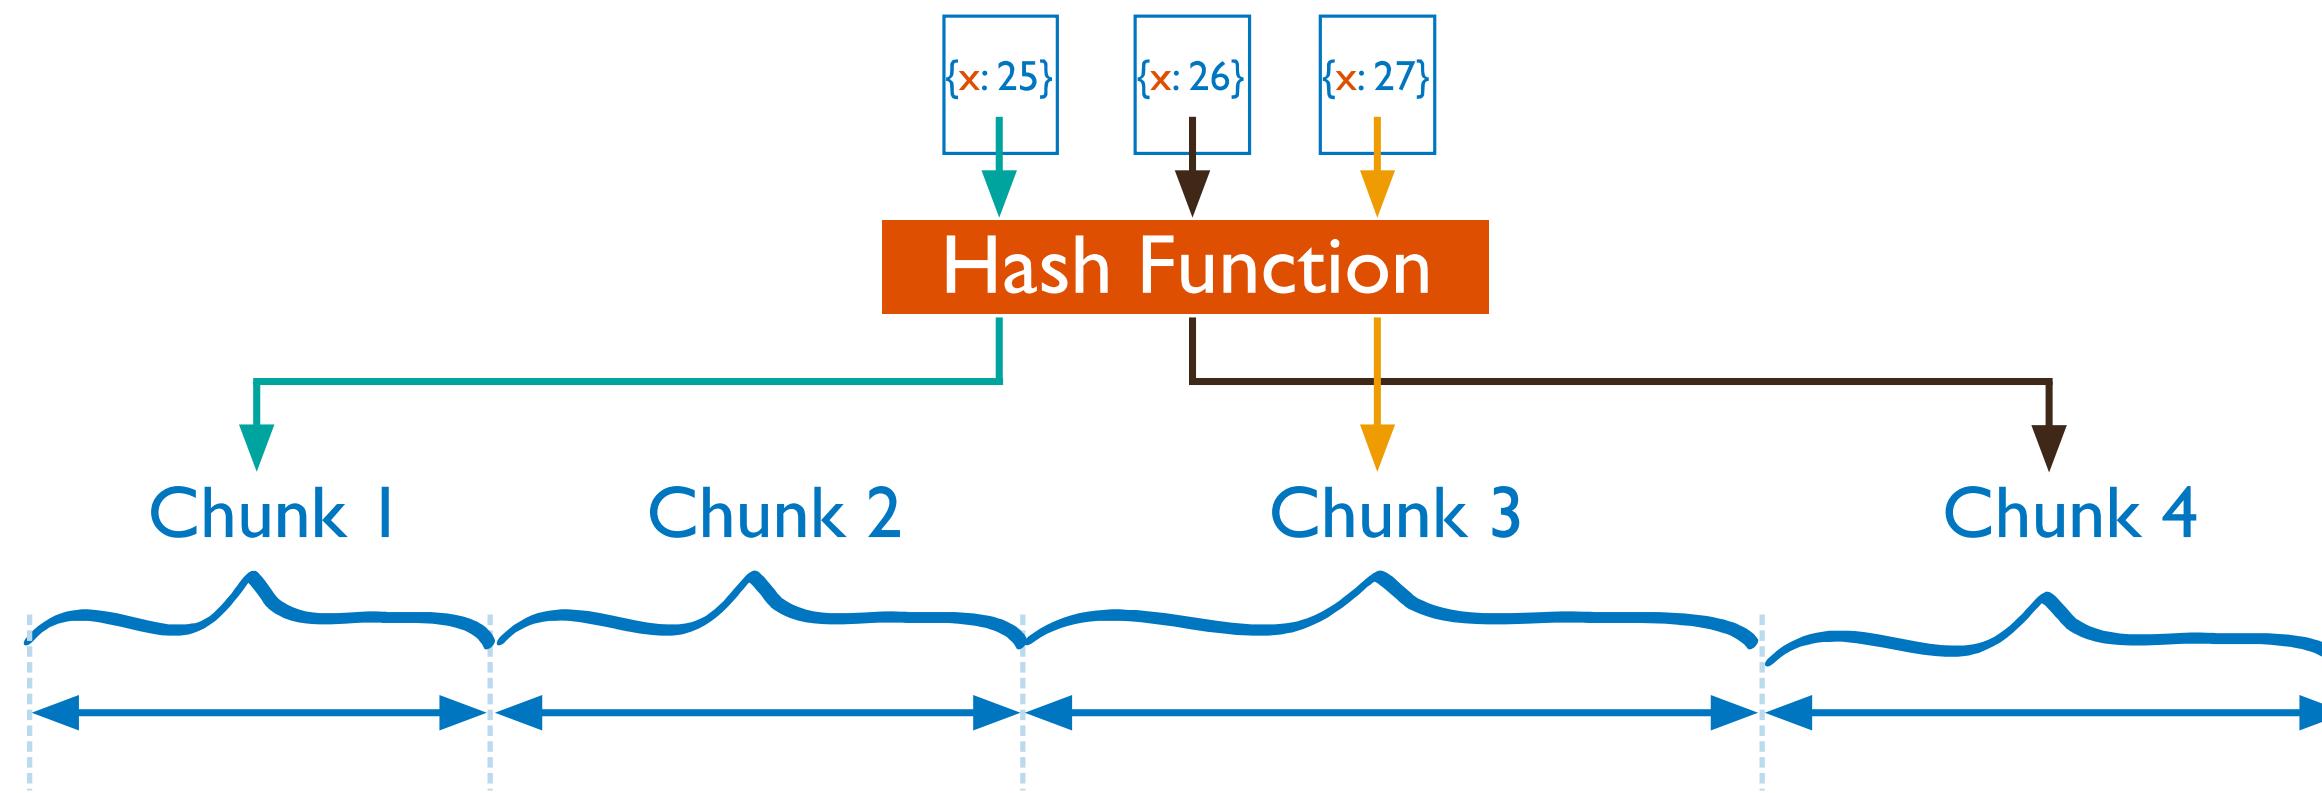
## Range Sharding

- Splits shards based on sub-range of a key (or also multiple keys combined)
  - Simple Shard Key: {deviceId}
  - Composite Shard Key: {deviceId, timestamp}



## Hash Sharding

- MongoDB applies a MD5 hash on the key when a hash shard key is used:
  - Hash Shard Key(deviceId) = MD5(deviceId)
  - Ensures data is distributed randomly within the range of MD5 values



## Tag Sharding

Tag-aware sharding allows subset of shards to be tagged, and assigned to a sub-range of the shard-key.

Example: Sharding User Data belong to users from 100 “regions”

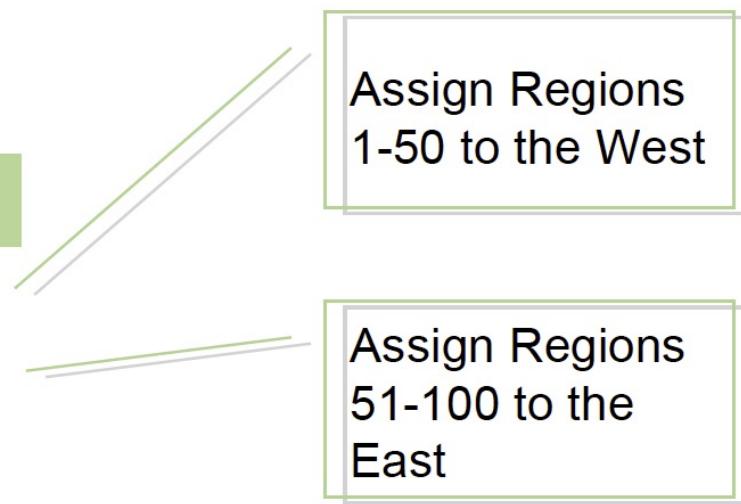
Collection: Users, Shard Key: {uld, regionCode}

Tag based on macro regions

## Tag Sharding Example

**Collection: Users, Shard Key: {uld, regionCode}**

Tag	Start	End
West	MinKey, MinKey	MaxKey, 50
East	MinKey, 50	MaxKey, MaxKey
Shard1, Tag=West	Shard2, Tag=West	Shard3, Tag=East
<b>Primary</b>	<b>Primary</b>	<b>Primary</b>
<b>Secondary</b>	<b>Secondary</b>	<b>Secondary</b>
<b>Secondary</b>	<b>Secondary</b>	<b>Secondary</b>



## Which Sharding to use?

### Usage

### Required Strategy

Scale

Range or Hash

Geo-Locality

Tag-aware

Hardware Optimization

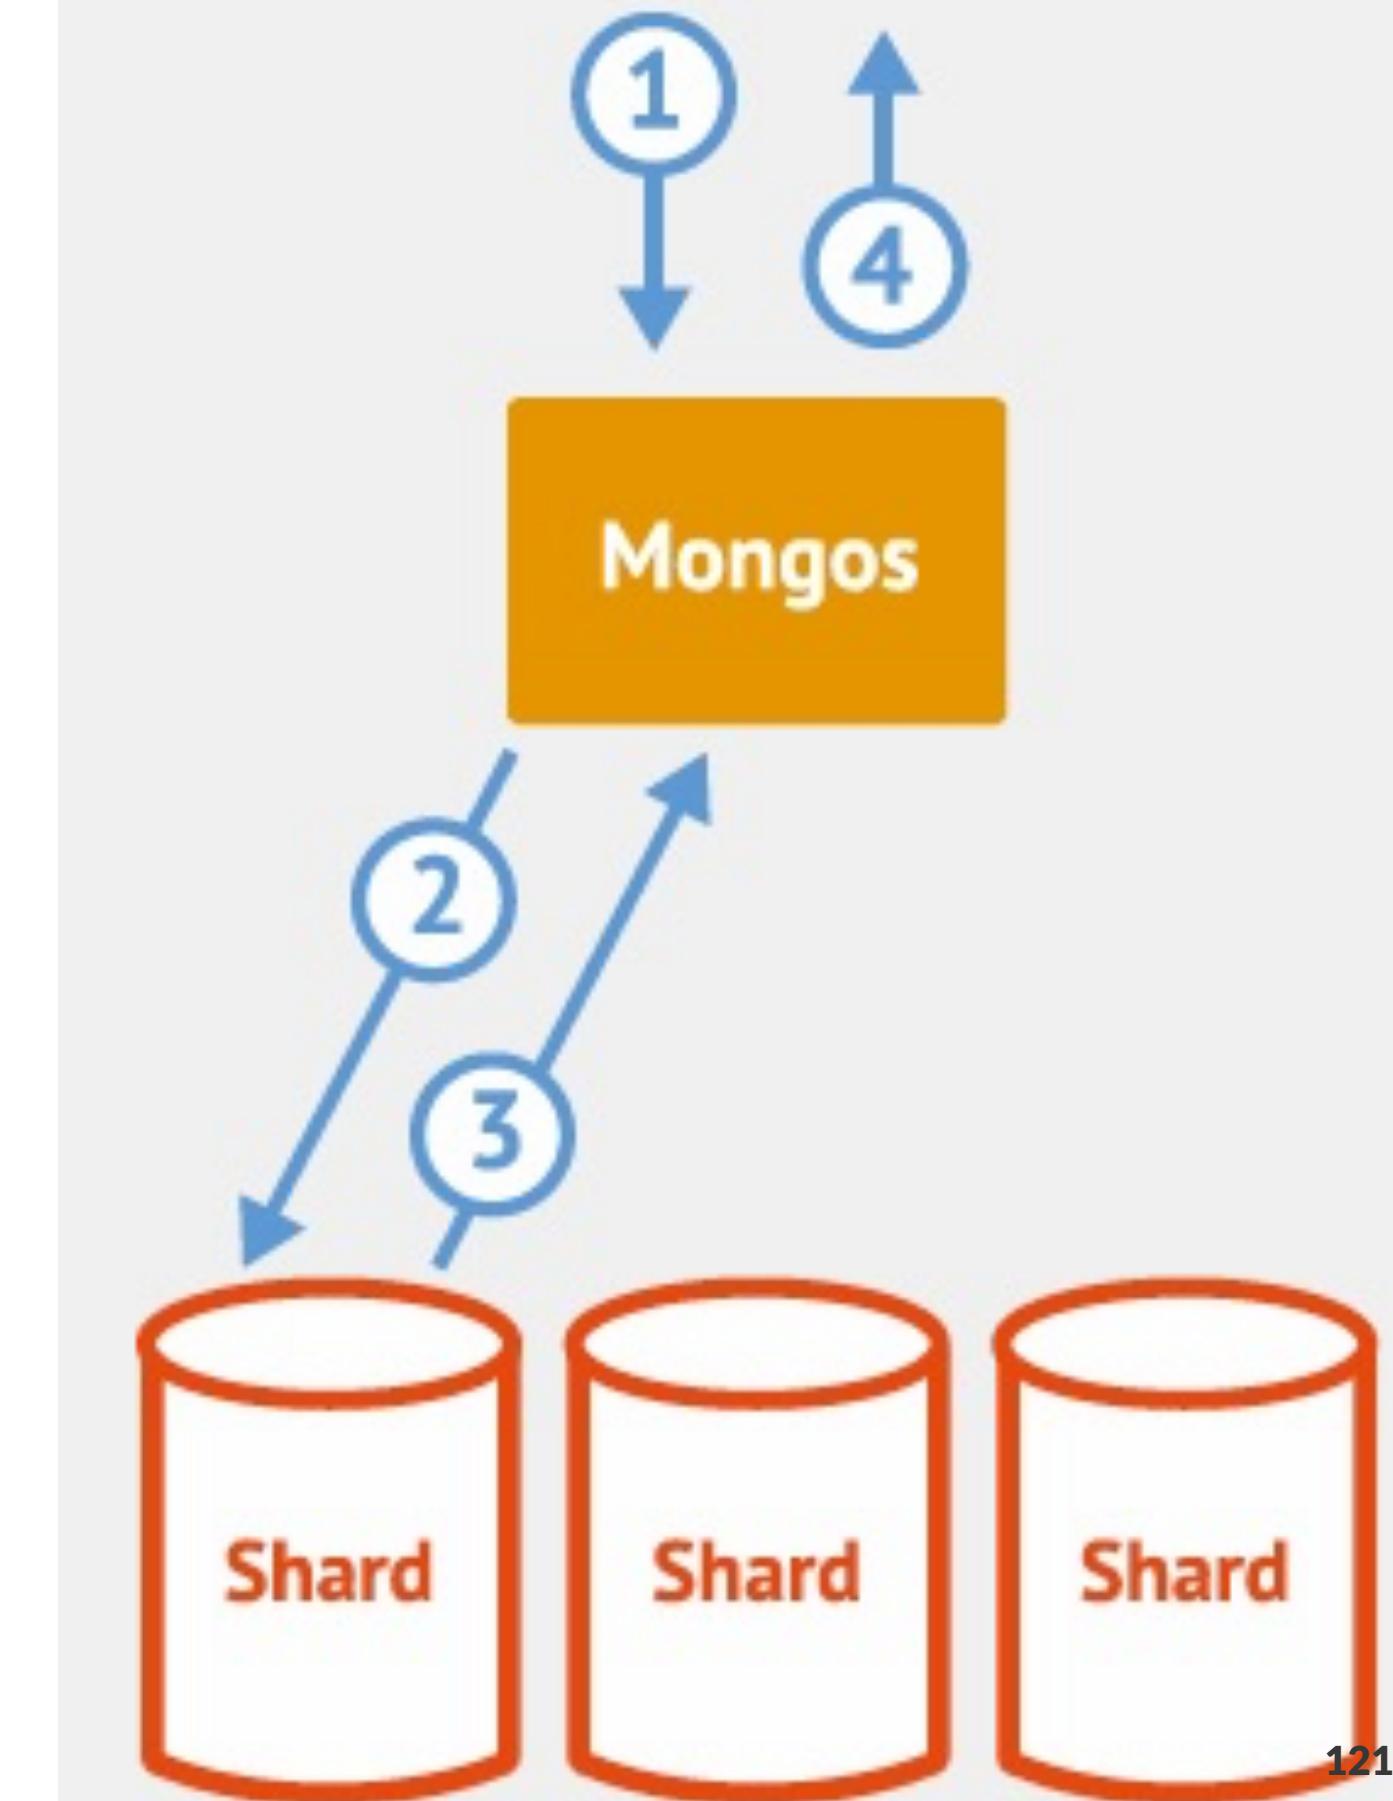
Tag-aware

Lower Recovery Times

Range or Hash

## Routing and Balancing

- Queries routed to specific shards
- MongoDB balances cluster
- MongoDB migrates data to new nodes



# MongoDB Security

- SSL
  - between client and server
  - Intra-cluster communication
- Authorization at the database level
  - Read Only/Read+Write/Administrator



## References

- [Mongodb.com](#)
- No SQL Distilled by P. Sadalage and M. Fowler
- MongoDB Applied Design Patterns by R. Copeland
- The Definitive Guide to MongoDB by Plugge, Membry and Hawkins

