



Data Engineering

LTAT.02.007

Ass Prof. Riccardo Tommasini
Fabiano Spiga, Hassan Eldeeb, Mohamed Ragab

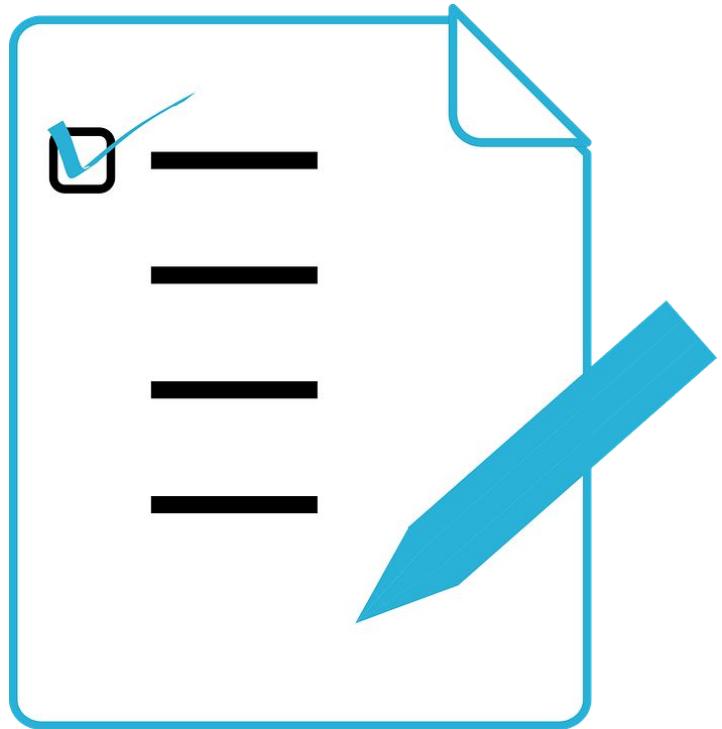
Key-Value DBS/Stores (**REDIS**)

Lab. 03



Lab Agenda

- SQL is Great but...
- NOSQL DBS
- KEY-VALUE Stores
- REDIS
 - What is REDIS?
 - Why We need Redis?
 - REDIS Caching
 - REDIS PUB/SUB
- REDIS Data-Types

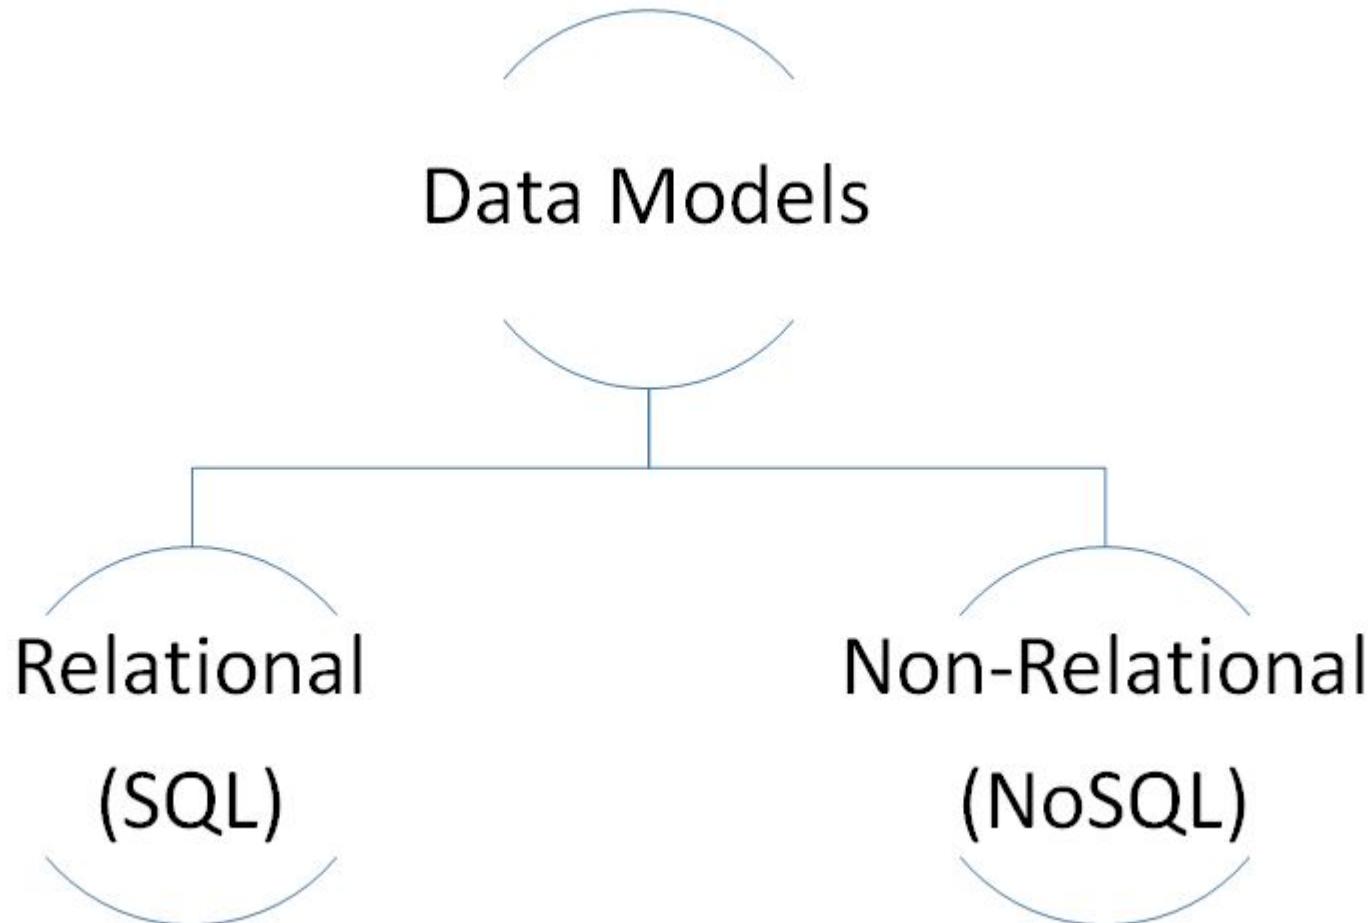


SQL is great but...

- Transactions, Integration, Reporting,..etc.
- Sometimes SQL is not the fit for what we are doing.
- One Size can't fit all!
- The data isn't always relational/tabular.
- With the rise of the WWW,
Semi-structured or unstructured data
can't fit in tables.



Data Models



What NoSQL means?

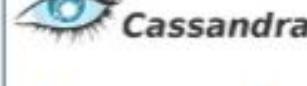
- Refers to "not-only-SQL" or "non-relational" database.
- NOSQL refer to databases that **don't support** the **relational** model and **not use SQL**.
- These non-relational DBMS are designed to solve problems where relational databases are not **adequate**.



NoSQL

Overview – NoSQL Family

- Data stored in **4 types**:
 - Key-Value pairs
 - Document (JSON, XML,...)
 - Graph
 - Wide-column

Document Database	Graph Databases
 Couchbase  mongoDB	 Neo4j  InfiniteGraph The Distributed Graph Database
Key-value Databases	Wide Column Stores
 redis  AEROSPIKE  riak	 ACCUMULO HYPERTABLE  Cassandra  APACHE HBASE Amazon SimpleDB

RDBs(SQL) VS. NoSQL (Key Differences)

- **Data Models and Schema**
 - Structured Tabular VS. unstructured data models .
 - Fixed/Predefined VS. dynamic schema
- **Data Structures**
 - Table-based VS. document-based, graph, key-value, wide columns.
- **Scaling**
 - Vertically scalable, typically expensive --->
 - Horizontally scalable, much cheaper
- **Development Model**
 - Typically closed-source with licencing fees VS. Open- Source community.



https://www.pinterest.com/pin/161355599127763483/?nic_v1=1a3Uvq1cr

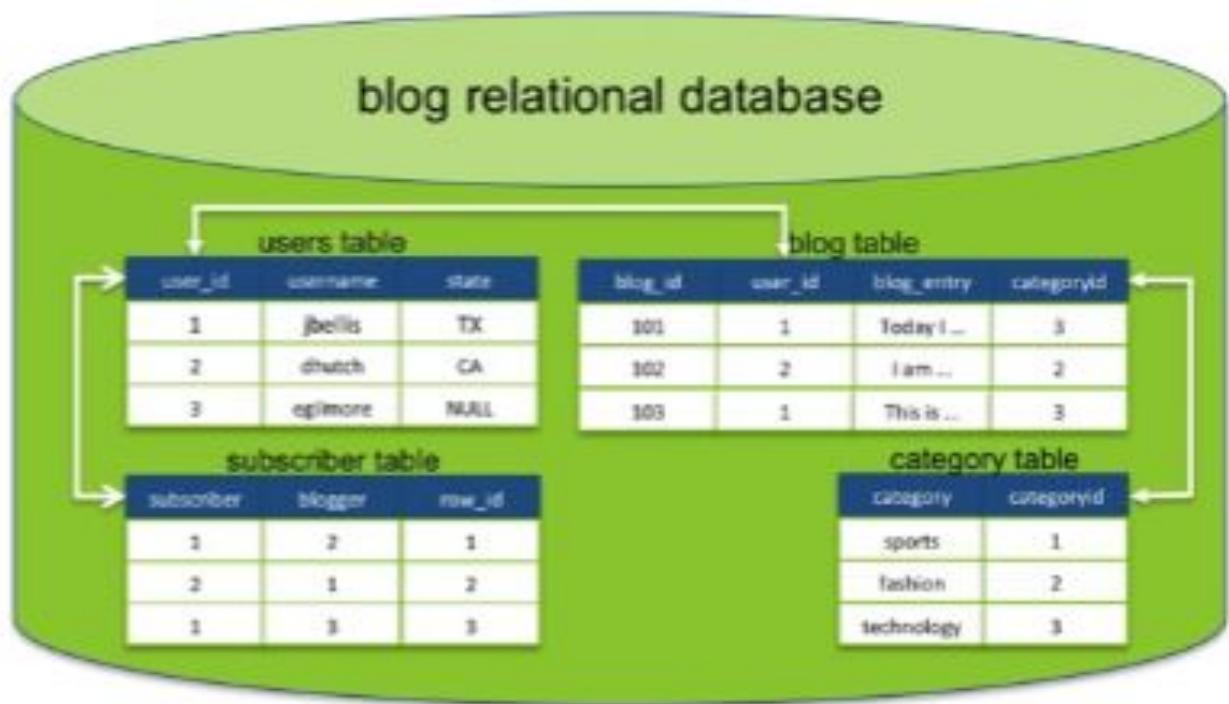
Key-Value Databases

Why Key-Value Stores?

- Essentially Key-Value stores **assign a key to a value**.
- Examples:
 - (Twitter.com) tweet_id → information about the tweet.
 - (amazon.com) item_number → information about the item
 - (yourbank.com) account_number → account details
 - ...etc.

Isn't that just a database?!

- Queried using SQL
- Key-based
- Foreign keys
- Indexes
- Joins



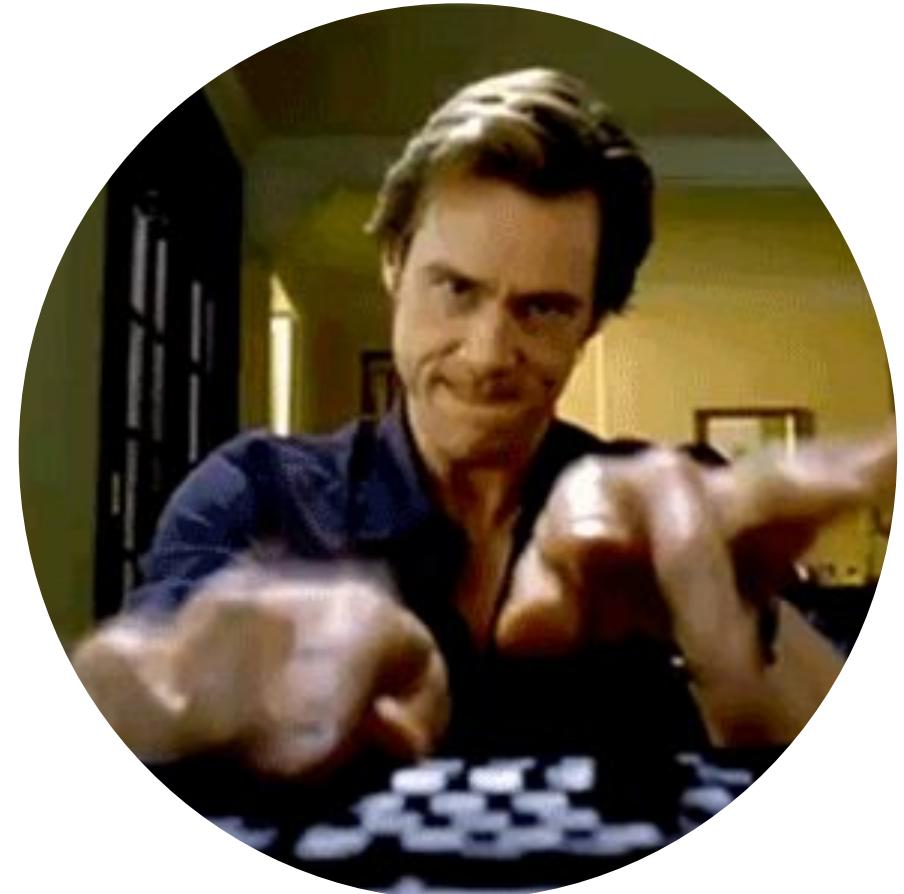
SQL queries: `SELECT user_id from users
WHERE
username = "jbellis"`

Mismatch with Today's Workloads

- **Data: Large and Unstructured**
 - Sometimes there's no predefined schema
- **Lots of random reads and writes**
 - RDBS are not optimized for **millions of concurrent** reads/writes.
- **Sometimes write-heavy**
 - RDBs are typically optimized for **read-heavy** workloads
- **Foreign-keys rarely needed, Joins infrequent as well**
 - Sometimes what we need is **much simpler** than this.

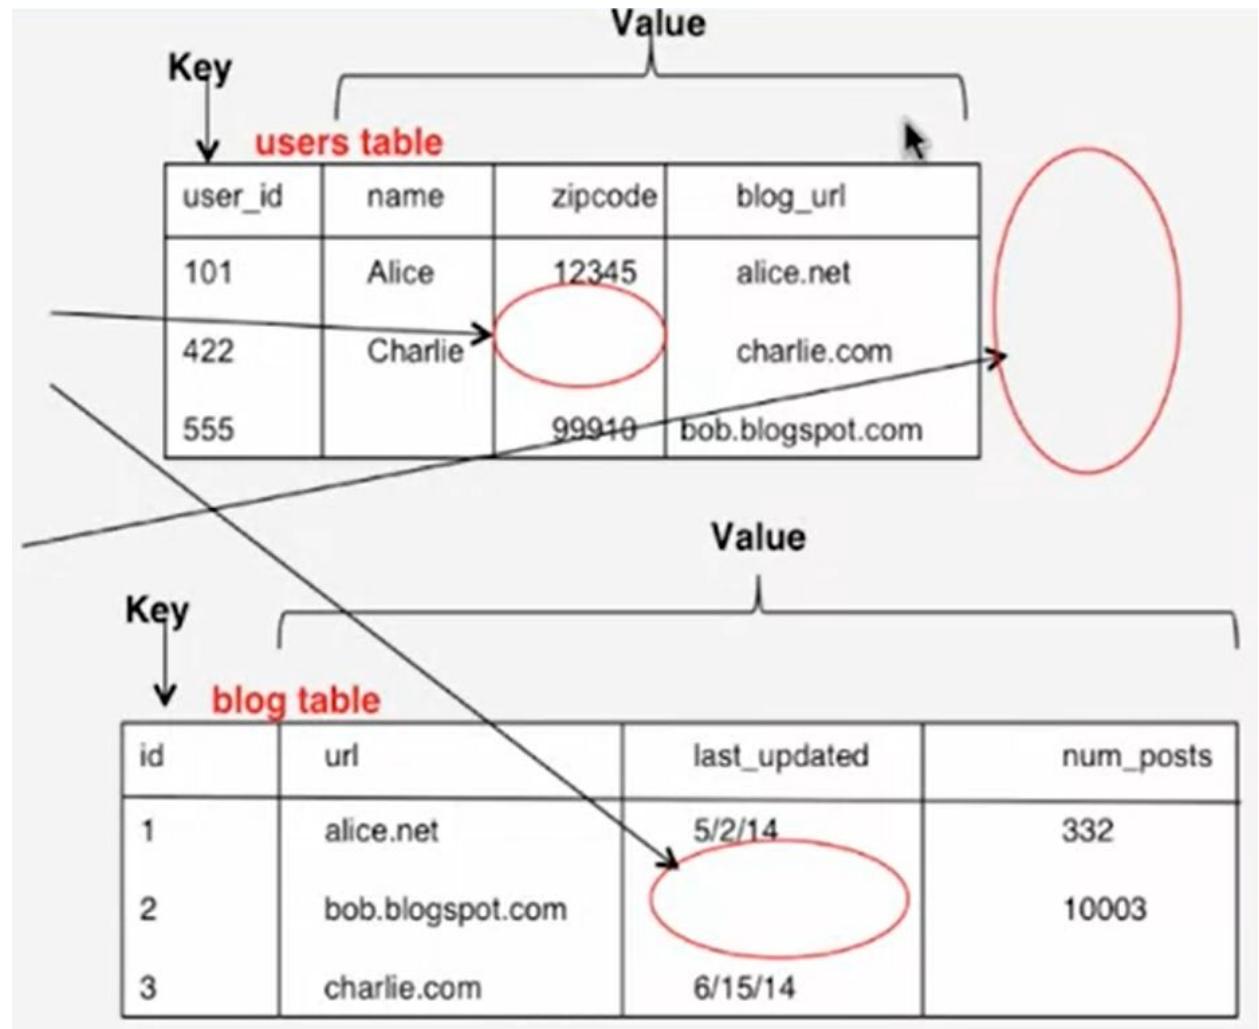
Needs of Today's Workloads

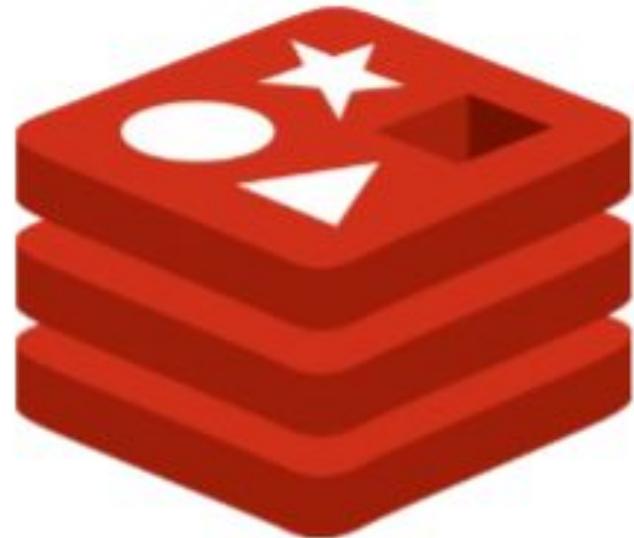
- SPEED
- Avoid Single Point of Failure (SPoF)
- Low TCO (Total Cost of Operation)
- Fewer System administrators
- Incremental Scalability
- Scale out, not up
 - What?



What is Key-Value Store?

- Unstructured
- No Schema imposed
- Columns missing from some rows
- No foreign keys
- Joins may not be supported as well





redis
REmote DIctionary
Server

What is Redis ?

- ✓ Redis, standing for **REmote DIctionary Server**,
- ✓ **Open Source, in-memory data structure store** which can be used as a database and/or cache and message broker.

Redis can persist data to the disk

Redis is fast

Redis is not only a key-value store

- ✓ Supports **multiple** data structures (Strings, Hashes, Lists, Sets,...)
- ✓ Built-in “**Master-Slave**” Replication.

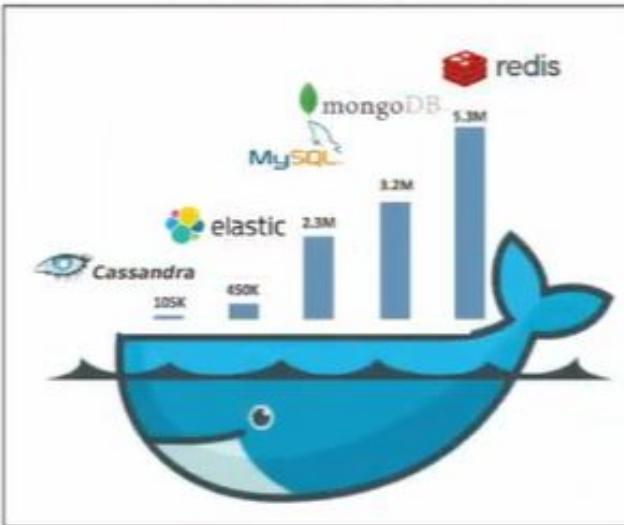
Why We need/use Redis?

- FAST
- NOT CPU Intensive
- Well Documented
- Simple and Versatile
- Scalable



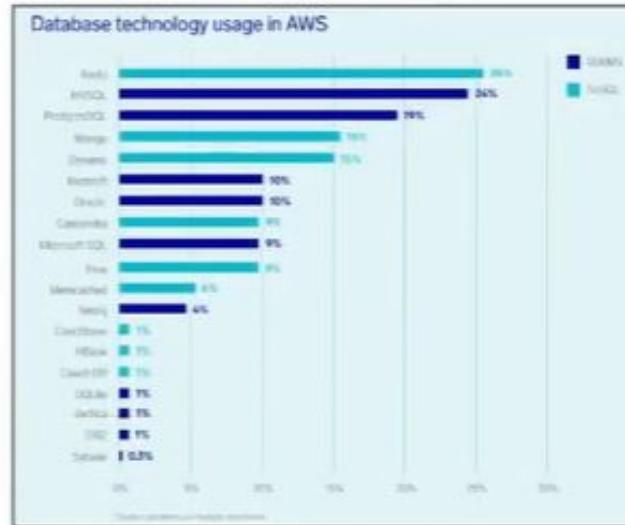
HOW Popular is REDIS?

MOST LAUNCHED



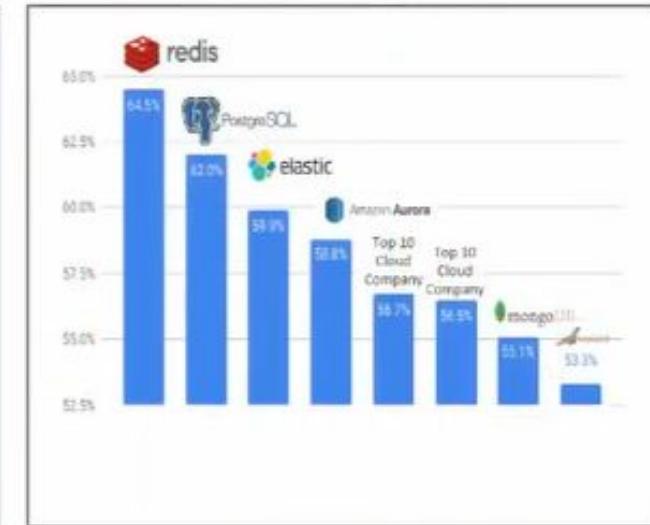
Launches/day: Docker Hub, Nov. 2019

MOST USED



Sumo Logic, Sept. 2019

MOST LOVED



Stack Overflow, 2017, 2018, 2019

Modern Models of REDIS



Streams



RediSearch



RedisGraph



RedisTimeSeries



RedisBloom

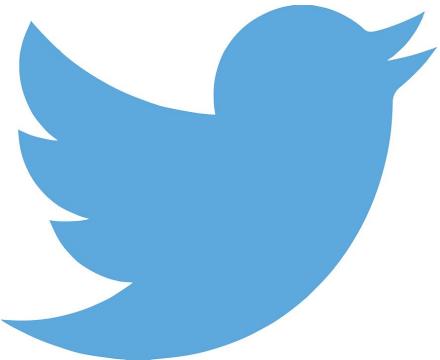


RedisJSON



RedisAI

Who Uses REDIS?



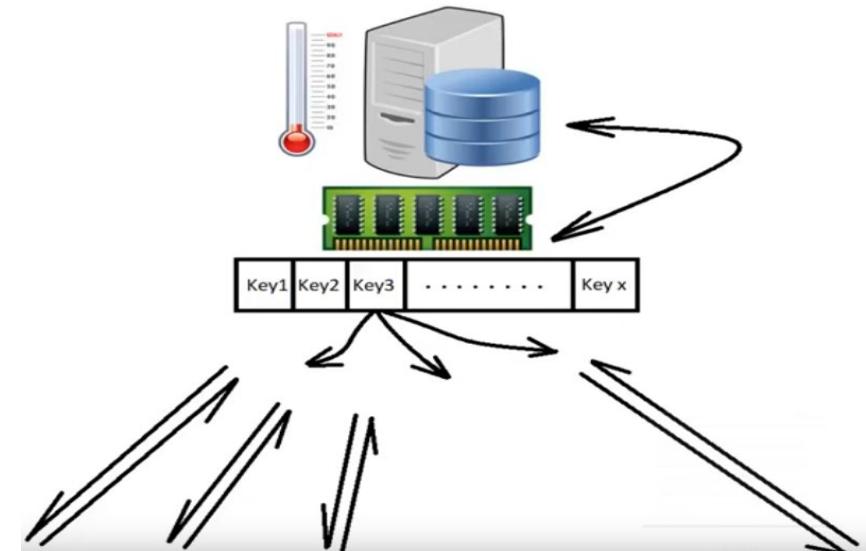
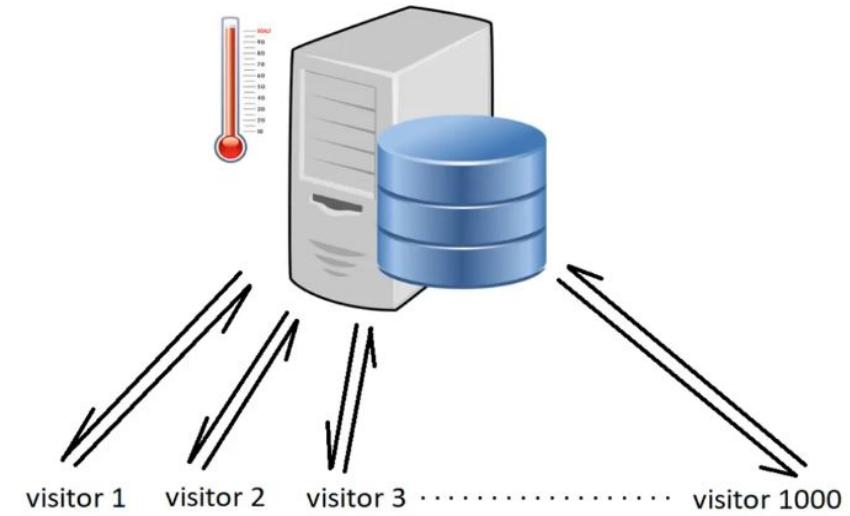
REDIS Common Use Cases

- Caching
- Analytics (e.g. Counting things)
- Leaderboard (voting applications)
- Messaging & Pub/Sub
- Queues
- Cookie storage & **Epiring/Vanishing** Data
- Distributing Locks
- High I/O Workloads



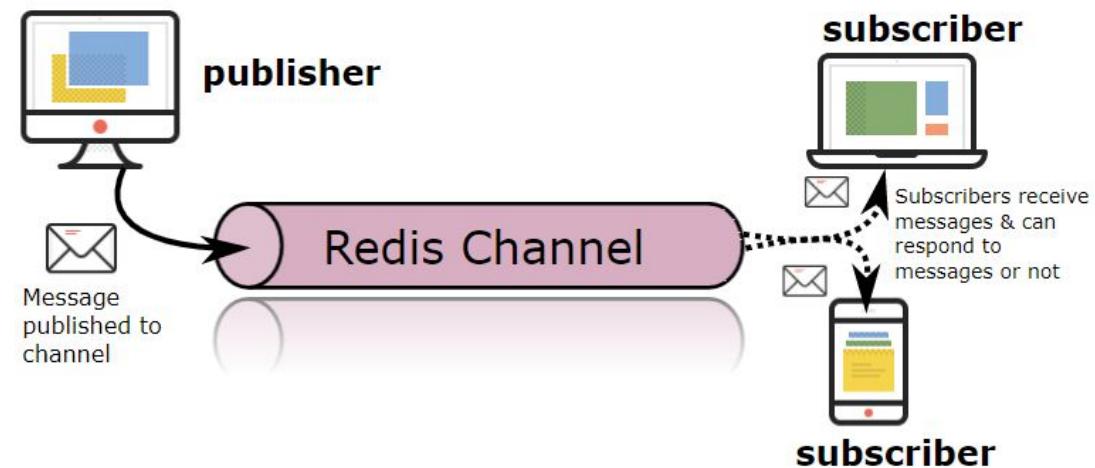
Client Side Caching

- A technique used to create high performance services.
- Normally, when some data is required, the application servers will ask the database about such information.
- When client side caching is used,
 - The application will store the reply of popular queries **directly inside the application memory(in REDIS)**.
 - so that it can reuse such replies later, without contacting the database again.



PUB/SUB in REDIS

- It is mainly characterized by listeners subscribing to channels, with publishers sending **binary string messages to channels**.
- The communication is processed by the **broker**, it helps publisher to deliver that information to the subscriber(s).



Other Caching Tools/DBs



<https://memcached.org/>
<https://varnish-cache.org/>

When to Consider REDIS?

- ▶ Speed is critical
- ▶ More than just key-value pairs
- ▶ Dataset can fit in memory
- ▶ Dataset is not critical



How to access REDSI (Redis Interfaces)

- Try REDIS Online Service
 - [This link](https://redis.io/topics/quickstart)
- Using the terminal / any other CLIs.
- Using Redis from Applications
(python,Ruby,...)



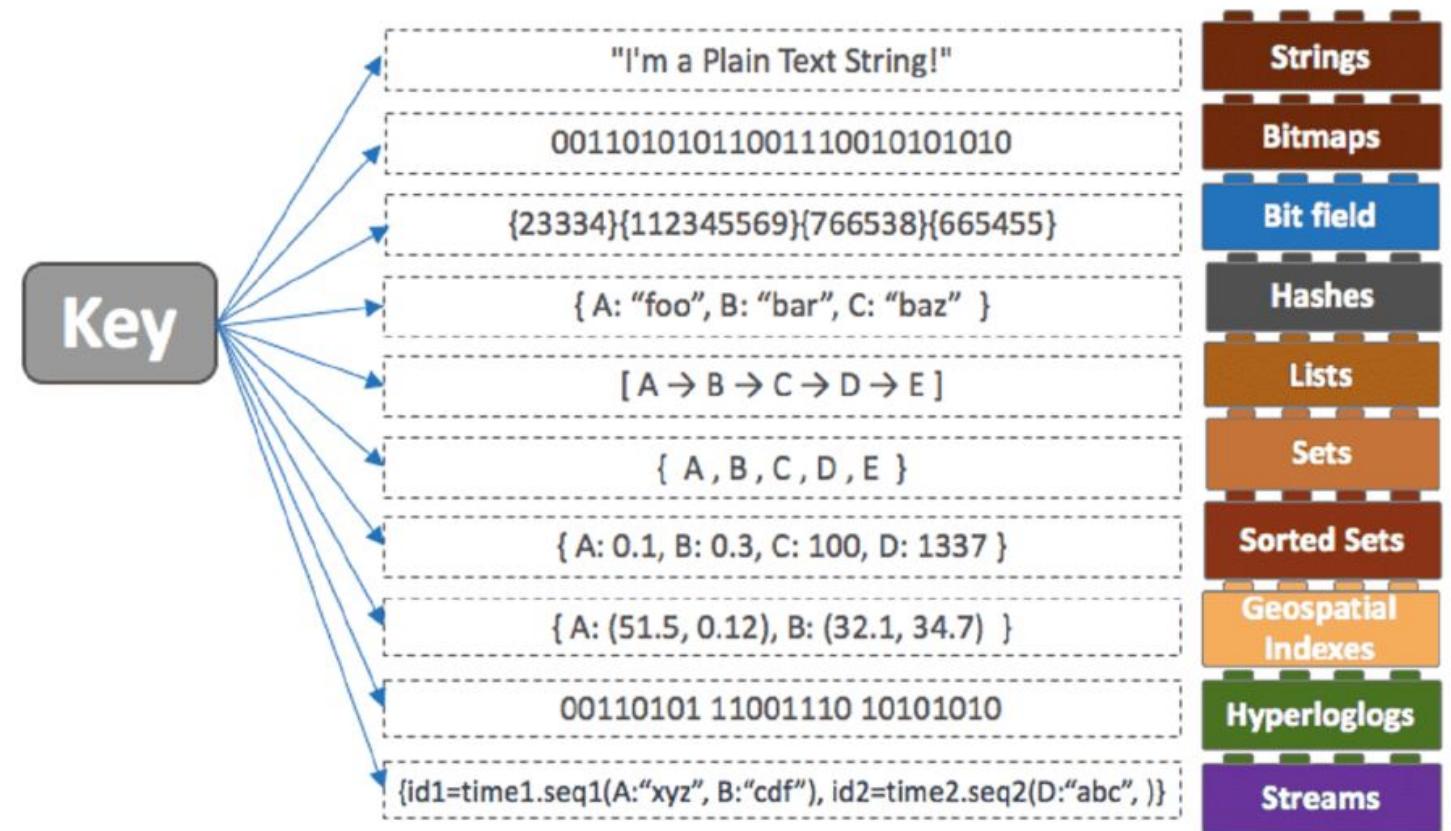
```
C:\Users\Mohamed Ragab>redis-server
[11120] 16 Sep 21:05:45.041 # Warning: no config file specified, using the de
file use redis-server /path/to/redis.conf

Redis 3.0.504 (00000000/0) 64 bit
Running in standalone mode
Port: 6379
PID: 11120

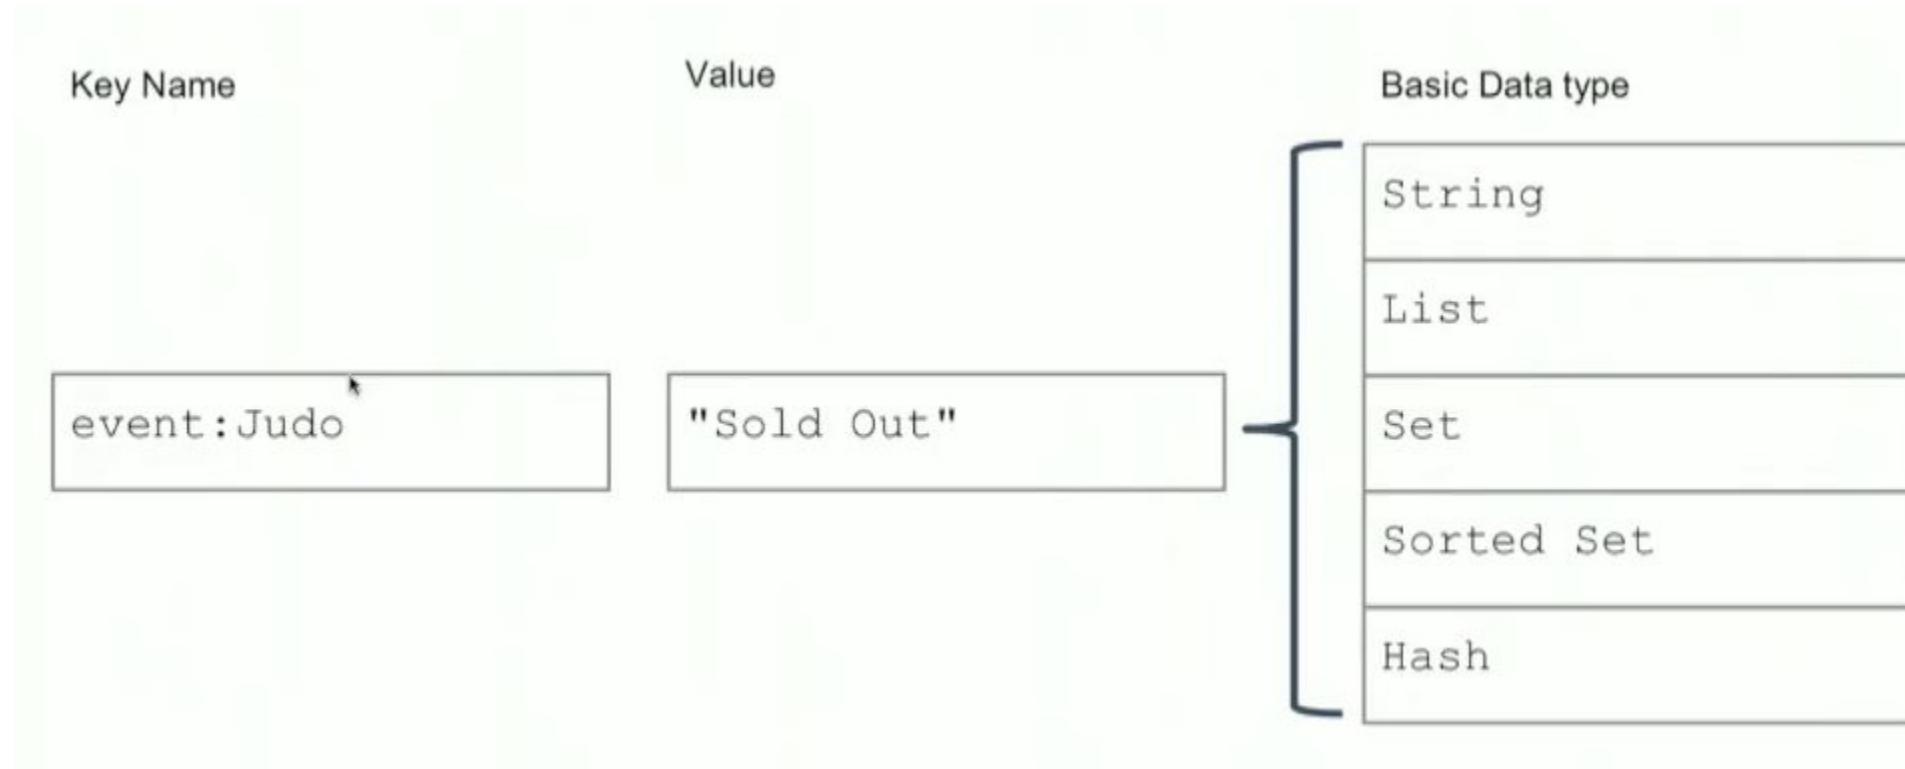
http://redis.io
```

```
C:\Users\Mohamed Ragab>redis-cli
127.0.0.1:6379> PING
PONG
127.0.0.1:6379>
```

REDIS data types



KEYs, VALUEs && Data Types



REDIS KEYS

Key Name

event:Judo

Keys

- Unique
- Binary Safe
- Up to 512MB in size

REDIS KEY Naming

- Up to you!
- Be consistent.
- Colon is a typical community convention
 - Domain Object Name
 - Colon
 - Unique ID
- Plurals are nice for Lists, Sets
- Case Sensitive

Key Name

event:Judo

event:venues

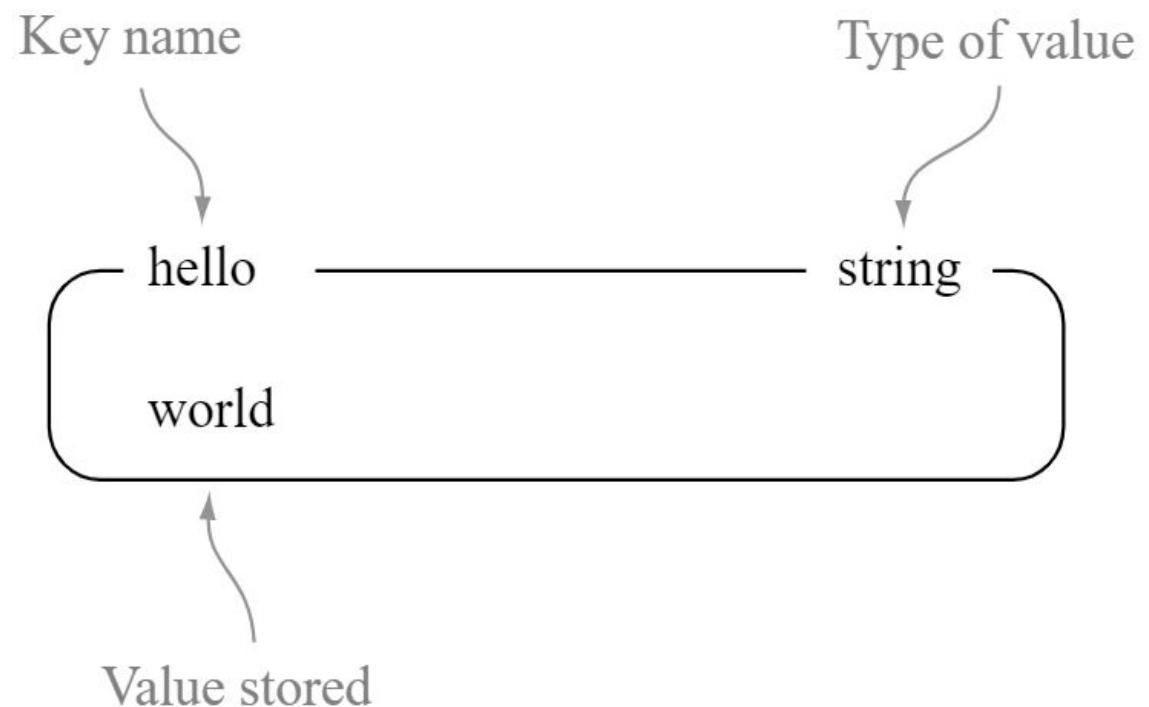
event:Fencing

event:fencing

Event:fencing

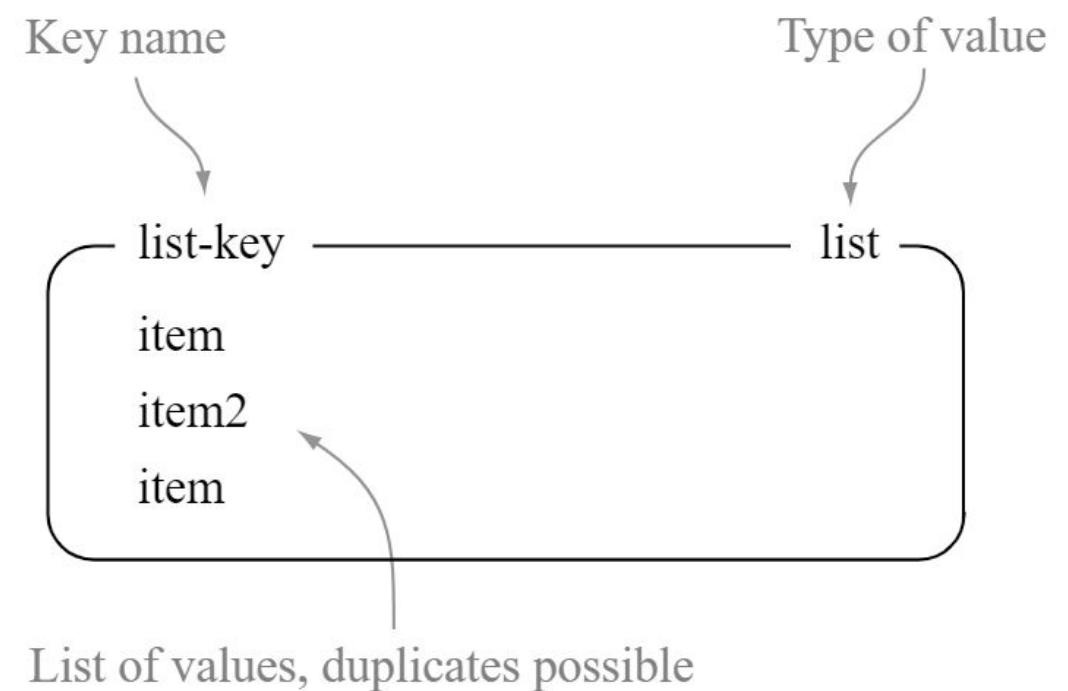
STRINGS

- The most important type.
- Store Strings, numerical values, serialized JSON, and binary.
 - **Serialization, and deserialization.**
- Commonly used for caching when combined with the **EXPIRE** command.
- Strings can handle up to size **512MB**.



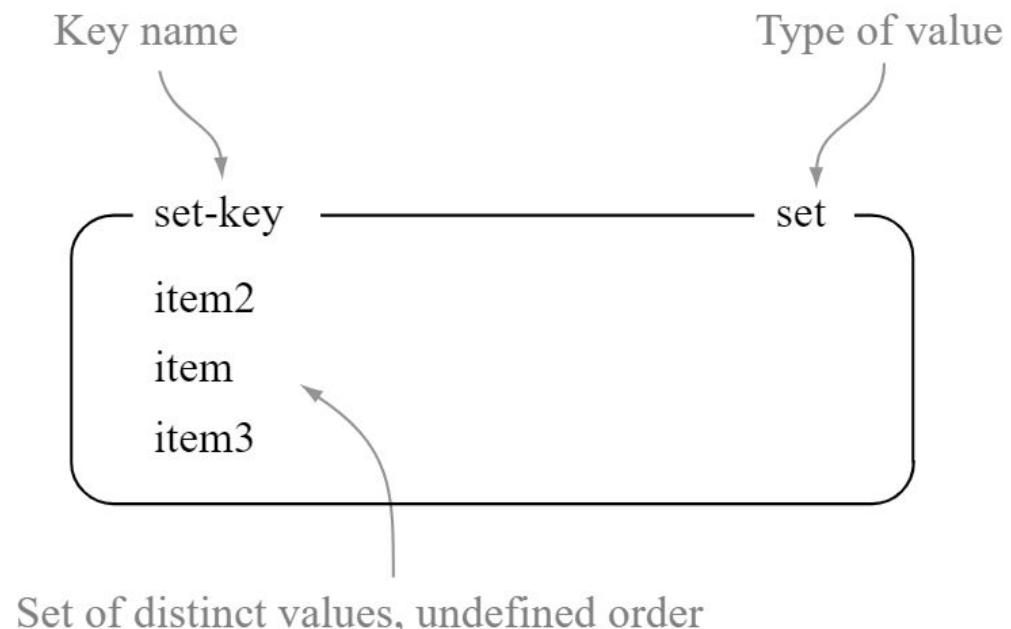
LISTS

- Think of Lists as an ordered sequence of strings like java ArrayList, Javascript array, or python n lists.
- We can use lists to implement stacks and queues.
- Lists accept duplicates.



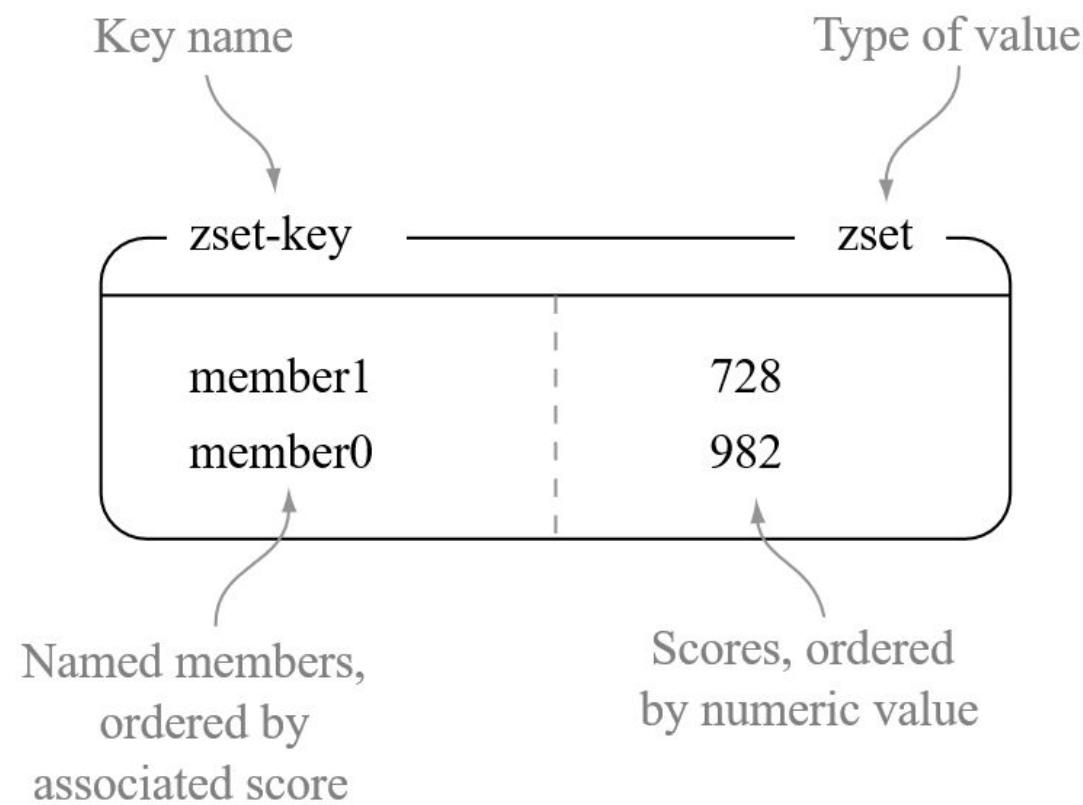
SETS

- Unordered collection of strings
- Contains no duplicates.
 - This makes Sets supernatural option for de-duplication applications.
- Questions that we can Answer using Sets:
 - Did I see this IP address in the last hour?
 - Is this user online?
 - Has this URL been blacklisted?
- Sets support standard operations:
 - Intersection
 - Difference
 - Union



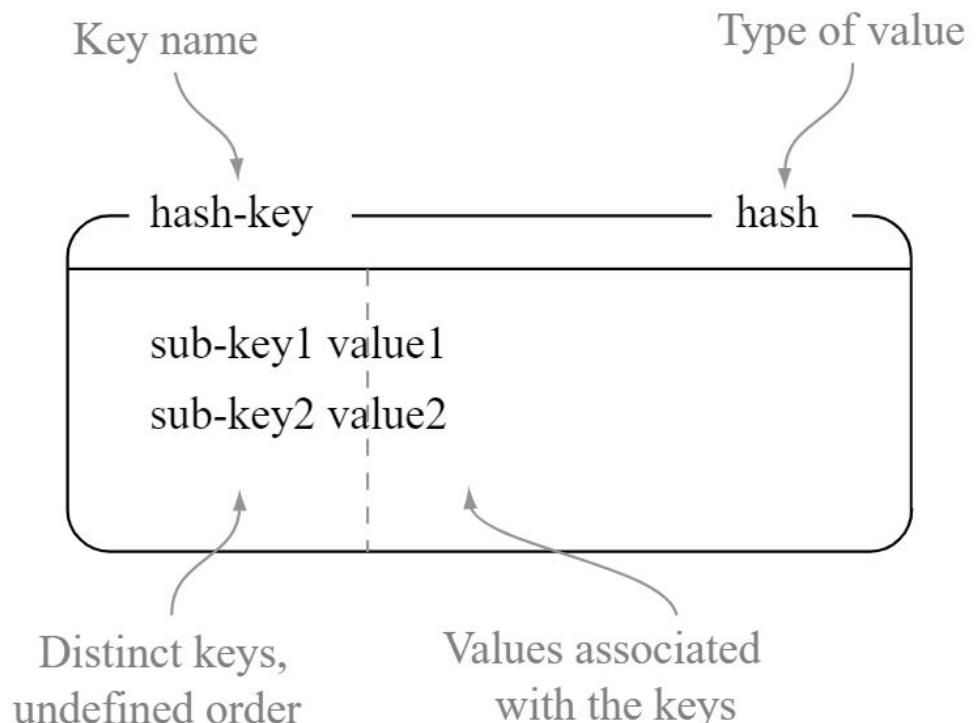
SORTED SETS

- Ordered collections of unique members.
- These members are ordered according to their associated score.
- Whenever you add to the sorted set, you are specifying a member and a score.
- Sorted Sets keep everything sorted from the beginning.
- Sorted sets are good choice for:
 - priority queues, Low-latency leaderboards, Secondary indexing
- Questions that we can Answer using Sorted Sets (e.g, in an online-game):
 - Who are the top 10 players?
 - what is the rank of a specific Player?
 - what is the current score of the player?



HASHES

- Collections of Keys-Value pairs.
- Similar to JSON object, Java HashMAPs, and Python Dictionary.
- They are Mutable:
 - ADD,CHANGE,INCREMENT,REMOVE field value-pairs anytime.
- Hashes are Schema-less.



Now, It's time to say ...

THANK YOU

**SEE YOU
NEXT TIME!**